

Minimalist Data Wrangling with Python

Marek Gagolewski

Prof. **Marek Gagolewski**

Warsaw University of Technology, Poland

Systems Research Institute, Polish Academy of Sciences

<https://www.gagolewski.com/>

Copyright (C) 2022–2026 by Marek Gagolewski. Some rights reserved.

This open-access textbook is an independent, non-profit project. It is published under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). Please spread the word about it.

This project received no funding, administrative, technical, or editorial support from Deakin University, Warsaw University of Technology, Polish Academy of Sciences, nor any other source.

A little peculiar is the world some people decided to immerse themselves in, so here is a message stating the obvious. Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is provided without warranty, either express or implied. The author will, of course, not be held liable for any damages caused or alleged to be caused directly or indirectly by this book.

Any bug reports/corrections/feature requests are welcome. To make this textbook even better, please file them at <https://github.com/gagolews/datawranglingpy>.

Typeset with Xe_{La}TeX. Please be understanding: it was an algorithmic process. Hence, the results are $\in$ [good enough, perfect).

Homepage: <https://datawranglingpy.gagolewski.com/>

Datasets: <https://github.com/gagolews/teaching-data>

Release: v1.1.0.9003 (2026-02-17T11:40:45UTC)

ISBN: 978-0-6455719-1-2 (v1.0; 2022; Melbourne: Marek Gagolewski)

DOI: [10.5281/zenodo.6451068](https://doi.org/10.5281/zenodo.6451068)

Contents

Preface	xiii
0.1 The art of data wrangling	xiii
0.2 Aims, scope, and design philosophy	xiv
0.2.1 We need maths	xv
0.2.2 We need some computing environment	xv
0.2.3 We need data and domain knowledge	xvi
0.3 Structure	xvii
0.4 The Rules	xix
0.5 About the author	xxi
0.6 Acknowledgements	xxi
0.7 You can make this book better	xxii
 I Introducing Python	 1
1 Getting started with Python	3
1.1 Installing Python	3
1.2 Working with Jupyter notebooks	4
1.2.1 Launching JupyterLab	5
1.2.2 First notebook	5
1.2.3 More cells	6
1.2.4 Edit vs command mode	7
1.2.5 Markdown cells	8
1.3 The best note-taking app	10
1.4 Initialising each session and getting example data	10
1.5 Exercises	12
 2 Scalar types and control structures in Python	 13
2.1 Scalar types	13
2.1.1 Logical values	13
2.1.2 Numeric values	13
Arithmetic operators	14
Creating named variables	15
2.1.3 Character strings	15
F-strings (formatted string literals)	16
2.2 Calling built-in functions	17
2.2.1 Positional and keyword arguments	17
2.2.2 Modules and packages	18

2.2.3	Slots and methods	18
2.3	Controlling program flow	19
2.3.1	Relational and logical operators	19
2.3.2	The <code>if</code> statement	20
2.3.3	The <code>while</code> loop	21
2.4	Defining functions	22
2.4.1	Lambda expressions	23
2.4.2	(*) Own modules	24
2.5	Exercises	24
3	Sequential and other types in Python	25
3.1	Sequential types	25
3.1.1	Lists	25
3.1.2	Tuples	26
3.1.3	Ranges	26
3.1.4	Strings (again)	27
3.2	Working with sequences	27
3.2.1	Extracting elements	27
3.2.2	Slicing	28
3.2.3	Modifying elements of mutable sequences	29
3.2.4	Searching for specific elements	30
3.2.5	Arithmetic operators	30
3.3	Dictionaries	30
3.4	Iterable types	32
3.4.1	The <code>for</code> loop	32
3.4.2	Tuple assignment	34
3.4.3	Argument unpacking (*)	36
3.4.4	Variadic arguments: <code>*args</code> and <code>**kwargs</code> (*)	37
3.5	Object references and copying (*)	37
3.5.1	Copying references	37
3.5.2	Pass by assignment	38
3.5.3	Object copies	38
3.5.4	Modify in place or return a modified copy?	39
3.6	Further reading	40
3.7	Exercises	41
II	Unidimensional data	43
4	Unidimensional numeric data and their empirical distribution	45
4.1	Creating vectors in <code>numpy</code>	46
4.1.1	Enumerating elements	47
4.1.2	Arithmetic progressions	48
4.1.3	Repeating values	49
4.1.4	<code>numpy.r_</code> (*)	49
4.1.5	Generating pseudorandom variates	50
4.1.6	Loading data from files	50
4.2	Some mathematical notation	51

4.3	Inspecting the data distribution with histograms	52
4.3.1	heights: A bell-shaped distribution	52
4.3.2	income: A right-skewed distribution	53
4.3.3	How many bins?	55
4.3.4	peds: A bimodal distribution (already binned)	57
4.3.5	matura: A bell-shaped distribution (almost)	59
4.3.6	marathon (truncated – fastest runners): A left-skewed distribution	59
4.3.7	Log-scale and heavy-tailed distributions	60
4.3.8	Cumulative probabilities and the empirical cumulative distribution function	63
4.4	Exercises	64
5	Processing unidimensional data	67
5.1	Aggregating numeric data	67
5.1.1	Measures of location	68
	Arithmetic mean and median	68
	Sensitive to outliers vs robust	69
	Sample quantiles	70
5.1.2	Measures of dispersion	72
	Standard deviation (and variance)	73
	Interquartile range	73
5.1.3	Measures of shape	74
5.1.4	Box (and whisker) plots	75
5.1.5	Other aggregation methods (*)	76
5.2	Vectorised mathematical functions	78
5.2.1	Logarithms and exponential functions	79
5.2.2	Trigonometric functions	80
5.3	Arithmetic operators	81
5.3.1	Vector-scalar case	82
5.3.2	Application: Feature scaling	82
	Standardisation and z-scores	83
	Min-max scaling and clipping	84
	Normalisation (l_2 ; dividing by magnitude)	85
	Normalisation (l_1 ; dividing by sum)	86
5.3.3	Vector-vector case	86
5.4	Indexing vectors	88
5.4.1	Integer indexing	88
5.4.2	Logical indexing	89
5.4.3	Slicing	90
5.5	Other operations	91
5.5.1	Cumulative sums and iterated differences	91
5.5.2	Sorting	91
5.5.3	Dealing with tied observations	92
5.5.4	Determining the ordering permutation and ranking	93
5.5.5	Searching for certain indexes (argmin, argmax)	94
5.5.6	Dealing with round-off and measurement errors	94

5.5.7	Vectorising scalar operations with list comprehensions	96
5.6	Exercises	98

6 Continuous probability distributions 101

6.1	Normal distribution	102
6.1.1	Estimating parameters	102
6.1.2	Data models are useful	103
6.2	Assessing goodness-of-fit	105
6.2.1	Comparing cumulative distribution functions	105
6.2.2	Comparing quantiles	107
6.2.3	Kolmogorov–Smirnov test (*)	109
6.3	Other noteworthy distributions	111
6.3.1	Log-normal distribution	111
6.3.2	Pareto distribution	115
6.3.3	Uniform distribution	118
6.3.4	Distribution mixtures (*)	120
6.4	Generating pseudorandom numbers	122
6.4.1	Uniform distribution	122
6.4.2	Not exactly random	122
6.4.3	Sampling from other distributions	123
6.4.4	Natural variability	124
6.4.5	Adding jitter (white noise)	126
6.4.6	Independence assumption	127
6.5	Further reading	127
6.6	Exercises	128

III Multidimensional data 129

7 From uni- to multidimensional numeric data 131

7.1	Creating matrices	131
7.1.1	Reading CSV files	131
7.1.2	Enumerating elements	133
7.1.3	Repeating arrays	133
7.1.4	Stacking arrays	134
7.1.5	<code>numpy.r_</code> revisited (*)	135
7.1.6	Other functions	135
7.2	Reshaping matrices	136
7.3	Mathematical notation	137
7.3.1	Transpose	138
7.3.2	Row and column vectors	139
7.3.3	Identity and other diagonal matrices	140
7.4	Visualising multidimensional data	140
7.4.1	2D Data	141
7.4.2	3D data and beyond	142
7.4.3	Scatter plot matrix (pairs plot)	145
7.5	Exercises	146

8	Processing multidimensional data	149
8.1	Extending vectorised operations to matrices	149
8.1.1	Vectorised mathematical functions	149
8.1.2	Componentwise aggregation	149
8.1.3	Arithmetic, logical, and relational operations	150
	Matrix vs scalar	151
	Matrix vs matrix	151
	Matrix vs any vector	153
	Row vector vs column vector (*)	154
8.1.4	Other row and column transforms (*)	155
8.2	Indexing matrices	156
8.2.1	Slice-based indexing	157
8.2.2	Scalar-based indexing	157
8.2.3	Mixed logical/integer vector and scalar/slice indexers	158
8.2.4	Two vectors as indexers (*)	158
8.2.5	Views of existing arrays (*)	160
8.2.6	Adding and modifying rows and columns	160
8.3	Matrix multiplication, dot products, and Euclidean norm (*)	161
8.4	Pairwise distances and related methods (*)	163
8.4.1	Euclidean metric (*)	163
8.4.2	Centroids (*)	166
8.4.3	Multidimensional dispersion and other aggregates (**)	166
8.4.4	Fixed-radius and k -nearest neighbour search (**)	167
8.4.5	Spatial search with multidimensional binary search trees (**)	169
8.5	Exercises	170
9	Exploring relationships between variables	173
9.1	Measuring correlation	174
9.1.1	Pearson linear correlation coefficient	174
	Perfect linear correlation	175
	Strong linear correlation	176
	No linear correlation does not imply independence	177
	False correlations	178
	Correlation is not causation	180
9.1.2	Correlation heat map	180
9.1.3	Linear correlation coefficients on transformed data	183
9.1.4	Spearman rank correlation coefficient	184
9.2	Regression tasks (*)	185
9.2.1	K -nearest neighbour regression (*)	185
9.2.2	From data to (linear) models (*)	188
9.2.3	Least squares method (*)	189
9.2.4	Analysis of residuals (*)	192
9.2.5	Multiple regression (*)	195
9.2.6	Variable transformation and linearisable models (**)	196
9.2.7	Descriptive vs predictive power (**)	198
9.2.8	Fitting regression models with <code>scikit-learn</code> (*)	204
9.2.9	Ill-conditioned model matrices (**)	205

9.3	Finding interesting combinations of variables (*)	209
9.3.1	Dot products, angles, collinearity, and orthogonality (*)	209
9.3.2	Geometric transformations of points (*)	210
9.3.3	Matrix inverse (*)	213
9.3.4	Singular value decomposition (*)	214
9.3.5	Dimensionality reduction with SVD (*)	215
9.3.6	Principal component analysis (*)	219
9.4	Further reading	221
9.5	Exercises	222

IV Heterogeneous data 225

10 Introducing data frames 227

10.1	Creating data frames	228
10.1.1	Data frames are matrix-like	229
10.1.2	Series	230
10.1.3	Index	232
10.2	Aggregating data frames	235
10.3	Transforming data frames	237
10.4	Indexing Series objects	239
10.4.1	Do not use [...] directly (in the current version of pandas)	240
10.4.2	loc[...]	241
10.4.3	iloc[...]	242
10.4.4	Logical indexing	243
10.5	Indexing data frames	243
10.5.1	loc[...] and iloc[...]	243
10.5.2	Adding rows and columns	245
10.5.3	Modifying items	246
10.5.4	Pseudorandom sampling and splitting	246
10.5.5	Hierarchical indexes (*)	248
10.6	Further operations on data frames	250
10.6.1	Sorting	250
10.6.2	Stacking and unstacking (long/tall and wide forms)	254
10.6.3	Joining (merging)	255
10.6.4	Set-theoretic operations and removing duplicates	257
10.6.5	...and (too) many more	259
10.7	Exercises	260

11 Handling categorical data 261

11.1	Representing and generating categorical data	261
11.1.1	Encoding and decoding factors	262
11.1.2	Binary data as logical and probability vectors	264
11.1.3	One-hot encoding (*)	265
11.1.4	Binning numeric data (revisited)	266
11.1.5	Generating pseudorandom labels	268
11.2	Frequency distributions	268
11.2.1	Counting	268

11.2.2	Two-way contingency tables: Factor combinations	269
11.2.3	Combinations of even more factors	270
11.3	Visualising factors	272
11.3.1	Bar plots	272
11.3.2	Political marketing and statistics	274
11.3.3		275
11.3.4	Pareto charts (*)	275
11.3.5	Heat maps	278
11.4	Aggregating and comparing factors	279
11.4.1	Mode	279
11.4.2	Binary data as logical vectors	280
11.4.3	Pearson chi-squared test (*)	281
11.4.4	Two-sample Pearson chi-squared test (*)	282
11.4.5	Measuring association (*)	284
11.4.6	Binned numeric data	286
11.4.7	Ordinal data (*)	286
11.5	Exercises	287
12	Processing data in groups	289
12.1	Basic methods	290
12.1.1	Aggregating data in groups	292
12.1.2	Transforming data in groups	293
12.1.3	Manual splitting into subgroups (*)	294
12.2	Plotting data in groups	295
12.2.1	Series of box plots	296
12.2.2	Series of bar plots	296
12.2.3	Semitransparent histograms	297
12.2.4	Scatter plots with group information	297
12.2.5	Grid (trellis) plots	299
12.2.6	Kolmogorov–Smirnov test for comparing ECDFs (*)	300
12.2.7	Comparing quantiles (*)	302
12.3	Classification tasks (*)	303
12.3.1	K -nearest neighbour classification (*)	305
12.3.2	Assessing prediction quality (*)	308
12.3.3	Splitting into training and test sets (*)	311
12.3.4	Validating many models (parameter selection) (**)	311
12.4	Clustering tasks (*)	313
12.4.1	K -means method (*)	313
12.4.2	Solving k -means is hard (*)	316
12.4.3	Lloyd algorithm (*)	317
12.4.4	Local minima (*)	318
12.4.5	Random restarts (*)	320
12.5	Further reading	324
12.6	Exercises	324
13	Accessing databases	325
13.1	Example database	325

13.2	Exporting data to a database	327
13.3	Exercises on SQL vs pandas	328
13.3.1	Filtering	329
13.3.2	Ordering	330
13.3.3	Removing duplicates	331
13.3.4	Grouping and aggregating	332
13.3.5	Joining	333
13.3.6	Solutions to exercises	334
13.4	Closing the database connection	337
13.5	Common data serialisation formats for the Web	338
13.6	Working with many files	339
13.6.1	File paths	339
13.6.2	File search	340
13.6.3	Exception handling	340
13.6.4	File connections (*)	340
13.7	Further reading	341
13.8	Exercises	341

V Other data types 343

14 Text data 345

14.1	Basic string operations	345
14.1.1	Unicode as the universal encoding	346
14.1.2	Normalising strings	346
14.1.3	Substring searching and replacing	347
14.1.4	Locale-aware services in ICU (*)	348
14.1.5	String operations in pandas	349
14.1.6	String operations in numpy (*)	351
14.2	Working with string lists	353
14.3	Formatted outputs for reproducible report generation	354
14.3.1	Formatting strings	355
14.3.2	str and repr	355
14.3.3	Aligning strings	355
14.3.4	Direct Markdown output in Jupyter	356
14.3.5	Manual Markdown file output (*)	356
14.4	Regular expressions (*)	358
14.4.1	Regex matching with re (*)	358
14.4.2	Regex matching with pandas (*)	360
14.4.3	Matching individual characters (*)	361
	Matching anything (almost) (*)	362
	Defining character sets (*)	362
	Complementing sets (*)	363
	Defining code point ranges (*)	363
	Using predefined character sets (*)	363
14.4.4	Alternating and grouping subexpressions (*)	364
	Alternation operator (*)	364
	Grouping subexpressions (*)	364

	Non-grouping parentheses (*)	364
14.4.5	Quantifiers (*)	365
14.4.6	Capture groups and references thereto (**)	366
	Extracting capture group matches (**)	366
	Replacing with capture group matches (**)	368
	Back-referencing (**)	368
14.4.7	Anchoring (*)	369
	Matching at the beginning or end of a string (*)	369
	Matching at word boundaries (*)	369
	Looking behind and ahead (**)	369
14.5	Exercises	370
15	Missing, censored, and questionable data	371
15.1	Missing data	371
15.1.1	Representing and detecting missing values	372
15.1.2	Computing with missing values	372
15.1.3	Missing at random or not?	374
15.1.4	Discarding missing values	374
15.1.5	Mean imputation	375
15.1.6	Imputation by classification and regression (*)	376
15.2	Censored and interval data (*)	377
15.3	Incorrect data	377
15.4	Outliers	379
15.4.1	The 3/2 IQR rule for normally-distributed data	379
15.4.2	Unidimensional density estimation (*)	380
15.4.3	Multidimensional density estimation (*)	382
15.5	Exercises	385
16	Time series	387
16.1	Temporal ordering and line charts	387
16.2	Working with date-times and time-deltas	389
16.2.1	Representation: The UNIX epoch	389
16.2.2	Time differences	390
16.2.3	Date-times in data frames	390
16.3	Basic operations	394
16.3.1	Iterated differences and cumulative sums revisited	394
16.3.2	Smoothing with moving averages	397
16.3.3	Detecting trends and seasonal patterns	398
16.3.4	Imputing missing values	401
16.3.5	Plotting multidimensional time series	402
16.3.6	Candlestick plots (*)	404
16.4	Further reading	406
16.5	Exercises	406

Changelog	409
------------------	------------

References	413
-------------------	------------

Minimalist Data Wrangling with Python is envisaged as a student's first **introduction to data science**, providing a high-level overview as well as discussing key concepts in detail. We explore methods for cleaning data gathered from different sources, transforming, selecting, and extracting features, performing exploratory data analysis and dimensionality reduction, identifying naturally occurring data clusters, modelling patterns in data, comparing data between groups, and reporting the results.

For many students around the world, educational resources are hardly affordable. Therefore, I have decided that this book should **remain an independent, non-profit, open-access project** (available both in [PDF](#)¹ and [HTML](#)² forms). Whilst, for some people, the presence of a “designer tag” from a major publisher might still be a proxy for quality, it is my hope that this publication will prove useful to those who seek knowledge for knowledge's sake.

Any bug/typo reports/fixes are appreciated. Please submit them via this project's [GitHub repository](#)³. Thank you.

Citation: Gagolewski M. (2026), *Minimalist Data Wrangling with Python*, Melbourne, DOI:[10.5281/zenodo.6451068](#)⁴, ISBN:978-0-6455719-1-2, URL:[https://datawranglingpy.gagolewski.com/](#).

Make sure to check out *Deep R Programming*⁵ [36] too.

¹ [https://datawranglingpy.gagolewski.com/datawranglingpy.pdf](#)

² [https://datawranglingpy.gagolewski.com/](#)

³ [https://github.com/gagolews/datawranglingpy/issues](#)

⁴ [https://dx.doi.org/10.5281/zenodo.6451068](#)

⁵ [https://deepr.gagolewski.com/](#)



Preface

0.1 The art of data wrangling

*Data science*⁶ aims at making sense of and generating predictions from data that have⁷ been collected in copious quantities from various sources, such as physical sensors, surveys, online forms, access logs, and (pseudo)random number generators, to name a few. They can take diverse forms, e.g., be given as vectors, matrices, or other tensors, graphs/networks, audio/video streams, or text.

Researchers in psychology, economics, sociology, agriculture, engineering, cybersecurity, biotechnology, pharmacy, sports science, medicine, and genetics, amongst many others, need statistical methods to make new discoveries and confirm or falsify existing hypotheses. What is more, with the increased availability of open data, everyone can do remarkable work for the common good, e.g., by volunteering for non-profit NGOs or debunking false news and overzealous acts of wishful thinking on any side of the political spectrum.

Data scientists, machine learning engineers, statisticians, and business analysts are among *the most well-paid specialists*⁸. This is because data-driven decision-making, modelling, and prediction proved themselves especially effective in many domains, including healthcare, food production, pharmaceuticals, transportation, financial services (banking, insurance, investment funds), real estate, and retail.

Overall, data science (and its assorted flavours, including operational research, machine learning, pattern recognition, business and artificial intelligence) can be applied wherever we have some information repository at hand and there is a need to describe, understand, model, or improve the underlying processes.

Exercise 0.1 *Miniaturisation, increased computing power, cheaper storage, and the popularity of various internet services all caused data to become ubiquitous. Think about how much information people consume and generate when they interact with news feeds or social media on their phones.*

Data usually do not come in a *tidy* and *tamed* form. *Data wrangling* is the very broad process of appropriately curating raw information chunks and then exploring the underlying data structure so that they become *analysable*.

⁶ Traditionally known as *statistics*.

⁷ Yes, *data* are plural (*datum* is singular).

⁸ <https://survey.stackoverflow.co/2024/work/>

0.2 Aims, scope, and design philosophy

This course is envisaged as a student's first exposure to data science⁹, providing a high-level overview as well as discussing key concepts at a healthy level of detail.

By no means do we have the ambition to be comprehensive with regard to any topic we cover. Time for that will come later in separate lectures on calculus, matrix algebra, probability, mathematical statistics, continuous and combinatorial optimisation, information theory, stochastic processes, statistical/machine learning, algorithms and data structures, take a deep breath, databases and big data analytics, operational research, graphs and networks, differential equations and dynamical systems, time series analysis, signal processing, etc.

Instead, we lay solid groundwork for the aforementioned by introducing all the objects at an appropriate level of generality, and building the most crucial connections between them. We provide the necessary intuitions behind the more advanced methods and concepts. This way, further courses do not need to waste our time introducing the most elementary definitions and answering the metaphysical questions like “but why do we need that (e.g., matrix multiplication) in the first place”.

For those reasons, in this book, we explore methods for:

- performing exploratory data analysis (e.g., aggregation and visualisation),
- working with varied types of data (e.g., numerical, categorical, text, time series),
- cleaning data gathered from structured and unstructured sources, e.g., by identifying outliers, normalising strings, extracting numbers from text, imputing missing data,
- transforming, selecting, and extracting features, dimensionality reduction,
- identifying naturally occurring data clusters,
- discovering patterns in data via approximation/modelling approaches using the most popular probability distributions and the easiest to understand statistical/machine learning algorithms,
- testing whether two data distributions are significantly different,
- reporting the results of data analysis.

We primarily focus on methods and algorithms that have stood the test of time and that continue to inspire researchers and practitioners. They all meet the reality check comprised of the three undermentioned properties, which we believe are essential in practice:

⁹ We might have entitled it *Introduction to Data Science (with Python)*.

- simplicity (and thus interpretability, being equipped with no or only a few underlying tunable parameters; being based on some sensible intuitions that can be explained in our own words),
- mathematical analysability (at least to some extent; so that we can understand their strengths and limitations),
- implementability (not too abstract on the one hand, but also not requiring any advanced computer-y hocus-pocus on the other).

Note Many *more complex* algorithms are merely variations on or clever combinations of the more basic ones. This is why we need to study the foundations in great detail. We might not see it now, but this will become evident as we progress.

0.2.1 We need maths

The maths we introduce is the most elementary possible, in a good sense. Namely, we do not go beyond:

- simple analytic functions (affine maps, logarithms, cosines),
- the natural linear ordering of points on the real line (and the lack thereof in the case of multidimensional data),
- the sum of squared differences between things (including the Euclidean distance between points),
- linear vector/matrix algebra, e.g., to represent the most useful geometric transforms (rotation, scaling, translation),
- the frequentist interpretation (as in: *in samples of large sizes, we expect that...*) of some common objects from probability theory and statistics.

This is the kind of toolkit that we believe is a *sine qua non* requirement for every prospective data scientist. We cannot escape falling in love with it.

0.2.2 We need some computing environment

We no longer practice data analysis solely using a piece of paper and a pencil¹⁰. Over the years, dedicated computer programs that solve the *most common* problems arising in the most straightforward scenarios were developed, e.g., spreadsheet-like click-here-click-there standalone statistical packages. Still, we need a tool that will enable us to respond to *any* challenge in a manner that is scientifically rigorous, i.e., well organised and reproducible.

¹⁰ We acknowledge that some more theoretically inclined readers might ask the question: *but why do we need programming at all?* Unfortunately, some mathematicians have forgotten that probability and statistics are deeply rooted in the so-called real world. Theory beautifully supplements practice and provides us with very deep insights, but we still need to get our hands dirty from time to time.

This course uses the Python language which we shall introduce from scratch. Consequently, we do not require any prior programming experience.

The 2024 [StackOverflow Developer Survey](https://survey.stackoverflow.co/2024)¹¹ lists Python as the second most popular programming language (slightly behind JavaScript, whose primary use is in Web development). Over the last couple of years, it has proven to be a quite robust choice for learning and applying data wrangling techniques. This is possible thanks to the devoted community of open-source programmers who wrote the famous high-quality packages such as `numpy`, `scipy`, `matplotlib`, `pandas`, `seaborn`, and `scikit-learn`.

Nevertheless, Python and its third-party packages are amongst *many* software tools which can help extract knowledge from data. Certainly, this ecosystem is not ideal for all the applications, nor is it the most polished. The `R`¹² environment [36, 65, 96, 102] is one¹³ of the recommended alternatives worth considering.

Important We will focus on developing *transferable skills*: most of what we learn here can be applied (using different syntax but the same kind of reasoning) in other environments. Thus, this is a course on data wrangling (*with* Python), not a course *on* Python (with examples in data wrangling).

We want the reader to become an *independent* user of this computing environment. Somebody who is not overwhelmed when they are faced with any intermediate-level data analysis problem. A user whose habitual response to a new challenge is not to look everything up on the internet even in the simplest possible scenarios. Someone who will not be replaced by stupid artificial “intelligence” in the future.

We believe we have found a healthy trade-off between the minimal set of tools that need to be mastered and the less frequently used ones that can later be found in the documentation or online. In other words, the reader will discover the joy of programming and using logical reasoning to tinker with things.

0.2.3 We need data and domain knowledge

There is no data science or machine learning without *data*, and data’s purpose is to represent a given problem domain. Mathematics allows us to study different processes at a healthy level of abstractness/specificity. Still, we need to be familiar with the reality behind the numbers we have at hand, for example, by working closely with various experts or pursuing our own research in the relevant field

Courses such as this one, out of necessity, must use some generic datasets that are familiar to most readers (e.g., life expectancy and GDP by country, time to finish a marathon, yearly household incomes).

¹¹ <https://survey.stackoverflow.co/2024>

¹² <https://www.r-project.org/>

¹³ Julia also deserves a mention. There are also some commercial solutions available on the market, but we believe that ultimately all software should be free. Consequently, we are not going to talk about them here at all.

Regrettably, many textbooks introduce statistical concepts using carefully fabricated datasets where everything runs smoothly, and all models work out of the box. This gives a false sense of security and builds a too cocky a level of confidence. In practice, however, most datasets are not only unpolished; they are dull, even after some careful treatment. Such is life. We will not be avoiding the *more difficult and less attractive* problems during our journey.

0.3 Structure

This book is a whole course. We recommend reading it from the beginning to the end. The material has been divided into five parts.

1. Introducing Python:

- [Chapter 1](#) discusses how to set up the Python environment, including Jupyter Notebooks which are a flexible tool for the reproducible generation of reports from data analyses.
- [Chapter 2](#) introduces the elementary scalar types in base Python, ways to call existing and to compose our own functions, and control a code chunk's execution flow.
- [Chapter 3](#) mentions sequential and other iterable types in base Python. The more advanced data structures (vectors, matrices, data frames) will build upon these concepts.

2. Unidimensional data:

- [Chapter 4](#) introduces vectors from `numpy`, which we use for storing data on the real line (think: individual columns in a tabular dataset). Then, we look at the most common types of empirical distributions of data, e.g., bell-shaped, right-skewed, heavy-tailed ones.
- In [Chapter 5](#), we list the most basic ways for processing sequences of numbers, including methods for data aggregation, transformation (e.g., standardisation), and filtering. We also mention that a computer's floating-point arithmetic is imprecise and what we can do about it.
- [Chapter 6](#) reviews the most common probability distributions (normal, log-normal, Pareto, uniform, and mixtures thereof), methods for assessing how well they fit empirical data. It also covers pseudorandom number generation which is crucial in experiments based on simulations.

3. Multidimensional data:

- [Chapter 7](#) introduces matrices from `numpy`. They are a convenient means of storing multidimensional quantitative data, i.e., many points described by possibly many numerical features. We also present some methods for their

visualisation (and the problems arising from our being three-dimensional creatures).

- [Chapter 8](#) is devoted to operations on matrices. We will see that some of them simply extend upon what we have learnt in [Chapter 5](#), but there is more: for instance, we discuss how to determine the set of each point's nearest neighbours.
- [Chapter 9](#) discusses ways to explore the most basic relationships between the variables in a dataset: the Pearson and Spearman correlation coefficients (and what it means that correlation is not causation), k -nearest neighbour and linear regression (including sad cases where a model matrix is ill-conditioned), and finding noteworthy combinations of variables that can help reduce the dimensionality of a problem (via principal component analysis).

4. Heterogeneous data:

- [Chapter 10](#) introduces `Series` and `DataFrame` objects from `pandas`, which we can think of as vectors and matrices on steroids. For instance, they allow rows and columns to be labelled and columns to be of different types. We emphasise that most of what we learnt in the previous chapters still applies, but now we can do even more: run methods for joining (merging) many datasets, converting between long and wide formats, etc.
- In [Chapter 11](#), we introduce the ways to represent and handle categorical data as well as how (not) to lie with statistics.
- [Chapter 12](#) covers the case of aggregating, transforming, and visualising data in groups defined by one or more qualitative variables. It introduces an approach to data classification using the k -nearest neighbours scheme, which is useful when we are asked to fill the gaps in a categorical variable. We will also discover naturally occurring partitions using the k -means method, which is an example of a computationally hard optimisation problem that needs to be tackled with some imperfect heuristics.
- [Chapter 13](#) is an interlude where we solve some pleasant exercises on data frames and learn the basics of SQL. This will come in handy when our datasets do not fit in a computer's memory.

5. Other data types:

- [Chapter 14](#) discusses ways to handle text data and extract information from them, e.g., through regular expressions. We also briefly mention the challenges related to the processing of non-English text, including phrases like *pozdrowienia z Bródna*, *Viele Grüße und viel Spaß*, and *χαίρετε*.
- [Chapter 15](#) emphasises that some data may be missing or be questionable (e.g., censored, incorrect, rare) and what we can do about it.
- In [Chapter 16](#), we cover the most rudimentary methods for the processing of

time series because, ultimately, everything changes, and we should be able to track the evolution of things.

Note (*) Parts marked with a single or double asterisk (e.g., some sections or examples) can be skipped on first reading for they are of lesser importance or greater difficulty.

0.4 The Rules

Our goal here, in the long run, is for you, dear reader, to become a skilled expert who is independent, ethical, and capable of critical thinking; one who hopefully will make some contribution towards making this world a slightly better place. To guide you through this challenging journey, we have a few tips.

1. *Follow the rules.*
2. *Technical textbooks are not belletristic – purely for shallow amusement.* Sometimes a single page will be very meaning-intense. Do not try to consume too much all at once. Go for a walk, reflect on what you learnt, and build connections between different concepts. In case of any doubt, go back to the previous sections. Learning is an iterative process, not a linear one.
3. *Solve all the suggested exercises.* We might be introducing ideas or developing crucial intuitions there as well. Also, try implementing most of the methods you learn about instead of looking for copy-paste solutions on the internet. How else will you be able to master the material and develop the necessary programming skills?
4. *Code is an integral part of the text.* Each piece of good code is worth 1234 words (on average). Do not skip it. On the contrary, you are encouraged to play and experiment with it. Run every major line of code, inspect the results generated, and read more about the functions you use in the official documentation. What is the type (class) of the object returned? If it is an array or a data frame, what is its shape? What would happen if we replaced X with Y? Do not fret; your computer will not blow up.
5. *Harden up*¹⁴. Your journey towards expertise will take years, there are no shortcuts, but it will be fairly enjoyable every now and then, so don't give up. Still, sitting all day in front of your computer is unhealthy. Exercise and socialise between 28 and 31 times per month for you're not, nor will ever be, a robot.
6. *Learn maths.* Our field has a very long history and stands on the shoulders of many giants; many methods we use these days are merely minor variations on the classical, fundamental results that date back to Newton, Leibniz, Gauss, and Laplace.

¹⁴ Cyclists know.

Eventually, you will need some working knowledge of mathematics to understand them (linear algebra, calculus, probability and statistics). Remember that software products/APIs seem to change frequently, but they are just a facade, a flashy wrapping around the methods we were using for quite a while.

7. *Use only methods that you can explain.* You ought to refrain from working with algorithms/methods/models whose definitions (pseudocode, mathematical formulae, objective functions they are trying to optimise) and properties you do not know, understand, or cannot rephrase in your own words. That they might be accessible or easy to use should not make any difference to you. Also, prefer simple models over black boxes.
8. *Compromises are inevitable*¹⁵. There will never be a single best metric, algorithm, or way to solve all the problems. Even though some solutions might be superior to others with regard to certain criteria, this will only be true under very specific assumptions (if they fit a theoretical model). Beware that focusing too much on one aspect leads to undesirable consequences with respect to other factors, especially those that cannot be measured easily. Refraining from improving things might sometimes be better than pushing too hard. Always apply common sense.
9. *Be scientific and ethical.* Make your reports reproducible, your toolkit well-organised, and all the assumptions you make explicit. Develop a dose of scepticism and impartiality towards everything, from marketing slogans, through your ideological biases, to all hotly debated topics. Most data analysis exercises end up with conclusions like: “it’s too early to tell”, “data don’t show it’s either way”, “there is a difference, but it is hardly significant”, “yeah, but our sample is not representative for the entire population” – and there is nothing wrong with this. Communicate in a precise manner [88]. Remember that it is highly unethical to use statistics to tell lies [98]; this includes presenting only one side of the overly complex reality and totally ignoring all others (compare Rule#8). Using statistics for doing dreadful things (tracking users to find their vulnerabilities, developing products and services which are addictive) is a huge no-no!
10. *The best things in life are free.* These include the open-source software and open-access textbooks (such as this one) we use in our journey. Spread the good news about them and – if you can – don’t only be a taker: contribute something valuable yourself (even as small as reporting typos in their documentation or helping others in different forums when they are stuck). After all, it is our shared responsibility.

¹⁵ Some people would refer to this rule as *There is no free lunch*, but in our – overall friendly – world, many things are actually free (see Rule #10). Therefore, this name is misleading.

0.5 About the author

I, [Marek Gagolewski](https://www.gagolewski.com/)¹⁶ (pronounced like Maa'rek (Mark) Gong-o-leaf-ski), am currently an Associate Professor in Data Science at the Faculty of Mathematics and Information Science, Warsaw University of Technology.

My research interests are related to data science, in particular: modelling complex phenomena, developing usable, general-purpose algorithms, studying their analytical properties, and finding out how people use, misuse, understand, and misunderstand methods of data analysis in research, commercial, and decision-making settings. I am an author of ~100 publications, including journal papers in outlets such as *Proceedings of the National Academy of Sciences (PNAS)*, *Journal of Statistical Software*, *The R Journal*, *Journal of Classification*, *Information Fusion*, *International Journal of Forecasting*, *Statistical Modelling*, *Physica A: Statistical Mechanics and its Applications*, *Information Sciences*, *Knowledge-Based Systems*, *IEEE Transactions on Fuzzy Systems*, and *Journal of Informetrics*.

In my “spare” time, I write books for my students: check out my *Deep R Programming*¹⁷ [36]. I also [develop](https://github.com/gagolewski)¹⁸ open-source software for data analysis, such as [stringi](https://github.com/gagolewski/stringi)¹⁹ (one of the most often downloaded R packages) and [genieclust](https://github.com/gagolewski/genieclust)²⁰ (a fast and robust clustering algorithm in both Python and R).

0.6 Acknowledgements

Minimalist Data Wrangling with Python is based on my experience as an author of a quite successful textbook *Przetwarzanie i analiza danych w języku Python* [37] that I wrote with my former (successful) PhD students, Maciej Bartoszek and Anna Cena – thanks! Even though the current book is an entirely different work, its predecessor served as an excellent testbed for many ideas conveyed here.

The teaching style exercised in this book has proven successful in many similar courses that yours truly has been responsible for, including at Warsaw University of Technology, Data Science Retreat (Berlin), and Deakin University (Melbourne). I thank all my students and colleagues for the feedback given over the last 10 or so years.

A thank-you to all the authors and contributors of the Python packages that we use throughout this course: [numpy](#) [48], [scipy](#) [97], [matplotlib](#) [54], [pandas](#) [66], and [seaborn](#) [99], amongst others (as well as the many C/C++/Fortran libraries they provide wrappers for). Their version numbers are given in [Section 1.4](#).

¹⁶ <https://www.gagolewski.com/>

¹⁷ <https://deepr.gagolewski.com/>

¹⁸ <https://github.com/gagolewski>

¹⁹ <https://stringi.gagolewski.com/>

²⁰ <https://genieclust.gagolewski.com/>

This book was prepared in a Markdown superset called [MyST](#)²¹, [Sphinx](#)²², and TeX (XeLaTeX). Python code chunks were processed with the R (sic!) package [knitr](#) [106]. A little help from Makefiles, custom shell scripts, and [Sphinx](#) plugins ([sphinxcontrib-bibtex](#)²³, [sphinxcontrib-proof](#)²⁴) dotted the j's and crossed the f's. The [Ubuntu Mono](#)²⁵ font is used for the display of code. The typesetting of the main text relies on the [Alegreya](#)²⁶ typeface.

This work received no funding, administrative, technical, or editorial support from Deakin University, Warsaw University of Technology, Polish Academy of Sciences, or any other source.

o.7 You can make this book better

When it comes to quality assurance, open, non-profit projects have to resort to the generosity of the readers' community.

If you find a typo, a bug, or a passage that could be rewritten or extended for better readability/clarity, do not hesitate to report it via the *Issues* tracker available at <https://github.com/gagolews/datawranglingpy>. New feature requests are welcome as well.

²¹ <https://myst-parser.readthedocs.io/en/latest/index.html>

²² <https://www.sphinx-doc.org/>

²³ <https://pypi.org/project/sphinxcontrib-bibtex>

²⁴ <https://pypi.org/project/sphinxcontrib-proof>

²⁵ <https://design.ubuntu.com/font>

²⁶ <https://www.huertatipografica.com/en>

Part I

Introducing Python

1

Getting started with Python

1.1 Installing Python

Python¹ was designed and implemented by the Dutch programmer Guido van Rossum in the late 1980s. It is an *immensely popular*² object-orientated programming language. Over the years, it proved particularly suitable for rapid prototyping. Its name is a tribute to the funniest British comedy troupe ever. We will surely be having a jolly good laugh³ along our journey.

In this course, we will be relying on the reference implementation of the Python language (called **CPython**⁴), version 3.11 (or any later one).

Users of UNIX-like operating systems (GNU/Linux⁵, FreeBSD, etc.) may download Python via their native package manager (e.g., `sudo apt install python3` in Debian and Ubuntu). Then, additional Python packages (see [Section 1.4](#)) can be *installed*⁶ by the said manager or directly from the Python Package Index (**PyPI**⁷) via the **pip** tool.

Users of other operating systems can download Python from the project's website or some other distribution available on the market, e.g., Anaconda or Miniconda.

Exercise 1.1 *Install Python on your computer.*

¹ <https://www.python.org/>

² <https://survey.stackoverflow.co/2023/#most-popular-technologies-language>

³ When we are all in tears because of mathematics and programming, those that we shed are often tears of joy.

⁴ (*) **CPython** was written in the C programming language. Many Python packages are just convenient wrappers around code written in C, C++, or Fortran.

⁵ GNU/Linux is the operating system of choice for machine learning engineers and data scientists both on the desktop and in the cloud. Switching to a free system at some point cannot be recommended highly enough.

⁶ <https://packaging.python.org/en/latest/tutorials/installing-packages>

⁷ <https://pypi.org/>

1.2 Working with Jupyter notebooks

Jupyter⁸ brings a web browser-based development environment supporting numerous⁹ programming languages. Even though, in the long run, it is not the most convenient space for exercising data science in Python (writing standalone scripts in some more advanced editors is the preferred option), we chose it here because of its educative advantages (interactivity, beginner-friendliness, etc.).




Figure 1.1. An example Jupyter notebook.

In Jupyter, we can work with:

- **Jupyter notebooks**¹⁰ — .ipynb documents combining code, text, plots, tables, and other rich outputs; importantly, code chunks can be created, modified, and run, which makes it a fine reporting tool for our common data science needs; see [Figure 1.1](#);

⁸ <https://jupyterlab.readthedocs.io/en/stable>

⁹ <https://github.com/jupyter/jupyter/wiki/Jupyter-kernels>

¹⁰ <https://jupyterlab.readthedocs.io/en/stable/user/notebook.html>

- code consoles — terminals where we evaluate code chunks in an interactive manner (a read-eval-print loop);
- source files in a variety of programming languages — with syntax highlighting and the ability to send code to the associated consoles;

and many more.

Exercise 1.2 Head to the official [documentation](#)¹¹ of the Jupyter project. Watch the introductory video linked in the Get Started section.

Note (*) More advanced students might consider, for example, [jupytertext](#)¹² as a means to create `.ipynb` files directly from Markdown documents.

1.2.1 Launching JupyterLab

How we launch JupyterLab (or its lightweight version, Jupyter Notebook) will vary from system to system. We all need to determine the best way to do it by ourselves.

Some users will be able to start JupyterLab via their start menu/application launcher. Alternatively, we can open the system terminal (**bash**, **zsh**, etc.) and type:

```
cd our/favourite/directory # change directory
jupyter lab # or jupyter-lab, depending on the system
```

This should launch the JupyterLab server and open the corresponding app in our favourite web browser.

Note Some commercial cloud-hosted instances or forks of the open-source JupyterLab project are available on the market, but we endorse none of them; even though they might be provided gratis, there are always strings attached. It is best to run our applications locally, where we are [free](#)¹³ to be in *full* control over the software environment. Contrary to the former solution, we do not have to remain on-line to use it.

1.2.2 First notebook

Follow the undermentioned steps to create your first notebook.

1. From JupyterLab, create a new notebook running a Python 3 kernel, for example, by selecting File → New → Notebook from the application menu.
2. Select File → Rename Notebook and change the filename to `HelloWorld.ipynb`.

¹¹ <https://jupyterlab.readthedocs.io/en/stable/index.html>

¹² <https://jupytertext.readthedocs.io/en/latest>

¹³ https://www.youtube.com/watch?v=Ag1AKIL_2GM

Important The file is stored relative to the running JupyterLab server instance's current working directory. Make sure you can locate `HelloWorld.ipynb` on your disk using your file explorer. On a side note, `.ipynb` is just a JSON file that can also be edited using ordinary text editors.

3. In the first code cell, input:

```
print("G'day!")
```

4. Press `Ctrl+Enter` (or `Cmd+Return` on `m**OS`) to execute the code cell and display the result; see Figure 1.2.




Figure 1.2. “Hello, World” in a Jupyter notebook.

1.2.3 More cells

We are on fire. We cannot stop now.

1. By pressing `Enter`, we enter the *Edit mode*. Modify the current cell's contents so that it becomes:

```
# My first code cell (this is a comment)
print("G'day!") # prints a message (this is a comment too)
print(2+5) # prints a number
```

2. Press `Ctrl+Enter` to execute whole code chunk and replace the previous outputs with the updated ones.
3. Enter another command that prints a message that you would like to share with the world. Note that character strings in Python must be enclosed either in double quotes or in apostrophes.
4. Press `Shift+Enter` to execute the current code cell, create a new one below it, and then enter the edit mode.
5. In the new cell, input and then execute:

```
import matplotlib.pyplot as plt # the main plotting library
plt.bar(
```

(continues on next page)

(continued from previous page)

```
[ "Python", "JavaScript", "HTML", "CSS" ], # a list of strings
[ 80, 30, 10, 15 ] # a list of integers (the respective bar heights)
)
plt.title("What makes you happy?")
plt.show()
```

6. Add three more code cells that display some text or create other bar plots.

Exercise 1.3 Change `print(2+5)` to `PRINT(2+5)`. Run the corresponding code cell and see what happens.

Note In the *Edit* mode, JupyterLab behaves like an ordinary text editor. Most keyboard shortcuts known from elsewhere are available, for example:

- Shift+LeftArrow, DownArrow, UpArrow, or RightArrow – select text,
 - Ctrl+c – copy,
 - Ctrl+x – cut,
 - Ctrl+v – paste,
 - Ctrl+z – undo,
 - Ctrl+] – indent,
 - Ctrl+[– dedent,
 - Ctrl+/ – toggle comment.
-

1.2.4 Edit vs command mode

By pressing ESC, we can enter the *Command mode*.

1. Use the arrow DownArrow and UpArrow keys to move between the code cells.
2. Press d,d (d followed by another d) to delete the currently selected cell.
3. Press z to undo the last operation.
4. Press a and b to insert a new blank cell, respectively, above and below the current one.
5. Note a simple drag and drop can relocate cells.

Important ESC and Enter switch between the *Command* and *Edit* modes, respectively.

Example 1.4 In Jupyter notebooks, the linear flow of chunks' execution is not strongly enforced. For instance:

```
## In [2]:
    x = [1, 2, 3]

## In [10]:
    sum(x)
## Out [10]:
##      18

## In [7]:
    sum(y)
## Out [7]:
##      6

## In [6]:
    x = [5, 6, 7]

## In [5]:
    y = x
```

The chunk IDs reveal the true order in which the author has executed them. By editing cells in a rather frivolous fashion, we may end up with matter that makes little sense when it is read from the beginning to the end. It is thus best to always select *Restart Kernel and Run All Cells* from the *Kernel* menu to ensure that evaluating content step by step renders results that meet our expectations.

1.2.5 Markdown cells

So far, we have only been playing with *code* cells. Notebooks are not just about writing code, though. They are meant to be read by humans too. Thus, we need some means to create formatted text.

Markdown is lightweight yet powerful enough markup (pun indented) language widely used on many popular platforms (e.g., on *Stack Overflow* and *GitHub*). We can convert the current cell to a Markdown block by pressing `m` in the *Command mode* (note that by pressing `y` we can turn it back to a *code* cell).

1. In a new Markdown cell, enter:

```
# Section
```

```
## Subsection
```

```
This is the first paragraph. It ~~was~~ *is* **very** nice.
Great success.
```

```
This is the second paragraph. Note that a blank line separates
it from the previous one. And now for something completely different;
a bullet list:
```

(continues on next page)

(continued from previous page)

```
* one,

* two,

    1. aaa,

    2. bbbb,

* [three](https://en.wikipedia.org/wiki/3).

---
```

And now some `2+2` in Python:

```
```python
some code to display (but not execute)
2+2
```
```

An image:

```
![Python logo](https://www.python.org/static/img/python-logo.png)
```

```
(\*) An equation (LaTeX):  $x_i = \frac{\sqrt{\pi}}{2}$ .
```

And a table:

| |
|---------|
| A B |
| -- -- |
| 1 3 |
| 2 4 |

2. Press Ctrl+Enter to display the formatted text.
3. Notice that Markdown cells can be modified in the *Edit mode* as usual (the Enter key).

Exercise 1.5 Read the [official introduction](https://daringfireball.net/projects/markdown/syntax)¹⁴ to the Markdown syntax.

Exercise 1.6 Follow [this](https://commonmark.org/help/tutorial)¹⁵ interactive Markdown tutorial.

Exercise 1.7 Apply what you have learnt by making the current Jupyter notebook more readable. At the beginning of the report, add a header featuring your name and your email address. Before and after each code cell, explain its purpose and how to interpret the results obtained.

¹⁴ <https://daringfireball.net/projects/markdown/syntax>

¹⁵ <https://commonmark.org/help/tutorial>

1.3 The best note-taking app

Our learning will not be effective if we do not take *good* note of the concepts that we come across during this course, especially if they are new to us. More often than not, we will find ourselves in a need to write down the definitions and crucial properties of the methods we discuss, draw simple diagrams and mind maps to build connections between different topics, verify the validity of some results, or derive simple mathematical formulae ourselves.

Let's not waste our time finding the best app for our computers, phones, or tablets. One versatile note-taking solution is an ordinary piece of A4 paper and a pen or a pencil. Loose sheets, 5 mm grid-ruled for graphs and diagrams, work nicely. They can be held together using a cheap landscape clip folder (the one with a clip on the long side). This way, it can be browsed through like an ordinary notebook. Also, new pages can be added anywhere, and their ordering altered arbitrarily.

1.4 Initialising each session and getting example data

From now on, we assume that the following commands are issued at the beginning of each Python session.

```
# import key packages (required):
import numpy as np
import scipy.stats
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# further settings (optional):
pd.set_option("display.notebook_repr_html", False) # disable "rich" output

import os
os.environ["COLUMNS"] = "74" # output width, in characters
np.set_printoptions(
    linewidth=74, # output width
    legacy="1.25", # print scalars without type information
)
pd.set_option("display.width", 74)

import sklearn
sklearn.set_config(display="text")
```

(continues on next page)

(continued from previous page)

```
plt.style.use("seaborn-v0_8") # plot style template

_colours = [ # the "R4" palette from R
    "#000000f0", "#DF536Bf0", "#61D04Ff0", "#2297E6f0",
    "#28E2E5f0", "#CD0B8Cf0", "#F5C710f0", "#999999f0"
]

_linestyles = [
    "solid", "dashed", "dashdot", "dotted"
]

plt.rcParams["axes.prop_cycle"] = plt.cycler(
    # each plotted line will have a different plotting style
    color=_colours, linestyle=_linestyles*2
)
plt.rcParams["patch.facecolor"] = _colours[0]

np.random.seed(123) # initialise the pseudorandom number generator
```

First, we imported the most frequently used packages (together with their usual aliases, we will get to that later). Then, we set up some further options that yours truly is particularly fond of. On a side note, [Section 6.4.2](#) discusses the issues in reproducible pseudorandom number generation.

Open-source software regularly enjoys feature extensions, API changes, and bug fixes. It is worthwhile to know which version of the Python environment was used to execute all the code listed in this book:

```
import sys
print(sys.version)
## 3.13.9 (main, Nov 23 2025, 01:32:31) [GCC]
```

Given beneath are the versions of the packages that we will be relying on. This information can usually be accessed by calling `print(package.__version__)`.

| Package | Version |
|----------------------------|-----------|
| numpy | 2.3.4 |
| scipy | 1.16.2 |
| matplotlib | 3.10.7 |
| pandas | 2.3.3 |
| seaborn | 0.13.2 |
| sklearn (scikit-learn) (*) | 1.7.1 |
| icu (PyICU) (*) | 2.15.3 |
| IPython (*) | 9.4.0 |
| mplfinance (*) | 0.12.10bo |

We expect 99% of our code to work in the (near-)future versions of the environment. If the diligent reader discovers that this is not the case, for the benefit of other students, filing a bug report at <https://github.com/gagolews/datawranglingpy> will be much appreciated.

Important All example datasets that we use throughout this course are available for download at <https://github.com/gagolews/teaching-data>.

Exercise 1.8 *Ensure you know how to access the files from our teaching data repository. Chose any file, e.g., `nhanes_adult_female_height_2020.txt` in the `marek` folder, and then click Raw. It is the URL where you have been redirected to, not the previous one, that includes the link to be used from within your Python session. Also, note that each dataset starts with several comment lines explaining its structure, the meaning of the variables, etc.*

1.5 Exercises

Exercise 1.9 *What is the difference between the Edit and the Command modes in Jupyter?*

Exercise 1.10 *How can we format a table in Markdown? How can we insert an image, a hyper-link, and an enumerated list?*

2

Scalar types and control structures in Python

In this part, we introduce the basics of the Python language itself. Being a general-purpose tool, various packages supporting data wrangling operations are provided as third-party extensions. In further chapters, extending upon the concepts discussed here, we will be able to use `numpy`, `scipy`, `matplotlib`, `pandas`, `seaborn`, and other packages with a healthy degree of confidence.

2.1 Scalar types

Scalars are *single* or *atomic* values. Their five ubiquitous types are:

- `bool` – logical,
- `int`, `float`, `complex` – numeric,
- `str` – character.

Let's discuss them in detail.

2.1.1 Logical values

There are only two possible logical (Boolean) values: `True` and `False`. By typing:

```
True  
## True
```

we instantiated the former. This is a dull exercise unless we have fallen into the under-mentioned pitfall.

Important Python is a case-sensitive language. Writing “`TRUE`” or “`true`” instead of “`True`” is an error.

2.1.2 Numeric values

The three numeric scalar types are:

- `int` – integers, e.g., `1`, `-42`, `1_000_000`;

- float – floating-point (real) numbers, e.g., -1.0, 3.14159, 1.23e-4;
- (*) complex – complex numbers, e.g., 1+2j.

In practice, numbers of the type `int` and `float` often interoperate seamlessly. We usually do not have to think about them as being of distinctive types. On the other hand, complex numbers are rather infrequently used in data science applications (but see [Section 4.1.4](#)).

Exercise 2.1 *1.23e-4 and 9.8e5 are examples of numbers in scientific notation, where “e” stands for “... times 10 to the power of...”. Additionally, 1_000_000 is a decorated (more human-readable) version of 1000000. Use the `print` function to check out their values.*

Arithmetic operators

Here is the list of available arithmetic operators:

```
1 + 2      # addition
## 3
1 - 7      # subtraction
## -6
4 * 0.5    # multiplication
## 2.0
7 / 3      # float division (results are always of the type float)
## 2.3333333333333335
7 // 3     # integer division
## 2
7 % 3      # division remainder
## 1
2 ** 4     # exponentiation
## 16
```

The precedence of these operators is [quite predictable](#)¹, e.g., exponentiation has higher priority than multiplication and division, which in turn bind more strongly than addition and subtraction. Thus,

```
1 + 2 * 3 ** 4
## 163
```

is the same as $1 + (2 * (3 ** 4))$ and is different from, e.g., $((1 + 2) * 3) ** 4$.

Note Keep in mind that computers’ floating-point arithmetic is precise only up to a dozen or so significant digits. As a consequence, the result of $7/3$ is only approximate; hence the 2.3333333333333335 above. We will discuss this topic in [Section 5.5.6](#).

¹ <https://docs.python.org/3/reference/expressions.html#operator-precedence>

Creating named variables

A named variable can be introduced through the *assignment operator*, `=`. It can store an arbitrary Python object which we can recall at any later time. Names of variables can include any lower- and uppercase letters, underscores, and (except at the beginning) digits.

To make our code easier to understand for humans, it is best to use names that are self-explanatory, like:

```
x = 7 # read: let `x` from now on be equal to 7 (or: `x` becomes 7)
```

“x” is great name: it means *something of general interest* in mathematics. Let’s print out the value it is bound to:

```
print(x) # or just `x`
## 7
```

New variables can be created easily based on existing ones:

```
my_2nd_variable = x/3 - 2 # creates `my_2nd_variable`
print(my_2nd_variable)
## 0.3333333333333335
```

Existing variables may be rebound to any other value freely:

```
x = x/3 # let the new `x` be equal to the old `x` (7) divided by 3
print(x)
## 2.3333333333333335
```

Exercise 2.2 Define two named variables *height* (in centimetres) and *weight* (in kilograms). Determine the corresponding body mass index (BMI^2).

Note (*) Augmented assignments are also available. For example:

```
x *= 3
print(x)
## 7.0
```

In this context, the foregoing is equivalent to `x = x*3`. In other words, it creates a new object. Nevertheless, in some scenarios, augmented assignments may modify the objects they act upon *in place*; compare [Section 3.5](#).

2.1.3 Character strings

Character strings (objects of the type `str`) store text data. They are created using apostrophes or double quotes:

² https://en.wikipedia.org/wiki/Body_mass_index

```
print("spam, spam, #, bacon, and spam")
## spam, spam, #, bacon, and spam
print('Cześć! ¿Qué tal?')
## Cześć! ¿Qué tal?
print('G\'day, how\'s it goin\''," he asked.\\n"All good," she responded.')
## "G'day, how's it goin'," he asked.
## "All good," she responded.
```

We see some examples of *escape sequences*³ here:

- “\ '” is a way to include an apostrophe in an apostrophe-delimited string,
- “\ \” enters a backslash,
- “\ n” inputs a newline character.

Multiline strings are created using three apostrophes or double quotes:

```
"""
spam|spam
tasty|t"spam"
lovely|t'spam'
"""
## '|nspam|spam|ntasty|t"spam"|nlovely|t|'spam|'|n'
```

Exercise 2.3 Call the `print` function on the above objects to reveal the meaning of the included escape sequences.

Important Many string operations are available, e.g., for formatting and pattern searching. They are especially important in the art of data wrangling as information often arrives in textual form. [Chapter 14](#) covers this topic in detail.

F-strings (formatted string literals)

F-strings are formatted string literals:

```
x = 2
f"x is equal to {x}"
## 'x is equal to 2'
```

Notice the “f” prefix. The “{x}” part was replaced with the value stored in the `x` variable.

The formatting of items can be fine-tuned. As usual, it is best to study the [documentation](#)⁴ in search of noteworthy features. Here, let’s just mention that we will frequently be referring to placeholders like “{value:width}” and “{value:width.precision}”,

³ https://docs.python.org/3/reference/lexical_analysis.html#string-and-bytes-literals

⁴ https://docs.python.org/3/reference/lexical_analysis.html#f-strings

which specify the field width and the number of fractional digits of a number. This way, we can output a series of values aesthetically aligned one beneath another.

```
n = 3.14159265358979323846
e = 2.71828182845904523536
print(f"""
n   = {n:10.8f}
e   = {e:10.8f}
ne2 = {(n*e**2):10.8f}
""")
##
## n   = 3.14159265
## e   = 2.71828183
## ne2 = 23.21340436
```

“10.8f” means that a value should be formatted as a float, be of width at least ten characters (text columns), and use eight fractional digits.

2.2 Calling built-in functions

We have a few base functions at our disposal. For instance, to round the Euler constant e to two decimal digits, we can call:

```
e = 2.718281828459045
round(e, 2)
## 2.72
```

Exercise 2.4 Call `help("round")` to access the function’s manual. Note that the second argument, called `ndigits`, which we set to 2, defaults to `None`. Check what happens when we omit it during the call.

2.2.1 Positional and keyword arguments

The `round` function has two parameters, `number` and `ndigits`. Thus, the following calls are equivalent:

```
print(
    round(e, 2),           # two arguments matched positionally
    round(e, ndigits=2),   # positional and keyword argument
    round(number=e, ndigits=2), # two keyword arguments
    round(ndigits=2, number=e) # the order does not matter for keyword args
)
## 2.72 2.72 2.72 2.72
```

Verifying that no other call scheme is permitted is left as an exercise, i.e., positionally matched arguments must be listed before the keyword ones.

2.2.2 Modules and packages

Python modules and packages (which are collections of modules) define thousands of additional functions. For example, `math` features the most common mathematical routines:

```
import math # the math module must be imported before we can use it
print(math.log(2.718281828459045)) # the natural logarithm (base e)
## 1.0
print(math.floor(-7.33)) # the floor function
## -8
print(math.sin(math.pi)) # sin(pi) equals 0 (with small numeric error)
## 1.2246467991473532e-16
```

See the [official documentation](#)⁵ for the comprehensive list of objects available. On a side note, all floating-point computations in any programming language are subject to round-off errors and other inaccuracies. This is why the result of $\sin \pi$ is not exactly 0, but some value very close thereto. We will elaborate on this topic in [Section 5.5.6](#).

Packages can be given aliases, for the sake of code readability or due to our being lazy. For instance, in [Chapter 4](#) we will get used to importing the `numpy` package under the `np` alias:

```
import numpy as np
```

And now, instead of writing, for example, `numpy.random.rand()`, we can call:

```
np.random.rand() # a pseudorandom value in [0.0, 1.0)
## 0.6964691855978616
```

2.2.3 Slots and methods

Python is an object-orientated programming language. Each object is an instance of some *class* whose name we can reveal by calling the `type` function:

```
x = 1+2j
type(x)
## <class 'complex'>
```

Important Classes define two kinds of *attributes*:

- *slots* – associated data,
 - *methods* – associated functions.
-

Exercise 2.5 Call `help("complex")` to reveal that the `complex` class defines, amongst others, the *conjugate* method and the *real* and *imag* slots.

⁵ <https://docs.python.org/3/library/math.html>

Here is how we can read the two slots:

```
print(x.real) # access slot `real` of object `x` of the class `complex`
## 1.0
print(x.imag)
## 2.0
```

And here is an example of a method call:

```
x.conjugate() # equivalently: complex.conjugate(x)
## (1-2j)
```

Notably, the documentation of this function can be accessed by typing `help("complex.conjugate")` (class name – dot – method name).

2.3 Controlling program flow

2.3.1 Relational and logical operators

We have several operators which return a single logical value:

```
1 == 1.0 # is equal to?
## True
2 != 3 # is not equal to?
## True
"spam" < "egg" # is less than? (with respect to the lexicographic order)
## False
```

Some more examples:

```
math.sin(math.pi) == 0.0 # well, numeric error...
## False
abs(math.sin(math.pi)) <= 1e-9 # is close to 0?
## True
```

Logical results can be combined using **and** (*conjunction*; for testing if both operands are true) and **or** (*alternative*; for determining whether at least one operand is true). Likewise, **not** stands for *negation*.

```
3 <= math.pi and math.pi <= 4 # is it between 3 and 4?
## True
not (1 > 2 and 2 < 3) and not 100 <= 3
## True
```

Notice that `not 100 <= 3` is equivalent to `100 > 3`. Also, based on the de Morgan laws, `not (1 > 2 and 2 < 3)` is true if and only if `1 <= 2 or 2 >= 3` holds.

Exercise 2.6 Assuming that p, q, r are logical and a, b, c, d are variables of the type *float*, simplify the following expressions:

- *not not* p ,
- *not* p and *not* q ,
- *not* (*not* p or *not* q or *not* r),
- *not* $a == b$,
- *not* ($b > a$ and $b < c$),
- *not* ($a >= b$ and $b >= c$ and $a >= c$),
- ($a > b$ and $a < c$) or ($a < c$ and $a > d$).

2.3.2 The if statement

The *if* statement executes a chunk of code *conditionally*, based on whether the provided expression is true or not. For instance, given some variable:

```
x = np.random.rand() # a pseudorandom value in [0.0, 1.0)
```

we can react enthusiastically to its being less than 0.5:

```
if x < 0.5: print("spam!") # note the colon after the tested condition
```

Actually, we remained cool as a cucumber (nothing was printed) because x is equal to:

```
print(x)
## 0.6964691855978616
```

Multiple *elif* (*else-if*) parts can be added. They are inspected one by one, until one of the tests turns out to be successful. At the end, we can include an optional *else* part. It is executed when all of the tested conditions turn out to be false.

```
if x < 0.25: print("spam!")
elif x < 0.5: print("ham!") # i.e., x in [0.25, 0.5)
elif x < 0.75: print("bacon!") # i.e., x in [0.5, 0.75)
else: print("eggs!") # i.e., x >= 0.75
## bacon!
```

Note that if we wrote the second condition as $x \geq 0.25$ and $x < 0.5$, we would introduce some redundancy; when it is being considered, we already know that $x < 0.25$ (the first test) is *not* true. Similarly, the *else* part is only executed when all the tests fail, which in our case happens if neither $x < 0.25$, $x < 0.5$, nor $x < 0.75$ is true, i.e., if $x \geq 0.75$.

Whenever more than one statement is to be executed conditionally, an *indented code block* can be introduced.

```

if x >= 0.25 and x <= 0.75:
    print("bacon!")
    print("I love it!")
else:
    print("I'd rather eat spam!")
print("more spam!") # executed regardless of the condition's state
## bacon!
## I love it!
## more spam!

```

Important The indentation must be neat and consistent. We recommend using *four spaces*. Note the kind of error generated when we try executing:

```

if x < 0.5:
    print("spam!")
    print("ham!") # :(
IndentationError: unindent does not match any outer indentation level

```

Exercise 2.7 For a given BMI, print out the corresponding category as defined by the WHO (underweight if less than 18.5 kg/m², normal range up to 25.0 kg/m², etc.). Bear in mind that the BMI is a simplistic measure. Both the medical and statistical communities pointed out its inherent limitations. Read the Wikipedia article thereon for more details (and appreciate the amount of data wrangling required for its preparation: tables, charts, calculations; something that we will be able to perform quite soon, given quality reference data, of course).

Exercise 2.8 (*) Check if it is easy to find on the internet (in reliable sources) some raw datasets related to the body mass studies, e.g., measuring subjects' height, weight, body fat and muscle mass, etc.

2.3.3 The while loop

The **while** loop executes a given statement or a series of statements *as long as* a given condition is true. For example, here is a simple simulator determining how long we have to wait *until* drawing the first value *not* greater than 0.01 whilst generating numbers in the unit interval:

```

count = 0
while np.random.rand() > 0.01:
    count = count + 1
print(count)
## 117

```

Exercise 2.9 Using the **while** loop, determine the arithmetic mean of 100 randomly generated numbers (i.e., the sum of the numbers divided by 100).

2.4 Defining functions

As a means for *code reuse*, we can define *our own* functions. For instance, below is a procedure that computes the minimum (with respect to the `<` relation) of three given objects:

```
def min3(a, b, c):
    """
    A function to determine the minimum of three given inputs.

    By the way, this is a docstring (documentation string);
    call help("min3") later to view it.
    """
    if a < b:
        if a < c:
            return a
        else:
            return c
    else:
        if b < c:
            return b
        else:
            return c
```

Example calls:

```
print(min3(10, 20, 30),
      min3(10, 30, 20),
      min3(20, 10, 30),
      min3(20, 30, 10),
      min3(30, 10, 20),
      min3(30, 20, 10))
## 10 10 10 10 10 10
```

Note that `min3` returns a value. The result it yields can be consumed in further computations:

```
x = min3(np.random.rand(), 0.5, np.random.rand()) # minimum of 3 numbers
x = round(x, 3) # transform the result somehow
print(x)
## 0.5
```

Exercise 2.10 Write a function named `bmi` which computes and returns a person's BMI, given their weight (in kilograms) and height (in centimetres). As documenting functions constitutes a good development practice, do not forget about including a docstring.

New variables can be introduced inside a function's body. This can help the function perform its duties.

```
def min3(a, b, c):
    """
    A function to determine the minimum of three given inputs
    (alternative version).
    """
    m = a # a local (temporary/auxiliary) variable
    if b < m:
        m = b
    if c < m: # be careful! no `else` or `elif` here – it's a separate `if`
        m = c
    return m
```

Example call:

```
m = 7
n = 10
o = 3
min3(m, n, o)
## 3
```

All *local variables* cease to exist after the function is called. Notice that `m` inside the function is a variable independent of `m` in the global (calling) scope.

```
print(m) # this is still the global `m` from before the call
## 7
```

Exercise 2.11 Implement a function `max3` which determines the maximum of three given values.

Exercise 2.12 Write a function `med3` which defines the median of three given values (the value that is in-between two other ones).

Exercise 2.13 (*) Indite a function `min4` to compute the minimum of four values.

2.4.1 Lambda expressions

Lambda expressions give us an uncomplicated way to define functions using a single line of code. They are defined using the syntax `lambda argument_name: return_expression`.

```
square = lambda x: x**2 # i.e., def square(x): return x**2
square(4)
## 16
```

Objects generated through lambda expressions do not have to be assigned a name: they can remain anonymous. This is useful when calling a method which takes another function as its argument. With lambdas, the latter can be generated on the fly.

```
def print_x_and_fx(x, f):
    """
    Arguments: x - some object; f - a function to be called on x
    """
    print(f"x = {x} and f(x) = {f(x)}")

print_x_and_fx(4, lambda x: x**2)
## x = 4 and f(x) = 16
print_x_and_fx(math.pi/4, lambda x: round(math.cos(x), 5))
## x = 0.7853981633974483 and f(x) = 0.70711
```

2.4.2 (*) Own modules

Definitions of functions and other Python objects can be placed in a separate source file. This way, they can be referred to from within multiple projects. For instance, in the current working directory, if we create a file `module.py` featuring the definition of the above `square` function, we will be able to call it like:

```
import module
module.square(4)
## 16
```

Unfortunately, once a module is loaded, any changes thereto will not be reflected until the Python session is restarted. Thus, in an interactive environment (such as when working with Jupyter notebooks), we may find the `importlib.reload` function useful.

2.5 Exercises

Exercise 2.14 What does `import xxxxxx as x` mean?

Exercise 2.15 What is the difference between `if` and `while`?

Exercise 2.16 Name the scalar types we introduced in this chapter.

Exercise 2.17 What is a function's docstring and how can we create and access it?

Exercise 2.18 What are keyword arguments of a function?

3

Sequential and other types in Python

3.1 Sequential types

Sequential objects store data items that can be accessed by index (position). The three main sequential types are: lists, tuples, and ranges.

As a matter of fact, strings (which we often treat as scalars) can also be considered of this kind. Therefore, amongst sequential objects are such diverse classes as:

- lists,
- tuples,
- ranges, and
- strings.

Nobody expected that.

3.1.1 Lists

Lists consist of arbitrary Python objects. They can be created using standalone square brackets:

```
x = [True, "two", 3, [4j, 5, "six"], None]
print(x)
## [True, 'two', 3, [4j, 5, 'six'], None]
```

The preceding is an example list featuring objects of the types: `bool`, `str`, `int`, `list` (yes, it is possible to have a list inside another list), and `None` (the `None` object is the only of this kind, it represents a placeholder for nothingness).

Note We will often be relying on lists when creating vectors in **numpy** or data frame columns in **pandas**. Furthermore, lists of lists of equal lengths can be used to create matrices.

Each list is *mutable*. Consequently, its state may freely be changed. For instance, we can append a new object at its end:

```
x.append("spam")
print(x)
## [True, 'two', 3, [4j, 5, 'six'], None, 'spam']
```

The call to the `list.append` method modified `x` in place.

3.1.2 Tuples

Next, *tuples* are like lists, but they are *immutable* (read-only): once created, they cannot be altered.

```
("one", [], (3j, 4))
## ('one', [], (3j, 4))
```

This gave us a *triple* (a 3-tuple) carrying a string, an empty list, and a *pair* (a 2-tuple). Let's stress that we can drop the round brackets and still get a tuple:

```
1, 2, 3 # the same as `(1, 2, 3)`
## (1, 2, 3)
```

Also:

```
42, # equivalently: `(42, )`
## (42,)
```

Note the trailing comma; we defined a *singleton* (a 1-tuple). It is not the same as the scalar 42 or (42), which is an object of the type `int`.

Note Having a separate data type representing an immutable sequence makes sense in certain contexts. For example, a data frame's *shape* is its inherent property that should not be tinkered with. If a tabular dataset has 10 rows and 5 columns, we disallow the user to set the former to 15 (without making further assumptions, providing extra data, etc.).

When creating collections of items, we usually prefer lists, as they are more flexible a data type. Yet, [Section 3.4.2](#) will mention that many functions return tuples. We are thus expected to be able to handle them with confidence.

3.1.3 Ranges

Objects defined by calling `range(from, to)` or `range(from, to, by)` represent arithmetic progressions of integers.

```
list(range(0, 5)) # i.e., range(0, 5, 1) - from 0 to 5 (exclusive) by 1
## [0, 1, 2, 3, 4]
list(range(10, 0, -1)) # from 10 to 0 (exclusive) by -1
## [10, 9, 8, 7, 6, 5, 4, 3, 2, 1]
```


We converted the two ranges to ordinary lists as otherwise their display is not particularly spectacular. Let's point out that the rightmost boundary (to) is *exclusive* and that by default to 1.

3.1.4 Strings (again)

Recall that we have already discussed character strings in [Section 2.1.3](#).

```
print("lovely\nspam")  
## lovely  
## spam
```

Strings are most often treated as scalars (atomic entities, as in: a string as a whole). However, we will soon find out that their individual characters can also be accessed by index. Furthermore, [Chapter 14](#) will discuss a plethora of operations on *parts* of strings.

3.2 Working with sequences

3.2.1 Extracting elements

The *index operator*, `[...]`, can be applied on any sequential object to extract an element at a position specified by a single integer.

```
x = ["one", "two", "three", "four", "five"]  
x[0] # the first element  
## 'one'  
x[1] # the second element  
## 'two'  
x[len(x)-1] # the last element  
## 'five'
```

The valid indexes are $0, 1, \dots, n-2, n-1$, where n is the length (size) of the sequence, which can be fetched by calling `len`.

Important Think of an index as the distance from the start of a sequence. For example, `x[3]` means “three items away from the beginning”, i.e., the fourth element.

Negative indexes count from the end:

```
x[-1] # the last element (ultimate)  
## 'five'  
x[-2] # the next to last (the last but one, penultimate)  
## 'four'
```

(continues on next page)

(continued from previous page)

```
x[-len(x)] # the first element
## 'one'
```

The index operator can be applied on any sequential object:

```
"string"[3]
## 'i'
```

More examples:

```
range(0, 10)[-1] # the last item in an arithmetic progression
## 9
(1, ) [0] # extract from a 1-tuple
## 1
```

Important The same “thing” can have different meanings in different contexts. Therefore, we must always remain vigilant.

For instance, raw square brackets are used to create a list (e.g., `[1, 2, 3]`) whereas their presence after a sequential object indicates some form of indexing (e.g., `x[1]` or even `[1, 2, 3][1]`). Similarly, `(1, 2)` creates a 2-tuple and `f(1, 2)` denotes a call to a function `f` with two arguments.

3.2.2 Slicing

We can also use slices of the form `from:to` or `from:to:by` to select a subsequence of a given sequence. Slices are similar to ranges, but ``:`` can only be used within square brackets.

```
x = ["one", "two", "three", "four", "five"]
x[1:4] # from the second to the fifth (exclusive)
## ['two', 'three', 'four']
x[-1:0:-2] # from the last to first (exclusive) by every second backwards
## ['five', 'three']
```

In fact, the `from` and `to` parts of a slice are optional. When omitted, they default to one of the sequence boundaries.

```
x[3:] # from the third element to the end
## ['four', 'five']
x[:2] # the first two
## ['one', 'two']
x[:0] # none (the first zero)
## []
x[::2] # every second element from the start
```

(continues on next page)

(continued from previous page)

```
## ['one', 'three', 'five']  
x[::-1] # the elements in reverse order  
## ['five', 'four', 'three', 'two', 'one']
```

Slicing can be applied on other sequential objects as well:

```
"spam, bacon, spam, and eggs"[13:17] # fetch a substring  
## 'spam'
```

Knowing the difference between element extraction and subsetting a sequence (creating a subsequence) is crucial. For example:

```
x[0] # extraction (indexing with a single integer)  
## 'one'
```

It gave the object *at* that index. Moreover:

```
x[0:1] # subsetting (indexing with a slice)  
## ['one']
```

It returned the object of the same type as *x* (here, a list), even though, in this case, only one object was fetched. However, a slice can potentially select any number of elements, including zero.

pandas data frames and **numpy** arrays will behave similarly, but there will be many more indexing options; see [Section 5.4](#), [Section 8.2](#), and [Section 10.5](#).

3.2.3 Modifying elements of mutable sequences

Lists are *mutable*: their state can be changed. The index operator can replace the elements at given indexes.

```
x = ["one", "two", "three", "four", "five"]  
x[0] = "spam" # replace the first element  
x[-3:] = ["bacon", "eggs"] # replace last three with given two  
print(x)  
## ['spam', 'two', 'bacon', 'eggs']
```

Exercise 3.1 There are quite a few methods that modify list elements: not only the aforementioned **append**, but also **insert**, **remove**, **pop**, etc. Invoke `help("list")`, read their descriptions, and call them on a few example lists.

Exercise 3.2 Verify that similar operations cannot be performed on tuples, ranges, and strings. In other words, check that these types are immutable.

3.2.4 Searching for specific elements

The `in` operator and its negation, `not in`, determine whether an element exists in a given sequence:

```
7 in range(0, 10)
## True
[2, 3] in [ 1, [2, 3], [4, 5, 6] ]
## True
```

For strings, `in` tests whether a string includes a specific *substring*:

```
"spam" in "lovely spams"
## True
```

Exercise 3.3 In the documentation of the `list` and other classes, check out the `count` and `index` methods.

3.2.5 Arithmetic operators

Some arithmetic operators were *overloaded* for certain sequential types. However, they carry different meanings from those for integers and floats. In particular, ``+`` joins (concatenates) strings, lists, and tuples:

```
"spam" + " " + "bacon"
## 'spam bacon'
[1, 2, 3] + [4]
## [1, 2, 3, 4]
```

Moreover, ``*`` duplicates (recycles) a given sequence:

```
"spam" * 3
## 'spamspamspam'
(1, 2) * 4
## (1, 2, 1, 2, 1, 2, 1, 2)
```

In each case, a new object has been returned.

3.3 Dictionaries

Dictionaries are sets of `key:value` pairs, where the `value` (any Python object) can be accessed by `key` (usually¹ a string). In other words, they map `keys` to `values`.

¹ Overall, *hashable* data types can be used as dictionary keys, e.g., integers, floats, strings, tuples, and ranges; see `hash`. It is required that hashable objects be immutable.

```
x = {
    "a": [1, 2, 3],
    "b": 7,
    "z": "spam!"
}
print(x)
## {'a': [1, 2, 3], 'b': 7, 'z': 'spam!'}
```

We can also create a dictionary with string keys using the `dict` function which accepts any keyword arguments:

```
dict(a=[1, 2, 3], b=7, z="spam!")
## {'a': [1, 2, 3], 'b': 7, 'z': 'spam!'}
```

The index operator extracts a specific element from a dictionary, uniquely identified by a given key:

```
x["a"]
## [1, 2, 3]
```

In this context, `x[0]` is not valid and raises an error: a dictionary is not an object of sequential type; a key of `0` does not exist in `x`. If we are unsure whether a specific key is defined, we can use the `in` operator:

```
"a" in x, 0 not in x, "z" in x, "w" in x # a tuple of four tests' results
## (True, True, True, False)
```

There is also a method called `get`, which returns an element associated with a given key, or something else (by default, `None`) if we have a mismatch:

```
x.get("a")
## [1, 2, 3]
x.get("c") # if missing, returns None by default
x.get("c") is None # indeed
## True
x.get("c", "unknown")
## 'unknown'
```

We can also add new elements to a dictionary using the index operator:

```
x["f"] = "more spam!"
print(x)
## {'a': [1, 2, 3], 'b': 7, 'z': 'spam!', 'f': 'more spam!'}
```

Example 3.4 (*) In practice, we often import JSON files (which is a popular data exchange format on the internet) exactly in the form of Python dictionaries. Let's demo it briefly:

```
import requests
x = requests.get("https://api.github.com/users/gagolews/starred").json()
```

Now *x* is a sequence of dictionaries giving the information on the repositories starred by yours truly on GitHub. As an exercise, the reader is encouraged to inspect its structure.

3.4 Iterable types

All the objects we discussed here are *iterable*. In other words, we can iterate through each element contained therein. In particular, the **list** and **tuple** functions take any iterable object and convert it to a sequence of the corresponding type. For instance:

```
list("spam")
## ['s', 'p', 'a', 'm']
tuple(range(0, 10, 2))
## (0, 2, 4, 6, 8)
list({ "a": 1, "b": ["spam", "bacon", "spam"] })
## ['a', 'b']
```

Exercise 3.5 Take a look at the documentation of the **extend** method for the **list** class. The manual page suggests that this operation takes any iterable object. Feed it with a list, tuple, range, and a string and see what happens.

The notion of iterable objects is essential, as they appear in many contexts. There exist other iterable types that are, for example, non-sequential: we cannot access their elements at random using the index operator.

Exercise 3.6 (*) Check out the **enumerate**, **zip**, and **reversed** functions and what kind of iterable objects they return.

3.4.1 The for loop

The **for** loop allows to perform a specific action on each element in an iterable object. For instance, we can access consecutive items in a list as follows:

```
x = [1, "two", ["three", 3j, 3], False] # some iterable object
for el in x: # for each element in `x`, let's call it `el`...
    print(el) # ... do something on `el`
## 1
## two
## ['three', 3j, 3]
## False
```

Another common pattern is to traverse a sequential object by means of element indexes:

```

for i in range(len(x)): # for i = 0, 1, ..., len(x)-1
    print(i, x[i], sep=": ") # sep (label separator) defaults to " "
## 0: 1
## 1: two
## 2: ['three', 3j, 3]
## 3: False

```

Example 3.7 Let's compute the elementwise multiplication of two vectors of equal lengths, i.e., the product of their corresponding elements:

```

x = [1, 2, 3, 4, 5] # for testing
y = [1, 10, 100, 1000, 10000] # just a test
z = [] # result list - start with an empty one
for i in range(len(x)):
    tmp = x[i] * y[i]
    print(f"The product of {x[i]:6} and {y[i]:6} is {tmp:6}")
    z.append(tmp)
## The product of      1 and      1 is      1
## The product of      2 and     10 is     20
## The product of      3 and    100 is    300
## The product of      4 and   1000 is   4000
## The product of      5 and  10000 is  50000

```

The items were printed with a little help of f-strings; see [Section 2.1.3](#). Here is the resulting list:

```

print(z)
## [1, 20, 300, 4000, 50000]

```

Example 3.8 A dictionary may be useful for recoding lists of labels:

```

map = dict( # from=to
    apple="red",
    pear="yellow",
    kiwi="green",
)

```

And now:

```

x = ["apple", "pear", "apple", "kiwi", "apple", "kiwi"]
recoded_x = []
for fruit in x:
    recoded_x.append(map[fruit]) # or, e.g., map.get(fruit, "unknown")

print(recoded_x)
## ['red', 'yellow', 'red', 'green', 'red', 'green']

```

Exercise 3.9 Here is a function that determines the minimum of a given iterable object (compare the built-in `min` function, see `help("min")`).

```

import math
def mymin(x):
    """
    Fetches the smallest element in an iterable object x.
    We assume that x consists of numbers only.
    """
    curmin = math.inf # infinity is greater than any other number
    for e in x:
        if e < curmin:
            curmin = e # a better candidate for the minimum
    return curmin

mymin([0, 5, -1, 100])
## -1
mymin(range(5, 0, -1))
## 1
mymin((1,))
## 1

```

Note that due to the use of `math.inf`, the function operates under the assumption that all elements in `x` are numeric. Rewrite it so that it will work correctly, e.g., in the case of lists of strings.

Exercise 3.10 Using the `for` loop, author some basic versions of the built-in `max`, `sum`, `any`, and `all` functions.

Exercise 3.11 (*) The `glob` function in the `glob` module lists all files in a given directory whose names match a specific wildcard, e.g., `glob.glob("~/Music/*.mp3")` gives the list of MP3 files in the current user's home directory; see [Section 13.6.1](#). Moreover, `getsize` from the `os.path` module returns the size of a file, in bytes. Compose a function that determines the total size of all the files in a given directory.

3.4.2 Tuple assignment

We can create many variables in one line of code by using the syntax `tuple_of_ids = iterable_object`, which unpacks the iterable object on the right side of the assignment operator:

```

a, b, c = [1, "two", [3, 3j, "three"]]
print(a)
## 1
print(b)
## two
print(c)
## [3, 3j, 'three']

```

This is useful, for example, when the swapping of two elements is needed:

```

a, b = 1, 2 # the same as (a, b) = (1, 2) - parentheses are optional
a, b = b, a # swap a and b

```

(continues on next page)

(continued from previous page)

```
print(a)
## 2
print(b)
## 1
```

Another use case is where we fetch outputs of functions that return many objects at once. For instance, later we will learn about `numpy.unique` which (depending on arguments passed) may return a tuple of arrays:

```
import numpy as np
result = np.unique([1, 2, 1, 2, 1, 1, 3, 2, 1], return_counts=True)
print(result)
## (array([1, 2, 3]), array([5, 3, 1]))
```

That this is a tuple of length two can be verified² as follows:

```
type(result), len(result)
## (<class 'tuple'>, 2)
```

Now, instead of:

```
values = result[0]
counts = result[1]
```

we can write:

```
values, counts = np.unique([1, 2, 1, 2, 1, 1, 3, 2, 1], return_counts=True)
```

This gives two separate variables, each storing a different array:

```
print(values)
## [1 2 3]
print(counts)
## [5 3 1]
```

If only the second item is of our interest, we can write:

```
counts = np.unique([1, 2, 1, 2, 1, 1, 3, 2, 1], return_counts=True)[1]
print(counts)
## [5 3 1]
```

because a tuple is a sequential object.

Example 3.12 (*) *The `dict.items` method generates an iterable object that can be used to traverse through all the (key, value) pairs:*

² We should have already been able to tell that by merely looking at the result: note the round brackets and the two objects separated by a comma.

```
x = { "a": 1, "b": ["spam", "bacon", "spam"] }
print(list(x.items())) # just a demo
## [('a', 1), ('b', ['spam', 'bacon', 'spam'])]
```

We can thus utilise tuple assignments in contexts such as:

```
for k, v in x.items(): # or: for (k, v) in x.items()...
    print(k, v, sep=" ")
## a: 1
## b: ['spam', 'bacon', 'spam']
```

Note (*) If there are more values to unpack than then number of identifiers, we can use the notation like `*name` inside the `tuple_of_identifiers` on the left side of the assignment operator. Such a placeholder gathers all the *surplus* objects in the form of a list:

```
for a, b, *c, d in [range(4), range(10), range(3)]:
    print(a, b, c, d, sep="; ")
## 0; 1; [2]; 3
## 0; 1; [2, 3, 4, 5, 6, 7, 8]; 9
## 0; 1; []; 2
```

3.4.3 Argument unpacking (*)

Sometimes we will need to call a function with many parameters or call a series of functions with similar arguments, e.g., when plotting many objects using the same plotting style like colour, shape, font. In such scenarios, it may be convenient to *pre-prepare* the data to be passed as their inputs before making the actual call.

Consider a function that takes four arguments and prints them out obtusely:

```
def test(a, b, c, d):
    "It is just a test - print the given arguments"
    print("a = ", a, ", b = ", b, ", c = ", c, ", d = ", d, sep="")
```

Arguments to be *matched positionally* can be wrapped inside any *iterable* object and then unpacked using the *asterisk operator*:

```
args = [1, 2, 3, 4] # merely an example
test(*args) # just like test(1, 2, 3, 4)
## a = 1, b = 2, c = 3, d = 4
```

Keyword arguments can be wrapped inside a *dictionary* and unpacked with a *double asterisk*:

```
kwargs = dict(a=1, c=3, d=4, b=2)
```

(continues on next page)

(continued from previous page)

```
test(**kwargs)
## a = 1, b = 2, c = 3, d = 4
```

The unpackings can be intertwined. For this reason, the following calls are equivalent:

```
test(1, *range(2, 4), 4)
## a = 1, b = 2, c = 3, d = 4
test(1, **dict(d=4, c=3, b=2))
## a = 1, b = 2, c = 3, d = 4
test(*range(1, 3), **dict(d=4, c=3))
## a = 1, b = 2, c = 3, d = 4
```

3.4.4 Variadic arguments: `*args` and `**kwargs` (*)

We can also construct a function that takes any number of positional or keyword arguments by including `*args` or `**kwargs` (those are customary names) in their parameter list:

```
def test(a, b, *args, **kwargs):
    "simply prints the arguments passed"
    print(
        "a = ", a, ", b = ", b,
        ", args = ", args, ", kwargs = ", kwargs, sep=""
    )
```

For example:

```
test(1, 2, 3, 4, 5, spam=6, eggs=7)
## a = 1, b = 2, args = (3, 4, 5), kwargs = {'spam': 6, 'eggs': 7}
```

We see that `*args` gathers all the positionally matched arguments (except `a` and `b`, which were set explicitly) into a tuple. On the other hand, `**kwargs` is a dictionary that stores all keyword arguments that are not mentioned in the function's parameter list.

Exercise 3.13 *From time to time, we will be coming across `*args` and `**kwargs` in various contexts. Study what `matplotlib.pyplot.plot` uses them for (by calling `help(plt.plot)`).*

3.5 Object references and copying (*)

3.5.1 Copying references

It is important to always keep in mind that when writing:

```
x = [1, 2, 3]
y = x
```

the assignment operator does *not* create a copy of `x`; both `x` and `y` refer to the same object in the computer's memory.

Important If `x` is mutable, any change made to it will affect `y` (as, again, they are two different means to access the *same* object). This will also be true for **numpy** arrays and **pandas** data frames.

For example:

```
x.append(4)
print(y)
## [1, 2, 3, 4]
```

3.5.2 Pass by assignment

Arguments are passed to functions *by assignment* too. In other words, they behave as if `=` was used: what we get is another reference to the existing object.

```
def myadd(z, i):
    z.append(i)
```

And now:

```
myadd(x, 5)
myadd(y, 6)
print(x)
## [1, 2, 3, 4, 5, 6]
```

3.5.3 Object copies

If we find the foregoing behaviour undesirable, we can always make a copy of a fragile object. It is customary for the mutable types to be equipped with a relevant method:

```
x = [1, 2, 3]
y = x.copy()
x.append(4)
print(y)
## [1, 2, 3]
```

This did not change the object referred to as `y` because it is now a different entity.

3.5.4 Modify in place or return a modified copy?

We now know that we *can* have functions or methods that change the state of a given object. Consequently, for all the functions we apply, it is important to read their documentation to determine if they modify their inputs *in place* or if they return an entirely new object.

In particular, the `sorted` function returns a sorted version of an iterable object:

```
x = [5, 3, 2, 4, 1]
print(sorted(x)) # returns a sorted copy of x (does not change x)
## [1, 2, 3, 4, 5]
print(x) # unchanged
## [5, 3, 2, 4, 1]
```

The `list.sort` method modifies the object it is applied on in place:

```
x = [5, 3, 2, 4, 1]
x.sort() # modifies x in place and returns nothing
print(x)
## [1, 2, 3, 4, 5]
```

Additionally, `random.shuffle` is a *function* (not: a method) that changes the state of the argument:

```
x = [5, 3, 2, 4, 1]
import random
random.shuffle(x) # modifies x in place, returns nothing
print(x)
## [5, 2, 4, 3, 1]
```

Later we will learn about the `Series` class in `pandas`, which represents data frame columns. It has the `sort_values` method which, by default, returns a sorted copy of the object it acts upon:

```
import pandas as pd
x = pd.Series([5, 3, 2, 4, 1])
print(list(x.sort_values())) # inplace=False
## [1, 2, 3, 4, 5]
print(list(x)) # unchanged
## [5, 3, 2, 4, 1]
```

This behaviour can, however, be altered:

```
x = pd.Series([5, 3, 2, 4, 1])
x.sort_values(inplace=True) # note the argument now
print(list(x)) # changed
## [1, 2, 3, 4, 5]
```

Important We are always advised to study the *official*³ documentation of every function we call. Although surely some patterns arise (such as: a method is more *likely* to modify an object in place whereas a similar standalone function will be returning a copy), ultimately, the functions' developers are free to come up with some exceptions to them if they deem it more sensible or convenient.

3.6 Further reading

Our overview of the Python language is by no means exhaustive. Still, it touches upon the most important topics from the perspective of data wrangling.

We will mention a few additional standard library features later in this course: list comprehensions in [Section 5.5.7](#), exception handling in [Section 13.6.3](#), file connection in [Section 13.6.4](#), string formatting in [Section 14.3.1](#), pattern searching with regular expressions in [Section 14.4](#), etc.

We have deliberately decided *not* to introduce some language constructs which we can easily manage without (e.g., **else** clauses on **for** and **while** loops, the **match** statement) or are perhaps too technical for an introductory course (**yield**, **iter** and **next**, sets, name binding scopes, deep copying of objects, defining new classes, overloading operators, function factories and closures).

Also, we skipped the constructs that do not work well with the third-party packages we will soon be using (e.g., a notation like $x < y < z$ is not valid if the three involved variables are **numpy** vectors of lengths greater than one).

The said simplifications were brought in so the student is not overwhelmed. We strongly advocate for minimalism in software development. Python is the basis for one of many possible programming environments for exercising data science. In the long run, it is best to focus on developing the most *transferable* skills, as other software solutions might not enjoy all the Python's syntactic sugar, and vice versa.

The reader is encouraged to skim through at least the following chapters of the [official Python 3 tutorial](#)⁴:

- [3. An Informal Introduction to Python](#)⁵,
- [4. More Control Flow Tools](#)⁶,
- [5. Data Structures](#)⁷.

³ And not some random tutorial on the internet displaying numerous ads.

⁴ <https://docs.python.org/3/tutorial/index.html>

⁵ <https://docs.python.org/3/tutorial/introduction.html>

⁶ <https://docs.python.org/3/tutorial/controlflow.html>

⁷ <https://docs.python.org/3/tutorial/datastructures.html>

3.7 Exercises

Exercise 3.14 *Name the sequential objects we introduced.*

Exercise 3.15 *Is every iterable object sequential?*

Exercise 3.16 *Is `dict` an instance of a sequential type?*

Exercise 3.17 *What is the meaning of ``+`` and ``*`` operations on strings and lists?*

Exercise 3.18 *Given a list `x` of numeric scalars, how can we create a new list of the same length giving the squares of all the elements in the former?*

Exercise 3.19 (*) *How can we make an object copy and when should we do so?*

Exercise 3.20 *What is the difference between `x[0]`, `x[1]`, `x[:0]`, and `x[:1]`, where `x` is a sequential object?*

Part II

Unidimensional data

Unidimensional numeric data and their empirical distribution

Our data wrangling adventure starts the moment we get access to loads of data points representing some measurements, such as industrial sensor readings, patient body measures, employee salaries, city sizes, etc.

For instance, consider the heights of adult females (in centimetres) in the longitudinal study called National Health and Nutrition Examination Survey (NHANES¹) conducted by the US Centres for Disease Control and Prevention.

```
heights = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +  
                        "teaching-data/master/marek/nhanes_adult_female_height_2020.txt")
```

Let's preview a few observations:

```
heights[:6] # the first six values  
## array([160.2, 152.7, 161.2, 157.4, 154.6, 144.7])
```

This is an example of *quantitative* (numeric) data. They are in the form of a series of numbers. It makes sense to apply various mathematical operations on them, including subtraction, division, taking logarithms, comparing, and so forth.

Most importantly, here, all the observations are *independent* of each other. Each value represents a different person. Our data sample consists of 4 221 points on the real line: a *bag* of points whose actual ordering does not matter. We depicted them in Figure 4.1. However, we see that merely looking at the *raw* numbers themselves tells us nothing. They are too plentiful.

This is why we are interested in studying a multitude of methods that can bring some insight into the reality behind the numbers. For example, inspecting their distribution.

¹ <https://wwwn.cdc.gov/nchs/nhanes/search/datapage.aspx>




Figure 4.1. The heights dataset is comprised of independent points on the real line. We added some jitter on the y-axis for dramatic effects only.

4.1 Creating vectors in numpy

In this chapter, we introduce the ways to create **numpy** vectors, which are an efficient data structure for storing and operating on numeric data.

numpy² [48] is an open-source add-on for numerical computing written by Travis Oliphant and other developers in 2005. Still, the project has a much longer **history**³ and stands on the shoulders of many giants (e.g., concepts from the APL and Fortran languages).

numpy implements multidimensional arrays similar to those available in R, S, GNU Octave, Scilab, Julia, Perl (via **Perl Data Language**), and some numerical analysis libraries such as **LAPACK**, **GNU GSL**, etc.

Many other Python packages discussed in this book are built on top of **numpy**, including: **scipy** [97], **pandas** [66], and **scikit-learn** [75]. This is why we will study it in such a great detail. In particular, whatever we learn about vectors will be beautifully transferable to the case of processing columns in data frames⁴.

It is customary to import the **numpy** package under the **np** alias:

² <https://numpy.org/doc/stable/reference/index.html>

³ https://scipy.github.io/old-wiki/pages/History_of_SciPy

⁴ And very similar packages for tensor processing on GPUs and other specialised computational units; e.g., **TensorFlow**, **PyTorch**, **Theano**, or **tinygrad**.

```
import numpy as np
```

Our code can now refer to the objects defined therein as `np.spam`, `np.bacon`, or `np.spam`.

4.1.1 Enumerating elements

One way to create a vector is by calling the `numpy.array` function:

```
x = np.array([10, 20, 30, 40, 50, 60])
x
## array([10, 20, 30, 40, 50, 60])
```

Here, the vector elements were specified by means of an ordinary list. Ranges and tuples can also be used as content providers. The earnest readers are encouraged to check it now themselves.

A vector of length (size) n is often used to represent a point in an n -dimensional space (for example, GPS coordinates of a place on Earth assume $n = 2$) or n readings of some one-dimensional quantity (e.g., recorded heights of n people). The said length can either be read using the previously-mentioned `len` function:

```
len(x)
## 6
```

A vector is a *one-dimensional array*. Accordingly, its *shape* is stored as a tuple of length 1 (the number of dimensions is given by querying `x.ndim`).

```
x.shape
## (6,)
```

We can therefore fetch its length also by accessing `x.shape[0]`. On a side note, matrices – two-dimensional arrays discussed in [Chapter 7](#) – will be of shape like `(number_of_rows, number_of_columns)`.

Important Recall that Python lists, e.g., `[1, 2, 3]`, represent simple sequences of objects of any kind. Their use cases are very broad, which is both an advantage and something quite the opposite. *Vectors* in `numpy` are like lists, but on steroids. They are powerful in scientific computing because of the underlying assumption that each object they store is of the same type⁵. In most scenarios, we will be dealing with vectors of logical values, integers, and floating-point numbers. Thanks to this, a wide range of

⁵ (*) Vectors are directly representable as simple arrays in the C programming language, in which the `numpy` methods are written. Operations on vectors are very fast provided that we rely on functions that process them *as a whole*. The readers with some background in other lower-level languages will need to get out of the habit of processing *individual* elements using `for`-like loops.

methods for performing the most popular mathematical operations could have been defined.

And so, above we created a sequence of integers:

```
x.dtype # data type
## dtype('int64')
```

To show that other element types are also available, we can convert it to a vector with elements of the type `float`:

```
x.astype(float) # or np.array(x, dtype=float)
## array([10., 20., 30., 40., 50., 60.] )
```

Let's emphasise this vector is printed differently from its `int`-based counterpart: note the decimal separators. Furthermore:

```
np.array([True, False, False, True])
## array([ True, False, False,  True])
```

gives a logical vector, for the array constructor detected that the common type of all the elements is `bool`. Also:

```
np.array(["spam", "spam", "bacon", "spam"])
## array(['spam', 'spam', 'bacon', 'spam'], dtype='<U5')
```

yields an array of strings in Unicode (i.e., capable of storing any character in any alphabet, emojis, mathematical symbols, etc.), each of no more than five⁶ code points in length.

4.1.2 Arithmetic progressions

`numpy`'s `arange` is similar to the built-in `range` function, but outputs a vector:

```
np.arange(0, 10, 2) # from 0 to 10 (exclusive) by 2
## array([0, 2, 4, 6, 8])
```

`numpy.linspace` creates a sequence of equidistant points on the linear scale in a given interval:

```
np.linspace(0, 1, 5) # from 0 to 1 (inclusive), 5 equispaced values
## array([0. , 0.25, 0.5 , 0.75, 1.  ])
```

Exercise 4.1 Call `help(np.linspace)` or `help("numpy.linspace")` to study the meaning of the endpoint argument. Find the same documentation page on the **numpy** project's website⁷. Another way is to use your favourite search engine such as DuckDuckGo and query “lin-

⁶ Chapter 14 will point out that replacing any element with new content results in the too-long strings' being truncated. We shall see that this can be remedied by calling `x.astype("<U10")`.

⁷ <https://numpy.org/doc/stable/reference/index.html>

`space site:numpy.org`⁸. Always remember to gather information from first-hand sources. You should become a frequent visitor to this page (and similar ones). In particular, every so often it is advisable to check out for significant updates at <https://numpy.org/news>.

4.1.3 Repeating values

`numpy.repeat` repeats *each* given value a specified number of times:

```
np.repeat(3, 6) # six 3s
## array([3, 3, 3, 3, 3, 3])
np.repeat([1, 2], 3) # three 1s, three 2s
## array([1, 1, 1, 2, 2, 2])
np.repeat([1, 2], [3, 5]) # three 1s, five 2s
## array([1, 1, 1, 2, 2, 2, 2, 2])
```

In the last case, every element from the list passed as the first argument was repeated the *corresponding* number of times, as defined by the second argument.

`numpy.tile`, on the other hand, repeats a whole sequence with *recycling*:

```
np.tile([1, 2], 3) # repeat [1, 2] three times
## array([1, 2, 1, 2, 1, 2])
```

Notice the difference between the above and the result of `numpy.repeat([1, 2], 3)`.

See also⁹ `numpy.zeros` and `numpy.ones` for some specialised versions of the foregoing functions.

4.1.4 `numpy.r_` (*)

`numpy.r_` gives perhaps the most flexible means for creating vectors involving quite a few of the aforementioned scenarios, albeit its syntax is quirky. For example:

```
np.r_[1, 2, 3, np.nan, 5, np.inf]
## array([ 1.,  2.,  3., nan,  5., inf])
```

Here, `nan` stands for a *not-a-number* and is used as a placeholder for missing values (discussed in Section 15.1) or *wrong* results, such as the square root of -1 in the domain of reals. The `inf` object, on the other hand, denotes the *infinity*, ∞ . We can think of it as a value that is too large to be represented in the set of floating-point numbers.

We see that `numpy.r_` uses square brackets instead of round ones. This is smart for we mentioned in Section 3.2.2 that slices (``:``) can only be created inside the index operator. And so:

⁸ DuckDuckGo also supports *search bangs* like “!numpy linspace”. They redirect us to the official documentation automatically.

⁹ When we write “see also”, it means that this is an exercise for the reader (Rule #3), in this case: to look something up in the official documentation.

```
np.r_[0:10:2] # like np.arange(0, 10, 2)
## array([0, 2, 4, 6, 8])
```

What is more, `numpy.r_` accepts the following syntactic sugar:

```
np.r_[0:1:5j] # like np.linspace(0, 1, 5)
## array([0. , 0.25, 0.5 , 0.75, 1.  ])
```

Here, `5j` does not have its literal meaning (a complex number). By an arbitrary convention, and only in this context, it designates the output length of the sequence to be generated. Could the `numpy` authors do that? Well, they could, and they did. End of story.

We can also combine many types of sequences into one:

```
np.r_[1, 2, [3]*2, 0:3, 0:3:3j]
## array([1. , 2. , 3. , 3. , 0. , 1. , 2. , 0. , 1.5, 3.  ])
```

4.1.5 Generating pseudorandom variates

The automatically-attached `numpy.random` module defines many functions to generate pseudorandom numbers. We will be discussing the reason for our using the *pseudo* prefix in Section 6.4, so now let's only take note of a way to sample from the uniform distribution on the unit interval:

```
np.random.rand(5) # five pseudorandom observations in [0, 1]
## array([0.49340194, 0.41614605, 0.69780667, 0.45278338, 0.84061215])
```

and to pick a few values from a given set with replacement (selecting the same value multiple times is allowed):

```
np.random.choice(np.arange(1, 10), 20) # replace=True
## array([7, 7, 4, 6, 6, 2, 1, 7, 2, 1, 8, 9, 5, 5, 9, 8, 1, 2, 6, 6])
```

4.1.6 Loading data from files

We will usually be reading whole heterogeneous tabular datasets using `pandas.read_csv`, being the topic we shall cover in Chapter 10. It is worth knowing, though, that arrays with elements of the same type can be read efficiently from text files (e.g., CSV) using `numpy.genfromtxt`. The code chunk at the beginning of this chapter serves as an example.

Exercise 4.2 Read the *population_largest_cities_unnamed*¹⁰ dataset directly from GitHub (click Raw to get access to its contents and use the URL you were redirected to, not the original one).

¹⁰ https://github.com/gagolews/teaching-data/blob/-/marek/population_largest_cities_unnamed.txt

4.2 Some mathematical notation

Mathematically, we will be denoting a number sequence like:

$$\mathbf{x} = (x_1, x_2, \dots, x_n),$$

where x_i is its i -th element, and n is its length (size). Using the programming syntax, x_i is `x[i - 1]` (because the first element is at index 0), and n corresponds to `len(x)` or, equivalently, `x.shape[0]`.

The bold font (hopefully clearly visible) is to emphasise that $\mathbf{x}$ is not an atomic entity (x), but rather a collection of items. For brevity, instead of writing “let $\mathbf{x}$ be a real-valued sequence of length n ”, we can state “let $\mathbf{x} \in \mathbb{R}^n$ ”. Here:

- the “ $\in$ ” symbol stands for “is in” or “is a member of”,
- $\mathbb{R}$ denotes the set of real numbers (the very one that includes, 0, -358745.2394 , 42 and π , amongst uncountably many others), and
- $\mathbb{R}^n$ is the set of real-valued sequences of length n (i.e., n such numbers considered at a time); e.g., $\mathbb{R}^2$ includes pairs such as $(1, 2)$, $(\pi/3, \sqrt{2}/2)$, and $(1/3, 10^3)$.

Overall, if $\mathbf{x} \in \mathbb{R}^n$, then we often say that $\mathbf{x}$ is a *sequence of n numbers*, a *(numeric) n -tuple*, an *n -dimensional real vector*, a *point in an n -dimensional real space*, or an *element of a real n -space*, etc. In many contexts, they are synonymic (although, mathematically, the devil is in the detail).

Note Mathematical notation is pleasantly *abstract* (general) in the sense that $\mathbf{x}$ can be anything, e.g., data on the incomes of households, sizes of the largest cities in some country, or heights of participants in some longitudinal study. At first glance, such a representation of objects from the so-called *real world* might seem overly simplistic, especially if we wish to store information on very complex entities. Nonetheless, in most cases, expressing them as *vectors* (i.e., establishing a set of numeric attributes that best describe them in a task at hand) is not only natural but also perfectly sufficient for achieving whatever we aim for.

Exercise 4.3 Consider the following problems:

- How would you represent patients in a medical clinic (for the purpose of conducting a research study in dietetics and sport sciences)?
- How would you represent cars in an insurance company’s database (to determine how much drivers should pay annually for the mandatory policy)?
- How would you represent students in a university (to grant them scholarships)?

In each case, list a few numeric features that best describe the reality at hand. On a side note, descriptive (categorical, qualitative) labels can always be encoded as numbers, e.g., female = 1, male = 2, but this will be the topic of [Chapter 11](#).

By $x_{(i)}$ (notice the bracket¹¹) we denote the i -th *order statistic*, that is, the i -th smallest value in x . In particular, $x_{(1)}$ is the sample minimum and $x_{(n)}$ is the maximum. Here is the same in Python:

```
x = np.array([5, 4, 2, 1, 3]) # just an example
x_sorted = np.sort(x)
x_sorted[0], x_sorted[-1] # the minimum and the maximum
## (1, 5)
```

To avoid the clutter of notation, in certain formulae (e.g., in the definition of the type-7 quantiles in Section 5.1.1), we will be assuming that $x_{(0)}$ is the same as $x_{(1)}$ and $x_{(n+1)}$ is equivalent to $x_{(n)}$.

4.3 Inspecting the data distribution with histograms

Histograms are one of the most intuitive tools for depicting the empirical distribution of a data sample. We will be drawing them using the classic plotting library `matplotlib-lib`¹² [54] (originally developed by John D. Hunter). Let's import it under its traditional alias:

```
import matplotlib.pyplot as plt
```

4.3.1 heights: A bell-shaped distribution

Let's draw a histogram of the heights dataset; see Figure 4.2.

```
counts, bins, __not_important = plt.hist(heights, bins=11,
    color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.show()
```

The data were split into 11 bins. Then, they were plotted so that the bar heights are proportional to the number of observations falling into each of the 11 intervals. The bins are non-overlapping, adjacent to each other, and of equal lengths. We can read their coordinates by looking at the bottom side of each rectangular bar. For example, circa 1200 observations fall into the interval [158, 163] (more or less) and roughly 400 into [168, 173] (approximately). To get more precise information, we can query the return objects:

```
bins # 12 interval boundaries; give 11 bins
## array([131.1          , 136.39090909, 141.68181818, 146.97272727,
##        152.26363636, 157.55454545, 162.84545455, 168.13636364,
```

(continues on next page)

¹¹ Some textbooks denote the i -th order statistic by $x_{i:n}$, but we will not.

¹² <https://matplotlib.org/>




Figure 4.2. A histogram of the heights dataset: the empirical distribution is nicely bell-shaped.

(continued from previous page)

```
##      173.42727273, 178.71818182, 184.00909091, 189.3      ] )
counts # the corresponding 11 counts
## array([  2.,  11., 116., 409., 992., 1206., 948., 404., 110.,
##        20.,   3.] )
```

This distribution is nicely symmetrical around about 160 cm. Traditionally, we are used to saying that it is in the shape of a *bell*. The most typical (*normal*, common) observations are somewhere in the middle, and the probability mass decreases quickly on both sides.

As a matter of fact, in [Chapter 6](#), we will model this dataset using a *normal* distribution and obtain an excellent fit. In particular, we will mention that observations outside the interval [139, 181] are very rare (probability less than 1%; via the 3σ rule, i.e., expected value ± 3 standard deviations).

4.3.2 income: A right-skewed distribution

For some of us, a normal distribution is often a *prototypical* one: we might expect (wish-fully think) that many phenomena enjoy similar regularities. And that is approximately the case¹³, e.g., in standardised testing (IQ and other ability scores or personality traits), physiology (the above heights), or when quantifying physical objects' attributes

¹³ In fact, we have a proposition stating that the sum or average of many observations or otherwise simpler components of some more complex entity, assuming that they are independent and follow the same (any!) distribution with finite variance, is approximately normally distributed. This is called the Central Limit Theorem and it is a very strong mathematical result.

with not-so-precise devices (distribution of measurement errors). We might be tempted to think now that *everything* is normally distributed, but this is very much untrue.

Consider another dataset. Figure 4.3 depicts the distribution of a simulated¹⁴ sample of disposable income of 1 000 randomly chosen UK households, in British Pounds, for the financial year ending 2020.

```
income = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/uk_income_simulated_2020.txt")
plt.hist(income, bins=20, color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.show()
```




Figure 4.3. A histogram of the `income` dataset: the distribution is right-skewed.

We notice that the probability density quickly increases, reaches its peak at around £15 500–£35 000, and then slowly goes down. We say that it has a *long tail on the right*, or that it is *right- or positive-skewed*. Accordingly, there are several people earning a decent amount of money. It is quite a non-normal distribution. Most people are rather unwealthy: their income is way below the typical per-capita revenue (being the average income for the whole population).

In Section 6.3.1, we will note that such a distribution is frequently encountered in biology and medicine, social sciences, or technology. For instance, the number of sexual partners or weights of humans are believed to be aligned this way.

Note Looking at Figure 4.3, we might have taken note of the relatively higher bars, as

¹⁴ For privacy and other reasons, the UK Office for National Statistics does not detail the individual taxpayers' incomes. This is why we needed to guesstimate them based on more coarse-grained data from a report published at <https://www.ons.gov.uk/peoplepopulationandcommunity>.

compared to their neighbours, at c. £100 000 and £120 000. Even though we might be tempted to try to invent a *story* about why there can be some difference in the relative probability mass, we ought to refrain from it. As our data sample is small, they might merely be due to some natural variability (Section 6.4.4). Of course, there might be some reasons behind it (theoretically), but we cannot read this only by looking at a single histogram. In other words, a histogram is a tool that we use to identify some rather general features of the data distribution (like the overall shape), not the specifics.

Exercise 4.4 *There is also the `nhanes_adult_female_weight_2020`¹⁵ dataset in our data repository, giving the weights (in kilograms) of the NHANES study participants. Draw its histogram. Does its shape resemble the `income` or `heights` distribution more?*

4.3.3 How many bins?

Unless some stronger assumptions about the data distribution are made, choosing the right number of bins is more art than science:

- too many will result in a rugged histogram,
- too few might cause us to miss some important details.

Figure 4.4 illustrates this.

```
plt.subplot(1, 2, 1) # one row, two columns; the first plot
plt.hist(income, bins=5, color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.subplot(1, 2, 2) # one row, two columns; the second plot
plt.hist(income, bins=200, color="lightgray", edgecolor="black")
plt.ylabel(None)
plt.show()
```

For example, in the histogram with five bins, we miss the information that the c. £20 000 income is more popular than the c. £10 000 one. (as given by the first two bars in Figure 4.3). On the other hand, the histogram with 200 bins seems to be too fine-grained already.

Important Usually, the “truth” is probably somewhere in-between. When preparing histograms for publication (e.g., in a report or on a webpage), we might be tempted to think “one must choose one and only one bin count”. In fact, we do not have to (despite that some will complain about it). Remember that it is we who are responsible for the data’s being presented in the most unambiguous a fashion possible. Providing two or three histograms can sometimes be a much better idea.

Further, we should be aware that someone might want to trick us by choosing the number of bins that depict the reality in a good light, when the object matter is slightly

¹⁵ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_weight_2020.txt




Figure 4.4. Too few and too many histogram bins (the income dataset).

more nuanced. For instance, the histogram on the left side of [Figure 4.4](#) hides the poorest households inside the first bar: the first income bracket is very wide. If we cannot request access to the original data, the best we can do is to simply ignore such a data visualisation instance and warn others not to trust it. A true data scientist must be sceptical.

Also, note that in the right histogram, we exactly know what is the income of the wealthiest person. From the perspective of privacy, this might be a bit unsensitive.

The documentation of `matplotlib.pyplot.hist` states that the `bins` argument is passed to `numpy.histogram_bin_edges` to determine the intervals into which our data are to be split. `numpy.histogram` uses the same function, but it returns the corresponding bin counts without plotting them.

```
counts, bins = np.histogram(income, 20)
counts
## array([131, 238, 238, 147, 95, 55, 29, 23, 10, 12, 5, 7, 4,
##        3, 2, 0, 0, 0, 0, 1])
bins
## array([ 5750. , 15460.95, 25171.9 , 34882.85, 44593.8 , 54304.75,
##        64015.7 , 73726.65, 83437.6 , 93148.55, 102859.5 , 112570.45,
##        122281.4 , 131992.35, 141703.3 , 151414.25, 161125.2 , 170836.15,
##        180547.1 , 190258.05, 199969.  ])
```

Thus, there are 238 observations both in the [15 461, 25 172) and [25 172, 34 883) intervals.

Note A table of ranges and the corresponding counts can be effective for data report-

ing. It is more informative and takes less space than a series of raw numbers, especially if we present them like in the table below.

Table 4.1. Incomes of selected British households; the bin edges are pleasantly round numbers

| income bracket [£1000s] | count |
|-------------------------|-------|
| 0–20 | 236 |
| 20–40 | 459 |
| 40–60 | 191 |
| 60–80 | 64 |
| 80–100 | 26 |
| 100–120 | 11 |
| 120–140 | 10 |
| 140–160 | 2 |
| 160–180 | 0 |
| 180–200 | 1 |

Reporting data in tabular form can also increase the privacy of the subjects (making subjects less identifiable, which is good) or hide some uncomfortable facts (which is not so good; “there are ten people in our company earning *more* than £200 000 p.a.” – this can be as much as £10 000 000, but shush).

Exercise 4.5 Find out how we can provide the `matplotlib.pyplot.hist` and `numpy.histogram` functions with custom bin breaks. Plot a histogram where the bin edges are 0, 20 000, 40 000, etc. (just like in the above table). Also let’s highlight the fact that bins do not have to be of equal sizes: set the last bin to [140 000, 200 000].

Exercise 4.6 (*) There are quite a few heuristics to determine the number of bins automatically, see `numpy.histogram_bin_edges` for a few formulae. Check out how different values of the `bins` argument (e.g., “sturges”, “fd”) affect the histogram shapes on both `income` and `heights` datasets. Each has its limitations, none is perfect, but some might be a sensible starting point for further fine-tuning.

We will get back to the topic of manual data binning in Section 11.1.4.

4.3.4 peds: A bimodal distribution (already binned)

Here are the December 2019 hourly average `pedestrian counts`¹⁶ near the Southern Cross Station in Melbourne:

```
peds = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
                    "teaching-data/master/marek/southern_cross_station_peds_2019_dec.txt")
peds
```

(continues on next page)

¹⁶ <http://www.pedestrian.melbourne.vic.gov.au/>

(continued from previous page)

```
## array([ 31.22580645,  18.38709677,  11.77419355,   8.48387097,
##         8.58064516,  58.70967742, 332.93548387, 1121.96774194,
##        2061.87096774, 1253.41935484,  531.64516129,  502.35483871,
##         899.06451613,  775.         ,  614.87096774,  825.06451613,
##        1542.74193548, 1870.48387097,  884.38709677,  345.83870968,
##        203.48387097, 150.4516129 , 135.67741935,  94.03225806])
```

This time, data have already been binned by somebody else. Consequently, we cannot use `matplotlib.pyplot.hist` to depict them. Instead, we can rely on a more low-level function, `matplotlib.pyplot.bar`; see [Figure 4.5](#).

```
plt.bar(np.arange(24)+0.5, width=1, height=peds,
        color="lightgray", edgecolor="black")
plt.xticks([0, 6, 12, 18, 24])
plt.show()
```




Figure 4.5. A histogram of the `peds` dataset: a bimodal (trimodal?) distribution.

This is an example of a bimodal (or even trimodal) distribution. There is a morning peak and an evening peak (and some analysts probably would distinguish a lunchtime one too).

4.3.5 matura: A bell-shaped distribution (almost)

Figure 4.6 depicts a histogram of another interesting dataset which also comes in a presummarised form. It gives the [distribution](#)¹⁷ of the 2019 Matura exam (secondary school exit diploma) scores in Poland (in %) – Polish literature¹⁸ at the basic level.

```
matura = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/matura_2019_polish.txt")
plt.bar(np.arange(0, 71), width=1, height=matura,
    color="lightgray", edgecolor="black")
plt.show()
```




Figure 4.6. A histogram of the *matura* dataset: a bell-shaped distribution... almost.

We probably expected the distribution to be bell-shaped. However, it is clear that someone tinkered with it. Still, knowing that:

- the examiners are good people: we teachers love our students,
- 20 points were required to pass,
- 50 points were given for an essay and beauty is in the eye of the beholder,

it all starts to make sense. Without graphically depicting this dataset, we would not know that some *anomalies* occurred: some students got a “lucky” pass grade.

4.3.6 marathon (truncated – fastest runners): A left-skewed distribution

Next, let’s consider the 37th Warsaw Marathon (2015) results.

¹⁷ https://cke.gov.pl/images/_EGZAMIN_MATURALNY_OD_2015/Informacje_o_wynikach/2019/sprawozdanie/Sprawozdanie%202019%20-%20J%C4%99zyk%20polski.pdf

¹⁸ Gombrowicz, Nałkowska, Miłosz, Tuwim, etc.; I recommend.

```
marathon = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/37_pzu_warsaw_marathon_mins.txt")
```

Here are the top five gun times (in minutes):

```
marathon[:5] # preview the first five (data are already sorted increasingly)
## array([129.32, 130.75, 130.97, 134.17, 134.68])
```

Figure 4.7 gives the histogram for the participants who finished the 42.2 km run in less than three hours, i.e., a *truncated* version of this dataset (more information about subsetting vectors using logical indexers will be given in Section 5.4).

```
plt.hist(marathon[marathon < 180], color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.show()
```




Figure 4.7. A histogram of a truncated version of the `marathon` dataset: the distribution is left-skewed.

We revealed that the data are highly *left-skewed*. This was not unexpected. There are only a few elite runners in the game, but, boy, are they fast. Yours truly wishes his personal best would be less than 180 minutes someday. We shall see. Running is fun, and so is walking; why not take a break for an hour and go outside?

Exercise 4.7 Plot the histogram of the untruncated (complete) version of this dataset.

4.3.7 Log-scale and heavy-tailed distributions

Consider the dataset on the populations of US cities in 2000:

```
cities = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/other/us_cities_2000.txt")
```

Let's focus only on the cities whose population is not less than 10 000 (another instance of truncating, this time on the other side of the distribution). Even though they constitute roughly 14% of all the US settlements, they are home to as much as about 84% of all the US citizens.

```
large_cities = cities[cities >= 10000]
len(large_cities) / len(cities)
## 0.13863320820692138
np.sum(large_cities) / np.sum(cities) # more on aggregation functions later
## 0.8351248599553305
```

Here are the populations of the five largest cities (can we guess which ones are they?):

```
large_cities[-5:] # preview last five - data are sorted increasingly
## array([1517550., 1953633., 2896047., 3694742., 8008654.])
```

The histogram is depicted in Figure 4.8.

```
plt.hist(large_cities, bins=20, color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.show()
```




Figure 4.8. A histogram of the `large_cities` dataset: the distribution is extremely heavy-tailed.

The histogram is virtually unreadable because the distribution is not just right-skewed; it is extremely *heavy-tailed*. Most cities are small, and those that are crowded,

such as New York, are *really* enormous. Had we plotted the whole dataset (cities instead of `large_cities`), the results' intelligibility would be even worse. For this reason, we should rather draw such a distribution on the *logarithmic* scale; see [Figure 4.9](#).

```
logbins = np.geomspace(np.min(large_cities), np.max(large_cities), 21)
plt.hist(large_cities, bins=logbins, color="lightgray", edgecolor="black")
plt.xscale("log")
plt.ylabel("Count")
plt.show()
```




Figure 4.9. Another histogram of the same `large_cities` dataset: the distribution is right-skewed even on a logarithmic scale.

The log-scale causes the x-axis labels not to increase linearly anymore: it is no longer based on steps of equal sizes, giving 0, 1 000 000, 2 000 000, ..., and so forth. Instead, the increases are now *geometrical*: 10 000, 100 000, 1 000 000, etc.

The current dataset enjoys a right-skewed distribution even on the logarithmic scale. Many real-world datasets behave alike; e.g., the frequencies of occurrences of words in books. On a side note, [Chapter 6](#) will discuss the Pareto distribution family which yields similar histograms.

Note (*) We relied on `numpy.geomspace` to generate bin edges manually:

```
np.round(np.geomspace(np.min(large_cities), np.max(large_cities), 21))
## array([ 10001.,  13971.,  19516.,  27263.,  38084.,  53201.,
##         74319., 103818., 145027., 202594., 283010., 395346.,
##         552272., 771488., 1077717., 1505499., 2103083., 2937867.,
##         4104005., 5733024., 8008654.])
```

Due to the fact that the natural logarithm is the inverse of the exponential function and vice versa (compare [Section 5.2](#)), equidistant points on a logarithmic scale can also be generated as follows:

```
np.round(np.exp(
    np.linspace(
        np.log(np.min(large_cities)),
        np.log(np.max(large_cities)),
        21
    )))
## array([ 10001.,  13971.,  19516.,  27263.,  38084.,  53201.,
##         74319., 103818., 145027., 202594., 283010., 395346.,
##         552272., 771488., 1077717., 1505499., 2103083., 2937867.,
##         4104005., 5733024., 8008654.] )
```

Exercise 4.8 Draw the histogram of `income` on the logarithmic scale. Does it resemble a bell-shaped distribution? We will get back to this topic in [Section 6.3.1](#).

4.3.8 Cumulative probabilities and the empirical cumulative distribution function

The histogram of the heights dataset in [Figure 4.2](#) revealed that, amongst others, 28.6% (1 206 of 4 221) of women are approximately 160.2 ± 2.65 cm tall. However, sometimes we might be more concerned with *cumulative* probabilities. Their interpretation is different; from [Figure 4.10](#), we can read that, e.g., 80% of all women are *no more than* roughly 166 cm tall (or that only 20% are taller than this height).

```
plt.hist(heights, bins=20, cumulative=True, density=True,
        color="lightgray", edgecolor="black")
plt.ylabel("Cumulative probability")
plt.show()
```

Very similar is the plot of the *empirical cumulative distribution function* (ECDF), which for a sample $\mathbf{x} = (x_1, \dots, x_n)$ we denote by $\hat{F}_n$. And so, at any given point $t \in \mathbb{R}$, $\hat{F}_n(t)$ is a step function¹⁹ that gives the proportion of observations in our sample that are not greater than t :

$$\hat{F}_n(t) = \frac{|\{i = 1, \dots, n : x_i \leq t\}|}{n}.$$

We read $|\{i = 1, \dots, n : x_i \leq t\}|$ as the number of indexes like i such that the corresponding x_i is less than or equal to t . Given the ordered inputs $x_{(1)} < x_{(2)} < \dots < x_{(n)}$, we have:

$$\hat{F}_n(t) = \begin{cases} 0 & \text{for } t < x_{(1)}, \\ k/n & \text{for } x_{(k)} \leq t < x_{(k+1)}, \\ 1 & \text{for } t \geq x_{(n)}. \end{cases}$$

¹⁹ We cannot see the steps in [Figure 4.11](#) for the points are too plentiful.




Figure 4.10. A cumulative histogram of the heights dataset.

Let's underline the fact that drawing the ECDF does not involve binning; we only need to arrange the observations in an ascending order. Then, assuming that all observations are unique (there are no ties), the arithmetic progression $1/n, 2/n, \dots, n/n$ is plotted against them; see Figure 4.11²⁰.

```
n = len(heights)
heights_sorted = np.sort(heights)
plt.plot(heights_sorted, np.arange(1, n+1)/n, drawstyle="steps-post")
plt.xlabel("$t$") # LaTeX maths
plt.ylabel("$\hat{F}_n(t)$, i.e., Prob(height $\leq t$)")
plt.show()
```

Thus, for example, the height of 150 cm is not exceeded by 10% of the women.

Note (*) Quantiles (which we introduce in Section 5.1.1) can be considered a generalised inverse of the ECDF.

4.4 Exercises

Exercise 4.9 What is the difference between `numpy.arange` and `numpy.linspace`?

²⁰ (*) We are using (La)TeX maths typesetting within "\$\dots\$" to obtain nice plot labels, see [72] for a comprehensive introduction.




Figure 4.11. The empirical cumulative distribution function for the heights dataset.

Exercise 4.10 (*) What happens when we convert a logical vector to a numeric one using the `astype` method? And what about when we convert a numeric vector to a logical one? We will discuss that later, but you might want to check it yourself now.

Exercise 4.11 Check what happens when we try to create a vector storing a mix of logical, integer, and floating-point values.

Exercise 4.12 Answer the following questions:

- What is a bell-shaped distribution?
- What is a right-skewed distribution?
- What is a heavy-tailed distribution?
- What is a multi-modal distribution?

Exercise 4.13 (*) When does logarithmic binning make sense?

Processing unidimensional data

Seldom will our datasets bring valid and valuable insights out of the box. The ones we are using for illustrational purposes in the first part of our book have already been curated. After all, it is an introductory course. We need to build the necessary skills up slowly, minding not to overwhelm the tireless reader with too much information all at once. We learn simple things first, learn them well, and then we move to more complex matters with a healthy level of confidence.

In real life, various *data cleansing* and *feature engineering* techniques will need to be applied. Most of them are based on the simple operations on vectors that we cover in this chapter:

- summarising data (for example, computing the median or sum),
- transforming values (applying mathematical operations on each element, such as subtracting a scalar or taking the natural logarithm),
- filtering (selecting or removing observations that meet specific criteria, e.g., those that are larger than the arithmetic mean ± 3 standard deviations).

Important Chapter 10 will be applying the same operations on individual data frame columns.

5.1 Aggregating numeric data

Recall that in the previous chapter, we discussed the heights and income datasets whose histograms we gave in [Figure 4.2](#) and [Figure 4.3](#), respectively. Such graphical summaries are based on binned data. They provide us with snapshots of how much probability mass is allocated in different parts of the data domain.

Instead of dealing with large datasets, we obtained a few counts. The process of binning and its textual or visual depictions is valuable in determining whether the distribution is unimodal or multimodal, skewed or symmetric around some point, what range of values contains most of the observations, and how small or large extreme values are.

Too fine a level of information granularity may sometimes be overwhelming. Besides,

revealing too much might not be a clever idea for privacy or confidentiality reasons¹. Consequently, we might be interested in even more coarse descriptions: data aggregates which reduce the whole dataset into a *single* number reflecting some of its characteristics. Such summaries can provide us with a kind of bird's-eye view of some of the dataset's aspects.

In this part, we discuss a few noteworthy *measures* of:

- *location*; e.g., central tendency measures such as the arithmetic mean and median;
- *dispersion*; e.g., standard deviation and interquartile range;
- *distribution shape*; e.g., skewness.

We also introduce *box-and-whisker plots*.

5.1.1 Measures of location

Arithmetic mean and median

Two main measures of *central tendency* are:

- *the arithmetic mean* (sometimes for simplicity called the average or simply *the mean*), defined as the sum of all observations divided by the sample size:

$$\bar{x} = \frac{(x_1 + x_2 + \dots + x_n)}{n} = \frac{1}{n} \sum_{i=1}^n x_i,$$

- *the median*, being the middle value in a sorted version of the sample if its length is odd, or the arithmetic mean of the two middle values otherwise:

$$m = \begin{cases} x_{((n+1)/2)} & \text{if } n \text{ is odd,} \\ \frac{x_{(n/2)} + x_{(n/2+1)}}{2} & \text{if } n \text{ is even.} \end{cases}$$

They can be computed using the `numpy.mean` and `numpy.median` functions.

```
heights = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_adult_female_height_2020.txt")
np.mean(heights), np.median(heights)
## (160.13679222932953, 160.1)

income = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/uk_income_simulated_2020.txt")
np.mean(income), np.median(income)
## (35779.994, 30042.0)
```

Let's underline what follows:

- for symmetric distributions, the arithmetic mean and the median are expected to be more or less equal;

¹ Nevertheless, we strongly advocate for all information of concern to the public to be openly available, so that *experienced* statisticians can put them to good use.

- for skewed distributions, the arithmetic mean will be biased towards the heavier tail.

Exercise 5.1 *Get the arithmetic mean and median for the `37_pzu_warsaw_marathon_mins` dataset mentioned in [Chapter 4](#).*

Exercise 5.2 (*) *Write a function that computes the median based on its mathematical definition and `numpy.sort`.*

Note (*) Technically, the arithmetic mean can also be computed using the `mean` method for the `numpy.ndarray` class. It will sometimes be the case that we have many ways to perform the same operation. We can even compose it manually using the `sum` function. Thus, all the following expressions are equivalent:

```
print(
    np.mean(income),
    income.mean(),
    np.sum(income)/len(income),
    income.sum()/income.shape[0]
)
## 35779.994 35779.994 35779.994 35779.994
```

Unfortunately, the `median` method for vectors is not available. As functions are more universal in `numpy`, we prefer sticking with them.

Sensitive to outliers vs robust

The arithmetic mean is strongly influenced by very large or very small observations which, in some contexts, we might refer to as *outliers*; see [Section 15.4](#). Let's invite a billionaire to swim in our income stream and check how the mean changes:

```
income2 = np.append(income, [1_000_000_000])
print(np.mean(income), np.mean(income2))
## 35779.994 1034745.2487512487
```

Comparing this new result to the previous one, oh we all feel much richer now, don't we? In fact, the arithmetic mean reflects the income each of us would get if all the wealth were gathered inside a single Santa Claus's (Robin Hood's or Joseph Stalin's) sack and then distributed equally amongst all of us. A noble idea provided that everyone contributes equally to the society which, sadly, is not the case.

On the other hand, the median is the value such that 50% of the observations are less than or equal to it and 50% of the remaining ones are not less than it. Hence, it is not at all sensitive to most of the data points on both the left and the right side of the distribution:

```
print(np.median(income), np.median(income2))
## 30042.0 30076.0
```

We cannot generally say that one measure is preferred to the other. It depends on the context (the nature of data, the requirements, etc.). Certainly, for symmetrical distributions with no outliers (e.g., heights), the mean will be better as it uses *all* data (and its efficiency can be proven for certain statistical models). For skewed distributions (e.g., income), the median has a nice interpretation, as it gives the value in the middle of the ordered sample. Remember that these data summaries allow us to look at a *single* data *aspect* only, and there can be many different, valid perspectives. The reality is complex.

Sample quantiles

Quantiles generalise the notion of the sample median and of the inverse of the empirical cumulative distribution function (Section 4.3.8). They provide us with the value that is not exceeded by the elements in a given sample with a predefined probability. Before proceeding with their formal definition, which is quite technical, let's point out that for larger sample sizes, we have the following rule of thumb.

Important For any p between 0 and 1, the p -quantile, denoted q_p , is a value dividing the sample in such a way that approximately $100p\%$ of observations are not greater than q_p , and the remaining c. $100(1 - p)\%$ are not less than q_p .

Quantiles appear under many different names, but they all refer to the same concept. In particular, we can speak about the $100p$ -th *percentiles*, e.g., the 0.5-quantile is the same as the 50th percentile. Furthermore:

- 0-quantile (q_0) is the minimum (also: `numpy.min`),
- 0.25-quantile ($q_{0.25}$) equals to the first quartile (denoted Q_1),
- 0.5-quantile ($q_{0.5}$) is the second quartile a.k.a. median (denoted Q_2 or m),
- 0.75-quantile ($q_{0.75}$) is the third quartile (denoted Q_3),
- 1.0-quantile (q_1) is the maximum (also: `numpy.max`).

Here are these five aggregates for our two example datasets:

```
np.quantile(heights, [0, 0.25, 0.5, 0.75, 1])
## array([131.1, 155.3, 160.1, 164.8, 189.3])
np.quantile(income, [0, 0.25, 0.5, 0.75, 1])
## array([ 5750. , 20669.75, 30042. , 44123.75, 199969. ])
```

Example 5.3 Let's print the aggregates neatly using *f-strings*; see Section 2.1.3:

```
wh = [0, 0.25, 0.5, 0.75, 1]
qheights = np.quantile(heights, wh)
qincome = np.quantile(income, wh)
print("      heights      income")
for i in range(len(wh)):
    print(f"wh[{i}] = {qheights[i]:.1f}, {qincome[i]:.1f}")
```

(continues on next page)

(continued from previous page)

```
print(f"q_{wh[i]:<4g} {qheights[i]:10.2f} {qincome[i]:10.2f}")
##           heights      income
## q_0         131.10      5750.00
## q_0.25      155.30     20669.75
## q_0.5       160.10     30042.00
## q_0.75      164.80     44123.75
## q_1         189.30    199969.00
```

Exercise 5.4 What is the income bracket for 95% of the most typical UK taxpayers? In other words, determine the 2.5th and 97.5th percentiles.

Exercise 5.5 Compute the midrange of *income* and *heights*, being the arithmetic mean of the minimum and the maximum (this measure is extremely sensitive to outliers).

Note (*) As we do not like the *approximately* part in the above “asymptotic definition”, in this course, we shall assume that for any $p \in [0, 1]$, the p -quantile is given by:

$$q_p = x_{(\lfloor k \rfloor)} + (k - \lfloor k \rfloor)(x_{(\lfloor k \rfloor + 1)} - x_{(\lfloor k \rfloor)}),$$

where $k = (n - 1)p + 1$ and $\lfloor k \rfloor$ is the floor function, i.e., the greatest integer less than or equal to k (e.g., $\lfloor 2.0 \rfloor = \lfloor 2.001 \rfloor = \lfloor 2.999 \rfloor = 2$, $\lfloor 3.0 \rfloor = \lfloor 3.999 \rfloor = 3$, $\lfloor -3.0 \rfloor = \lfloor -2.999 \rfloor = \lfloor -2.001 \rfloor = -3$, and $\lfloor -2.0 \rfloor = \lfloor -1.001 \rfloor = -2$).

q_p is a function that linearly interpolates between the points featuring the consecutive order statistics, $((k - 1)/(n - 1), x_{(k)})$ for $k = 1, \dots, n$. For instance, for $n = 5$, we connect the points $(0, x_{(1)})$, $(0.25, x_{(2)})$, $(0.5, x_{(3)})$, $(0.75, x_{(4)})$, $(1, x_{(5)})$. For $n = 6$, we do the same for $(0, x_{(1)})$, $(0.2, x_{(2)})$, $(0.4, x_{(3)})$, $(0.6, x_{(4)})$, $(0.8, x_{(5)})$, $(1, x_{(6)})$; see Figure 5.1.

Notice that for $p = 0.5$ we get the median regardless of whether n is even or not.

Note (**) There are many possible definitions of quantiles used in statistical software packages; see the `method` argument to `numpy.quantile`. They were nicely summarised in [56] as well as in the corresponding Wikipedia² article. They are all approximately equivalent for large sample sizes (i.e., asymptotically), but the best practice is to be explicit about which variant we are using in the computations when reporting data analysis results. Accordingly, in our case, we say that we are relying on the type-7 quantiles as described in [56]; see also [47].

In fact, simply mentioning that our computations are done with `numpy` version 1.xx (as specified in Section 1.4) implicitly implies that the default method parameters are used everywhere, unless otherwise stated. In many contexts, that is good enough.

² <https://en.wikipedia.org/wiki/Quantile>




Figure 5.1. q_p as a function of p for example vectors of length 5 (left subfigure) and 6 (right).

5.1.2 Measures of dispersion

Measures of central tendency quantify the location of the most *typical* value (whatever that means, and we already know it is complicated). That of dispersion (spread), on the other hand, will tell us how much *variability* or *diversity* is in our data. After all, when we say that the height of a group of people is 160 cm (on average) ± 14 cm (here, 2 standard deviations), the latter piece of information is a valuable addition (and is very different from the imaginary ± 4 cm case).

Some degree of variability might be welcome in certain contexts, and can be disastrous in others. A bolt factory should keep the variance of the fasteners' diameters as low as possible: this is how we define quality products (assuming that they all meet, on average, the required specification). Nevertheless, too much diversity in human behaviour, where everyone feels that they are special, is not really sustainable (but lack thereof would be extremely boring).

Let's explore the different ways in which we can quantify this data aspect.

Standard deviation (and variance)

The standard deviation³, is the root mean square distance to the arithmetic mean:

$$s = \sqrt{\frac{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + \dots + (x_n - \bar{x})^2}{n}} = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}.$$

Let's compute it with `numpy`:

```
np.std(heights), np.std(income)
## (7.062021850008261, 22888.77122437908)
```

The standard deviation quantifies the typical amount of spread around the arithmetic mean. It is overall adequate for making comparisons across different samples measuring similar things (e.g., heights of males vs of females, incomes in the UK vs in South Africa).

However, without further assumptions, it can be difficult to express the meaning of a particular value of s : for instance, the statement that the standard deviation of income is £22 900 is hard to interpret. This measure makes therefore most sense for data distributions that are symmetric around the mean.

Note (*) For bell-shaped data such as heights (more precisely: for normally-distributed samples; see the next chapter), we sometimes report $\bar{x} \pm 2s$. By the so-called 2σ rule, the theoretical expectancy is that roughly 95% of data points fall into the $[\bar{x} - 2s, \bar{x} + 2s]$ interval.

Further, the *variance* is the square of the standard deviation, s^2 . Mind that if data are expressed in centimetres, then the variance is in centimetres *squared*, which is not very intuitive. The standard deviation does not have this drawback. For many reasons, mathematicians find the square root in the definition of s annoying, though; it is why we will come across the s^2 measure every now and then too.

Interquartile range

The interquartile range (IQR) is another popular way to quantify data dispersion. It is defined as the difference between the third and the first quartile:

$$\text{IQR} = q_{0.75} - q_{0.25} = Q_3 - Q_1.$$

Computing it is effortless:

³ (**) We mean the one based on the so-called *uncorrected for statistical bias* version of the sample variance. We prefer it here for didactical reasons (simplicity, interpretability). Plus, it is the default one in `numpy`. Passing `ddof=1` (*delta degrees of freedom*) to `numpy.std` will apply division by $n - 1$ instead of by n (we will note later that the `std` methods in `pandas` have it activated by default). When used as an estimator of the distribution's standard deviation, the latter has slightly better statistical properties that we normally explore in a course on mathematical statistics, which this one is not.

```
np.quantile(heights, 0.75) - np.quantile(heights, 0.25)
## 9.5
np.quantile(income, 0.75) - np.quantile(income, 0.25)
## 23454.0
```

The IQR has an appealing interpretation: it is the range comprised of the 50% *most typical* values. It is a quite robust measure, as it ignores the 25% smallest and 25% largest observations. Standard deviation, on the other hand, is much more sensitive to outliers.

Furthermore, by *range* (or support) we will mean a measure based on extremal quantiles: it is the difference between the maximal and minimal observation.

5.1.3 Measures of shape

From a histogram, we can easily read whether a dataset is symmetric or skewed. It turns out that such a data characteristic can be easily quantified. Namely, the *skewness* is given by:

$$g = \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^3}{\left(\sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2} \right)^3}.$$

For symmetric distributions, skewness is approximately zero. Positive and negative skewness indicates a heavier right and left tail, respectively.

For example, heights are an instance of an almost-symmetric distribution:

```
import scipy.stats
scipy.stats.skew(heights)
## 0.0811184528074054
```

and income is right-skewed:

```
scipy.stats.skew(income)
## 1.9768735693998942
```

Note (*) It is worth stressing that $g > 0$ does not necessarily imply that the sample mean is greater than the median. As an alternative measure of skewness, the practitioners sometimes use:

$$g' = \frac{\bar{x} - m}{s}.$$

Yule's coefficient is an example of a robust skewness measure:

$$g'' = \frac{Q_3 + Q_1 - 2m}{Q_3 - Q_1}.$$

The computation thereof on our example datasets is left as an exercise.

Furthermore, *kurtosis* (Fisher's excess kurtosis) describes whether an empirical distribution is heavy- or thin-tailed; compare `scipy.stats.kurtosis`.

5.1.4 Box (and whisker) plots

A *box-and-whisker* plot (*box plot* for short) depicts many noteworthy features of a data sample all at the same time.

```
plt.subplot(2, 1, 1) # two rows, one column; the first subplot
plt.boxplot(heights, vert=False)
plt.yticks([1], ["heights"]) # label at y=1
plt.subplot(2, 1, 2) # two rows, one column; the second subplot
plt.boxplot(income, vert=False)
plt.yticks([1], ["income"]) # label at y=1
plt.show()
```




Figure 5.2. Example box plots.

Each box plot (compare [Figure 5.2](#)) consists of:

- the box, which spans between the first and the third quartile:
 - the median is clearly marked by a vertical bar inside the box;
 - the width of the box corresponds to the IQR;
- the whiskers, which span⁴ between:

⁴ The 1.5IQR rule is the most popular in the statistical literature, but some plotting software may use different whisker ranges. See [Section 15.4.1](#) for further discussion.

- the smallest observation (the minimum) or $Q_1 - 1.5\text{IQR}$ (the left side of the box minus 3/2 of its width), whichever is larger, and
- the largest observation (the maximum) or $Q_3 + 1.5\text{IQR}$ (the right side of the box plus 3/2 of its width), whichever is smaller.

Additionally, all observations that are less than $Q_1 - 1.5\text{IQR}$ (if any) or greater than $Q_3 + 1.5\text{IQR}$ (if any) are separately marked.

Note We are used to referring to the individually marked points as *outliers*, but it does not automatically mean there is anything *anomalous* about them. They are *atypical* in the sense that they are considerably farther away from the box. It might indicate some problems in data quality (e.g., when someone made a typo entering the data), but not necessarily. Actually, box plots are calibrated (via the nicely round magic constant 1.5) in such a way that we expect there to be no or only few outliers if the data are normally distributed. For skewed distributions, there will naturally be many outliers on either side; see [Section 15.4](#) for more details.

Box plots are based solely on sample quantiles. Most statistical packages *do not* draw the arithmetic mean. If they do, it is marked with a distinctive symbol.

Exercise 5.6 Call `matplotlib.pyplot.plot(numpy.mean(..data..), 0, "bX")` to mark the arithmetic mean with a blue cross. Alternatively, pass `showmeans=True` (amongst others) to `matplotlib.pyplot.boxplot`.

Box plots are particularly beneficial for comparing data samples with each other (e.g., body measures of men and women separately), both in terms of the relative shift (location) as well as spread and skewness; see, e.g., [Figure 12.1](#).

Example 5.7 (*) A violin plot, see [Figure 5.3](#), represents a kernel density estimator, which is a smoothened version of a histogram; see [Section 15.4.2](#).

```
plt.subplot(2, 1, 1) # two rows, one column; the first subplot
plt.violinplot(heights, vert=False, showextrema=False)
plt.boxplot(heights, vert=False)
plt.yticks([1], ["heights"])
plt.subplot(2, 1, 2) # two rows, one column; the second subplot
plt.violinplot(income, vert=False, showextrema=False)
plt.boxplot(income, vert=False)
plt.yticks([1], ["income"])
plt.show()
```

5.1.5 Other aggregation methods (*)

We said that the arithmetic mean is overly sensitive to extreme observations. The sample median is an example of a *robust* aggregate: it ignores all but 1–2 middle observations (we say it has a high *breakdown point*). Some measures of central tendency that are in-between the mean–median extreme include:




Figure 5.3. Example violin plots.

- *trimmed means* – the arithmetic mean of all the observations except several, say p , smallest and greatest ones,
- *winsorised means* – the arithmetic mean with p smallest and p greatest observations replaced with the $(p + 1)$ -th smallest and the $(p + 1)$ -th greatest one, respectively.

The two other famous means are the *geometric* and *harmonic* ones. The former is more meaningful for averaging growth rates and speedups whilst the latter can be used for computing the average speed from speed measurements at sections of identical lengths; see also the notion of the F measure in [Section 12.3.2](#). Also, the *quadratic* mean is featured in the definition of the standard deviation (it is the quadratic mean of the distances to the mean).

As far as spread measures are concerned, the interquartile range (IQR) is a robust statistic. If necessary, the standard deviation might be replaced with:

- mean absolute deviation from the mean: $\frac{1}{n} \sum_{i=1}^n |x_i - \bar{x}|$,
- mean absolute deviation from the median: $\frac{1}{n} \sum_{i=1}^n |x_i - m|$,
- median absolute deviation from the median: the median of $(|x_1 - m|, |x_2 - m|, \dots, |x_n - m|)$.

The *coefficient of variation*, being the standard deviation divided by the arithmetic mean, is an example of a *relative* (or normalised) spread measure. It can be appropriate for comparing data on different scales, as it is unitless (think how standard deviation changes when you convert between metres and centimetres).

The *Gini index*, widely used in economics, can also serve as a measure of relative dis-

persion, but assumes that all data points are nonnegative:

$$G = \frac{\sum_{i=1}^n \sum_{j=1}^n |x_i - x_j|}{2(n-1)n\bar{x}} = \frac{\sum_{i=1}^n (n-2i+1)x_{(n-i+1)}}{(n-1)\sum_{i=1}^n x_i}.$$

It is normalised so that it takes values in the unit interval. An index of 0 reflects the situation where all values in a sample are the same (0 variance; perfect equality). If there is a single entity in possession of all the “wealth”, and the remaining ones are 0, then the index is equal to 1.

For a more generic (algebraic) treatment of aggregation functions for unidimensional data; see, e.g., [13, 32, 33, 45]. Some measures might be preferred under certain (often strict) assumptions usually explored in a course on mathematical statistics, e.g., [41].

Overall, numerical aggregates can be used in cases where data are unimodal. For multimodal mixtures or data in groups, they should rather be applied to summarise each cluster/class separately; compare [Chapter 12](#). Also, [Chapter 8](#) will extend some of the summaries for the case of multidimensional data.

5.2 Vectorised mathematical functions

numpy, just like any other comprehensive numerical computing environment (e.g., R, GNU Octave, Scilab, and Julia), gives access to many mathematical functions:

- absolute value: `numpy.abs`,
- square and square root: `numpy.square` and `numpy.sqrt`, respectively,
- (natural) exponential function: `numpy.exp`,
- logarithms: `numpy.log` (the natural logarithm, i.e., base e) and `numpy.log10` (base 10),
- trigonometric functions: `numpy.cos`, `numpy.sin`, `numpy.tan`, etc., and their inverses: `numpy.arccos`, etc.
- rounding and truncating: `numpy.round`, `numpy.floor`, `numpy.ceil`, `numpy.trunc`.

Important The classical handbook of mathematical functions and the (in)equalities related to them is [1]; see [73] for its updated version.

Each of the aforementioned functions is *vectorised*. Applying, say, f , on a vector like $\mathbf{x} = (x_1, \dots, x_n)$, we obtain a sequence of the same length with all elements appropriately transformed:

$$f(\mathbf{x}) = (f(x_1), f(x_2), \dots, f(x_n)).$$

In other words, f operates *element by element* on the whole array. Vectorised operations are frequently used for making adjustments to data, e.g., as in [Figure 6.8](#), where we discover that the *logarithm* of the UK incomes has a bell-shaped distribution.

Here is an example call to the vectorised version of the rounding function:

```
np.round([-3.249, -3.151, 2.49, 2.51, 3.49, 3.51], 1)
## array([-3.2, -3.2, 2.5, 2.5, 3.5, 3.5])
```

The input list has been automatically converted to a **numpy** vector.

Important Thanks to the vectorised functions, our code is not only more readable, but also runs faster: we do not have to employ the generally slow Python-level **while** or **for** loops to traverse through each element in a given sequence.

5.2.1 Logarithms and exponential functions

Let's list some well-known properties of the natural logarithm and its inverse, the exponential function. By convention, Euler's number $e \simeq 2.718$, $\log x = \log_e x$, and $\exp(x) = e^x$.

- $\log 1 = 0, \log e = 1$,
- $\log x^y = y \log x$ and hence $\log e^x = x$,
- $\log(xy) = \log x + \log y$ and thus $\log(x/y) = \log x - \log y$,
- $e^0 = 1, e^1 = e$,
- $e^{\log x} = x$,
- $e^{x+y} = e^x e^y$ and so $e^{x-y} = e^x / e^y$,
- $e^{xy} = (e^x)^y$.

Logarithms are only defined for $x > 0$. Both functions are strictly increasing. For $x \geq 1$, the logarithm grows very slowly whereas the exponential function increases very rapidly; see [Figure 5.4](#). In the limit as $x \rightarrow 0$, we have that $\log x \rightarrow -\infty$. On the other hand, $e^x \rightarrow 0$ as $x \rightarrow -\infty$.

```
plt.subplot(1, 2, 1)
x = np.linspace(np.exp(-2), np.exp(3), 1001)
plt.plot(x, np.log(x), label="$y=\\log x$")
plt.legend()
plt.subplot(1, 2, 2)
x = np.linspace(-2, 3, 1001)
plt.plot(x, np.exp(x), label="$y=\\exp(x)$")
plt.legend()
plt.show()
```




Figure 5.4. The natural logarithm (left) and the exponential function (right).

Logarithms of different bases and non-natural exponential functions are also available. In particular, when drawing plots, we used the base-10 logarithmic scales on the axes. We have $\log_{10} x = \frac{\log x}{\log 10}$ and its inverse is $10^x = e^{x \log 10}$. For example:

```
10.0**np.array([-1, 0, 1, 2]) # exponentiation; see below
## array([ 0.1,  1. , 10. , 100. ])
np.log10([-1, 0.01, 0.1, 1, 2, 5, 10, 100, 1000, 10000])
## <string>:1: RuntimeWarning: invalid value encountered in log10
## array([ nan, -2. , -1. ,  0. ,  0.30103,  0.69897,
##        1. ,  2. ,  3. ,  4. ])
```

Take note of the warning and the not-a-number (NaN) generated.

Exercise 5.8 Check that when using the log-scale on the x-axis (`plt.xscale("log")`), the plot of the logarithm (of any base) is a straight line. Similarly, the log-scale on the y-axis (`plt.yscale("log")`) makes the exponential function linear.

5.2.2 Trigonometric functions

The trigonometric functions in **numpy** accept angles in radians. If x is the degree in angles, then to compute its cosine, we should instead write $\cos(x\pi/180)$.

Figure 5.5 depicts the cosine and sine functions.

```
x = np.linspace(-2*np.pi, 4*np.pi, 1001)
plt.plot(x, np.cos(x), label="cos") # black solid line (default style)
plt.plot(x, np.sin(x), 'r:', label="sin") # red dotted line
```

(continues on next page)

(continued from previous page)

```
plt.xticks(
    [-2*np.pi, -np.pi, 0, np.pi/2, np.pi, 3*np.pi/2, 2*np.pi, 4*np.pi],
    ["$-2\\pi$", "$-\\pi$", "$0$", "$\\pi/2$", "$\\pi$",
     "$3\\pi/2$", "$2\\pi$", "$4\\pi$"]
)
plt.legend()
plt.show()
```




Figure 5.5. The cosine and the sine functions.

Some identities worth memorising, which we shall refer to later:

- $\sin x = \cos(\pi/2 - x)$,
- $\cos(-x) = \cos x$,
- $\cos^2 x + \sin^2 x = 1$, where $\cos^2 x = (\cos x)^2$,
- $\cos(x + y) = \cos x \cos y - \sin x \sin y$,
- $\cos(x - y) = \cos x \cos y + \sin x \sin y$.

5.3 Arithmetic operators

We can apply the standard binary (two-argument) arithmetic operators `+`, `-`, `*`, `/`, `**`, `%`, and `//` on vectors too. Beneath we discuss the possible cases of the operands' lengths.

5.3.1 Vector-scalar case

Often, we will be referring to the binary operators in contexts where one operand is a vector and the other is a single value (scalar). For example:

```
np.array([-2, -1, 0, 1, 2, 3])**2
## array([4, 1, 0, 1, 4, 9])
(np.array([-2, -1, 0, 1, 2, 3])+2)/5
## array([0. , 0.2, 0.4, 0.6, 0.8, 1. ])
```

Each element was transformed (e.g., squared, divided by 5) and we got a vector of the same length in return. In these cases, the operators work just like the aforementioned vectorised mathematical functions.

Mathematically, we commonly assume that the scalar multiplication is performed in this way. In this book, we will also extend this behaviour to a scalar addition. Thus:

$$c\mathbf{x} + t = (cx_1 + t, cx_2 + t, \dots, cx_n + t).$$

We will also become used to writing $(\mathbf{x} - t)/c$, which is equivalent to $(1/c)\mathbf{x} + (-t/c)$.

5.3.2 Application: Feature scaling

Vector-scalar operations and aggregation functions are the basis for the popular *feature scalers* that we discuss in the sequel:

- standardisation,
- min-max scaling and clipping,
- normalisation.

They can increase the interpretability of data points by bringing them onto a common, unitless scale. They might also be essential when computing pairwise distances between multidimensional points; see [Section 8.4](#).

These transformations are of the form $\mathbf{y} = c\mathbf{x} + t$, i.e., they are *affine*. We can interpret them geometrically as scaling (stretching or shrinking) and shifting (translating); see [Figure 5.6](#) for an illustration.

Note Let $\mathbf{y} = c\mathbf{x} + t$ and let $\bar{x}, \bar{y}, s_x, s_y$ denote the vectors' arithmetic means and standard deviations. The following properties hold.

- The arithmetic mean is *equivariant* with respect to translation and scaling; we have $\bar{y} = c\bar{x} + t$. This is also true for all the quantiles (including, of course, the median).
- The standard deviation is *invariant* to translations, and *equivariant* to scaling: $s_y = cs_x$. The same happens for the interquartile range and the range.

As a byproduct, for the variance, we get $s_y^2 = c^2 s_x^2$.




Figure 5.6. Scaled and shifted versions of an example vector.

Standardisation and z-scores

A *standardised* version of a vector $\mathbf{x} = (x_1, \dots, x_n)$ consists in subtracting, from each element, the sample arithmetic mean (which we call *centring*) and then dividing it by the standard deviation, i.e., $\mathbf{z} = (\mathbf{x} - \bar{x})/s$. In other words, we transform each x_i to obtain:

$$z_i = \frac{x_i - \bar{x}}{s}.$$

Consider again the female heights dataset:

```
heights[-5:] # preview
## array([157. , 167.4, 159.6, 168.5, 147.8])
```

whose mean $\bar{x}$ and standard deviation s are equal to:

```
np.mean(heights), np.std(heights)
## (160.13679222932953, 7.062021850008261)
```

With **numpy**, standardisation is as simple as applying two aggregation functions and two arithmetic operations:

```
heights_std = (heights - np.mean(heights)) / np.std(heights)
heights_std[-5:] # preview
## array([-0.44417764,  1.02848843, -0.07601113,  1.18425119, -1.74692071])
```

What we obtained is sometimes referred to as the *z-scores*. They are nicely interpretable:

- z-score of 0 corresponds to an observation equal to the sample mean (perfectly average);
- z-score of 1 is obtained for a datum 1 standard deviation above the mean;
- z-score of -2 means that it is a value 2 standard deviations below the mean;

and so forth.

Because of the way they emerge, the mean of the z-scores is always 0 and their standard deviation is 1 (up to a tiny numerical error, as usual; see [Section 5.5.6](#)):

```
np.mean(heights_std), np.std(heights_std)
## (1.8920872660373198e-15, 1.0)
```

Even though the original heights were measured in centimetres, the z-scores are *unitless* (centimetres divided by centimetres).

Important Standardisation enables the comparison of measurements on different scales (think: height in centimetres vs weight in kilograms or apples vs oranges). It makes the most sense for bell-shaped distributions, in particular normally-distributed ones. [Section 6.1.2](#) will introduce the 2σ rule for the normal family (but not necessarily other models!). We will learn that we can *expect* that 95% of observations have z-scores between -2 and 2. Further, z-scores less than -3 and greater than 3 are highly unlikely.

Exercise 5.9 We have a patient whose height z-score is 1 and weight z-score is -1. How can we interpret this piece of information?

What about a patient whose weight z-score is 0 but BMI z-score is 2?

On a side note, sometimes we might be interested in performing some form of *robust* standardisation (e.g., for data with outliers or skewed). In such a case, we can replace the mean with the median and the standard deviation with the IQR.

Min-max scaling and clipping

A less frequently but still noteworthy transformation is called *min-max scaling* and involves subtracting the minimum and then dividing by the range, $(x - x_{(1)}) / (x_{(n)} - x_{(1)})$.

```
x = np.array([-1.5, 0.5, 3.5, -1.33, 0.25, 0.8])
(x - np.min(x)) / (np.max(x) - np.min(x))
## array([0. , 0.4 , 1. , 0.034, 0.35 , 0.46 ])
```

Here, the smallest value is mapped to 0 and the largest becomes equal to 1. Let's stress that, in this context, 0.5 does not represent the value which is equal to the mean (unless we are incredibly lucky).

Also, *clipping* can be used to replace all values less than 0 with 0 and those greater than 1 with 1.

```
np.clip(x, 0, 1)
## array([0. , 0.5 , 1. , 0. , 0.25, 0.8 ])
```

The function is, of course, flexible: another popular choice involves clipping to $[-1, 1]$. Note that this operation can also be composed by means of the vectorised pairwise minimum and maximum functions:

```
np.minimum(1, np.maximum(0, x))
## array([0. , 0.5 , 1. , 0. , 0.25, 0.8 ])
```

Normalisation (l_2 ; dividing by magnitude)

Normalisation is the scaling of a given vector so that it is of *unit length*. Usually, by *length* we mean the square root of the sum of squares, i.e., the Euclidean (l_2) norm also known as the *magnitude*:

$$\|(x_1, \dots, x_n)\| = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2} = \sqrt{\sum_{i=1}^n x_i^2}.$$

Its special case for $n = 2$ we know well from school: the length of a vector (a, b) is $\sqrt{a^2 + b^2}$, e.g., $\|(1, 2)\| = \sqrt{5} \approx 2.236$. Also, we are advised to program our brains so that when we see $\|x\|^2$ next time, we immediately think of the *sum of squares*.

Consequently, a normalised vector:

$$\frac{\mathbf{x}}{\|\mathbf{x}\|} = \left(\frac{x_1}{\|\mathbf{x}\|}, \frac{x_2}{\|\mathbf{x}\|}, \dots, \frac{x_n}{\|\mathbf{x}\|} \right),$$

can be determined via:

```
x = np.array([1, 5, -4, 2, 2.5]) # example vector
x/np.sqrt(np.sum(x**2)) # x divided by the Euclidean norm of x
## array([ 0.13834289,  0.69171446, -0.55337157,  0.27668579,  0.34585723])
```

Exercise 5.10 *Normalisation is similar to standardisation if data are already centred (when the mean was subtracted). Show that we can obtain one from the other via the scaling by $\sqrt{n}$.*

Important A common confusion is that normalisation is supposed to make data *more normally* distributed. This is not the case⁵, as we only scale (stretch or shrink) the observations here.

⁵ (*) A Box–Cox transformation [12] can help achieve this in some datasets. Chapter 6 will apply its particular case: it will turn out that the logarithm of incomes follow a normal distribution (hence, incomes follow a log-normal distribution). Generally, there is nothing “wrong” or “bad” about data’s not being normally distributed. It is just a nice feature to have in certain contexts.

Normalisation (l_1 ; dividing by sum)

At other times, we might be interested in considering the Manhattan (l_1) norm:

$$\|(x_1, \dots, x_n)\|_1 = |x_1| + |x_2| + \dots + |x_n| = \sum_{i=1}^n |x_i|,$$

being the sum of absolute values.

```
x / np.sum(np.abs(x))
## array([ 0.06896552,  0.34482759, -0.27586207,  0.13793103,  0.17241379])
```

l_1 normalisation is frequently applied on vectors of nonnegative values, whose normalised versions can be interpreted as *probabilities* or *proportions*: values between 0 and 1 which sum to 1 (or, equivalently, 100%).

Example 5.11 *Given some binned data:*

```
c, b = np.histogram(heights, [-np.inf, 150, 160, 170, np.inf])
print(c) # counts
## [ 306 1776 1773  366]
```

We can convert the counts to empirical probabilities:

```
p = c/np.sum(c) # np.abs is not needed here
print(p)
## [0.07249467 0.42075338 0.42004264 0.08670931]
```

We did not apply `numpy.abs` because the values were already nonnegative.

5.3.3 Vector-vector case

We have been applying ``*`, `+``, etc., so far on vectors and scalars only. All arithmetic operators can also be applied on two vectors of equal lengths. In such a case, they will act *elementwisely*: taking each element from the first operand and combining it with the *corresponding* element from the second argument:

```
np.array([2, 3, 4, 5]) * np.array([10, 100, 1000, 10000])
## array([ 20,  300, 4000, 50000])
```

We see that the first element in the left operand (2) was multiplied by the first element in the right operand (10). Then, we multiplied 3 by 100 (the second corresponding elements), and so forth. Such a behaviour of the binary operators is inspired by the usual convention in vector algebra where applying $+$ (or $-$) on $\mathbf{x} = (x_1, \dots, x_n)$ and $\mathbf{y} = (y_1, \dots, y_n)$ means exactly:

$$\mathbf{x} + \mathbf{y} = (x_1 + y_1, x_2 + y_2, \dots, x_n + y_n).$$

Using other operators this way (elementwisely) is less standard in mathematics (for instance multiplication might denote the dot product), but in `numpy` it is really convenient.

Example 5.12 Let's compute the value of the expression $h = -(p_1 \log p_1 + \dots + p_n \log p_n)$, i.e., $h = -\sum_{i=1}^n p_i \log p_i$ (the entropy):

```
p = np.array([0.1, 0.3, 0.25, 0.15, 0.12, 0.08]) # example vector
-np.sum(p*np.log(p))
## 1.6790818544987114
```

It involves the use of a unary vectorised minus (change sign), an aggregation function (sum), a vectorised mathematical function (log), and an elementwise multiplication of two vectors of the same lengths.

Example 5.13 Assume we would like to plot two mathematical functions: the sine, $f(x) = \sin x$, and a polynomial of degree 7, $g(x) = x - x^3/6 + x^5/120 - x^7/5040$ for x in the interval $[-\pi, 3\pi/2]$. To do this, we can probe the values of f and g at sufficiently many points using the vectorised operations, and then use the `matplotlib.pyplot.plot` function to draw what we see in Figure 5.7.

```
x = np.linspace(-np.pi, 1.5*np.pi, 1001) # many points in the said interval
yf = np.sin(x)
yg = x - x**3/6 + x**5/120 - x**7/5040
plt.plot(x, yf, 'k-', label="f(x)") # black solid line
plt.plot(x, yg, 'r:', label="g(x)") # red dotted line
plt.legend()
plt.show()
```




Figure 5.7. With vectorised functions, it is easy to generate plots like this one. We used different line styles so that the plot is readable also when printed in black and white.

Decreasing the number of points in x will reveal that the plotting function merely draws a series of straight-line segments. Computer graphics is essentially discrete.

Exercise 5.14 Using a single line of code, compute the vector of BMIs of all people in the *nhanes_adult_female_height_2020*⁶ and *nhanes_adult_female_weight_2020*⁷ datasets. It is assumed that the *i*-th elements therein both refer to the same person.

5.4 Indexing vectors

Recall from Section 3.2.1 and Section 3.2.2 that sequential objects in Python (lists, tuples, strings, ranges) support indexing with scalars and slices:

```
x = [10, 20, 30, 40, 50]
x[1] # scalar index: extract
## 20
x[1:2] # slice index: subset
## [20]
```

numpy vectors support two *additional* indexer types: integer and boolean sequences.

5.4.1 Integer indexing

First, indexing with a single integer *extracts* a particular element:

```
x = np.array([10, 20, 30, 40, 50])
x[0] # the first
## 10
x[1] # the second
## 20
x[-1] # the last
## 50
```

Second, we can also use lists or vectors of integer indexes. They return a subvector with elements at the specified indexes:

```
x[ [0] ]
## array([10])
x[ [0, 1, -1, 0, 1, 0, 0] ]
## array([10, 20, 50, 10, 20, 10, 10])
x[ [] ]
## array([], dtype=int64)
```

Spaces between the square brackets were added only for readability, as `x[[0]]` looks slightly more obscure. (What are these double square brackets? Nah, it is a list inside the index operator.)

⁶ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_height_2020.txt

⁷ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_weight_2020.txt

To return the vector with the elements *except* at the given indexes, we call the `numpy.delete` function. Its name is slightly misleading for the input vector is not modified in place.

```
np.delete(x, [0, -1]) # except the first and the last element
## array([20, 30, 40])
```

5.4.2 Logical indexing

Subsetting using a logical vector of the same length as the indexed vector is possible too:

```
x[ [True, False, True, True, False] ]
## array([10, 30, 40])
```

It returned the first, third, and fourth element (select the first, omit the second, choose the third, pick the fourth, skip the fifth).

Such type of indexing is particularly useful as a *data filtering* technique. Knowing that the relational vector operators `<`, `<=`, `=`, `!=`, `>=`, and `>` are performed elementwisely, just like `+`, `*`, etc., for instance:

```
x >= 30 # elementwise comparison
## array([False, False, True, True, True])
```

we can write:

```
x[ x >= 30 ] # indexing by a logical vector
## array([30, 40, 50])
```

to mean “select the elements in `x` which are not less than 30”. Of course, the indexed vector and the vector specifying the *filter* do not⁸ have to be the same:

```
y = (x/10) % 2 # whatever
y # equal to 0 if a number is a multiply of 10 times an even number
## array([1., 0., 1., 0., 1.])
x[ y == 0 ]
## array([20, 40])
```

Important Sadly, if we wish to combine many logical vectors, we cannot use the **and**, **or**, and **not** operators. They are not vectorised (this is a limitation at the language level). Instead, in `numpy`, we use the `&`, `|`, and `~` operators. Alas, they have a higher order of precedence than `<`, `<=`, `=`, etc. Therefore, the bracketing of the comparisons is obligatory.

⁸ (*) The indexer is computed first, and its *value* is passed as an argument to the index operator. Python neither is a symbolic programming language, nor does it feature any nonstandard evaluation techniques. In other words, `[...]` does not care how the indexer was obtained.

For example:

```
x[ (20 <= x) & (x <= 40) ] # check what happens if we skip the brackets
## array([20, 30, 40])
```

means “elements in x between 20 and 40” (greater than or equal to 20 and less than or equal to 40). Also:

```
len(x[ (x < 15) | (x > 35) ])
## 3
```

computes the number of elements in x which are less than 15 or greater than 35 (are not between 15 and 35).

Exercise 5.15 Compute the BMIs only of the women whose height is between 150 and 170 cm.

5.4.3 Slicing

Just as with ordinary lists, slicing with ``:`` fetches the elements at indexes in a given range like `from:to` or `from:to:by`.

```
x[:3] # the first three elements
## array([10, 20, 30])
x[::2] # every second element
## array([10, 30, 50])
x[1:4] # from the second (inclusive) to the fifth (exclusive)
## array([20, 30, 40])
```

Important For efficiency reasons, slicing returns a *view* of existing data. It does not have to make an independent copy of the subsetted elements: by definition, sliced ranges are *regular*.

In other words, both x and its sliced version share the same memory. This is important when we apply operations which modify a given vector in place, such as the `sort` method that we mention in the sequel.

```
def zilchify(x):
    x[:] = 0 # re-sets all values in x

y = np.array([6, 4, 8, 5, 1, 3, 2, 9, 7])
zilchify(y[::2]) # modifies parts of y in place
y # has changed
## array([0, 4, 0, 5, 0, 3, 0, 9, 0])
```

It zeroed every second element in y . On the other hand, indexing with an integer or logical vector always returns a copy.


```

zilchify(y[ [1, 3, 5, 7] ]) # modifies a new object and then forgets about it
y # has not changed since the last modification
## array([0, 4, 0, 5, 0, 3, 0, 9, 0])

```

The original vector has *not* been modified, because we applied the function on a *new* (temporary) object. However, we note that compound operations such as `+=` work differently, because setting elements at specific indexes is always possible:

```

y[ [0, -1] ] += 7 # the same as y[ [0, 7] ] = y[ [0, 7] ] + 7
y
## array([7, 4, 0, 5, 0, 3, 0, 9, 7])

```

5.5 Other operations

5.5.1 Cumulative sums and iterated differences

Recall that the ``+`` operator acts on two vectors elementwisely and that the `numpy.sum` function aggregates all values into a single one. We have a similar function, but vectorised in a slightly different fashion: `numpy.cumsum` returns the vector of *cumulative sums*:

```

np.cumsum([5, 3, -4, 1, 1, 3])
## array([5, 8, 4, 5, 6, 9])

```

It gave, in this order: the first element, the sum of the first two elements, the sum of the first three elements, ..., the sum of all elements.

Iterated differences are a somewhat inverse operation:

```

np.diff([5, 8, 4, 5, 6, 9])
## array([ 3, -4,  1,  1,  3])

```

It returned the difference between the second and the first element, then the difference between the third and the second, and so forth. The resulting vector is one element shorter than the input one.

We often make use of cumulative sums and iterated differences when processing time series, e.g., stock exchange data (e.g., by how much the price changed since the previous day?; [Section 16.3.1](#)) or determining cumulative distribution functions ([Section 4.3.8](#)).

5.5.2 Sorting

The `numpy.sort` function returns a sorted copy of a given vector, i.e., determines the order statistics.

```
x = np.array([50, 30, 10, 40, 20, 30, 50])
np.sort(x)
## array([10, 20, 30, 30, 40, 50, 50])
```

The `sort` method, on the other hand, sorts the vector in place (and returns nothing).

```
x # before
## array([50, 30, 10, 40, 20, 30, 50])
x.sort()
x # after
## array([10, 20, 30, 30, 40, 50, 50])
```

Exercise 5.16 Readers concerned more with chaos than bringing order should give *numpy.random.permutation* a try: it shuffles the elements in a given vector.

5.5.3 Dealing with tied observations

Some statistical methods, especially for continuous data⁹, assume that all observations in a vector are unique, i.e., there are no *ties*. In real life, however, some values might be recorded multiple times. For instance, two marathoners might finish their runs in exactly the same time, data can be rounded up to a certain number of fractional digits, or it just happens that observations are inherently integer. Therefore, we should be able to detect duplicated entries.

`numpy.unique` is a workhorse for dealing with tied observations.

```
x = np.array([40, 10, 20, 40, 40, 30, 20, 40, 50, 10, 10, 70, 30, 40, 30])
np.unique(x)
## array([10, 20, 30, 40, 50, 70])
```

It returned a *sorted*¹⁰ version of a given vector with duplicates removed.

We can also get the corresponding counts:

```
np.unique(x, return_counts=True) # returns a tuple of length 2
## (array([10, 20, 30, 40, 50, 70]), array([3, 2, 3, 5, 1, 1]))
```

It can help determine if all the values in a vector are unique:

```
np.all(np.unique(x, return_counts=True)[1] == 1)
## False
```

Exercise 5.17 Play with the `return_index` argument to `numpy.unique`. It permits pinpointing the indexes of the first occurrences of each unique value.

⁹ Where, theoretically, the probability of obtaining a tie is equal to 0.

¹⁰ Later we will mention `pandas.unique` which lists the values in the order of appearance.

5.5.4 Determining the ordering permutation and ranking

`numpy.argsort` returns a sequence of indexes that lead to an ordered version of a given vector (i.e., an ordering permutation).

```
x = np.array([50, 30, 10, 40, 20, 30, 50])
np.argsort(x)
## array([2, 4, 5, 1, 3, 0, 6])
```

Which means that the smallest element is at index 2, then the second smallest is at index 4, the third smallest at index 1, etc. Therefore:

```
x[np.argsort(x)]
## array([10, 20, 30, 30, 40, 50, 50])
```

is equivalent to `numpy.sort(x)`.

Note (**) If there are tied observations in `x`, `numpy.argsort(x, kind="stable")` will use a *stable* sorting algorithm ([timsort](#)¹¹, a variant of mergesort), which guarantees that the ordering permutation is unique: tied elements are placed in the order of appearance.

Next, `scipy.stats.rankdata` returns a vector of *ranks*.

```
x = np.array([50, 30, 10, 40, 20, 30, 50])
scipy.stats.rankdata(x)
## array([6.5, 3.5, 1. , 5. , 2. , 3.5, 6.5])
```

Element 10 is the smallest (“the winner”, say, the quickest racer). Hence, it ranks first. The two tied elements equal to 30 are the third/fourth on the podium (*ex aequo*). Thus, they receive the average rank of 3.5. And so on.

On a side note, there are many methods in *nonparametric* statistics (where we do not make any particular assumptions about the underlying data distribution) that are based on ranks. In particular, [Section 9.1.4](#) will cover the Spearman correlation coefficient.

Exercise 5.18 Consult the manual page of `scipy.stats.rankdata` and test various methods for dealing with ties.

Exercise 5.19 What is the interpretation of a rank divided by the length of the sample?

Note (**) Calling `numpy.argsort` on a vector representing a permutation of indexes generates its inverse. In particular, `np.argsort(np.argsort(x, kind="stable"))+1` is equivalent to `scipy.stats.rankdata(x, method="ordinal")`.

¹¹ <https://github.com/python/cpython/blob/3.12/Objects/listsort.txt>

5.5.5 Searching for certain indexes (argmin, argmax)

`numpy.argmin` and `numpy.argmax` return the index at which we can find the smallest and the largest observation in a given vector.

```
x = np.array([50, 30, 10, 40, 20, 30, 50])
np.argmin(x), np.argmax(x)
## (2, 0)
```

Using mathematical notation, the former is denoted by:

$$i = \arg \min_j x_j,$$

and we read it as “let i be the index of the smallest element in the sequence”. Alternatively, it is the *argument of the minimum*, whenever:

$$x_i = \min_j x_j,$$

i.e., the i -th element is the smallest.

If there are multiple minima or maxima, the leftmost index is returned.

We can use `numpy.flatnonzero` to fetch the indexes where a logical vector has elements equal to True (Section 11.1.2 mentions that a value equal to zero is treated as the logical False, and as True in all other cases). For example:

```
np.flatnonzero(x == np.max(x))
## array([0, 6])
```

It is a version of `numpy.argmax` that lets us decide what we would like to do with the tied maxima (there are two).

Exercise 5.20 Let x be a vector with possible ties. Create an expression that returns a randomly chosen index pinpointing one of the sample maxima.

5.5.6 Dealing with round-off and measurement errors

Mathematics tells us (the easy proof is left as an exercise for the reader) that a centred version of a given vector $\mathbf{x}$, $\mathbf{y} = \mathbf{x} - \bar{x}$, has the arithmetic mean of 0, i.e., $\bar{y} = 0$. Of course, it is also true on a computer. But is it?

```
heights_centred = (heights - np.mean(heights))
np.mean(heights_centred) == 0
## False
```

The average is actually equal to:

```
np.mean(heights_centred)
## 1.3359078775153175e-14
```

which is *almost* zero (0.000000000000134), but not *exactly* zero (it is zero for an engineer, not a mathematician). We saw a similar result in [Section 5.3.2](#), when we standardised a vector (which involves centring).

Important All floating-point operations on a computer¹² (not only in Python) are performed with *finite* precision of 15–17 decimal digits.

We know it from school. For example, some fractions cannot be represented as decimals. When asked to add or multiply them, we will always have to apply some rounding that ultimately leads to precision loss. We know that $1/3 + 1/3 + 1/3 = 1$, but using a decimal representation with one fractional digit, we get $0.3 + 0.3 + 0.3 = 0.9$. With two digits of precision, we obtain $0.33 + 0.33 + 0.33 = 0.99$. And so on. This sum will never be equal exactly to 1 when using a finite precision.

Note Our data are often imprecise by nature. When asked about people's heights, rarely will they provide a non-integer answer (assuming they know how tall they are and are not lying about it, but it is a different story). We will most likely get data rounded to 0 decimal digits. In our heights dataset, the precision is a bit higher:

```
heights[:6] # preview
## array([160.2, 152.7, 161.2, 157.4, 154.6, 144.7])
```

But still, we expect there to be some inherent *observational error*.

Moreover, errors induced at one stage will propagate onto further operations. For instance, that the heights data are not necessarily accurate, makes their aggregates such as the mean *approximate* as well. Most often, the errors should more or less cancel out, but in extreme cases, they can lead to undesirable consequences (like for some model matrices in linear regression; see [Section 9.2.9](#)).

Exercise 5.21 Compute the BMIs of all females in the NHANES study. Determine their arithmetic mean. Compare it to the arithmetic mean computed for BMIs rounded to 1, 2, 3, 4, etc., decimal digits.

There is no reason to panic, though. The rule to remember is as follows.

Important As the floating-point values are precise up to a few decimal digits, we must refrain from comparing them using the `==` operator, which tests for *exact* equality.

When a comparison is needed, we need to take some error margin $\varepsilon > 0$ into account. Ideally, instead of testing $x == y$, we either inspect the *absolute error*:

$$|x - y| \leq \varepsilon,$$

¹² Double precision float64 format as defined by the IEEE Standard for Floating-Point Arithmetic (IEEE 754).

or, assuming $y \neq 0$, the *relative error*:

$$\frac{|x - y|}{|y|} \leq \varepsilon.$$

For instance, `numpy.allclose(x, y)` checks (by default) if for all corresponding elements in both vectors, we have `numpy.abs(x-y) <= 1e-8 + 1e-5*numpy.abs(y)`, which is a combination of both tests.

```
np.allclose(np.mean(heights_centred), 0)
## True
```

To avoid sorrow surprises, even the testing of inequalities like `x >= 0` should rather be performed as, say, `x >= 1e-8`.

Note (*) Another problem is related to the fact that floats on a computer use the binary base, not the decimal one. Therefore, some fractional numbers that we *believe* to be representable exactly, require an infinite number of bits. As a consequence, they are subject to rounding.

```
0.1 + 0.1 + 0.1 == 0.3 # obviously
## False
```

This is because `0.1`, `0.1+0.1+0.1`, and `0.3` are literally represented as, respectively:

```
print(f"{0.1:.19f}, {0.1+0.1+0.1:.19f}, and {0.3:.19f}.")
## 0.1000000000000000056, 0.3000000000000000444, and 0.29999999999999989.
```

A suggested introductory reference to the topic of numerical inaccuracies is [43]; see also [51, 59] for a more comprehensive treatment of numerical analysis.

5.5.7 Vectorising scalar operations with list comprehensions

List comprehensions of the form `[ expression for name in iterable ]` are part of base Python. They create lists based on transformed versions of individual elements in a given iterable object. Hence, they might work in cases where a task at hand cannot be solved by means of vectorised `numpy` functions.

For example, here is a way to generate the squares of a few positive natural numbers:

```
[ i**2 for i in range(1, 11) ]
## [1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
```

The result can be passed to `numpy.array` to convert it to a vector. Further, given an example vector:

```
x = np.round(np.random.rand(9)*2-1, 2)
x
## array([ 0.86, -0.37, -0.63, -0.59,  0.14,  0.19,  0.93,  0.31,  0.5 ])
```

if we wish to filter out all elements that are not positive, we can write:

```
[ e for e in x if e > 0 ]
## [0.86, 0.14, 0.19, 0.93, 0.31, 0.5]
```

We can also use the ternary operator of the form `x_true if cond else x_false` to return either `x_true` or `x_false` depending on the truth value of `cond`.

```
e = -2
e**0.5 if e >= 0 else (-e)**0.5
## 1.4142135623730951
```

Combined with a list comprehension, we can write, for instance:

```
[ round(e**0.5 if e >= 0 else (-e)**0.5, 2) for e in x ]
## [0.93, 0.61, 0.79, 0.77, 0.37, 0.44, 0.96, 0.56, 0.71]
```

This gave the square root of absolute values.

There is also a tool which vectorises a scalar function so that it can be used on **numpy** vectors:

```
def clip01(x):
    """clip to the unit interval"""
    if x < 0: return 0
    elif x > 1: return 1
    else: return x

clip01s = np.vectorize(clip01) # returns a function object
clip01s([0.3, -1.2, 0.7, 4, 9])
## array([0.3, 0. , 0.7, 1. , 1. ])
```

Overall, vectorised **numpy** functions lead to faster, more readable code. However, if the corresponding operations are unavailable (e.g., string processing, reading many files), list comprehensions can serve as their reasonable replacement.

Exercise 5.22 Write equivalent versions of the above expressions using vectorised **numpy** functions. Moreover, implement them using base Python lists, the **for** loop and the **list.append** method (start from an empty list that will store the result). Use the **timeit** module to compare the run times of different approaches on large vectors.

5.6 Exercises

Exercise 5.23 What are some benefits of using a *numpy* vector over an ordinary Python list? What are the drawbacks?

Exercise 5.24 How can we interpret the possibly different values of the arithmetic mean, median, standard deviation, interquartile range, and skewness, when comparing between heights of men and women?

Exercise 5.25 There is something scientific and magical about numbers that make us approach them with some kind of respect. However, taking into account that there are many possible data aggregates, there is a risk that a party may be cherry-picking: report the value that portrays the analysed entity in a good or bad light, e.g., choose the mean instead of the median or vice versa. Is there anything that can be done about it?

Exercise 5.26 Even though, mathematically speaking, all measures can be computed on all data, it does not mean that it always makes sense to do so. For instance, some distributions will have skewness of 0. However, we should not automatically assume that they are delightfully symmetric and bell-shaped (e.g., this can be a bimodal distribution). We always ought to visualise our data. Give some examples of datasets where we need to be critical of the obtained aggregates.

Exercise 5.27 Give the mathematical definitions, use cases, and interpretations of standardisation, normalisation, and min-max scaling.

Exercise 5.28 How are *numpy.log* and *numpy.exp* related to each other? What about *numpy.log* vs *numpy.log10*, *numpy.cumsum* vs *numpy.diff*, *numpy.min* vs *numpy.argmin*, *numpy.sort* vs *numpy.argsort*, and *scipy.stats.rankdata* vs *numpy.argsort*?

Exercise 5.29 What is the difference between *numpy.trunc*, *numpy.floor*, *numpy.ceil*, and *numpy.round*?

Exercise 5.30 What happens when we apply ``+`` on two vectors of different lengths?

Exercise 5.31 List the four general ways to index a vector.

Exercise 5.32 What is wrong with the expression `x[ x >= 0 and x <= 1 ]`, where *x* is a numeric vector? What about `x[ x >= 0 & x <= 1 ]`?

Exercise 5.33 What does it mean that slicing returns a view of existing data?

Exercise 5.34 Given a numeric vector *x*, for instance:

```
np.random.seed(123)
x = np.round(np.random.randn(20), 2)
```

perform what follows.

1. Print out all values in *x* that belong to the set $[-2, -1] \cup [1, 2]$.
2. Print out the number and the fraction of nonnegative values in *x*.
3. Find the position (index) of the first value greater than 2.

4. Create two vectors x_o and x_e consisting of the items at odd and even indices, respectively.
5. Compute the range, i.e., the difference between the greatest and the smallest value.
6. Compute the midrange, i.e., the arithmetic mean of the maximum and the minimum.
7. Compute the mean of absolute values.
8. Find the values closest to and farthest away from 0.
9. Find the values closest to and farthest away from 2.

Exercise 5.35 Write a function to compute the k -winsorised mean, given a numeric vector x and $k \leq \frac{n-1}{2}$, i.e., the arithmetic mean of a version of x , where the k smallest and k greatest values were replaced by the $(k + 1)$ -th smallest and greatest value, respectively.

Exercise 5.36 (**) Reflect on the famous¹³ saying: not everything that can be counted counts, and not everything that counts can be counted.

Exercise 5.37 (**) Being a data scientist can be a frustrating job, especially when you care for some causes. Reflect on: some things that count can be counted, but we will not count them because we ran over budget or because our stupid corporate manager simply doesn't care.

Exercise 5.38 (**) Being a data scientist can be a frustrating job, especially when you care for the truth. Reflect on: some things that count can be counted, but we will not count them for some people might be offended or find it unpleasant.

Exercise 5.39 (**) Assume you were to become the benevolent dictator of a nation living on some remote island. How would you measure if your people are happy or not? Assume that you need to come up with three quantitative measures (key performance indicators). What would happen if your policy-making was solely focused on optimising those KPIs? What about the same problem but in the context your company and employees? Think about what can go wrong in other areas of life.

¹³ <https://quoteinvestigator.com/2010/05/26/everything-counts-einstein>

Continuous probability distributions

Successful data analysts deal with hundreds or thousands of datasets in their lifetimes. In the long run, at some level, most of them will be deemed *boring* (datasets, not analysts). This is because only a few common patterns will be occurring over and over again. In particular, the previously mentioned bell-shapedness and right-skewness are prevalent in the so-called real world. Surprisingly, however, this is exactly when things become scientific and interesting, allowing us to study various phenomena at an *appropriate level of generality*.

Mathematically, such idealised patterns in the histogram shapes can be formalised using the notion of a *probability density function* (PDF) of a *continuous, real-valued random variable*. Intuitively¹, a PDF is a smooth curve that would arise if we drew a histogram for the entire *population* (e.g., all women living currently on Earth and beyond or otherwise an extremely large data sample obtained by independently querying the same underlying data generating process) in such a way that the total area of all the bars is equal to 1 and the bin sizes are very small. On the other hand, a *real-valued random variable* is a theoretical process that generates quantitative data. From this perspective, a *sample* at hand is assumed to be *drawn* from a given distribution; it is a *realisation* of the underlying process.

We do not intend ours to be a course in probability theory and mathematical statistics. Rather, a one that precedes and motivates them (e.g., [23, 40, 41, 82, 83]). Therefore, our definitions must be simplified so that they are digestible. We will thus consider the following characterisation.

Important (*) We call an integrable function $f : \mathbb{R} \rightarrow \mathbb{R}$ a *probability density function*, if $f(x) \geq 0$ for all x and $\int_{-\infty}^{\infty} f(x) dx = 1$. In other words, f is nonnegative and normalised in such a way that the total area under the whole curve is 1.

For any $a < b$, we treat $\int_a^b f(x) dx$, being the area under the fragment of the $f(x)$ curve for x between a and b , as the probability of the underlying real-valued random variable's falling into the $[a, b]$ interval.

Some distributions arise more frequently than others and appear to fit empirical data or their parts particularly well [29]. In this chapter, we review a few noteworthy prob-

¹ (*) This intuition is, of course, theoretically grounded and is based on the asymptotic behaviour of the histograms as the estimators of the underlying probability density function; see, e.g., [30] and the many references therein.

ability distributions: the normal, log-normal, Pareto, and uniform families (we will also mention the chi-squared, Kolmogorov, and exponential ones in this course).

6.1 Normal distribution

A *normal (Gaussian) distribution* has a prototypical, nicely symmetric, bell-shaped density. It is described by two parameters:

- $\mu \in \mathbb{R}$ being its expected value, at which the PDF is centred,
- $\sigma > 0$ is the standard deviation, saying how much the distribution is dispersed around μ .

The probability density function of $N(\mu, \sigma)$, compare Figure 6.1, is given by:

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right).$$




Figure 6.1. The probability density functions of some normal distributions $N(\mu, \sigma)$. Note that μ is responsible for shifting and σ affects scaling/stretching of the probability mass.

6.1.1 Estimating parameters

A course in mathematical statistics, may tell us that the sample arithmetic mean $\bar{x}$ and standard deviation s are natural, statistically well-behaving *estimators* of the said para-

meters. If all observations are really drawn independently from $N(\mu, \sigma)$ each, then we will *expect* $\bar{x}$ and s to be equal to, more or less, μ and σ . Furthermore, the larger the sample size, the smaller the error.

Recall the heights (females from the NHANES study) dataset and its bell-shaped histogram in Figure 4.2.

```
heights = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_adult_female_height_2020.txt")
n = len(heights)
n
## 4221
```

Let's estimate the said parameters of the normal distribution:

```
mu = np.mean(heights)
sigma = np.std(heights, ddof=1)
mu, sigma
## (160.13679222932953, 7.062858532891359)
```

Mathematically, we will denote these two by $\hat{\mu}$ and $\hat{\sigma}$ (mu and sigma with a hat) to emphasise that they are merely guesstimates of the unknown theoretical parameters μ and σ describing the whole population. On a side note, in this context, the requested `ddof=1` estimator has slightly better statistical properties.

Figure 6.2 shows the fitted density function, i.e., the PDF of $N(160.1, 7.06)$, which we computed using `scipy.stats.norm.pdf`, on top of a histogram.

```
plt.hist(heights, density=True, color="lightgray", edgecolor="black")
x = np.linspace(np.min(heights), np.max(heights), 1000)
plt.plot(x, scipy.stats.norm.pdf(x, mu, sigma), "r--",
    label=f"PDF of N({mu:.1f}, {sigma:.2f})")
plt.ylabel("Density")
plt.legend()
plt.show()
```

We passed `density=True` to `matplotlib.pyplot.hist` to normalise the bars' heights so that their total area is 1.

At first glimpse, the density matches the histogram nicely. Before proceeding with an overview of the ways to assess the goodness-of-fit more rigorously, we should heap praise on the potential benefits of getting access to idealised *models* of our datasets.

6.1.2 Data models are useful

If (provided that, assuming that, on condition that) our sample is *really* a realisation of the independent random variables following a given distribution, or a data analyst judges that such an approximation might be justified or beneficial, then we can *reduce* them to merely a few parameters.




Figure 6.2. A histogram for the heights dataset and the probability density function of the fitted normal distribution.

We can risk assuming that the heights data follow the normal distribution (assumption 1) with parameters $\mu = 160.1$ and $\sigma = 7.06$ (assumption 2). Note that the choice of the distribution family is one thing, and the way² we estimate the underlying parameters (in our case, we use the aforementioned $\hat{\mu}$ and $\hat{\sigma}$) is another.

Creating a data model only saves storage space and computational time, but also – based on what we can learn from a course in probability and statistics (by appropriately integrating the normal PDF) – we can imply the facts such as:

- c. 68% of (i.e., a *majority*) women are $\mu \pm \sigma$ tall (the 1σ rule),
- c. 95% of (i.e., *the most typical*) women are $\mu \pm 2\sigma$ tall (the 2σ rule),
- c. 99.7% of (i.e., *almost all*) women are $\mu \pm 3\sigma$ tall (the 3σ rule).

Also, if we knew that the distribution of heights of men is also normal with some other parameters (spoiler alert: $N(173.8, 7.66)$), we could make some comparisons between the two samples. For example, we could compute the probability that a passerby who is 155 cm tall is actually a man.

² (*) Sometimes we will have many point estimators to choose from, some being more suitable than others if data are not of top quality (e.g., contain outliers). For instance, in the normal model, we can also estimate μ and σ via the sample median and $\text{IQR}/1.349$.

(**) It might also be the case that we will have to obtain the estimates of a probability distribution's parameters by numerical optimisation because there are no known open-form formulae therefor. For example, in the case of the normal family, the maximum likelihood estimation problem involves minimising $\mathcal{L}(\mu, \sigma) = \sum_{i=1}^n \left(\frac{(x_i - \mu)^2}{\sigma^2} + \log \sigma^2 \right)$ with respect to μ and σ (here, we are lucky for its solution is exactly the sample mean and standard deviation).

Exercise 6.1 How different manufacturing industries (e.g., clothing) can make use of such models? Are simplifications necessary when dealing with complexity of the real world? What are the alternatives?

Furthermore, assuming a particular model gives us access to a range of *parametric* statistical methods (ones that are derived for the corresponding family of probability distributions), e.g., the t-test to compare the expected values.

Important We should always verify the assumptions of the tool at hand before we apply it in practice. In particular, we will soon discover that the UK annual incomes are not normally distributed. Therefore, we must not refer to the aforementioned 2σ rule in their case. A hammer neither barks nor can it serve as a screwdriver. Period.

6.2 Assessing goodness-of-fit

6.2.1 Comparing cumulative distribution functions

Bell-shaped histograms are encountered fairly frequently in real-world data: e.g., measurement errors in physical experiments and standardised tests' results (like IQ and other ability scores) tend to be distributed this way, at least approximately. If we yearn for more precision, there is a better way of assessing the extent to which a sample deviates from a hypothesised distribution. Namely, we can measure the discrepancy between some theoretical *cumulative distribution function* (CDF) and the empirical one (ECDF which we defined in [Section 4.3.8](#)).

Important If f is a PDF, then the corresponding theoretical CDF is defined as $F(x) = \int_{-\infty}^x f(t) dt$, i.e., the probability of the corresponding random variable's being less than or equal to x . By definition³, each CDF takes values in the unit interval $([0, 1])$ and is nondecreasing.

For the normal distribution family, the values of the theoretical CDF can be computed by calling `scipy.stats.norm.cdf`; compare [Figure 6.4](#) below.

[Figure 6.3](#) depicts the CDF of $N(160.1, 7.06)$ and the empirical CDF of the heights dataset. This *looks* like a superb match.

```
x = np.linspace(np.min(heights), np.max(heights), 1001)
probs = scipy.stats.norm.cdf(x, mu, sigma) # sample the CDF at many points
```

(continues on next page)

³ The probability distribution of any real-valued random variable X can be uniquely defined by means of a nondecreasing, right (upward) continuous function $F : \mathbb{R} \rightarrow [0, 1]$ such that $\lim_{x \rightarrow -\infty} F(x) = 0$ and $\lim_{x \rightarrow \infty} F(x) = 1$, in which case $\Pr(X \leq x) = F(x)$. The probability density function only exists for continuous random variables and is defined as the derivative of F .

(continued from previous page)

```
plt.plot(x, probs, "r--", label=f"CDF of N({mu:.1f}, {sigma:.2f})")
heights_sorted = np.sort(heights)
plt.plot(heights_sorted, np.arange(1, n+1)/n,
         drawstyle="steps-post", label="Empirical CDF")
plt.xlabel("$x$")
plt.ylabel("Prob(height $\leq x$)")
plt.legend()
plt.show()
```




Figure 6.3. The empirical CDF and the fitted normal CDF for the heights dataset: the fit is superb.

Example 6.2 $F(b) - F(a) = \int_a^b f(t) dt$ is the probability of generating a value in the interval $[a, b]$. Let's compute the probability related to the 3σ rule:

```
F = lambda x: scipy.stats.norm.cdf(x, mu, sigma)
F(mu+3*sigma) - F(mu-3*sigma)
## 0.9973002039367398
```

A common way to summarise the discrepancy between the empirical CDF $\hat{F}_n$ and a given theoretical CDF F is by computing the greatest absolute deviation:

$$\hat{D}_n = \sup_{t \in \mathbb{R}} |\hat{F}_n(t) - F(t)|,$$

where the *supremum* is a continuous version of the maximum. We have:

$$\hat{D}_n = \max \left\{ \max_{k=1, \dots, n} \left\{ \left| \frac{k-1}{n} - F(x_{(k)}) \right| \right\}, \max_{k=1, \dots, n} \left\{ \left| \frac{k}{n} - F(x_{(k)}) \right| \right\} \right\},$$

i.e., F needs to be probed only at the n points from the sorted input sample.


```
def compute_Dn(x, F): # equivalent to scipy.stats.kstest(x, F)[0]
    Fx = F(np.sort(x))
    n = len(x)
    k = np.arange(1, n+1) # 1, 2, ..., n
    Dn1 = np.max(np.abs((k-1)/n - Fx))
    Dn2 = np.max(np.abs(k/n - Fx))
    return max(Dn1, Dn2)
```

If the difference is *sufficiently*⁴ small, then we can assume that the normal model describes data quite well.

```
Dn = compute_Dn(heights, F)
Dn
## 0.010470976524201148
```

This is indeed the case here: we may estimate the probability of someone's being as tall as the given height with an error less than about 1.05%.

6.2.2 Comparing quantiles

A Q-Q plot (quantile-quantile or probability plot) is another graphical method for comparing two distributions. This time, instead of working with a cumulative distribution function F , we will be dealing with the related quantile function Q .

Important Given a continuous⁵ CDF F , the corresponding *quantile function* Q is exactly its inverse, i.e., we have $Q(p) = F^{-1}(p)$ for all $p \in (0, 1)$.

The theoretical quantiles can be generated by the `scipy.stats.norm.ppf` function; compare Figure 6.4. Here, *ppf* stands for the percent point function which is another (yet quite esoteric) name for the above Q .

Example 6.3 In our $N(160.1, 7.06)$ -distributed *heights* dataset, $Q(0.9)$ is the height not exceeded by 90% of the female population. In other words, only 10% of American women are taller than:

```
scipy.stats.norm.ppf(0.9, mu, sigma)
## 169.18820963937648
```

A Q-Q plot draws the sample quantiles *against* the corresponding theoretical quantiles. In Section 5.1.1, we mentioned that there are a few possible definitions thereof in the literature. Thus, we have some degree of flexibility. For simplicity, instead of using

⁴ The larger the sample size, the less tolerant regarding the size of this disparity we are; see Section 6.2.3.

⁵ More generally, for an arbitrary F , Q is its *generalised* inverse, defined for any $p \in (0, 1)$ as $Q(p) = \inf\{x : F(x) \geq p\}$, i.e., the smallest x such that the probability of drawing a value not greater than x is at least p .




Figure 6.4. The cumulative distribution functions (left) and the quantile functions (being the inverse of the CDF; right) of some normal distributions.

`numpy.quantile`, we will assume that the $i/(n+1)$ -quantile⁶ is equal to $x_{(i)}$, i.e., the i -th smallest value in a given sample $(x_1, x_2, \dots, x_n)$ and consider only $i = 1, 2, \dots, n$. This way, we mitigate the problem which arises when the 0- or 1-quantiles of the theoretical distribution, i.e., $Q(0)$ or $Q(1)$, are not finite (and this is the case for the normal distribution family).

```
def qq_plot(x, Q):
    """
    Draws a Q-Q plot, given:
    * x - a data sample (vector)
    * Q - a theoretical quantile function
    """
    n = len(x)
    q = np.arange(1, n+1)/(n+1)    # 1/(n+1), 2/(n+2), ..., n/(n+1)
    x_sorted = np.sort(x)          # sample quantiles
    quantiles = Q(q)               # theoretical quantiles
    plt.plot(quantiles, x_sorted, "o")
    plt.axline((x_sorted[n//2], x_sorted[n//2]), slope=1,
               linestyle=":", color="gray") # identity line
```

Figure 6.5 depicts the Q-Q plot for our example dataset.

```
qq_plot(heights, lambda q: scipy.stats.norm.ppf(q, mu, sigma))
plt.xlabel(f"Quantiles of N({mu:.1f}, {sigma:.2f})")
```

(continues on next page)

⁶ (*) `scipy.stats.probplot` uses a slightly different definition (there are many other ones in common use).

(continued from previous page)

```
plt.ylabel("Sample quantiles")
plt.show()
```




Figure 6.5. The Q-Q plot for the heights dataset. It is a nice fit.

Ideally, we wish that the points would be arranged on the $y = x$ line. In our case, there are small discrepancies⁷ in the tails (e.g., the smallest observation was slightly smaller than expected, and the largest one was larger than expected), although it is *common* a behaviour for small samples and certain distribution families. However, overall, we can say that we observe a fine fit.

6.2.3 Kolmogorov–Smirnov test (*)

To be more scientific, we can introduce a more formal method for assessing the quality of fit. It will enable us to test the null hypothesis stating that a given empirical distribution $\hat{F}_n$ does not differ *significantly* from the theoretical continuous CDF F :

$$\begin{cases} H_0 : \hat{F}_n = F & \text{(null hypothesis)} \\ H_1 : \hat{F}_n \neq F & \text{(two-sided alternative)} \end{cases}$$

The popular goodness-of-fit test by Kolmogorov and Smirnov can give us a conservative interval of the acceptable values of the largest deviation between the empirical and theoretical CDF, $\hat{D}_n$, as a function of n .

Namely, if the *test statistic* $\hat{D}_n$ is smaller than some *critical value* K_n , then we shall deem

⁷ (*) We can quantify (informally) the goodness of fit by using the Pearson linear correlation coefficient; see [Section 9.1.1](#).

the difference insignificant. This is to take into account the fact that reality might deviate from the ideal. Section 6.4.4 mentions that even for samples that truly come from a hypothesised distribution, there is some inherent variability. We need to be somewhat tolerant.

An authoritative textbook in mathematical statistics will tell us (and prove) that, under the assumption that $\hat{F}_n$ is the ECDF of a sample of n independent variables *really* generated from a continuous CDF F , the random variable $\hat{D}_n = \sup_{t \in \mathbb{R}} |\hat{F}_n(t) - F(t)|$ follows the Kolmogorov distribution with parameter n .

In other words, if we generate many samples of length n from F , and compute $\hat{D}_n$ s for each of them, we expect it to be distributed like in Figure 6.6, which we obtained by referring to `scipy.stats.kstwo`.




Figure 6.6. Densities (left) and cumulative distribution functions (right) of some Kolmogorov distributions. The greater the sample size, the smaller the acceptable deviations between the theoretical and empirical CDFs.

The choice of the critical value K_n involves a trade-off between our desire to:

- accept the null hypothesis when it is true (data *really* come from F), and
- reject it when it is false (data follow some other distribution, i.e., the difference is significant enough).

These two needs are, unfortunately, mutually exclusive.

In the framework of frequentist hypothesis testing, we assume some fixed upper bound (*significance level*) for making the former kind of mistake, which we call the *type-I error*. A nicely conservative (in a good way) value that we suggest employing is $\alpha = 0.001 = 0.1\%$, i.e., only 1 out of 1000 samples that really come from F will be rejected as not coming from F .

Such a K_n may be determined by considering the inverse of the CDF of the Kolmogorov distribution, Ξ_n . Namely, $K_n = \Xi_n^{-1}(1 - \alpha)$:

```
alpha = 0.001 # significance level
scipy.stats.kstwo.ppf(1-alpha, n)
## 0.029964456376393188
```

In our case $\hat{D}_n < K_n$ because $0.01047 < 0.02996$. We conclude that our empirical (heights) distribution does not differ significantly (at significance level 0.1%) from the assumed one, i.e., $N(160.1, 7.06)$. In other words, we do not have enough evidence against the statement that data are normally distributed. It is the presumption of innocence: they are normal enough.

We will return to this discussion in [Section 6.4.4](#) and [Section 12.2.6](#).

6.3 Other noteworthy distributions

6.3.1 Log-normal distribution

We say that a sample is *log-normally distributed*, if its logarithm is normally distributed. Such a behaviour is frequently observed in biology and medicine (size of living tissue), social sciences (number of sexual partners), or technology (file sizes). [Figure 6.7](#) suggests that it might also be true in the case of the UK taxpayers' incomes⁸.

```
income = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/uk_income_simulated_2020.txt")
plt.hist(np.log(income), bins=30, color="lightgray", edgecolor="black")
plt.ylabel("Count")
plt.show()
```

Let's thus proceed with the fitting of a log-normal model, $LN(\mu, \sigma)$. The procedure is similar to the normal case, but this time we determine the mean and standard deviation based on the logarithms of the observations:

```
lmu = np.mean(np.log(income))
lsigma = np.std(np.log(income), ddof=1)
lmu, lsigma
## (10.314409794364623, 0.5816585197803816)
```

⁸ Except for the few filthy rich, who are interesting on their own; see [Section 6.3.2](#) where we discuss the Pareto distribution.




Figure 6.7. A histogram of the logarithm of incomes.

Unintuitively, `scipy.stats.lognorm` identifies a distribution via the parameter s equal to σ and $scale$ equal to e^μ . Computing the PDF at different points must thus be done as follows:

```
x = np.linspace(np.min(income), np.max(income), 101)
fx = scipy.stats.lognorm.pdf(x, s=lsigma, scale=np.exp(lmu))
```

Figure 6.8 depicts the histograms on the log- and original scale together with the fitted probability density function. On the whole, the fit is not too bad; after all, we are only dealing with a sample of 1000 households. The original UK Office of National Statistics [data](#)⁹ could tell us more about the quality of this model in general, but it is beyond the scope of our simple exercise.

```
plt.subplot(1, 2, 1)
logbins = np.geomspace(np.min(income), np.max(income), 31)
plt.hist(income, bins=logbins, density=True,
        color="lightgray", edgecolor="black")
plt.plot(x, fx, "r--")
plt.xscale("log") # log-scale on the x-axis
plt.ylabel("Density")

plt.subplot(1, 2, 2)
plt.hist(income, bins=30, density=True, color="lightgray", edgecolor="black")
plt.plot(x, fx, "r--", label=f"PDF of LN({lmu:.1f}, {lsigma:.2f})")
plt.ylabel("Density")
```

(continues on next page)

⁹ <https://www.ons.gov.uk/peoplepopulationandcommunity/personalandhouseholdfinances/incomeandwealth/bulletins/householddisposableincomeandinequality/financialyear2020>

(continued from previous page)

```
plt.legend()
```

```
plt.show()
```




Figure 6.8. A histogram and the probability density function of the fitted log-normal distribution for the income dataset, on log- (left) and original (right) scale.

Next, the left side of [Figure 6.9](#) gives the quantile-quantile plot for the above log-normal model (note the double logarithmic scale). Additionally, on the right, we check the sensibility of the normality assumption (using a “normal” normal distribution, not its “log” version).

```
plt.subplot(1, 2, 1)
qq_plot( # see above for the definition
    income,
    lambda q: scipy.stats.lognorm.ppf(q, s=lsigma, scale=np.exp(lmu))
)
plt.xlabel(f"Quantiles of LN({lmu:.1f}, {lsigma:.2f})")
plt.ylabel("Sample quantiles")
plt.xscale("log")
plt.yscale("log")

plt.subplot(1, 2, 2)
mu = np.mean(income)
sigma = np.std(income, ddof=1)
qq_plot(income, lambda q: scipy.stats.norm.ppf(q, mu, sigma))
plt.xlabel(f"Quantiles of N({mu:.1f}, {sigma:.2f})")

plt.show()
```




Figure 6.9. The Q-Q plots for the income dataset vs the fitted log-normal (good fit; left) and normal (bad fit; right) distribution.

Exercise 6.4 Graphically compare the ECDF for *income* and the CDF of $\text{LN}(10.3, 0.58)$.

Exercise 6.5 (*) Perform the Kolmogorov–Smirnov goodness-of-fit test as in Section 6.2.3, to verify that the hypothesis of log-normality is not rejected at the $\alpha = 0.001$ significance level. At the same time, the income distribution significantly differs from a normal one.

The hypothesis that our data follow a normal distribution is most likely false. On the other hand, the log-normal model might be adequate. We can again reduce the whole dataset to merely two numbers, μ and σ , based on which (and probability theory), we may deduce that:

- the expected average (mean) income is $e^{\mu + \sigma^2/2}$,
- median is e^{μ} ,
- the most probable value (mode) in $e^{\mu - \sigma^2}$,

and so forth.

Note Recall again that for skewed a distribution such as the log-normal one, reporting the mean might be misleading. This is why *most* people are sceptical when they read the news about our prospering economy (“yeah, we’d like to see that kind of money in our pockets”). It is not only μ that matters, but also σ that quantifies the discrepancy between the rich and the poor.

For a normal distribution, the situation is vastly different. The mean, the median, and the most probable outcomes are the same: the distribution is symmetric around μ .

Exercise 6.6 What is the fraction of people with earnings below the mean in our $\text{LN}(10.3, 0.58)$ model? Hint: use `scipy.stats.lognorm.cdf` to get the answer.

6.3.2 Pareto distribution

Consider again the populations of the US cities:

```
cities = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/other/us_cities_2000.txt")
len(cities), sum(cities) # number of cities, total population
## (19447, 175062893.0)
```

Figure 6.10 gives the histogram of the city sizes on the log-scale. It looks like a log-normal distribution again, which the readers can fit themselves when they are feeling playful and have nothing better to do. (But, honestly, is there anything more delightful than doing stats?)

```
logbins = np.geomspace(np.min(cities), np.max(cities), 21)
plt.hist(cities, bins=logbins, color="lightgray", edgecolor="black")
plt.xscale("log")
plt.ylabel("Count")
plt.show()
```




Figure 6.10. A histogram of the unabridged `cities` dataset. Note the log-scale on the x-axis.

This time, however, we will be concerned with not what is *typical*, but what is in some sense *anomalous* or *extreme*. Just like in Section 4.3.7, let's look at the *truncated* version of the city size distribution by considering the cities with 10 000 or more inhabitants.

```
s = 10_000
large_cities = cities[cities >= s] # a right tail of the original dataset
len(large_cities), sum(large_cities) # number of cities, total population
## (2696, 146199374.0)
```

Plotting it on a double logarithmic scale can be performed by calling additionally `plt.yscale("log")`, which is left as an exercise. Doing so will lead to a picture similar to [Figure 6.11](#), which reveals something remarkable. The bar tops on the double log-scale are arranged more or less in a straight line. There are many datasets that exhibit this behaviour. We say that they follow a *power law* (power in the arithmetic sense, not the political one); see, e.g., [3].

The *Pareto distribution* (type I) family has a prototypical power law-like density. It is identified by two parameters:

- the (what `scipy` calls it) scale parameter $s > 0$ is equal to the shift from 0,
- the shape parameter, $\alpha > 0$, controls the slope of the said line on the double log-scale.

The probability density function of $P(\alpha, s)$ is given for $x \geq s$ by:

$$f(x) = \frac{\alpha s^\alpha}{x^{\alpha+1}},$$

and $f(x) = 0$ if $x < s$.

s is usually taken as the sample minimum (i.e., 10 000 in our case). α can be estimated through the reciprocal of the mean of the scaled logarithms of our observations:

```
alpha = 1/np.mean(np.log(large_cities/s))
alpha
## 0.9496171695997675
```

The left side of [Figure 6.11](#) compares the theoretical density and an empirical histogram on the double log-scale. The right part gives the corresponding Q-Q plot on a double logarithmic scale. We see that the populations of the largest cities are overestimated. The model could be better, but the cities are still growing, aren't they?

```
plt.subplot(1, 2, 1)
logbins = np.geomspace(s, np.max(large_cities), 21) # bin boundaries
plt.hist(large_cities, bins=logbins, density=True,
        color="lightgray", edgecolor="black")
plt.plot(logbins, scipy.stats.pareto.pdf(logbins, alpha, scale=s),
        "r--", label=f"PDF of P({alpha:.3f}, {s})")
plt.ylabel("Density")
plt.xscale("log")
plt.yscale("log")
plt.legend()
```

(continues on next page)

(continued from previous page)

```
plt.subplot(1, 2, 2)
qq_plot(large_cities, lambda q: scipy.stats.pareto.ppf(q, alpha, scale=s))
plt.xlabel(f"Quantiles of P({alpha:.3f}, {s})")
plt.ylabel("Sample quantiles")
plt.xscale("log")
plt.yscale("log")

plt.show()
```




Figure 6.11. A histogram (left) and a Q-Q plot (right) of the `large_cities` dataset vs the fitted density of a Pareto distribution on a double log-scale.

Example 6.7 (*) We might also be keen on verifying how accurately the probability of a randomly selected city's being at least of a given size can be predicted. Let's denote by $S(x) = 1 - F(x)$ the complementary cumulative distribution function (CCDF; sometimes referred to as the survival function), and by $\hat{S}_n(x) = 1 - \hat{F}_n(x)$ its empirical version. Figure 6.12 compares the empirical and the fitted CCDFs with probabilities on the linear- and log-scale.

```
x = np.geomspace(np.min(large_cities), np.max(large_cities), 1001)
probs = scipy.stats.pareto.cdf(x, alpha, scale=s)
n = len(large_cities)
for i in [1, 2]:
    plt.subplot(1, 2, i)
    plt.plot(x, 1-probs, "r--", label=f"CCDF of P({alpha:.3f}, {s})")
    plt.plot(np.sort(large_cities), 1-np.arange(1, n+1)/n,
             drawstyle="steps-post", label="Empirical CCDF")
    plt.xlabel("$x$")
    plt.xscale("log")
```

(continues on next page)

(continued from previous page)

```
plt.yscale(["linear", "log"][i-1])
if i == 1:
    plt.ylabel("Prob(city size > x)")
    plt.legend()
plt.show()
```




Figure 6.12. The empirical and theoretical complementary cumulative distribution functions for the `large_cities` dataset with probabilities on the linear- (left) and log-scale (right) and city sizes on the log-scale.

In terms of the maximal absolute distance between the two functions, $\hat{D}_n$, from the left plot we see that the fit seems acceptable. Still, let's stress that the log-scale overemphasises the relatively minor differences in the right tail and should not be used for judging the value of $\hat{D}_n$.

However, that the Kolmogorov–Smirnov goodness-of-fit test rejects the hypothesis of Paretianity (at a significance level 0.1%) is left as an exercise for the reader.

6.3.3 Uniform distribution

In the Polish *Lotto* lottery, six numbered balls $\{1, 2, \dots, 49\}$ are drawn without replacement from an urn. Here is a dataset that summarises the number of times each ball has been drawn in all the games in the period 1957–2016:

```
lotto = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/lotto_table.txt")
lotto
## array([720., 720., 714., 752., 719., 753., 701., 692., 716., 694., 716.,
##        668., 749., 713., 723., 693., 777., 747., 728., 734., 762., 729.,
```

(continues on next page)

(continued from previous page)

```
##      695., 761., 735., 719., 754., 741., 750., 701., 744., 729., 716.,
##      768., 715., 735., 725., 741., 697., 713., 711., 744., 652., 683.,
##      744., 714., 674., 654., 681.]
```

All events seem to occur more or less with the same probability. Of course, the numbers on the balls are integer, but in our idealised scenario, we may try modelling this dataset using a continuous *uniform distribution* $U(a, b)$, which yields arbitrary real numbers on a given interval (a, b) , i.e., between some a and b . Its probability density function is given for $x \in (a, b)$ by:

$$f(x) = \frac{1}{b-a},$$

and $f(x) = 0$ otherwise. Notice that `scipy.stats.uniform` uses parameters a and $scale$, the latter being equal to our $b - a$.

In the Lotto case, it makes sense to set $a = 1$ and $b = 50$ and interpret an outcome like 49.1253 as representing the 49th ball (compare the notion of the floor function, $\lfloor x \rfloor$).

```
x = np.linspace(1, 50, 1001)
plt.bar(np.arange(1, 50), width=1, height=lotto/np.sum(lotto),
        color="lightgray", edgecolor="black", alpha=0.8, align="edge")
plt.plot(x, scipy.stats.uniform.pdf(x, 1, scale=49), "r--",
        label="PDF of U(1, 50)")
plt.ylim(0, 0.025)
plt.legend()
plt.show()
```




Figure 6.13. A histogram of the lotto dataset.

Visually, see Figure 6.13, this model makes much sense, but again, some more rigorous statistical testing would be required to determine if someone has not been tampering with the lottery results, i.e., if data do not deviate from the uniform distribution significantly. Unfortunately, we cannot use the Kolmogorov–Smirnov test in the foregoing version as data are not continuous. See, however, Section 11.4.3 for the Pearson chi-squared test which is applicable here.

Exercise 6.8 *Does playing lotteries and engaging in gambling make rational sense at all, from the perspective of an individual player? Well, we see that 16 is the most frequently occurring outcome in Lotto, maybe there's some magic in it? Also, some people sometimes became millionaires, don't they?*

Note In data modelling (e.g., Bayesian statistics), sometimes a uniform distribution is chosen as a placeholder for “we know nothing about a phenomenon, so let's just assume that every event is equally likely”. Nonetheless, it is fascinating that in the end, the real world tends to be structured. Patterns that emerge are plentiful, and most often they are far from being uniformly distributed. Even more strikingly, they are subject to quantitative analysis.

6.3.4 Distribution mixtures (*)

Certain datasets may fail to fit through simple probabilistic models. It may sometimes be due to their non-random behaviour: statistics gives one of many means to create data idealisations, but in data science we can also employ partial differential equations, graphs and complex networks, agent-based modelling, cellular automata, amongst many others. They all might be worth giving a study (and then try).

It may also happen that what we observe is, in fact, a *mixture* of simpler processes. The dataset representing the December 2021 hourly averages pedestrian counts near the Southern Cross Station in Melbourne is a likely instance of such a scenario; compare Figure 4.5.

```
peds = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/southern_cross_station_peds_2019_dec.txt")
```

It might not be a bad idea to try to fit a probabilistic (convex) combination of three normal distributions f_1, f_2, f_3 , corresponding to the morning, lunchtime, and evening pedestrian count peaks. This yields the PDF:

$$f(x) = w_1 f_1(x) + w_2 f_2(x) + w_3 f_3(x),$$

for some coefficients $w_1, w_2, w_3 \geq 0$ such that $w_1 + w_2 + w_3 = 1$.

Figure 6.14 depicts a mixture of $N(8.4, 0.9)$, $N(14.2, 4)$, and $N(17.3, 0.9)$ with the corresponding weights of 0.27, 0.53, and 0.2. This dataset is coarse-grained: we only have

24 bar heights at our disposal. Consequently, the estimated¹⁰ coefficients should be taken with a pinch of Sichuan pepper.

```
plt.bar(np.arange(24)+0.5, width=1, height=peds/np.sum(peds),
        color="lightgray", edgecolor="black", alpha=0.8)
x = np.linspace(0, 24, 101)
p1 = scipy.stats.norm.pdf(x, 8.4, 0.9)
p2 = scipy.stats.norm.pdf(x, 14.2, 4.0)
p3 = scipy.stats.norm.pdf(x, 17.3, 0.9)
p = 0.27*p1 + 0.53*p2 + 0.20*p3 # a weighted combination of three densities
plt.plot(x, p, "r--", label="PDF of a normal mixture")
plt.xticks([0, 6, 12, 18, 24])
plt.legend()
plt.show()
```




Figure 6.14. A histogram of the peds dataset and an estimated mixture of three normal distributions.

Important More complex entities (models, methods) frequently arise as combinations of simpler (primitive) components. This is why we ought to spend a great deal of time studying the *fundamentals*.

Note Some data clustering techniques (in particular, the k -means algorithm that we briefly discuss later in this course) can be used to split a data sample into disjoint

¹⁰ The estimates were obtained by running `mixtools::normalmixEM` in R (expectation-maximisation for mixtures of univariate normals; [7]). Note that to turn peds, which is a table of counts, to a real-valued sample, we had to call `numpy.repeat(numpy.arange(24)+0.5, peds)`.

chunks corresponding to different mixture components. Also, it might be the case that the subpopulations are identified by another categorical variable that divides the dataset into natural groups; compare [Chapter 12](#).

6.4 Generating pseudorandom numbers

A probability distribution is useful not only for describing a dataset. It also enables us to perform many experiments on data that we do not currently have, but we might obtain in the future, to test various scenarios and hypotheses. Let's thus discuss some methods for generating random samples of independent (not related to each other) observations.

6.4.1 Uniform distribution

When most people say *random*, they implicitly mean *uniformly distributed*. For example:

```
np.random.rand(5)
## array([0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897])
```

gives five observations sampled independently from the uniform distribution on the unit interval, i.e., $U(0, 1)$. Here is the same with `scipy`, but this time the support is $(-10, 15)$.

```
scipy.stats.uniform.rvs(-10, scale=25, size=5) # from -10 to -10+25
## array([ 0.5776615 , 14.51910496,  7.12074346,  2.02329754, -0.19706205])
```

Alternatively, we could have shifted and scaled the output of the random number generator on the unit interval using the formula `numpy.random.rand(5)*25-10`.

6.4.2 Not exactly random

We generate numbers using a computer, which is a purely deterministic machine. Albeit they are indistinguishable from truly random when subject to rigorous tests for randomness, we refer to them as *pseudorandom* or random-like ones.

To prove that they are not random-random, let's set a specific initial state of the generator (the *seed*) and inspect what values are produced:

```
np.random.seed(123) # set the seed (the ID of the initial state)
np.random.rand(5)
## array([0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897])
```

Now, let's set the same seed and see how “random” the next values are:


```
np.random.seed(123) # set seed
np.random.rand(5)
## array([0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897])
```

Nobody expected that. Such a behaviour is very welcome, though. It enables us to perform completely *reproducible* numerical experiments, and truly scientific inquiries tend to nourish identical results under the same conditions.

Note If we do not set the seed manually, it will be initialised based on the current wall time, which is different every... time. As a result, the numbers will *seem* random to us, but only because we are slightly ignorant.

Many Python packages that we refer to in the sequel, including **pandas** and **scikit-learn**, rely on **numpy**'s random number generator. To harness them, we will have to become used to calling **numpy.random.seed**. Additionally, some of them (e.g., **sklearn.model_selection.train_test_split** or **pandas.DataFrame.sample**) will be equipped with the **random_state** argument, which behaves as if we *temporarily* changed the seed (for just one call to that function). For instance:

```
scipy.stats.uniform.rvs(size=5, random_state=123)
## array([0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897])
```

We obtained the same sequence again.

6.4.3 Sampling from other distributions

Generating data from other distributions is possible too; there are many **rvs** methods implemented in **scipy.stats**. For example, here is a sample from $N(160.1, 7.06)$:

```
scipy.stats.norm.rvs(160.1, 7.06, size=3, random_state=50489)
## array([166.01775384, 136.7107872 , 185.30879579])
```

Pseudorandom deviates from the *standard* normal distribution, i.e., $N(0, 1)$, can also be generated using **numpy.random.randn**. As $N(160.1, 7.06)$ is a scaled and shifted version thereof, the preceding is equivalent to:

```
np.random.seed(50489)
np.random.randn(3)*7.06 + 160.1
## array([166.01775384, 136.7107872 , 185.30879579])
```

Important Conclusions based on simulated data are trustworthy for they cannot be manipulated. Or can they?

The above pseudorandom number generator's seed, 50489, is a bit suspicious. It may suggest that someone wanted to *prove* some point (in this case, the violation of the 3σ rule). This is why we recommend sticking to only one seed, e.g., 123, or – when

performing simulations – setting the consecutive natural seeds in each iteration of the for loop: 1, 2,

Exercise 6.9 Generate 1000 pseudorandom numbers from the log-normal distribution and draw its histogram.

Note (*) Having a reliable pseudorandom number generator from the uniform distribution on the unit interval is crucial as sampling from other distributions usually involves transforming the independent $U(0, 1)$ variates. For instance, realisations of random variables following any continuous cumulative distribution function F can be constructed through the *inverse transform sampling*:

1. Generate a sample $x_1, \dots, x_n$ independently from $U(0, 1)$.
2. Transform each x_i by applying the quantile function, $y_i = F^{-1}(x_i)$.

Now $y_1, \dots, y_n$ follows the CDF F .

Exercise 6.10 (*) Generate 1000 pseudorandom numbers from the log-normal distribution using inverse transform sampling.

Exercise 6.11 (**) Generate 1000 pseudorandom numbers from the distribution mixture discussed in Section 6.3.4.

6.4.4 Natural variability

Even a sample truly generated from a specific distribution will deviate from it, sometimes considerably. Such effects will be especially visible for small sample sizes, but they usually dissolve¹¹ when the availability of data increases.

For example, Figure 6.15 depicts the histograms of nine different samples of size 100, all drawn independently from the standard normal distribution.

```
plt.figure(figsize=plt.rcParams["figure.figsize"]*[0], )*2) # width=height
for i in range(9):
    plt.subplot(3, 3, i+1, xlim=(-4, 4), ylim=(0, 0.6), ylabel=None)
    sample = scipy.stats.norm.rvs(size=100, random_state=i+1)
    plt.hist(sample, density=True, bins=10,
             color="lightgray", edgecolor="black")
plt.legend()
plt.show()
```

¹¹ Compare the Fundamental Theorem of Statistics (the Glivenko–Cantelli theorem).




Figure 6.15. All nine samples are normally distributed.

There is a certain ruggedness in the bar sizes that a naïve observer would try to interpret as something meaningful. Competent data scientists train their eyes to ignore such impurities. In this case, they are only due to random effects. Nevertheless, we must always be ready to detect *deviations* from the assumed model that are worth attention.

Exercise 6.12 Repeat the above experiment for samples of sizes 10, 1 000, and 10 000.

Example 6.13 (*) Using a simple Monte Carlo simulation, we can verify (approximately) that the Kolmogorov–Smirnov goodness-of-fit test introduced in Section 6.2.3 has been calibrated properly, i.e., that for samples that really follow the assumed distribution, the null hypothesis is indeed rejected in roughly 0.1% of the cases.

Assume we are interested in the null hypothesis referencing the standard normal distribution, $N(160.1, 7.06)$, and sample size $n = 4221$. We need to generate many (we assume 10 000

below) such samples. For each of them, we compute and store the maximal absolute deviation from the theoretical CDF, i.e., $\hat{D}_n$.

```
n = 4221
distrib = scipy.stats.norm(160.1, 7.06) # the assumed distribution
Dns = []
for i in range(10000): # increase this for better precision
    x = distrib.rvs(size=n, random_state=i+1) # really follows the distrib.
    Dns.append(compute_Dn(x, distrib.cdf))
Dns = np.array(Dns)
```

Now let's compute the proportion of cases which lead to $\hat{D}_n$ greater than the critical value K_n :

```
len(Dns[Dns >= scipy.stats.kstwo.ppf(1-0.001, n)]) / len(Dns)
## 0.0008
```

Its expected value is 0.001. But our approximation is necessarily imprecise because we rely on randomness. Increasing the number of trials from 10 000 to, say, 1 000 000 should make the above estimate closer to the theoretical expectation.

It is also worth checking that the density histogram of Dns resembles the Kolmogorov distribution that we can compute via `scipy.stats.kstwo.pdf`.

Exercise 6.14 (*) It might also be interesting to verify the test's power, i.e., the probability that when the null hypothesis is false, it will actually be rejected. Modify the above code in such a way that x in the `for` loop is not generated from $N(160.1, 7.06)$, but $N(140, 7.06)$, $N(141, 7.06)$, etc., and check the proportion of cases where we deem the sample distribution significantly different from $N(160.1, 7.06)$. Small deviations from the true location parameter μ are usually ignored, but this improves with sample size n .

6.4.5 Adding jitter (white noise)

We mentioned that measurements might be subject to observational error. Rounding can also occur as early as in the data collection phase. In particular, our heights dataset is precise up to 1 fractional digit. However, in statistics, when we say that data follow a continuous distribution, the probability of having two identical values in a sample is 0. Therefore, some data analysis methods might assume no ties in the input vector, i.e., that all values are unique.

The easiest way to deal with such numerical inconveniences is to add some white noise with the expected value of 0, either uniformly or normally distributed.

For example, for heights, it makes sense to add some jitter from $U[-0.05, 0.05]$:

```
heights_jitter = heights + (np.random.rand(len(heights))*0.1-0.05)
heights_jitter[:6] # preview
## array([160.21704623, 152.68870195, 161.24482407, 157.3675293 ,
##        154.61663465, 144.68964596])
```

Adding noise also might be performed for aesthetic reasons, e.g., when drawing scatter plots.

6.4.6 Independence assumption

Let's generate nine binary digits in a pseudorandom fashion:

```
np.random.choice([0, 1], 9)
## array([1, 1, 1, 1, 1, 1, 1, 1, 1])
```

We can consider ourselves very lucky; all numbers are the same. So, the next number *must* finally be a “zero”, right?

```
np.random.choice([0, 1], 1)
## array([1])
```

Wrong. The numbers we generate are *independent* of each other. There is no history. In the current model of randomness (Bernoulli trials; two possible outcomes with the same probability), there is a 50% chance of obtaining a “one” *regardless* of how many “ones” were observed previously.

We should not seek patterns where no regularities exist. Our brain forms expectations about the world, and overcoming them is hard work. This must be done as the reality could not care less about what we consider it to be.

6.5 Further reading

For an excellent general introductory course on probability and statistics, see [40, 41] and also [82, 83]. More advanced students are likely to enjoy other classics such as [5, 9, 19, 28, 46]. To go beyond the basics, check out [26]. Topics in random number generation are covered in [39, 59, 81].

For a more detailed introduction to exploratory data analysis, see the classical books by Tukey [92, 93], Tufte [91], and Wainer [98].

We took the logarithm of the log-normally distributed incomes and obtained a normally distributed sample. In statistical practice, it is not rare to apply different non-linear transforms of the input vectors at the data preprocessing stage (see, e.g., [Section 9.2.6](#)). In particular, the Box–Cox (power) transform [12] is of the form $x \mapsto (x^\lambda - 1)/\lambda$ for some λ . Interestingly, in the limit as $\lambda \rightarrow 0$, this formula yields $x \mapsto \log x$ which is exactly what we were applying in this chapter.

Newman et al. [16, 71] give a nice overview of the power-law-like behaviour of some “rich” or otherwise extreme datasets. It is worth noting that the logarithm of a Paretian sample divided by the minimum follows an exponential distribution (which we

discuss in [Chapter 16](#)). For a comprehensive catalogue of statistical distributions, their properties, and relationships between them, see [\[29\]](#).

6.6 Exercises

Exercise 6.15 *Why is the notion of the mean income confusing to the general public?*

Exercise 6.16 *When manually setting the seed of a pseudorandom number generator makes sense?*

Exercise 6.17 *Given a log-normally distributed sample x , how can we turn it to a normally distributed one, i.e., $y=f(x)$, with f being... what?*

Exercise 6.18 *What is the 3σ rule for normally distributed data?*

Exercise 6.19 *Can the 3σ rule be applied for log-normally distributed data?*

Exercise 6.20 (*) *How can we verify graphically if a sample follows a hypothesised theoretical distribution?*

Exercise 6.21 (*) *Explain the meaning of the type I error, significance level, and a test's power.*

Part III

Multidimensional data

From uni- to multidimensional numeric data

From the perspective of structured datasets, a vector often represents n independent measurements of the same quantitative property, e.g., heights of n different patients, incomes in n randomly chosen households, or finishing times of n marathoners. More generally, these are all instances of a bag of n points on the real line. By far¹, we should have become fairly fluent with the methods for processing such one-dimensional arrays.

Let's increase the level of complexity by describing the n entities by m features, for any $m \geq 1$. In other words, we will be dealing with n points in an m -dimensional space, $\mathbb{R}^m$.

We can arrange all the observations in a table with n rows and m columns (just like in spreadsheets). We can represent it with **numpy** as a two-dimensional array which we will refer to as a *matrix*. Thanks to matrices, we can keep the n tuples of length m together in a single object (or m tuples of length n , depending on how we want to look at them) and process them all at once. How convenient.

Important Just like vectors, matrices were designed to store data of the same type. [Chapter 10](#) will cover **pandas** *data frames*, which support mixed data types, e.g., numerical and categorical. Moreover, they let their rows and columns be named. **pandas** is built on top of **numpy**, and implements many recipes for the most *popular* data wrangling tasks. We, however, would like to be able to tackle *any* computational problem. It is worth knowing that many data analysis and machine learning algorithms automatically convert numerical parts of data frames to matrices so that **numpy** can do most of the mathematical heavy lifting.

7.1 Creating matrices

7.1.1 Reading CSV files

Tabular data are often stored and distributed in a very portable plain-text format called CSV (comma-separated values) or one of its variants. We can read them easily with `numpy.genfromtxt` (or later with `pandas.read_csv`).

¹ Assuming we solved all the suggested exercises, which we did (see Rule #3), didn't we?

```
body = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_adult_female_bmx_2020.csv",
    delimiter=",") [1:, :] # skip the first row (column names)
```

The file specifies column names in the first non-comment line (we suggest inspecting it in a web browser). Therefore, we had to omit it manually (more on matrix indexing later). Here is a preview of the first few rows:

```
body[:6, :] # the first six rows, all columns
## array([[ 97.1, 160.2,  34.7,  40.8,  35.8, 126.1, 117.9],
##        [ 91.1, 152.7,  33.5,  33. ,  38.5, 125.5, 103.1],
##        [ 73. , 161.2,  37.4,  38. ,  31.8, 106.2,  92. ],
##        [ 61.7, 157.4,  38. ,  34.7,  29. , 101. ,  90.5],
##        [ 55.4, 154.6,  34.6,  34. ,  28.3,  92.5,  73.2],
##        [ 62. , 144.7,  32.5,  34.2,  29.8, 106.7,  84.8]])
```

It is an excerpt from the National Health and Nutrition Examination Survey (NHANES²), where the consecutive columns give seven body measurements of adult females:

```
body_columns = np.array([
    "weight (kg)",
    "standing height (cm)", # we know 'heights' from the previous chapters
    "upper arm len. (cm)",
    "upper leg len. (cm)",
    "arm circ. (cm)",
    "hip circ. (cm)",
    "waist circ. (cm)",
])
```

We noted the column names down as **numpy** matrices give no means for storing column labels. It is only a minor inconvenience.

body is a **numpy** array:

```
type(body) # class of this object
## <class 'numpy.ndarray'>
```

but this time it is a two-dimensional one:

```
body.ndim # number of dimensions
## 2
```

which means that its shape slot is now a tuple of length two:

```
body.shape
## (4221, 7)
```

² <https://wwwn.cdc.gov/nchs/nhanes/search/datapage.aspx>

We obtained the total number of rows and columns, respectively.

7.1.2 Enumerating elements

`numpy.array` can take a sequence of vector-like objects of the same lengths specifying consecutive rows of a matrix. For example:

```
np.array([ # list of lists
    [ 1, 2, 3, 4 ], # the first row
    [ 5, 6, 7, 8 ], # the second row
    [ 9, 10, 11, 12 ] # the third row
])
## array([[ 1,  2,  3,  4],
##        [ 5,  6,  7,  8],
##        [ 9, 10, 11, 12]])
```

gives a 3×4 (3-by-4) matrix. Next:

```
np.array([ [1], [2], [3] ])
## array([[1],
##        [2],
##        [3]])
```

yields a 3×1 array. Such two-dimensional arrays with one column will be referred to as *column vectors* (they are matrices still). Moreover:

```
np.array([ [1, 2, 3, 4] ])
## array([[1, 2, 3, 4]])
```

produces a 1×4 array (a *row vector*).

Note An ordinary vector (a unidimensional array) only displays a single pair of square brackets:

```
np.array([1, 2, 3, 4])
## array([1, 2, 3, 4])
```

7.1.3 Repeating arrays

The previously-mentioned `numpy.tile` and `numpy.repeat` can also generate some nice matrices. For instance:

```
np.repeat([[1, 2, 3, 4]], 3, axis=0) # over the first axis
## array([[1, 2, 3, 4],
##        [1, 2, 3, 4],
##        [1, 2, 3, 4]])
```

repeats a row vector rowwisely, i.e., over the first axis (0). Replicating a column vector columnwisely is possible as well:

```
np.repeat([[1], [2], [3]], 4, axis=1) # over the second axis
## array([[1, 1, 1, 1],
##        [2, 2, 2, 2],
##        [3, 3, 3, 3]])
```

Exercise 7.1 *Generate matrices of the following kinds:*

$$\begin{bmatrix} 1 & 2 \\ 1 & 2 \\ 1 & 2 \\ 3 & 4 \\ 3 & 4 \\ 3 & 4 \\ 3 & 4 \end{bmatrix}, \quad \begin{bmatrix} 1 & 2 & 1 & 2 & 1 & 2 \\ 1 & 2 & 1 & 2 & 1 & 2 \end{bmatrix}, \quad \begin{bmatrix} 1 & 1 & 2 & 2 & 2 \\ 3 & 3 & 4 & 4 & 4 \end{bmatrix}.$$

7.1.4 Stacking arrays

`numpy.column_stack` and `numpy.vstack` take a tuple of array-like objects and bind them column- or rowwisely to form a new matrix:

```
np.column_stack(([10, 20], [30, 40], [50, 60])) # a tuple of lists
## array([[10, 30, 50],
##        [20, 40, 60]])
np.vstack(([10, 20], [30, 40], [50, 60]))
## array([[10, 20],
##        [30, 40],
##        [50, 60]])
np.column_stack((
    np.vstack(([10, 20], [30, 40], [50, 60])),
    [70, 80, 90]
))
## array([[10, 20, 70],
##        [30, 40, 80],
##        [50, 60, 90]])
```

Note the double round brackets: we called these functions on tuples.

Exercise 7.2 *Perform similar operations using `numpy.append`, `numpy.hstack`, `numpy.stack`, `numpy.concatenate`, and (*) `numpy.c_`. Are they worth taking note of, or are they redundant?*

Exercise 7.3 *Using `numpy.insert`, add a new row/column at the beginning, end, and in the middle of an array. Let's stress that this function returns a new array.*

7.1.5 `numpy.r_` revisited (*)

In Section 4.1.4, we introduced the `numpy.r_` object that simplifies vector creation. It turns out that its first argument can be a string that controls the way that the given items are merged:

```
np.r_['r', 1:5] # row vector
## matrix([[1, 2, 3, 4]])
np.r_['c', 1:5] # column vector
## matrix([[1],
##         [2],
##         [3],
##         [4]])
np.r_['0,2', [1, 2], [3, 4], [5, 6]] # concatenate along axis 0, ndim=2
## array([[1, 2],
##        [3, 4],
##        [5, 6]])
np.r_['1,2,0', [1, 2], [3, 4], [5, 6]] # along axis 1, make them column vecs
## array([[1, 3, 5],
##        [2, 4, 6]])
```

Furthermore, the last expression can be equivalently rewritten using `numpy.c_`:

```
np.c_[ [1, 2], [3, 4], [5, 6] ] # column stack
## array([[1, 3, 5],
##        [2, 4, 6]])
```

7.1.6 Other functions

Many built-in functions can generate arrays of arbitrary shapes (not only vectors). For example:

```
np.random.seed(123)
np.random.rand(2, 5) # not: rand((2, 5))
## array([[0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897],
##        [0.42310646, 0.9807642 , 0.68482974, 0.4809319 , 0.39211752]])
```

The same with `scipy`:

```
scipy.stats.uniform.rvs(0, 1, size=(2, 5), random_state=123)
## array([[0.69646919, 0.28613933, 0.22685145, 0.55131477, 0.71946897],
##        [0.42310646, 0.9807642 , 0.68482974, 0.4809319 , 0.39211752]])
```

The way we specify the output shapes might differ across functions and packages. Consequently, as usual, it is always best to refer to their documentation.

Exercise 7.4 Check out the documentation of the following functions: `numpy.eye`, `numpy.diag`, `numpy.zeros`, `numpy.ones`, and `numpy.empty`.

7.2 Reshaping matrices

Let's consider an example 3×4 matrix:

```
A = np.array([
    [ 1,  2,  3,  4 ],
    [ 5,  6,  7,  8 ],
    [ 9, 10, 11, 12 ]
])
```

Internally, a matrix is represented using a *long* flat vector where elements are stored in the row-major³ order:

```
A.size # the total number of elements
## 12
A.ravel() # the underlying flat array
## array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12])
```

It is the shape slot that is causing the 12 elements to be treated as if they were arranged on a 3×4 grid, for example in different algebraic operations and during the printing of the matrix. This virtual arrangement can be altered anytime without modifying the underlying array:

```
A.shape = (4, 3)
A
## array([[ 1,  2,  3],
##        [ 4,  5,  6],
##        [ 7,  8,  9],
##        [10, 11, 12]])
```

This way, we obtained a different *view* of the same data.

For convenience, the **reshape** method returns a modified version of the object it is applied on:

```
A.reshape(-1, 6) # A.reshape(don't make me compute this for you mate!, 6)
## array([[ 1,  2,  3,  4,  5,  6],
##        [ 7,  8,  9, 10, 11, 12]])
```

Here, the *placeholder* “-1” means that **numpy** must deduce by itself how many rows we want in the result. Twelve elements are supposed to be arranged in six columns, so the maths behind it is not rocket science. Thanks to this, generating row or column vectors is straightforward:

```
np.linspace(0, 1, 5).reshape(1, -1) # one row, guess the number of columns
                                     (continues on next page)
```

³ (*) Sometimes referred to as a C-style array, as opposed to the Fortran-style which is used in, e.g., R.

(continued from previous page)

```
## array([[0. , 0.25, 0.5 , 0.75, 1.  ]])
np.array([9099, 2537, 1832]).reshape(-1, 1) # one column, guess row count
## array([[9099],
##        [2537],
##        [1832]])
```

Note (*) Higher-dimensional arrays are also available. For example:

```
np.arange(24).reshape(2, 4, 3)
## array([[ 0,  1,  2],
##        [ 3,  4,  5],
##        [ 6,  7,  8],
##        [ 9, 10, 11]],
##       [[12, 13, 14],
##        [15, 16, 17],
##        [18, 19, 20],
##        [21, 22, 23]])
```

Is an array of “depth” 2, “height” 4, and “width” 3; we can see it as two 4×3 matrices stacked together.

Multidimensional arrays can be used for representing contingency tables for products of many factors, but we usually prefer working with *long* data frames instead (Section 10.6.2) due to their more aesthetic display and better handling of sparse data.

7.3 Mathematical notation

Mathematically, a matrix with n rows and m columns (an $n \times m$ matrix) $\mathbf{X}$ can be written as:

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n,1} & x_{n,2} & \cdots & x_{n,m} \end{bmatrix}.$$

We denote it by $\mathbf{X} \in \mathbb{R}^{n \times m}$. Spreadsheets display data in a similar fashion. We see that $x_{i,j} \in \mathbb{R}$ is the element in the i -th row (e.g., the i -th *observation* or *case*) and the j -th column (e.g., the j -th *feature* or *variable*).

Important Matrices can encode many different kinds of data:

- n points in an m -dimensional space, like n observations for which there are m

measurements/features recorded, where each row describes a different object; this is the most common scenario (in particular, if $\mathbf{X}$ represents the body dataset, then $x_{5,2}$ is the height of the fifth person);

- m time series sampled at n points in time (e.g., prices of m different currencies on n consecutive days; see [Chapter 16](#));
- a single kind of measurement for data in m groups, each consisting of n subjects (e.g., heights of n males and n females); here, the order of elements in each column does not usually matter as observations are not *paired*; there is no relationship between $x_{i,j}$ and $x_{i,k}$ for $j \neq k$; a matrix is used merely as a convenient container for storing a few unrelated vectors of identical sizes; we will be dealing with a more generic case of possibly nonhomogeneous groups in [Chapter 12](#);
- two-way contingency tables (see [Section 11.2.2](#)), where an element $x_{i,j}$ gives the number of occurrences of items at the i -th level of the first categorical variable and, at the same time, being at the j -th level of the second variable (e.g., blue-eyed *and* blonde-haired);
- graphs and other relationships between objects, e.g., $x_{i,j} = 0$ might mean that the i -th object is not connected⁴ with the j -th one, and $x_{k,l} = 1$ that there is a connection between k and l (e.g., who is a friend of whom, whether a user recommends a particular item);
- images, where $x_{i,j}$ represents the intensity of a colour component (e.g., red, green, blue or shades of grey or hue, saturation, brightness; compare [Section 16.4](#)) of a pixel in the $(n - i + 1)$ -th row and the j -th column.

In practice, more complex and less-structured data can often be mapped to a tabular form. For instance, a set of audio recordings can be described by measuring the overall loudness, timbre, and danceability of each song. Also, a collection of documents can be described by means of the degrees of belongingness to some automatically discovered topics (e.g., someone may claim that the *Lord of the Rings* is 80% travel literature, 70% comedy, and 50% heroic fantasy, but let's not take it for granted).

7.3.1 Transpose

The *transpose* of a matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$ is an $(m \times n)$ -matrix $\mathbf{Y}$ given by:

$$\mathbf{Y} = \mathbf{X}^T = \begin{bmatrix} x_{1,1} & x_{2,1} & \cdots & x_{m,1} \\ x_{1,2} & x_{2,2} & \cdots & x_{m,2} \\ \vdots & \vdots & \ddots & \vdots \\ x_{1,n} & x_{2,n} & \cdots & x_{m,n} \end{bmatrix},$$

i.e., it enjoys $y_{i,j} = x_{j,i}$. For example:

⁴ (*) Such matrices are usually sparse, i.e., have many elements equal to 0. We have special, memory-efficient data structures for handling such data; see `scipy.sparse`.


```

A # before
## array([[ 1,  2,  3],
##        [ 4,  5,  6],
##        [ 7,  8,  9],
##        [10, 11, 12]])
A.T # the transpose of A
## array([[ 1,  4,  7, 10],
##        [ 2,  5,  8, 11],
##        [ 3,  6,  9, 12]])

```

Rows became columns and vice versa. It is not the same as the aforementioned reshaping, which does not change the order of elements in the underlying array:

7.3.2 Row and column vectors

Additionally, we will sometimes use the following notation to emphasise that $\mathbf{X}$ consists of n rows:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_{1,\cdot} \\ \mathbf{x}_{2,\cdot} \\ \vdots \\ \mathbf{x}_{n,\cdot} \end{bmatrix}.$$

Here, $\mathbf{x}_{i,\cdot}$ is a *row vector* of length m , i.e., a $(1 \times m)$ -matrix:

$$\mathbf{x}_{i,\cdot} = [x_{i,1} \quad x_{i,2} \quad \cdots \quad x_{i,m}].$$

Alternatively, we can specify the m columns:

$$\mathbf{X} = [\mathbf{x}_{\cdot,1} \quad \mathbf{x}_{\cdot,2} \quad \cdots \quad \mathbf{x}_{\cdot,m}],$$

where $\mathbf{x}_{\cdot,j}$ is a *column vector* of length n , i.e., an $(n \times 1)$ -matrix:

$$\mathbf{x}_{\cdot,j} = [x_{1,j} \quad x_{2,j} \quad \cdots \quad x_{n,j}]^T = \begin{bmatrix} x_{1,j} \\ x_{2,j} \\ \vdots \\ x_{n,j} \end{bmatrix},$$

Thanks to the use of the matrix transpose, $\cdot^T$, we can save some vertical space (we want this enjoyable to be as long as possible, but maybe not this way).

Also, recall that we are used to denoting *vectors* of length m by $\mathbf{x} = (x_1, \dots, x_m)$. A vector is a one-dimensional array (not a two-dimensional one), hence the slightly different bold font which is crucial where any ambiguity could be troublesome.

Note To avoid notation clutter, we will often be implicitly promoting vectors like $\mathbf{x} =$

$(x_1, \dots, x_m)$ to row vectors $\mathbf{x} = [x_1 \dots x_m]$. This is the behaviour that `numpy`⁵ uses; see Chapter 8.

7.3.3 Identity and other diagonal matrices

`I` denotes the *identity matrix*, being a square $n \times n$ matrix (with n most often clear from the context) with 0s everywhere except on the main diagonal which is occupied by 1s.

```
np.eye(5) # I
## array([[1., 0., 0., 0., 0.],
##        [0., 1., 0., 0., 0.],
##        [0., 0., 1., 0., 0.],
##        [0., 0., 0., 1., 0.],
##        [0., 0., 0., 0., 1.]])
```

The identity matrix is a neutral element of the matrix multiplication (Section 8.3).

More generally, any diagonal matrix, $\text{diag}(a_1, \dots, a_n)$, can be constructed from a given sequence of elements by calling:

```
np.diag([1, 2, 3, 4])
## array([[1, 0, 0, 0],
##        [0, 2, 0, 0],
##        [0, 0, 3, 0],
##        [0, 0, 0, 4]])
```

7.4 Visualising multidimensional data

Let's go back to our body dataset:

```
body[:6, :] # preview
## array([[ 97.1, 160.2,  34.7,  40.8,  35.8, 126.1, 117.9],
##        [ 91.1, 152.7,  33.5,  33. ,  38.5, 125.5, 103.1],
##        [ 73. , 161.2,  37.4,  38. ,  31.8, 106.2,  92. ],
##        [ 61.7, 157.4,  38. ,  34.7,  29. , 101. ,  90.5],
##        [ 55.4, 154.6,  34.6,  34. ,  28.3,  92.5,  73.2],
##        [ 62. , 144.7,  32.5,  34.2,  29.8, 106.7,  84.8]])
body.shape
## (4221, 7)
```

It is an example of tabular (“structured”) data whose important property is that the elements in any chosen row describe the same person. We can freely reorder all the columns at the same time (change the order of participants), and this dataset will

⁵ Some textbooks assume that all vectors are *column* vectors, though.

make the same sense. However, sorting a single column and leaving others unchanged will be semantically invalid.

Mathematically, we can consider the above as a set of 4221 points in a seven-dimensional space, $\mathbb{R}^7$. Let's discuss how we can visualise its different natural *projections*.

7.4.1 2D Data

A *scatter plot* visualises one variable against another one.

```
plt.plot(body[:, 1], body[:, 3], "o", c="#00000022")
plt.xlabel(body_columns[1])
plt.ylabel(body_columns[3])
plt.show()
```




Figure 7.1. An example scatter plot.

Figure 7.1 depicts upper leg length (the y-axis) vs (versus; against; as a function of) standing height (the x-axis) in the form of a point cloud with (x, y) coordinates like $(\text{body}[i, 1], \text{body}[i, 3])$ for all $i = 1, \dots, 4221$.

Example 7.5 Here are the exact coordinates of the point corresponding to the person of the smallest height:

```
body[np.argmin(body[:, 1]), [1, 3]]
## array([131.1, 30.8])
```

Locate it in Figure 7.1. Also, pinpoint the one with the greatest upper leg length:

```
body[np.argmax(body[:, 3]), [1, 3]]
## array([168.9, 49.1])
```

As the points are abundant, normally we cannot easily see *where* most of them are located. As a simple remedy, we plotted the points using a semi-transparent colour. This gave a kind of the points' density estimate. The colour specifier was of the form `#rrggbaa`, giving the intensity of the red, green, blue, and alpha (opaqueness) channel in four series of two hexadecimal digits (between `00` = 0 and `ff` = 255).

Overall, the plot reveals that there is a *general tendency* for small heights and small upper leg lengths to occur frequently together. The taller the person, the longer her legs on average, and vice verse.

But there is some natural variability: for example, looking at people of height roughly equal to 160 cm, their upper leg length can be anywhere between 25 and 50 cm (range), yet we expect the majority to lie somewhere between 35 and 40 cm. [Chapter 9](#) will explore two measures of correlation that will enable us to quantify the degree (strength) of association between variable pairs.

7.4.2 3D data and beyond

With more variables to visualise, we might be tempted to use a three-dimensional scatter plot like the one in [Figure 7.2](#).

```
fig = plt.figure()
ax = fig.add_subplot(projection="3d", facecolor="#ffffff00")
ax.scatter(body[:, 1], body[:, 3], body[:, 0], color="#00000011")
ax.view_init(elev=30, azim=60, vertical_axis="y")
ax.set_xlabel(body_columns[1])
ax.set_ylabel(body_columns[3])
ax.set_zlabel(body_columns[0])
plt.show()
```

Infrequently will such a 3D plot provide us with readable results, though. We are projecting a three-dimensional reality onto a two-dimensional screen or a flat page. Some information must inherently be lost. What we see is relative to the position of the virtual camera and some angles can be more meaningful than others.

Exercise 7.6 (*) Try finding an interesting elevation and azimuth angle by playing with the arguments passed to the `mpl_toolkits.mplot3d.axes3d.Axes3D.view_init` function. Also, depict arm circumference, hip circumference, and weight on a 3D plot.

Note (*) We may have facilities for creating an *interactive* scatter plot (running the above from the Python's console enables this), where the virtual camera can be freely repositioned with a mouse/touch pad. This can give some more insight into our data. Also, there are means of creating animated sequences, where we can fly over the data scene. Some people find it cool, others find it annoying, but the biggest problem there-with is that they cannot be included in printed material. If we are only targeting the




Figure 7.2. A three-dimensional scatter plot reveals almost nothing.

display for the Web (this includes mobile devices), we can try [some Python libraries](#)⁶ that output HTML+CSS+JavaScript code which instructs the browser engine to create some more sophisticated interactive graphics, e.g., **bokeh** or **plotly**.

Example 7.7 *Instead of drawing a 3D plot, it might be better to play with a 2D scatter plot that uses different marker colours (or sometimes sizes: think of them as bubbles). Suitable [colour maps](#)⁷ can distinguish between low and high values of a third variable.*

```
plt.scatter(
    body[:, 4],    # x
    body[:, 5],    # y
    c=body[:, 0],  # "z" - colours
    cmap=plt.colormaps.get_cmap("copper"), # colour map
    alpha=0.5     # opaqueness level between 0 and 1
)
plt.xlabel(body_columns[4])
plt.ylabel(body_columns[5])
plt.axis("equal")
plt.rcParams["axes.grid"] = False
cbar = plt.colorbar()
plt.rcParams["axes.grid"] = True
cbar.set_label(body_columns[0])
plt.show()
```

⁶ <https://wiki.python.org/moin/NumericAndScientific/Plotting>

⁷ <https://matplotlib.org/stable/tutorials/colors/colormaps.html>




Figure 7.3. A two-dimensional scatter plot displaying three variables.

In Figure 7.3, we see some tendency for the weight to be greater as both the arm and the hip circumferences increase.

Exercise 7.8 Play around with different colour palettes. However, be wary that every 1 in 12 men (8%) and 1 in 200 women (0.5%) have colour vision deficiencies, especially in the red-green or blue-yellow spectrum. For this reason, some diverging colour maps might be worse than others.

A piece of paper is two-dimensional: it only has height and width. By looking around, we also perceive the notion of depth. So far so good. But with more-dimensional data, well, suffice it to say that we are three-dimensional creatures and any attempts towards visualising them will simply not work, don't even trip.

Luckily, it is where mathematics comes to our rescue. With some more knowledge and intuitions, and this book helps us develop them, it will be easy⁸ to consider a generic m -dimensional space, and then assume that, say, $m = 7$ or 42. This is exactly why data science relies on automated methods for knowledge/pattern discovery. Thanks to them, we can identify, describe, and analyse the structures that might be present in the data, but cannot be *experienced* with our imperfect senses.

Note Linear and nonlinear dimensionality reduction techniques can be applied to visualise some aspects of high-dimensional data in the form of 2D (or 3D) plots. In particular, the principal component analysis (PCA; Section 9.3) finds a potentially *note-worthy* angle from which we can try to look at the data.

⁸ This is an old funny joke that most funny mathematicians find funny. Ha.

7.4.3 Scatter plot matrix (pairs plot)

We can also try depicting all pairs of selected variables in the form of a scatter plot matrix.

```
def pairplot(X, labels, bins=21, alpha=0.1):
    """
    Draws a scatter plot matrix, given:
    * X - data matrix,
    * labels - list of column names
    """
    assert X.shape[1] == len(labels)

    k = X.shape[1]
    fig, axes = plt.subplots(nrows=k, ncols=k, sharex="col", sharey="row",
                             figsize=plt.rcParams["figure.figsize"][0], )*2)
    for i in range(k):
        for j in range(k):
            ax = axes[i, j]
            if i == j: # diagonal
                ax.text(0.5, 0.5, labels[i], transform=ax.transAxes,
                       ha="center", va="center", size="x-small")
            else:
                ax.plot(X[:, j], X[:, i], ".", color="black", alpha=alpha)
```

And now:

```
which = [1, 0, 4, 5]
pairplot(body[:, which], body_columns[which])
plt.show()
```

Plotting variables against themselves is rather silly (exercise: what would that be?). Therefore, on the main diagonal of [Figure 7.4](#), we printed out the variable names.

A scatter plot matrix can be a valuable tool for identifying noteworthy combinations of columns in our datasets. We see that some pairs of variables are more “structured” than others, e.g., hip circumference and weight are more or less aligned on a straight line. This is why [Chapter 9](#) will describe ways to model the possible relationships between the variables.

Exercise 7.9 Create a pairs plot where weight, arm circumference, and hip circumference are on the log-scale.

Exercise 7.10 (*) Call `seaborn.pairplot` to create a scatter plot matrix with histograms on the main diagonal, thanks to which you will be able to see how the marginal distributions are distributed. Note that the matrix must, unfortunately, be converted to a *pandas* data frame first.

Exercise 7.11 (**) Modify our `pairplot` function so that it displays the histograms of the marginal distributions on the main diagonal.




Figure 7.4. The scatter plot matrix for selected columns in the body dataset.

7.5 Exercises

Exercise 7.12 What is the difference between `[1, 2, 3]`, `[[1, 2, 3]]`, and `[[1], [2], [3]]` in the context of an array's creation?

Exercise 7.13 If A is a matrix with five rows and six columns, what is the difference between `A.reshape(6, 5)` and `A.T`?

Exercise 7.14 If A is a matrix with 5 rows and 6 columns, what is the meaning of: `A.reshape(-1)`, `A.reshape(3, -1)`, `A.reshape(-1, 3)`, `A.reshape(-1, -1)`, `A.shape = (3, 10)`, and `A.shape = (-1, 3)`?

Exercise 7.15 List some methods to add a new row or column to an existing matrix.

Exercise 7.16 *Give some ways to visualise three-dimensional data.*

Exercise 7.17 *How can we set point opaqueness/transparency when drawing a scatter plot? When would we be interested in this?*

8

Processing multidimensional data

8.1 Extending vectorised operations to matrices

The vector operations from [Chapter 5](#) are brilliant examples of the *write less, do more* principle in practice. Let's see how are they extended to matrices.

8.1.1 Vectorised mathematical functions

Applying vectorised functions such as `numpy.round`, `numpy.log`, and `numpy.exp` returns an array of the same shape, with all elements transformed accordingly:

$$f(\mathbf{X}) = \begin{bmatrix} f(x_{1,1}) & f(x_{1,2}) & \cdots & f(x_{1,m}) \\ f(x_{2,1}) & f(x_{2,2}) & \cdots & f(x_{2,m}) \\ \vdots & \vdots & \ddots & \vdots \\ f(x_{n,1}) & f(x_{n,2}) & \cdots & f(x_{n,m}) \end{bmatrix}.$$

```
A = np.array([
    [0.2, 0.6, 0.4, 0.4],
    [0.0, 0.2, 0.4, 0.7],
    [0.8, 0.8, 0.2, 0.1]
]) # example matrix
```

For instance, to take the square of every element, we can call:

```
np.square(A)
## array([[0.04, 0.36, 0.16, 0.16],
##        [0.  , 0.04, 0.16, 0.49],
##        [0.64, 0.64, 0.04, 0.01]])
```

8.1.2 Componentwise aggregation

Unidimensional aggregation functions (e.g., `numpy.mean`, `numpy.quantile`) can be applied to summarise:

- all data into a single number (`axis=None`, being the default),
- data in each column (`axis=0`), as well as
- data in each row (`axis=1`).

Here are the corresponding examples:

```
np.mean(A)
## 0.39999999999999997
np.mean(A, axis=0)
## array([0.33333333, 0.53333333, 0.33333333, 0.4       ])
np.mean(A, axis=1)
## array([0.4    , 0.325, 0.475])
```

Important Let's stress that `axis=1` does not mean that we get the column means (even though columns constitute the second axis, and we count starting at 0). It denotes the axis *along* which the matrix is sliced. Sadly, even yours truly sometimes does not get it right.

Exercise 8.1 Given the `nhanes_adult_female_bmx_2020`¹ dataset, compute the mean, standard deviation, the minimum, and the maximum of each body measure.

We will get back to the topic of the aggregation of multidimensional data in [Section 8.4](#).

8.1.3 Arithmetic, logical, and relational operations

Recall that for vectors, binary operators such as `+`, `*`, `==`, `<=`, and `&` as well as similar elementwise functions (e.g., `numpy.minimum`) can be applied if both inputs are of the same length. For example:

```
np.array([1, 10, 100, 1000]) * np.array([7, -6, 2, 8]) # elementwisely
## array([ 7, -60, 200, 8000])
```

Alternatively, one input can be a scalar:

```
np.array([1, 10, 100, 1000]) * -3
## array([-3, -30, -300, -3000])
```

More generally, a set of rules referred to in the `numpy` manual as *broadcasting*² describes how this package handles arrays of different shapes.

Important Generally, for two matrices, their column/row counts must match or be equal to 1. Also, if one operand is a one-dimensional array, it will be promoted to a row vector.

Let's explore all the possible scenarios.

¹ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_bmx_2020.csv

² <https://numpy.org/devdocs/user/basics.broadcasting.html>

Matrix vs scalar

If one operand is a scalar, then it is going to be propagated over all matrix elements. For example:

```
(-1)*A
## array([[ -0.2,  -0.6,  -0.4,  -0.4],
##        [ -0. ,  -0.2,  -0.4,  -0.7],
##        [ -0.8,  -0.8,  -0.2,  -0.1]])
```

It changed the sign of every element, which is, mathematically, an instance of multiplying a matrix $\mathbf{X}$ by a scalar c :

$$c\mathbf{X} = \begin{bmatrix} cx_{1,1} & cx_{1,2} & \cdots & cx_{1,m} \\ cx_{2,1} & cx_{2,2} & \cdots & cx_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ cx_{n,1} & cx_{n,2} & \cdots & cx_{n,m} \end{bmatrix}.$$

Furthermore, we can take the square³ of each element:

```
A**2
## array([[0.04, 0.36, 0.16, 0.16],
##        [0. , 0.04, 0.16, 0.49],
##        [0.64, 0.64, 0.04, 0.01]])
```

or compare each element to 0.25.

```
A >= 0.25
## array([[False,  True,  True,  True],
##        [False, False,  True,  True],
##        [ True,  True, False, False]])
```

Matrix vs matrix

For two matrices of identical sizes, we act on the *corresponding* elements:

```
B = np.tri(A.shape[0], A.shape[1]) # just an example
B # a lower triangular 0-1 matrix
## array([[1., 0., 0., 0.],
##        [1., 1., 0., 0.],
##        [1., 1., 1., 0.]])
```

And now:

```
A * B
## array([[0.2, 0. , 0. , 0. ],
##        [0. , 0.2, 0. , 0. ],
##        [0.8, 0.8, 0.2, 0. ]])
```

³ This is not the same as matrix-multiply by itself which we cover in [Section 8.3](#).

multiplies each $a_{i,j}$ by the corresponding $b_{i,j}$.

This behaviour extends upon the idea from algebra that given **A** and **B** with n rows and m columns each, the result of $+$ is:

$$\mathbf{A} + \mathbf{B} = \begin{bmatrix} a_{1,1} + b_{1,1} & a_{1,2} + b_{1,2} & \cdots & a_{1,m} + b_{1,m} \\ a_{2,1} + b_{2,1} & a_{2,2} + b_{2,2} & \cdots & a_{2,m} + b_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n,1} + b_{n,1} & a_{n,2} + b_{n,2} & \cdots & a_{n,m} + b_{n,m} \end{bmatrix}.$$

Thanks to the matrix-matrix and matrix-scalar operations, we can perform various tests on a per-element basis, e.g.,

```
(A >= 0.25) & (A <= 0.75) # logical matrix & logical matrix
## array([[False,  True,  True,  True],
##        [False, False,  True,  True],
##        [False, False, False, False]])
```

Example 8.2 (*) Figure 8.1 depicts a (filled) contour plot of Himmelblau's function, $f(x, y) = (x^2 + y - 11)^2 + (x + y^2 - 7)^2$, for $x \in [-5, 5]$ and $y \in [-4, 4]$. To draw it, we probe 250 points from these two intervals, and call `numpy.meshgrid` to generate two matrices, both of shape 250 by 250, giving the x - and y -coordinates of all the points on the corresponding two-dimensional grid. Thanks to this, we are able to use vectorised mathematical operations to compute the values off thereon.

```
x = np.linspace(-5, 5, 250)
y = np.linspace(-4, 4, 250)
xg, yg = np.meshgrid(x, y)
z = (xg**2 + yg - 11)**2 + (xg + yg**2 - 7)**2
plt.contourf(x, y, z, levels=20)
CS = plt.contour(x, y, z, levels=[1, 5, 10, 20, 50, 100, 150, 200, 250])
plt.clabel(CS, colors="black")
plt.show()
```

To understand the result generated by `numpy.meshgrid`, let's inspect its output for a smaller number of probe points:

```
x = np.linspace(-5, 5, 3)
y = np.linspace(-4, 4, 5)
xg, yg = np.meshgrid(x, y)
xg
## array([[-5.,  0.,  5.],
##        [-5.,  0.,  5.],
##        [-5.,  0.,  5.],
##        [-5.,  0.,  5.],
##        [-5.,  0.,  5.]])
```

Here, each column consists of the same values.




Figure 8.1. An example filled contour plot with additional labelled contour lines.

```
yg
## array([[ -4.,  -4.,  -4.],
##        [ -2.,  -2.,  -2.],
##        [  0.,   0.,   0.],
##        [  2.,   2.,   2.],
##        [  4.,   4.,   4.]])
```

In this case, each row is constant. Therefore, calling:

```
(xg**2 + yg - 11)**2 + (xg + yg**2 - 7)**2
## array([[116., 306., 296.],
##        [208., 178., 148.],
##        [340., 170., 200.],
##        [320.,  90., 260.],
##        [340., 130., 520.]])
```

gives a matrix $\mathbf{Z}$ such that $z_{i,j}$ is generated by considering the i -th element in y and the j -th item in x , which is exactly what we desired. We will provide an alternative implementation in [Example 8.5](#).

Matrix vs any vector

An $n \times m$ matrix can also be combined with an $n \times 1$ column vector:

```
A * np.array([1, 10, 100]).reshape(-1, 1)
## array([[ 0.2,  0.6,  0.4,  0.4],
##        [ 0. ,  2. ,  4. ,  7. ],
##        [80. , 80. , 20. , 10. ]])
```

It propagated the column vector over all columns (left to right). Similarly, combining a matrix with a $1 \times m$ row vector recycles the latter over all rows (top to bottom).

```
A + np.array([1, 2, 3, 4]).reshape(1, -1)
## array([[1.2, 2.6, 3.4, 4.4],
##        [1. , 2.2, 3.4, 4.7],
##        [1.8, 2.8, 3.2, 4.1]])
```

If one operand is a one-dimensional array or a list of length m , it will be treated as a row vector. For example, here is an instance of *centring* of each column:

```
np.round(A - np.mean(A, axis=0), 3) # matrix - vector
## array([[ -0.133,  0.067,  0.067, -0.   ],
##        [ -0.333, -0.333,  0.067,  0.3   ],
##        [ 0.467,  0.267, -0.133, -0.3   ]])
```

An explicit `.reshape(1, -1)` was not necessary.

Mathematically, although it is not necessarily a standard notation, we will allow adding and subtracting row vectors from matrices of compatible sizes:

$$\mathbf{X} + \mathbf{t} = \mathbf{X} + [t_1 \ t_2 \ \dots \ t_m] = \begin{bmatrix} x_{1,1} + t_1 & x_{1,2} + t_2 & \dots & x_{1,m} + t_m \\ x_{2,1} + t_1 & x_{2,2} + t_2 & \dots & x_{2,m} + t_m \\ \vdots & \vdots & \ddots & \vdots \\ x_{n,1} + t_1 & x_{n,2} + t_2 & \dots & x_{n,m} + t_m \end{bmatrix}.$$

This corresponds to shifting (translating) every row in the matrix.

Exercise 8.3 In the `nhanes_adult_female_bmx_2020`⁴ dataset, standardise, normalise, and min-max scale every column (compare Section 5.3.2). A single line of code will suffice in each case.

Row vector vs column vector (*)

A row vector combined with a column vector results in an operation's being performed on each *combination* of all pairs of elements in the two arrays (i.e., the cross-product; not just the *corresponding* pairs).

```
np.arange(1, 8).reshape(1, -1) * np.array([1, 10, 100]).reshape(-1, 1)
## array([[ 1,  2,  3,  4,  5,  6,  7],
##        [10, 20, 30, 40, 50, 60, 70],
##        [100, 200, 300, 400, 500, 600, 700]])
```

Exercise 8.4 Check out that `numpy.nonzero` relies on similar shape broadcasting rules as the binary operators we discussed here, but not with respect to all three arguments.

Example 8.5 (*) Himmelblau's function in Example 8.2 is defined by means of arithmetic operators only, and they all rely on the kind of shape broadcasting that we discuss in this section. Consequently, calling `numpy.meshgrid` to evaluate `f` on a point grid was not really necessary:

⁴ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_bmx_2020.csv


```
x = np.linspace(-5, 5, 3)
y = np.linspace(-4, 4, 5)
xg = x.reshape(1, -1)
yg = y.reshape(-1, 1)
(xg**2 + yg - 11)**2 + (xg + yg**2 - 7)**2
## array([[116., 306., 296.],
##        [208., 178., 148.],
##        [340., 170., 200.],
##        [320., 90., 260.],
##        [340., 130., 520.]])
```

See also the `sparse` parameter in `numpy.meshgrid`, and [Section 12.3.1](#) where this function turns out useful after all.

8.1.4 Other row and column transforms (*)

Some functions discussed in the previous part of this course are equipped with the `axis` argument, which supports processing each row or column independently. For example, to compute the ranks of elements in each column, we can call:

```
scipy.stats.rankdata(A, axis=0) # columnwisely (along the rows)
## array([[2. , 2. , 2.5, 2. ],
##        [1. , 1. , 2.5, 3. ],
##        [3. , 3. , 1. , 1. ]])
```

Some functions have the default argument `axis=-1` meaning that they are applied along the last⁵ axis (i.e., columns in the matrix case):

```
np.diff(A) # means axis=1 in this context (along the columns)
## array([[ 0.4, -0.2,  0. ],
##        [ 0.2,  0.2,  0.3],
##        [ 0. , -0.6, -0.1]])
```

Compare the foregoing to the iterated differences in each column separately (along the rows):

```
np.diff(A, axis=0)
## array([[ -0.2, -0.4,  0. ,  0.3],
##        [ 0.8,  0.6, -0.2, -0.6]])
```

If a vectorised function is not equipped with the `axis` argument, we can propagate it over all the rows or columns by calling `numpy.apply_along_axis`. For instance, here is another (did you solve [Exercise 8.3?](#)) way to compute the z-scores in each matrix column:

```
def standardise(x):
```

(continues on next page)

⁵ `numpy.mean` is amongst the many exceptions to this rule.

(continued from previous page)

```

return (x-np.mean(x))/np.std(x)

np.round(np.apply_along_axis(standardise, 0, A), 2) # round for readability
## array([[ -0.39,  0.27,  0.71, -0.  ],
##        [ -0.98, -1.34,  0.71,  1.22],
##        [  1.37,  1.07, -1.41, -1.22]])

```

But, of course, we prefer `(x-np.mean(x, axis=0))/np.std(x, axis=0)`.

Note (*) Matrices are iterable (in the sense of [Section 3.4](#)), but in an interesting way. Namely, an iterator traverses through each row in a matrix. Writing:

```
r1, r2, r3 = A # A has three rows
```

creates three variables, each representing a separate row in `A`, the second of which is:

```

r2
## array([0. , 0.2, 0.4, 0.7])

```

8.2 Indexing matrices

Recall that for unidimensional arrays, we have four possible indexer choices (i.e., when performing filtering like `x[i]`):

- scalar (extracts a single element),
- slice (selects a regular subsequence, e.g., every second element or the first six items; returns a *view* of existing data: it does not make an independent copy of the subsetted elements),
- integer vector (selects the elements at given indexes),
- logical vector (selects the elements that correspond to `True` in the indexer).

Matrices are two-dimensional arrays. Subsetting thus requires two indexes. By writing `A[i, j]`, we select rows given by `i` and columns given by `j`. Both `i` and `j` can be one of the four aforementioned types, so we have at ten different cases to consider (skipping the symmetric ones).

Important Generally:

- each scalar index reduces the dimensionality of the subsetted object by 1;
- slice-slice and slice-scalar indexing returns a *view* of the existing array, so we need to be careful when modifying the resulting object;

- usually, indexing returns a submatrix (subblock), which is a combination of elements at given rows and columns;
- indexing with two integer or logical vectors is performed elementwisely, and should be avoided if find the rules of shape broadcasting too complicated.

8.2.1 Slice-based indexing

Our favourite example matrix again:

```
A = np.array([
    [0.2, 0.6, 0.4, 0.4],
    [0.0, 0.2, 0.4, 0.7],
    [0.8, 0.8, 0.2, 0.1]
])
```

Indexing based on two slices selects a submatrix:

```
A[::2, 3:] # every second row, skip the first three columns
## array([[0.4],
##        [0.1]])
```

An empty slice selects all elements on the corresponding axis:

```
A[:, ::-1] # all rows, reversed columns
## array([[0.4, 0.4, 0.6, 0.2],
##        [0.7, 0.4, 0.2, 0. ],
##        [0.1, 0.2, 0.8, 0.8]])
```

Let's stress that the result is *always* in the form of a matrix.

8.2.2 Scalar-based indexing

Indexing by a scalar selects a given row or column, reducing the dimensionality of the output object:

```
A[:, 3] # one scalar: from two to one dimensions
## array([0.4, 0.7, 0.1])
```

It selected the fourth column and gave a flat vector (we can always use the `reshape` method to convert the resulting object back to a matrix). Furthermore:

```
A[0, -1] # two scalars: from two to zero dimensions
## 0.4
```

It yielded the element (scalar) in the first row and the last column.

8.2.3 Mixed logical/integer vector and scalar/slice indexers

A logical and integer vector-like object can also be employed for element selection. If the other indexer is a slice or a scalar, the result is quite predictable. For instance:

```
A[ [0, -1, 0], ::-1 ]
## array([[0.4, 0.4, 0.6, 0.2],
##        [0.1, 0.2, 0.8, 0.8],
##        [0.4, 0.4, 0.6, 0.2]])
```

It selected the first, the last, and the first row again. Then, it reversed the order of columns.

```
A[ A[:, 0] > 0.1, : ]
## array([[0.2, 0.6, 0.4, 0.4],
##        [0.8, 0.8, 0.2, 0.1]])
```

It chose the rows from A where the values in the first column of A are greater than 0.1.

```
A[np.mean(A, axis=1) > 0.35, : ]
## array([[0.2, 0.6, 0.4, 0.4],
##        [0.8, 0.8, 0.2, 0.1]])
```

It fetched the rows whose mean is greater than 0.35.

```
A[np.argsort(A[:, 0]), : ]
## array([[0. , 0.2, 0.4, 0.7],
##        [0.2, 0.6, 0.4, 0.4],
##        [0.8, 0.8, 0.2, 0.1]])
```

It ordered the matrix with respect to the values in the first column (all rows permuted in the same way, together).

Exercise 8.6 In the *nhanes_adult_female_bmx_2020*⁶ dataset, select all the participants whose heights are within their mean ± 2 standard deviations.

8.2.4 Two vectors as indexers (*)

Indexing based on two logical or integer vectors is a tad more horrible, as in this case not only some form of *shape broadcasting* comes into play but also all the headache-inducing exceptions listed in the perhaps not the most clearly written *Advanced Indexing*⁷ section of the **numpy** manual. Cheer up, though: [Section 10.5](#) points out that indexing in **pandas** is even more troublesome.

For the sake of our maintaining sanity, in practice, it is best to be extra careful when using two vector indexers and stick only to the scenarios discussed beneath. First, with two flat integer indexers, we pick elementwisely:

⁶ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_adult_female_bmx_2020.csv

⁷ <https://numpy.org/doc/stable/user/basics.indexing.html>

```
A[ [0, -1, 0, 2, 0], [1, 2, 0, 2, 1] ]
## array([0.6, 0.2, 0.2, 0.2, 0.6])
```

It yielded `A[0, 1]`, `A[-1, 2]`, `A[0, 0]`, `A[2, 2]`, and `A[0, 1]`.

Second, to select a submatrix (a *subblock*) using integer indexes, it is best to make sure that the first indexer is a column vector, and the second one is a row vector (or some objects like these, e.g., compatible lists of lists).

```
A[ [[0], [-1]], [[1, 3]] ] # column vector-like list, row vector-like list
## array([[0.6, 0.4],
##        [0.8, 0.1]])
```

Third, if indexing involves logical vectors, it is best to convert them to integer ones first (e.g., by calling `numpy.flatnonzero`).

```
A[ np.flatnonzero(np.mean(A, axis=1) > 0.35).reshape(-1, 1), [[0, 2, 3, 0]] ]
## array([[0.2, 0.4, 0.4, 0.2],
##        [0.8, 0.2, 0.1, 0.8]])
```

The necessary reshaping can be outsourced to `numpy.ix_` function:

```
A[ np.ix_( np.mean(A, axis=1) > 0.35, [0, 2, 3, 0] ) ] # np.ix_(rows, cols)
## array([[0.2, 0.4, 0.4, 0.2],
##        [0.8, 0.2, 0.1, 0.8]])
```

Alternatively, we can always apply indexing twice:

```
A[np.mean(A, axis=1) > 0.45, :][:, [0, 2, 3, 0]]
## array([[0.8, 0.2, 0.1, 0.8]])
```

This is only a mild inconvenience. We will be forced to apply such double indexing anyway in **pandas** whenever selecting rows *by position* and columns *by name* is required; see [Section 10.5](#).

Note (*) Interestingly, we can also index a vector using an integer matrix. This is like subsetting using a list of integer indexes, but the output's shape matches that of the indexer:

```
u = np.array(["a", "b", "c"])
V = np.array([ [0, 1], [1, 0], [2, 1] ])
u[V] # like u[V.ravel()].reshape(V.shape)
## array(['a', 'b'],
##        ['b', 'a'],
##        ['c', 'b']], dtype='<U1')
```

8.2.5 Views of existing arrays (*)

Only the indexing involving two slices or a slice and a scalar returns a [view](#)⁸ on an existing array. For example:

```
B = A[:, ::2]
B
## array([[0.2, 0.4],
##        [0. , 0.4],
##        [0.8, 0.2]])
```

Now B and A share memory. By modifying B in place, e.g.:

```
B *= -1
```

the changes will be visible in A as well:

```
A
## array([[ -0.2,  0.6, -0.4,  0.4],
##        [ -0. ,  0.2, -0.4,  0.7],
##        [ -0.8,  0.8, -0.2,  0.1]])
```

This is time and memory efficient, but might lead to some unexpected results if we are being rather absent-minded. The readers have been warned.

In all other cases, we get a copy of the subsetted array.

8.2.6 Adding and modifying rows and columns

With slice- and scalar-based indexers, rows, columns, or individual elements can be replaced by new content in a natural way:

```
A[:, 0] = A[:, 0]**2
```

With **numpy** arrays, however, new rows or columns cannot be added via the index operator. Instead, the whole array needs to be created from scratch using, e.g., one of the functions discussed in [Section 7.1.4](#). For example:

```
A = np.column_stack((A, np.sqrt(A[:, 0])))
A
## array([[ 0.04,  0.6 , -0.4 ,  0.4 ,  0.2 ],
##        [ 0. ,  0.2 , -0.4 ,  0.7 ,  0. ],
##        [ 0.64,  0.8 , -0.2 ,  0.1 ,  0.8 ]])
```

⁸ <https://numpy.org/devdocs/user/basics.copies.html>

8.3 Matrix multiplication, dot products, and Euclidean norm (*)

Matrix algebra is at the core of all the methods used in data analysis with the matrix multiplication being often the most crucial operation (e.g., [22, 42]). Given $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{B} \in \mathbb{R}^{p \times m}$, their *multiply* is a matrix $\mathbf{C} = \mathbf{AB} \in \mathbb{R}^{n \times m}$ such that $c_{i,j}$ is the sum of the i -th row in $\mathbf{A}$ and the j -th column in $\mathbf{B}$ multiplied elementwisely:

$$c_{i,j} = a_{i,1}b_{1,j} + a_{i,2}b_{2,j} + \dots + a_{i,p}b_{p,j} = \sum_{k=1}^p a_{i,k}b_{k,j},$$

for $i = 1, \dots, n$ and $j = 1, \dots, m$. For example:

```
A = np.array([
    [1, 0, 1],
    [2, 2, 1],
    [3, 2, 0],
    [1, 2, 3],
    [0, 0, 1],
])
B = np.array([
    [1, 0, 0, 0],
    [0, 4, 1, 3],
    [2, 0, 3, 1],
])
```

And now:

```
C = A @ B # or: A.dot(B)
C
## array([[ 3,  0,  3,  1],
##        [ 4,  8,  5,  7],
##        [ 3,  8,  2,  6],
##        [ 7,  8, 11,  9],
##        [ 2,  0,  3,  1]])
```

Mathematically, we can denote it by:

$$\begin{bmatrix} 1 & 0 & 1 \\ 2 & 2 & 1 \\ 3 & 2 & 0 \\ \mathbf{1} & \mathbf{2} & \mathbf{3} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & \mathbf{0} & 0 \\ 0 & 4 & \mathbf{1} & 3 \\ 2 & 0 & \mathbf{3} & 1 \end{bmatrix} = \begin{bmatrix} 3 & 0 & 3 & 1 \\ 4 & 8 & 5 & 7 \\ 3 & 8 & 2 & 6 \\ 7 & 8 & \mathbf{11} & 9 \\ 2 & 0 & 3 & 1 \end{bmatrix}.$$

For example, the element in the fourth row and the third column, $c_{4,3}$ takes the fourth row in the left matrix $\mathbf{a}_{4,\cdot} = [1 \ 2 \ 3]$ and the third column in the right matrix $\mathbf{b}_{\cdot,3} = [0 \ 1 \ 3]^T$ (emphasised with bold), multiplies the corresponding elements, and computes their sum, i.e., $c_{4,3} = 1 \cdot 0 + 2 \cdot 1 + 3 \cdot 3 = 11$.

Important Matrix multiplication can only be performed on two matrices of *compatible sizes*: the number of columns in the left matrix must match the number of rows in the right operand.

Another example:

```
A = np.array([
    [1, 2],
    [3, 4]
])
I = np.array([ # np.eye(2)
    [1, 0],
    [0, 1]
])
A @ I # or A.dot(I)
## array([[1, 2],
##        [3, 4]])
```

We matrix-multiplied **A** by the identity matrix **I**, which is the neutral element of the said operation. This is why the result is identical to **A**.

Important In most textbooks, just like in this one, **AB** always denotes the *matrix* multiplication. This is a very different operation from the *elementwise* multiplication.

Compare the above to:

```
A * I # elementwise multiplication (the Hadamard product)
## array([[1, 0],
##        [0, 4]])
```

Exercise 8.7 (*) Show that $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$. Also notice that, typically, matrix multiplication is not commutative, i.e., $\mathbf{AB} \neq \mathbf{BA}$.

Note By definition, matrix multiplication is convenient for denoting sums of products of corresponding elements in many pairs of vectors, which we refer to as dot products. More formally, given two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^p$, their *dot (or scalar) product* is:

$$\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^p x_i y_i.$$

In matrix multiplication terms, if $\mathbf{x}$ is a row vector and $\mathbf{y}^T$ is a column vector, then the above can be written as $\mathbf{xy}^T$. The result is a single number.

In particular, the dot product of a vector and itself:

$$\mathbf{x} \cdot \mathbf{x} = \sum_{i=1}^p x_i^2,$$

is the square of the Euclidean norm of $\mathbf{x}$ (simply the sum of squares), which is used to measure the *magnitude* of a vector (Section 5.3.2):

$$\|\mathbf{x}\| = \sqrt{\sum_{i=1}^p x_i^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}} = \sqrt{\mathbf{x}\mathbf{x}^T}.$$

It is worth pointing out that the Euclidean norm fulfils (amongst others) the condition that $\|\mathbf{x}\| = 0$ if and only if $\mathbf{x} = \mathbf{0} = (0, 0, \dots, 0)$. The same naturally holds for its square.

Exercise 8.8 Show that $\mathbf{A}^T \mathbf{A}$ gives the matrix that consists of the dot products of all the pairs of columns in $\mathbf{A}$ and $\mathbf{A}\mathbf{A}^T$ stores the dot products of all the pairs of rows.

Section 9.3.2 will note that matrix multiplication can be used as a way to express certain geometrical transformations of points in a dataset, e.g., scaling and rotating. Also, Section 9.3.3 briefly discusses the concept of the inverse of a matrix. Furthermore, Section 9.3.4 introduces its singular value decomposition.

8.4 Pairwise distances and related methods (*)

Many data analysis methods rely on the notion of *distances*, which quantify the extent to which two points (e.g., two rows in a matrix) differ from each other. Here, we will be dealing with the most natural⁹ distance called the Euclidean metric. We know it from school, where we measured it using a ruler.

8.4.1 Euclidean metric (*)

Given two points in $\mathbb{R}^m$, $\mathbf{u} = (u_1, \dots, u_m)$ and $\mathbf{v} = (v_1, \dots, v_m)$, the *Euclidean metric* is defined in terms of the corresponding Euclidean norm:

$$\|\mathbf{u} - \mathbf{v}\| = \sqrt{(u_1 - v_1)^2 + (u_2 - v_2)^2 + \dots + (u_m - v_m)^2} = \sqrt{\sum_{i=1}^m (u_i - v_i)^2},$$

i.e., the square root of the sum of squared differences between the corresponding coordinates.

In particular, for unidimensional data ($m = 1$), we have $\|\mathbf{u} - \mathbf{v}\| = |u_1 - v_1|$, i.e., the absolute value of the difference.

⁹ There are many possible distances, allowing to measure the similarity of points not only in $\mathbb{R}^m$, but also character strings (e.g., the Levenshtein metric), ratings (e.g., cosine dissimilarity), etc. There is even an encyclopedia of distances, [25].

Important Given two vectors of equal lengths $\mathbf{x}, \mathbf{y} \in \mathbb{R}^m$, the dot product of their difference:

$$(\mathbf{x} - \mathbf{y}) \cdot (\mathbf{x} - \mathbf{y}) = (\mathbf{x} - \mathbf{y})(\mathbf{x} - \mathbf{y})^T = \sum_{i=1}^m (x_i - y_i)^2,$$

is nothing else than the square of the Euclidean distance between them.

Exercise 8.9 Consider the following matrix $\mathbf{X} \in \mathbb{R}^{4 \times 2}$:

$$\mathbf{X} = \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ -3/2 & 1 \\ 1 & 1 \end{bmatrix}.$$

Calculate (by hand): $\|\mathbf{x}_{1,\cdot} - \mathbf{x}_{2,\cdot}\|$, $\|\mathbf{x}_{1,\cdot} - \mathbf{x}_{3,\cdot}\|$, $\|\mathbf{x}_{1,\cdot} - \mathbf{x}_{4,\cdot}\|$, $\|\mathbf{x}_{2,\cdot} - \mathbf{x}_{4,\cdot}\|$, $\|\mathbf{x}_{2,\cdot} - \mathbf{x}_{3,\cdot}\|$, $\|\mathbf{x}_{1,\cdot} - \mathbf{x}_{1,\cdot}\|$, and $\|\mathbf{x}_{2,\cdot} - \mathbf{x}_{1,\cdot}\|$.

`scipy.spatial.distance.cdist` computes the distances between all the possible pairs of rows in two matrices $\mathbf{X} \in \mathbb{R}^{n \times m}$ and $\mathbf{Y} \in \mathbb{R}^{k \times m}$. We need to be careful, though. It brings about a distance matrix of size $n \times k$, which can become large. For instance, for $n = k = 100\,000$, we need roughly 80 GB of RAM to store it.

Here are the distances between all the pairs of points in the same dataset.

```
X = np.array([
    [0, 0],
    [1, 0],
    [-1.5, 1],
    [1, 1]
])
import scipy.spatial.distance
D = scipy.spatial.distance.cdist(X, X)
D
## array([[0.          , 1.          , 1.80277564, 1.41421356],
##        [1.          , 0.          , 2.6925824 , 1.          ],
##        [1.80277564, 2.6925824 , 0.          , 2.5          ],
##        [1.41421356, 1.          , 2.5          , 0.          ]])
```

Hence, $d_{i,j} = \|\mathbf{x}_{i,\cdot} - \mathbf{x}_{j,\cdot}\|$. That we have zeros on the diagonal is due to the fact that $\|\mathbf{u} - \mathbf{v}\| = 0$ if and only if $\mathbf{u} = \mathbf{v}$. Furthermore, $\|\mathbf{u} - \mathbf{v}\| = \|\mathbf{v} - \mathbf{u}\|$, which implies the symmetry of $\mathbf{D}$, i.e., we have $\mathbf{D}^T = \mathbf{D}$.

Figure 8.2 illustrates the six non-trivial pairwise distances. In the left subplot, our perception of distance is disturbed because the aspect ratio (the ratio between the range of the x-axis to the range of the y-axis) is not 1:1. To be able to assess spatial relationships, it is thus very important to call `matplotlib.pyplot.axis("equal")`.

```

for s in range(2):
    plt.subplot(1, 2, s+1)
    if s == 1: plt.axis("equal") # right subplot
    plt.plot(X[:, 0], X[:, 1], "ko")
    for i in range(X.shape[0]-1):
        for j in range(i+1, X.shape[0]):
            plt.plot(X[[i,j], 0], X[[i,j], 1], "k-", alpha=0.2)
            plt.text(
                np.mean(X[[i,j], 0]),
                np.mean(X[[i,j], 1]),
                np.round(D[i, j], 2)
            )
plt.show()

```




Figure 8.2. Distances between four example points. In the left plot, their perception is disturbed because the aspect ratio is not 1:1.

Exercise 8.10 (*) Each metric also enjoys the triangle inequality: $\|\mathbf{u} - \mathbf{v}\| \leq \|\mathbf{u} - \mathbf{w}\| + \|\mathbf{w} - \mathbf{v}\|$ for all $\mathbf{u}, \mathbf{v}, \mathbf{w}$. Verify that this property holds by studying each triple of points in an example distance matrix.

Important A few popular data science techniques rely on pairwise distances, e.g.:

- multidimensional data aggregation (undermentioned),
- k -means clustering (Section 12.4),
- k -nearest neighbour regression (Section 9.2.1) and classification (Section 12.3.1),
- missing value imputation (Section 15.1),

- density estimation (which we can use outlier detection, see [Section 15.4](#)).

They assume that data have been appropriately preprocessed; compare, e.g., [2]. In particular, matrix columns should be on the same scale (e.g., standardised) as otherwise computing sums of their squared differences might not make sense at all.

8.4.2 Centroids (*)

So far we have been only discussing ways to aggregate unidimensional data, e.g., each matrix column separately. Some of the introduced summaries can be generalised to the multidimensional case.

For instance, the arithmetic mean of a vector $(x_1, \dots, x_n)$ is a point c that minimises the sum of the *squared* unidimensional distances between itself and all the x_i s, i.e., the minimiser of $\sum_{i=1}^n \|x_i - c\|^2 = \sum_{i=1}^n (x_i - c)^2$. More generally, we can define the *centroid* of a dataset $\mathbf{X} \in \mathbb{R}^{n \times m}$ as the point $\mathbf{c} \in \mathbb{R}^m$ to which the overall *squared* distance is the smallest:

$$\text{minimise } \sum_{i=1}^n \|\mathbf{x}_{i,\cdot} - \mathbf{c}\|^2 \quad \text{w.r.t. } \mathbf{c}.$$

Its solution is:

$$\mathbf{c} = \frac{1}{n} (\mathbf{x}_{1,\cdot} + \mathbf{x}_{2,\cdot} + \dots + \mathbf{x}_{n,\cdot}) = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_{i,\cdot},$$

which is the componentwise (columnwise) arithmetic mean. In other words, its j -th component is given by:

$$c_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}.$$

For example, the centroid of the dataset depicted in [Figure 8.2](#) is:

```
c = np.mean(X, axis=0)
c
## array([0.125, 0.5  ])
```

Centroids are the basis for the k -means clustering method that we discuss in [Section 12.4](#).

8.4.3 Multidimensional dispersion and other aggregates (**)

Furthermore, as a measure of multidimensional dispersion, we can consider the natural generalisation of the standard deviation:

$$s = \sqrt{\frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_{i,\cdot} - \mathbf{c}\|^2},$$

being the square root of the average squared distance to the centroid. Notice that s is a single number.

```
np.sqrt(np.mean(scipy.spatial.distance.cdist(X, c.reshape(1, -1))**2))
## 1.1388041973930374
```

Note (**) Generalising other aggregation functions is not a trivial task because, amongst others, there is no natural linear ordering relation in the multidimensional space (see, e.g., [78]). For instance, any point on the convex hull of a dataset could serve as an analogue of the minimal and maximal observation.

Furthermore, the componentwise median does not behave nicely (it may, for example, fall outside the convex hull). Instead, we usually consider a different generalisation of the median: the point $\mathbf{m}$ which minimises the sum of distances (not squared), $\sum_{i=1}^n \|\mathbf{x}_{i,\cdot} - \mathbf{m}\|$. Even though it does not have an analytic solution, it can be determined algorithmically.

Note (**) A bag plot [84] is one of the possible multidimensional generalisations of the box-and-whisker plot. Unfortunately, its use is not popular amongst practitioners.

8.4.4 Fixed-radius and k -nearest neighbour search (**)

Several data analysis techniques rely on aggregating information about what is happening in the *local neighbourhoods* of given points. Let $\mathbf{X} \in \mathbb{R}^{n \times m}$ be a dataset and $\mathbf{x}' \in \mathbb{R}^m$ be some point, not necessarily from $\mathbf{X}$. We have two options:

- *fixed-radius search*: for some radius $r > 0$, we seek the indexes of all the points in $\mathbf{X}$ whose distance to $\mathbf{x}'$ is not greater than r :

$$B_r(\mathbf{x}') = \{i : \|\mathbf{x}_{i,\cdot} - \mathbf{x}'\| \leq r\};$$

- *few nearest neighbour search*: for some (usually small) integer $k \geq 1$, we seek the indexes of the k points in $\mathbf{X}$ which are the closest to $\mathbf{x}'$:

$$N_k(\mathbf{x}') = (i_1, i_2, \dots, i_k),$$

such that for all $j \notin \{i_1, \dots, i_k\}$:

$$\|\mathbf{x}_{i_1,\cdot} - \mathbf{x}'\| \leq \|\mathbf{x}_{i_2,\cdot} - \mathbf{x}'\| \leq \dots \leq \|\mathbf{x}_{i_k,\cdot} - \mathbf{x}'\| \leq \|\mathbf{x}_{j,\cdot} - \mathbf{x}'\|.$$

Note The set $S_r(\mathbf{x}') = \{\mathbf{u} : \|\mathbf{u} - \mathbf{x}'\| \leq r\}$ is the m -dimensional Euclidean ball (a solid hypersphere) of radius r centred at $\mathbf{x}'$. In particular, in $\mathbb{R}^1$, $S_r(\mathbf{x}')$ is the interval of

length $2r$ centred at $\mathbf{x}'$, i.e., $[x'_1 - r, x'_1 + r]$. In $\mathbb{R}^2$, $S_r(\mathbf{x}')$ is the circle of radius r centred at (x'_1, x'_2) .

Here is an example dataset, consisting of some randomly generated points; compare Figure 8.3.

```
np.random.seed(777)
X = np.random.randn(25, 2)
```

Let's inspect the local neighbourhood of the point $\mathbf{x}' = (0, 0)$ by computing the distances to each point in $\mathbf{X}$.

```
x_test = np.array([0, 0])
import scipy.spatial.distance
D = scipy.spatial.distance.cdist(X, x_test.reshape(1, -1))
```

For instance, here are the indexes of the points in $B_{0.75}(\mathbf{x}')$:

```
r = 0.75
B = np.flatnonzero(D <= r)
B
## array([ 1, 11, 14, 16, 24])
```

And here are the 11 nearest neighbours, $N_{11}(\mathbf{x}')$:

```
k = 11
N = np.argsort(D.reshape(-1))[:k]
N
## array([14, 24, 16, 11, 1, 22, 7, 19, 0, 9, 15])
```

Note that to prepare Figure 8.3, we need to set the aspect ratio to 1:1 as otherwise the circle would look like an ellipse.

```
fig, ax = plt.subplots()
ax.add_patch(plt.Circle(x_test, r, color="red", alpha=0.1))
for i in range(k):
    plt.plot(
        [x_test[0], X[N[i], 0]],
        [x_test[1], X[N[i], 1]],
        "r:", alpha=0.4
    )
plt.plot(X[:, 0], X[:, 1], "bo", alpha=0.1)
for i in range(X.shape[0]):
    plt.text(X[i, 0], X[i, 1], str(i), va="center", ha="center")
plt.plot(x_test[0], x_test[1], "rx")
plt.text(x_test[0], x_test[1], "$\\mathbf{x}'$", va="center", ha="center")
plt.axis("equal")
plt.show()
```




Figure 8.3. Fixed-radius search.

8.4.5 Spatial search with multidimensional binary search trees (**)

For efficiency reasons, rather than computing all pairwise distances, it is better to rely on dedicated data structures, especially if we have a large number of neighbourhood-related queries. `scipy` implements a spatial search algorithm based on multidimensional binary search trees called *K-d trees*¹⁰.

Note (*) In *K-d trees*, the data space is partitioned into hyperrectangles along the axes of the Cartesian coordinate system (standard basis). Thanks to such a representation, all subareas too far from the query point can be pruned to speed up the search.

Let's create a data structure for searching relative to the **X** matrix.

```
import scipy.spatial
T = scipy.spatial.KDTree(X)
```

Assume we would like to make queries with regard to three pivot points:

```
X_test = np.array([
    [0, 0],
    [2, 2],
    [2, -2]
])
```

¹⁰ In our context, we would like to refer to them as *m-d trees*, but we decided to stick with the traditional name.

Here are the results for the fixed radius searches ($r = 0.75$):

```
T.query_ball_point(X_test, 0.75)
## array([list([1, 11, 14, 16, 24]), list([20]), list([])], dtype=object)
```

We see that the method is nicely vectorised. We made a query about three points at the same time. As a result, we received a list-like object storing three lists representing the indexes of interest. Note that in the case of the third point, there are no elements in X within the requested range (circle), hence the empty index list.

And here are the five nearest neighbours:

```
distances, indexes = T.query(X_test, 5) # returns a tuple of length two
```

We obtained both the distances to the nearest neighbours:

```
distances
## array([[0.31457701, 0.44600012, 0.54848109, 0.64875661, 0.71635172],
##        [0.20356263, 1.45896222, 1.61587605, 1.64870864, 2.04640408],
##        [1.2494805 , 1.35482619, 1.93984334, 1.95938464, 2.08926502]])
```

and the indexes:

```
indexes
## array([[14, 24, 16, 11, 1],
##        [20, 5, 13, 2, 9],
##        [17, 3, 21, 12, 22]])
```

Each is a matrix with three rows (corresponding to the number of pivot points) and five columns (the number of neighbours sought).

Note (*) We expect the K -d trees to be much faster than the brute-force approach (where we compute all pairwise distances) in low-dimensional spaces. Nonetheless, due to the phenomenon called the *curse of dimensionality*, sometimes already for $m \geq 5$ the speed gains might be very small; see, e.g., [11].

8.5 Exercises

Exercise 8.11 Does `numpy.mean(A, axis=0)` compute the rowwise or columnwise means of A ?

Exercise 8.12 How does shape broadcasting work? List the most common pairs of shape cases when performing arithmetic operations like addition or multiplication.

Exercise 8.13 What are the possible ways to index a matrix?

Exercise 8.14 Which kinds of matrix indexers return a view of an existing array?

Exercise 8.15 (*) How can we select a submatrix comprised of the first and the last row and the first and the last column?

Exercise 8.16 Why appropriate data preprocessing is required when computing the Euclidean distance between the points in a matrix?

Exercise 8.17 What is the relationship between the dot product, the Euclidean norm, and the Euclidean distance?

Exercise 8.18 What is a centroid? How is it defined by means of the Euclidean distance between the points in a dataset?

Exercise 8.19 What is the difference between the fixed-radius and the few nearest neighbour search?

Exercise 8.20 (*) When K-d trees or other spatial search data structures might be better than a brute-force search based on `scipy.spatial.distance.cdist`?

Exercise 8.21 (**) See what kind of vector and matrix processing capabilities are available in the following packages: *TensorFlow*, *PyTorch*, *Theano*, and *tinygrad*. Are their APIs similar to that of *numpy*?

Exploring relationships between variables

Recall that in [Section 7.4](#), we performed some graphical exploratory analysis of the body measures recorded by the National Health and Nutrition Examination Survey:

```
body = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_adult_female_bmx_2020.csv",
    delimiter=",")[1:, :] # skip the first row (column names)
body[:6, :] # preview: the first six rows, all columns
## array([[ 97.1, 160.2,  34.7,  40.8,  35.8, 126.1, 117.9],
##        [ 91.1, 152.7,  33.5,  33. ,  38.5, 125.5, 103.1],
##        [ 73. , 161.2,  37.4,  38. ,  31.8, 106.2,  92. ],
##        [ 61.7, 157.4,  38. ,  34.7,  29. , 101. ,  90.5],
##        [ 55.4, 154.6,  34.6,  34. ,  28.3,  92.5,  73.2],
##        [ 62. , 144.7,  32.5,  34.2,  29.8, 106.7,  84.8]])
body.shape
## (4221, 7)
```

We already know that $n = 4221$ adult female participants are described by seven different numerical features, in this order:

```
body_columns = np.array([
    "weight",      # weight (kg)
    "height",      # standing height (cm)
    "arm len.",    # upper arm length (cm)
    "leg len.",    # upper leg length (cm)
    "arm circ.",   # arm circumference (cm)
    "hip circ.",   # hip circumference (cm)
    "waist circ."  # waist circumference (cm)
])
```

The data in different columns are somewhat *related* to each other. [Figure 7.4](#) indicates that higher hip circumferences *tend to occur together* with higher arm circumferences, and that the latter metric does not really tell us much about heights. In this chapter, we discuss ways to describe the possible relationships between variables.

9.1 Measuring correlation

Let's explore some basic means for measuring (expressing as a single number) the degree of association between *two* features.

9.1.1 Pearson linear correlation coefficient

The Pearson *linear correlation* coefficient is given by:

$$r(\mathbf{x}, \mathbf{y}) = \frac{1}{n} \sum_{i=1}^n \frac{x_i - \bar{x}}{s_x} \frac{y_i - \bar{y}}{s_y},$$

with s_x, s_y denoting the standard deviations and $\bar{x}, \bar{y}$ being the means of the two sequences $\mathbf{x} = (x_1, \dots, x_n)$ and $\mathbf{y} = (y_1, \dots, y_n)$, respectively.

r is the mean of the pairwise products¹ of the standardised versions of two vectors. It is a normalised measure of how they *vary together* (co-variance). Here is how we can compute it manually on two example vectors:

```
x = body[:, 4] # arm circumference
y = body[:, 5] # hip circumference
x_std = (x - np.mean(x)) / np.std(x) # z-scores for x
y_std = (y - np.mean(y)) / np.std(y) # z-scores for y
np.mean(x_std * y_std)
## 0.8680627457873239
```

And here is the built-in function that implements the same formula:

```
scipy.stats.pearsonr(x, y)[0] # the function returns more than we need
## 0.8680627457873241
```

Important The properties of Pearson's r include:

1. $r(\mathbf{x}, \mathbf{y}) = r(\mathbf{y}, \mathbf{x})$ (it is symmetric);
2. $|r(\mathbf{x}, \mathbf{y})| \leq 1$ (it is bounded from below by -1 and from *above* by 1);
3. $r(\mathbf{x}, \mathbf{y}) = 1$ if and only if $\mathbf{y} = a\mathbf{x} + b$ for some $a > 0$ and any b (it reaches the maximum when one variable is an increasing *affine* function of the other one; their graph forms an ascending line);
4. $r(\mathbf{x}, -\mathbf{y}) = -r(\mathbf{x}, \mathbf{y})$ (negative scaling (reflection) of one variable changes its sign);
5. $r(\mathbf{x}, a\mathbf{y} + b) = r(\mathbf{x}, \mathbf{y})$ for any $a > 0$ and b (it is invariant to translation and scaling of inputs, as long as it does not change the sign of elements).

¹ Section 9.3.1 mentions that r is the cosine of the angle between the centred and normalised versions of the vectors.

To help develop some intuitions, let's illustrate a few noteworthy *correlations* using a helper function that draws a scatter plot and prints out Pearson's r (and Spearman's ρ discussed in [Section 9.1.4](#); let's ignore it by then):

```
def plot_corr(x, y, axes_eq=False):
    r = scipy.stats.pearsonr(x, y)[0]
    rho = scipy.stats.spearmanr(x, y)[0]
    plt.plot(x, y, "o")
    plt.title(f"$r = {r:.3}$, $\rho = {rho:.3}$",
             fontdict=dict(fontsize=10))
    if axes_eq: plt.axis("equal")
```

Perfect linear correlation

The aforementioned properties imply that $r(x, y) = -1$ if and only if $y = ax + b$ for some $a < 0$ and any b (reaches the minimum when variables are decreasing affine functions of one another). Furthermore, a variable is trivially perfectly correlated with itself, $r(x, x) = 1$. Consequently, we get perfect *linear correlation* (-1 or 1) when one variable is a scaled and shifted version of the other variable; see [Figure 9.1](#).

```
x = np.random.rand(100)
plt.subplot(1, 2, 1); plot_corr(x, -0.5*x+3, axes_eq=True) # negative slope
plt.subplot(1, 2, 2); plot_corr(x, 3*x+10, axes_eq=True) # positive slope
plt.show()
```




Figure 9.1. Perfect linear correlation (negative and positive).

Important A negative correlation means that when one variable increases, the other one decreases, and vice versa (like in: weight vs expected life span).

Strong linear correlation

Next, if two variables are *approximately* affine functions of one another, the correlations will be close to -1 or 1. The degree of association goes towards 0 as the linear relationship becomes less articulated; see Figure 9.2.

```
plt.figure(figsize=plt.rcParams["figure.figsize"][0], )*2) # width=height
x = np.random.rand(100) # random x (whatever)
y = 0.5*x # y is a linear function of x
e = np.random.randn(len(x)) # random Gaussian noise (expected value 0)
plt.subplot(2, 2, 1); plot_corr(x, y) # original y
plt.subplot(2, 2, 2); plot_corr(x, y+0.05*e) # some noise added to y
plt.subplot(2, 2, 3); plot_corr(x, y+0.1*e) # more noise
plt.subplot(2, 2, 4); plot_corr(x, y+0.25*e) # even more noise
plt.show()
```

Recall that the arm and hip circumferences enjoy high-ish positive degree of linear correlation ($r \simeq 0.868$). Their scatter plot (Figure 7.4) looks somewhat similar to one of the cases depicted here.

Exercise 9.1 Draw a series of similar plots but for the case of negatively correlated point pairs, e.g., $y = -2x + 5$.

Important As a rule of thumb, a linear correlation coefficient of 0.9 or more (or -0.9 or less) is fairly strong. Closer to -0.8 and 0.8, we should already start being sceptical about two variables' possibly being linearly correlated. Some statisticians are more lenient; in particular, it is not uncommon in the social sciences to consider 0.6 a decent degree of correlation, but this is like building your evidence-base on sand. No wonder why some fields are suffering from a reproducibility crisis these days (albeit there are beautiful exceptions to this general observation).

If the dataset at hand does not give us too strong an evidence, it is our ethical duty to refrain from making unjustified claims. We must not mislead the recipients of our data analysis exercises. Still, it can sometimes be interesting to discover that some factors popularly conceived as dependent are actually not correlated, for this is an instance of myth-busting.




Figure 9.2. Linear correlation coefficients for data on a line, but with different amounts of noise.

No linear correlation does not imply independence

For two *independent* variables, we expect the correlation coefficient to be approximately equal to 0. Nevertheless, correlations close to 0 do not necessarily mean that two variables are unrelated. Pearson's r is a *linear* correlation coefficient, so we are quantifying only² these types of relationships; see Figure 9.3 for an illustration.

```
plt.figure(figsize=(plt.rcParams["figure.figsize"][0], )*2) # width=height
plt.subplot(2, 2, 1)
plot_corr(x, np.random.rand(100)) # independent (not correlated)
```

(continues on next page)

² Note that in Section 6.2.3, we were also testing *one* very specific hypothesis: whether a distribution was normal, or whether it was anything else. We only know that if the data really follow that particular distribution, the null hypothesis will not be rejected in 0.1% of the cases. The rest is silence.

(continued from previous page)

```
plt.subplot(2, 2, 2)
plot_corr(x, (2*x-1)**2-1)      # quadratic dependence
plt.subplot(2, 2, 3)
plot_corr(x, np.abs(2*x-1))     # absolute value
plt.subplot(2, 2, 4)
plot_corr(x, np.sin(10*np.pi*x)) # sine
plt.show()
```




Figure 9.3. Are all of the variable pairs really uncorrelated?

False correlations

What is more, sometimes we can fall into the trap of *false* correlations. This happens when data are actually functionally dependent, the relationship is not affine, but the points are aligned not far from a line; see [Figure 9.4](#) for some examples.


```
plt.figure(figsize=(plt.rcParams["figure.figsize"][0], )*2) # width=height
plt.subplot(2, 2, 1)
plot_corr(x, np.sin(0.6*np.pi*x)) # sine
plt.subplot(2, 2, 2)
plot_corr(x, np.log(x+1)) # logarithm
plt.subplot(2, 2, 3);
plot_corr(x, np.exp(x**2)) # exponential of square
plt.subplot(2, 2, 4)
plot_corr(x, 1/(x/2+0.2)) # reciprocal
plt.show()
```




Figure 9.4. Example nonlinear relationships that look like linear, at least to Pearson's r .

A single measure cannot be perfect: we are trying to compress n data points into a single number here. It is obvious that many different datasets, sometimes remarkably diverse, will yield the same correlation degree.

Furthermore, even truly independent variables might return significant Pearson's r . Especially in datasets with many columns and very few observations, it will be possible to find paradoxical pairs of variables that our naïve tool will report as correlated. Tyler Vigen's *Spurious Correlations*³ project documents numerous hilarious instances of this kind, such as Divorce rates in the United Kingdom vs Disney movies released ($r = 0.93$) or Number of Google searches for *Why do I have green poop* vs Pirate attacks globally ($r = 0.853$). For every (small) feature, there is something that correlates with it.

Exercise 9.2 (*) Two independent, normally distributed samples, each of length n , yield Pearson's r which is approximately normally distributed with expectation 0 and standard deviation $1/\sqrt{n-3}$. Compute the probability of drawing an independent pair of samples that yields $|r| > 0.6$ for $n = 10, 25$, and 100.

Correlation is not causation

A high correlation degree (either positive or negative) does not mean that there is any *causal* relationship between the two variables. We cannot say that having large arm circumference affects hip size or vice versa. There might be some *latent* variable that influences these two (e.g., maybe also related to weight?).

Exercise 9.3 Quite often, medical advice is formulated based on correlations and similar association-measuring tools. We are expected to know how to interpret them, as it is never a true cause-effect relationship; rather, it is all about detecting common patterns in larger populations. For instance, in “obesity increases the likelihood of lower back pain and diabetes” we do not say that one necessarily implies another or that if you are not overweight, there is no risk of getting the two said conditions. It might also work the other way around, as lower back pain may lead to less exercise and then weight gain. Reality is complex. Find similar patterns in sets of health conditions.

Note Correlation analysis can aid in constructing regression models, where we would like to identify a transformation that expresses a variable as a function of one or more other features. For instance, when we say that y can be modelled approximately by $ax + b$, regression analysis can identify the best matching a and b coefficients; see [Section 9.2.3](#) for more details.

9.1.2 Correlation heat map

Calling `numpy.corrcoef(body, rowvar=False)` determines the linear correlation coefficients between all pairs of variables in our dataset. We can depict them nicely on a heat map based on a fancified call to `matplotlib.pyplot.imshow`.

```
def corrheatmap(R, labels):
    """
```

(continues on next page)

³ <https://tylervigen.com/spurious-correlations>

(continued from previous page)

```

Draws a correlation heat map, given:
* R - matrix of correlation coefficients for all variable pairs,
* labels - list of column names
"""
assert R.shape[0] == R.shape[1] and R.shape[0] == len(labels)
k = R.shape[0]

# plot the heat map using a custom colour palette
# (correlations are in [-1, 1])
plt.imshow(R, cmap=plt.colormaps.get_cmap("RdBu"), vmin=-1, vmax=1)

# add text labels
for i in range(k):
    for j in range(k):
        plt.text(i, j, f"{R[i, j]:.2f}", ha="center", va="center",
                 color="black" if np.abs(R[i, j]) < 0.5 else "white")

plt.xticks(np.arange(k), labels=labels, rotation=30)
plt.tick_params(axis="x", which="both",
                labelbottom=True, labeltop=True, bottom=False, top=False)

plt.yticks(np.arange(k), labels=labels)
plt.tick_params(axis="y", which="both",
                labelleft=True, labelright=True, left=False, right=False)

plt.grid(False)

```

See Figure 9.5 for the plot.

```

plt.figure(figsize=(plt.rcParams["figure.figsize"][0], )*2) # width=height
R = np.corrcoef(body, rowvar=False)
order = [4, 5, 6, 0, 2, 1, 3] # chiefly for aesthetics
corrheatmap(R[np.ix_(order, order)], body_columns[order])
plt.show()

```

Notice that we ordered⁴ the columns to reveal some naturally occurring variable clusters: for instance, arm, hip, waist circumference, and weight are all strongly correlated.

Of course, we have 1.0s on the main diagonal because a variable is trivially correlated with itself. This heat map is symmetric which is due to the property $r(\mathbf{x}, \mathbf{y}) = r(\mathbf{y}, \mathbf{x})$.

Example 9.4 (*) To fetch the row and column index of the most correlated pair of variables (either positively or negatively), we should first take the upper (or lower) triangle of the correlation matrix (see `numpy.triu` or `numpy.tril`) to ignore the irrelevant and repeating items:

⁴ (**) This can be done automatically by some hierarchical clustering algorithm applied onto the correlation matrix converted to a distance one, e.g., $1 - |\mathbf{R}|$ or $1 - \mathbf{R}^2$.




Figure 9.5. A correlation heat map.

```
Ru = np.triu(np.abs(R), 1)
np.round(Ru, 2)
## array([[0. , 0.35, 0.55, 0.19, 0.91, 0.95, 0.9 ],
##        [0. , 0. , 0.67, 0.66, 0.15, 0.2 , 0.13],
##        [0. , 0. , 0. , 0.48, 0.45, 0.46, 0.43],
##        [0. , 0. , 0. , 0. , 0.08, 0.1 , 0.03],
##        [0. , 0. , 0. , 0. , 0. , 0.87, 0.85],
##        [0. , 0. , 0. , 0. , 0. , 0. , 0.9 ],
##        [0. , 0. , 0. , 0. , 0. , 0. , 0. ]])
```

and then find the location of the maximum:

```
pos = np.unravel_index(np.argmax(Ru), Ru.shape)
pos # (row, column)
```

(continues on next page)

(continued from previous page)

```
## (0, 5)
body_columns[ list(pos) ] # indexing by a tuple has a different meaning
## array(['weight', 'hip circ.'], dtype='<U11')
```

Weight and hip circumference is the most strongly correlated pair.

Note that `numpy.argmax` returns an index in the flattened (unidimensional) array. We had to use `numpy.unravel_index` to convert it to a two-dimensional one.

Example 9.5 (*) Use `seaborn.heatmap` to draw the correlation heat map.

9.1.3 Linear correlation coefficients on transformed data

Pearson's coefficient can also be applied on nonlinearly transformed versions of variables, e.g., logarithms (remember incomes?), squares, square roots, etc.

Let's consider an excerpt from the 2020 CIA [World Factbook](https://www.cia.gov/the-world-factbook)⁵, where we have country-level data on gross domestic product per capita (based on purchasing power parity) and life expectancy at birth.

```
world = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/world_factbook_2020_subset1.csv",
    delimiter=",")[1:, :] # skip the first row (column names)
world[:6, :] # preview
## array([[ 2000. ,   52.8],
##        [12500. ,   79. ],
##        [15200. ,   77.5],
##        [11200. ,   74.8],
##        [49900. ,   83. ],
##        [ 6800. ,   61.3]])
```

Figure 9.6 depicts these data on a scatter plot.

```
plt.subplot(1, 2, 1)
plot_corr(world[:, 0], world[:, 1])
plt.xlabel("per capita GDP PPP")
plt.ylabel("life expectancy (years)")
plt.subplot(1, 2, 2)
plot_corr(np.log(world[:, 0]), world[:, 1])
plt.xlabel("log(per capita GDP PPP)")
plt.yticks()
plt.show()
```

If we compute Pearson's r between these two, we will note that the degree of linear correlation is rather small:

⁵ <https://www.cia.gov/the-world-factbook>




Figure 9.6. Scatter plots for life expectancy vs gross domestic product (purchasing power parity) on linear (left) and log-scale (right).

```
scipy.stats.pearsonr(world[:, 0], world[:, 1])[0]
## 0.656471945486374
```

However, already the logarithm of GDP is more strongly linearly correlated with life expectancy:

```
scipy.stats.pearsonr(np.log(world[:, 0]), world[:, 1])[0]
## 0.8066505089380016
```

which means that modelling our data via $y = a \log x + b$ could be an idea worth considering.

9.1.4 Spearman rank correlation coefficient

Sometimes we might be keen on measuring the degree of any kind of *monotonic* correlation – to what extent one variable is an increasing or decreasing function of another one (linear, logarithmic, quadratic over the positive domain, etc.). In such a scenario, the Spearman *rank correlation* coefficient is frequently used:

$$\rho(\mathbf{x}, \mathbf{y}) = r(R(\mathbf{x}), R(\mathbf{y})),$$

which is⁶ the Pearson linear coefficient computed over vectors of the corresponding ranks of all the elements in $\mathbf{x}$ and $\mathbf{y}$ (denoted by $R(\mathbf{x})$ and $R(\mathbf{y})$, respectively). Hence, the two following calls are equivalent:

⁶ If a method Y is nothing else than X on transformed data, we do not consider it a totally new method.

```

scipy.stats.spearmanr(world[:, 0], world[:, 1])[0]
## 0.8275220380818622
scipy.stats.pearsonr(
    scipy.stats.rankdata(world[:, 0]),
    scipy.stats.rankdata(world[:, 1])
)[0]
## 0.8275220380818619

```

Let's point out that this measure is invariant to monotone transformations of the input variables (up to the sign). This is because they do not change the observations' ranks (or only reverse them).

```

scipy.stats.spearmanr(np.log(world[:, 0]), -np.sqrt(world[:, 1]))[0]
## -0.8275220380818622

```

Exercise 9.6 We included the ρ s in all the outputs generated by our `plot_corr` function. Review all the preceding figures.

Exercise 9.7 Apply `numpy.corrcoef` and `scipy.stats.rankdata` (with the appropriate `axis` argument) to compute the Spearman correlation matrix for all the variable pairs in `body`. Draw it on a heat map.

Exercise 9.8 (*) Draw the scatter plots of the ranks of each column in the `world` and `body` data-sets.

9.2 Regression tasks (*)

Assume we are given a *training/reference* set of n points in an m -dimensional space represented as a matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$ and a set of n corresponding numeric outcomes $\mathbf{y} \in \mathbb{R}^n$. Regression aims to find a function between the m *independent/explanatory/predictor* variables and a chosen *dependent/response/predicted* variable that can be applied on any test point $\mathbf{x}' \in \mathbb{R}^m$:

$$\hat{y}' = f(x'_1, x'_2, \dots, x'_m),$$

and which *approximates* the reference outcomes in a *usable* way.

9.2.1 K-nearest neighbour regression (*)

A straightforward approach to regression relies on aggregating the reference outputs that are associated with a few nearest neighbours of the point $\mathbf{x}'$ tested; compare [Section 8.4.4](#).

In k -nearest neighbour regression, for a fixed $k \geq 1$ and any given $\mathbf{x}' \in \mathbb{R}^m$, $\hat{y} = f(\mathbf{x}')$ is computed as follows.

1. Find the indices $N_k(\mathbf{x}') = \{i_1, \dots, i_k\}$ of the k points from $\mathbf{X}$ closest to $\mathbf{x}'$, i.e., ones that fulfil for all $j \notin \{i_1, \dots, i_k\}$:

$$\|\mathbf{x}_{i_1} - \mathbf{x}'\| \leq \dots \leq \|\mathbf{x}_{i_k} - \mathbf{x}'\| \leq \|\mathbf{x}_{j_1} - \mathbf{x}'\|.$$

2. Return the arithmetic mean of $(y_{i_1}, \dots, y_{i_k})$ as the result.

Here is a straightforward implementation that generates the predictions for each point in X_{test} :

```
def knn_regress(X_test, X_train, y_train, k):
    t = scipy.spatial.KDTree(X_train.reshape(-1, 1))
    i = t.query(X_test.reshape(-1, 1), k)[1] # indices of NNs
    y_nn_pred = y_train[i] # corresponding reference outputs
    return np.mean(y_nn_pred, axis=1)
```

For example, let's express weight (the first column) as a function of hip circumference (the sixth column) in the body dataset:

$$\text{weight} = f_1(\text{hip circumference}) \quad (+\text{some error}).$$

We can also model the life expectancy at birth in different countries (world dataset) as a function of their GDP per capita (PPP):

$$\text{life expectancy} = f_2(\text{GDP per capita}) \quad (+\text{some error}).$$

Both are instances of the *simple* regression problem, i.e., where there is only one independent variable ($m = 1$). We can easily create its appealing visualisation by means of the following function:

```
def knn_regress_plot(x, y, K, num_test_points=1001):
    """
    x - 1D vector - reference inputs
    y - 1D vector - corresponding outputs
    K - numbers of near neighbours to test
    num_test_points - number of points to test at
    """
    plt.plot(x, y, "o", alpha=0.1)
    _x = np.linspace(x.min(), x.max(), num_test_points)
    for k in K:
        _y = knn_regress(_x, x, y, k) # see above
        plt.plot(_x, _y, label=f"$k={k}$")
    plt.legend()
```

Figure 9.7 depicts the fitted functions for a few different k s.

```
plt.subplot(1, 2, 1)
knn_regress_plot(body[:, 5], body[:, 0], [5, 25, 100])
plt.xlabel("hip circumference")
```

(continues on next page)

(continued from previous page)

```
plt.ylabel("weight")

plt.subplot(1, 2, 2)
knn_regress_plot(world[:, 0], world[:, 1], [5, 25, 100])
plt.xlabel("per capita GDP PPP")
plt.ylabel("life expectancy (years)")

plt.show()
```




Figure 9.7. k -nearest neighbour regression curves for example datasets. The greater the k , the more coarse-grained the approximation.

We obtained a *smoothened* version of the original dataset. The fact that we do not reproduce the reference data points in an exact manner is reflected by the (figurative) error term in the above equations. Its role is to emphasise the existence of some natural data variability; after all, one's weight is not purely determined by their hip size and life is not all about money.

For small k we adapt to the data points more closely. This can be worthwhile unless data are very noisy. The greater the k , the smoother the approximation at the cost of losing fine detail and restricted usability at the domain boundaries (here: in the left and right part of the plots).

Usually, the number of neighbours is chosen by trial and error (just like the number of bins in a histogram; compare [Section 4.3.3](#)).

Note (**) Some methods use weighted arithmetic means for aggregating the k reference outputs, with weights inversely proportional to the distances to the neighbours (closer inputs are considered more important).

Also, instead of few nearest neighbours, we can easily compose some form of fixed-radius search regression, by simply replacing $N_k(\mathbf{x}')$ with $B_r(\mathbf{x}')$; compare [Section 8.4.4](#). Yet, note that this way we might make the function undefined in sparsely populated regions of the domain.

9.2.2 From data to (linear) models (*)

Unfortunately, to generate predictions for new data points, k -nearest neighbours regression requires that the training sample is available at all times. It does not *synthesise* or *simplify* the inputs; instead, it works as a kind of a black box. If we were to provide a mathematical equation for the generated prediction, it would be disgustingly long and obscure.

In such cases, to emphasise that f is dependent on the training sample, we sometimes use the more explicit notation $f(\mathbf{x}'|\mathbf{X}, \mathbf{y})$ or $f_{\mathbf{X}, \mathbf{y}}(\mathbf{x}')$.

In many contexts we might prefer creating a data *model* instead, in the form of an easily interpretable mathematical function. A simple yet still flexible choice tackles regression problems via affine maps of the form:

$$y = f(x_1, x_2, \dots, x_m) = c_1 x_1 + c_2 x_2 + \dots + c_m x_m + c_{m+1},$$

or, in matrix multiplication terms:

$$y = \mathbf{c}\mathbf{x}^T + c_{m+1},$$

where $\mathbf{c} = [c_1 \ c_2 \ \dots \ c_m]$ and $\mathbf{x} = [x_1 \ x_2 \ \dots \ x_m]$.

For $m = 1$, the above simply defines a straight line, which we traditionally denote by:

$$y = ax + b,$$

i.e., where we mapped $x_1 \mapsto x$, $c_1 \mapsto a$ (slope), and $c_2 \mapsto b$ (intercept).

For $m > 1$, we obtain different hyperplanes (high-dimensional generalisations of the notion of a plane).

Important A separate intercept term “ $+c_{m+1}$ ” in the defining equation can be cumbersome. We will thus restrict ourselves to linear maps like:

$$y = \mathbf{c}\mathbf{x}^T,$$

but where we can possibly have an explicit constant-1 component somewhere *inside* $\mathbf{x}$. For instance:

$$\mathbf{x} = [x_1 \ x_2 \ \dots \ x_m \ 1].$$

Together with $\mathbf{c} = [c_1 \ c_2 \ \dots \ c_m \ c_{m+1}]$, as trivially $c_{m+1} \cdot 1 = c_{m+1}$, this new setting is equivalent to the original one.

Without loss of generality, from now on we assume that $\mathbf{x}$ is m -dimensional, regardless of its having a constant-1 inside or not.

9.2.3 Least squares method (*)

A linear model is uniquely⁷ encoded using only the coefficients $c_1, \dots, c_m$. To find them, for each point $\mathbf{x}_{i,\cdot}$ from the input (training) set, we typically desire the *predicted* value:

$$\hat{y}_i = f(x_{i,1}, x_{i,2}, \dots, x_{i,m}) = f(\mathbf{x}_{i,\cdot} | \mathbf{c}) = \mathbf{c} \mathbf{x}_{i,\cdot}^T,$$

to be as *close* to the corresponding reference y_i as possible.

There are many measures of *closeness*, but the most popular one⁸ uses the notion of the *sum of squared residuals* (true minus predicted outputs):

$$\text{SSR}(\mathbf{c} | \mathbf{X}, \mathbf{y}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - (c_1 x_{i,1} + c_2 x_{i,2} + \dots + c_m x_{i,m}))^2,$$

which is a function of $\mathbf{c} = (c_1, \dots, c_m)$ (for fixed $\mathbf{X}, \mathbf{y}$).

The *least squares* solution to the stated linear regression problem will be defined by the coefficient vector $\mathbf{c}$ that minimises the SSR. Based on what we said about matrix multiplication, this is equivalent to solving the optimisation task:

$$\text{minimise } (\mathbf{y} - \mathbf{c} \mathbf{X}^T) (\mathbf{y} - \mathbf{c} \mathbf{X}^T)^T \quad \text{w.r.t. } (c_1, \dots, c_m) \in \mathbb{R}^m,$$

because $\hat{\mathbf{y}} = \mathbf{c} \mathbf{X}^T$ gives the predicted values as a row vector (the diligent readers are encouraged to check that on a piece of paper now), $\mathbf{r} = \mathbf{y} - \hat{\mathbf{y}}$ computes all the n residuals, and $\mathbf{r} \mathbf{r}^T$ gives their sum of squares.

The method of least squares is one of the simplest and most natural approaches to regression analysis (curve fitting). Its theoretical foundations (calculus...) were developed more than 200 years ago by Gauss and then were polished by Legendre.

Note (*) Had the points lain on a hyperplane exactly (the interpolation problem), $\mathbf{y} = \mathbf{c} \mathbf{X}^T$ would have an exact solution, equivalent to solving the linear system of equations $\mathbf{y} - \mathbf{c} \mathbf{X}^T = \mathbf{0}$. However, in our setting we assume that there might be some measurement errors or other discrepancies between the reality and the theoretical model. To account for this, we are trying to solve a more general problem of finding a hyperplane for which $\|\mathbf{y} - \mathbf{c} \mathbf{X}^T\|^2$ is as small as possible.

This optimisation task can be solved analytically (compute the partial derivatives of SSR with respect to each $c_1, \dots, c_m$, equate them to 0, and solve a simple system of

⁷ To memorise the model for further reference, we only need to serialise its m coefficients, e.g., in a JSON or CSV file.

⁸ Due to computability and mathematical analysability, which we usually explore in more advanced courses on statistical data analysis such as [10, 24, 50].

linear equations). This spawns $\mathbf{c} = \mathbf{yX}(\mathbf{X}^T\mathbf{X})^{-1}$, where $\mathbf{A}^{-1}$ is the inverse of a matrix $\mathbf{A}$, i.e., the matrix such that $\mathbf{AA}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$; compare `numpy.linalg.inv`. As inverting larger matrices directly is not too robust numerically, we would rather rely on a more specialised algorithm.

The undermentioned `scipy.linalg.lstsq` function provides a fairly numerically stable (yet, see [Section 9.2.9](#)) procedure that is based on the singular value decomposition of the model matrix.

Let's go back to the NHANES study excerpt and express weight (the first column) as function of hip circumference (the sixth column) again, but this time using an affine map of the form⁹:

$$\text{weight} = a \cdot \text{hip circumference} + b \quad (+\text{some error}).$$

The *design (model) matrix* $\mathbf{X}$ and the reference *ys* are:

```
x_original = body[:, [5]]      # a column vector
X_train = x_original**[1, 0]    # hip circumference, 1s
y_train = body[:, 0]           # weight
```

We used the vectorised exponentiation operator to convert each x_i (the i -th hip circumference) to a pair $\mathbf{x}_{i\cdot} = (x_i^1, x_i^0) = (x_i, 1)$, which is a nice trick to append a column of 1s to a matrix. This way, we included the intercept term in the model (as discussed in [Section 9.2.2](#)). Here is a preview:

```
preview_indices = [4, 5, 6, 8, 12, 13]
X_train[preview_indices, :]
## array([[ 92.5,  1. ],
##        [106.7,  1. ],
##        [ 96.3,  1. ],
##        [102. ,  1. ],
##        [ 94.8,  1. ],
##        [ 97.5,  1. ]])
y_train[preview_indices]
## array([55.4, 62. , 66.2, 77.2, 64.2, 56.8])
```

Let's determine the least squares solution to our regression problem:

```
import scipy.linalg
res = scipy.linalg.lstsq(X_train, y_train)
```

That's it. The optimal coefficients vector (the one that minimises the SSR) is:

⁹ We sometimes explicitly list the error term that corresponds to the residuals. This is to assure the reader that we are not naïve and that we know what we are doing. We see from the scatter plot of the involved variables that the data do not lie on a straight line perfectly. Each model is merely an idealisation/simplification of the described reality. It is wise to remind ourselves about that every so often.

```
c = res[0]
c
## array([ 1.3052463 , -65.10087248])
```

The estimated model is:

$$\text{weight} = 1.305 \cdot \text{hip circumference} - 65.1 \quad (+\text{some error}).$$

The model is nicely interpretable. For instance, as hip circumference increases, we expect the weights to be greater and greater. As we said before, it does not mean that there is some *causal* relationship between the two (for instance, there can be some latent variables that affect both of them). Instead, there is some general tendency regarding how the data align in the sample space. For instance, that the “best guess” (according to the current model – there can be many; see below) weight for a person with hip circumference of 100 cm is 65.4 kg. Thanks to such models, we might get more insight into certain phenomena, or find some proxies for different variables (especially if measuring them directly is tedious, costly, dangerous, etc.).

Important What we have performed above is an instance of *machine learning*. The machine (the equation $y = ax + b$) is controlled by two parameters (memory; a, b). We have *learnt* their best values based on the knowledge about the world we had (the inputs x and the desired outputs y). Thus, this term is really an obscure marketing slogan designed to impress (or confuse) the laypeople.

Let’s determine the predicted weights for all of the participants:

```
y_pred = c @ X_train.T
np.round(y_pred[preview_indices], 2) # preview
## array([55.63, 74.17, 60.59, 68.03, 58.64, 62.16])
```

The scatter plot and the fitted regression line in Figure 9.8 indicates a fair fit but, of course, there is some natural variability.

```
plt.plot(x_original, y_train, "o", alpha=0.1) # scatter plot
_x = np.array([x_original.min(), x_original.max()]).reshape(-1, 1)
_y = c @ (_x**[1, 0]).T
plt.plot(_x, _y, "r-") # a line that goes through the two extreme points
plt.xlabel("hip circumference")
plt.ylabel("weight")
plt.show()
```

Exercise 9.9 The *Anscombe quartet*¹⁰ is a famous example dataset, where we have four pairs of variables that have almost identical means, variances, and linear correlation coefficients. Even though they can be approximated by the same straight line, their scatter plots are vastly different. Reflect upon this toy example.

¹⁰ <https://github.com/gagolews/teaching-data/raw/master/t/anscombe.csv>




Figure 9.8. The least squares line for weight vs hip circumference.

9.2.4 Analysis of residuals (*)

The residuals (i.e., the estimation errors – what we expected vs what we got), for the chosen six observations are visualised in Figure 9.9.

```
r = y_train - y_pred # residuals
np.round(r[preview_indices], 2) # preview
## array([-0.23, -12.17,  5.61,  9.17,  5.56, -5.36])
```

We wanted the squared residuals (on average – across all the points) to be as small as possible. The least squares method assures that this is the case *relative to the chosen model*, i.e., a linear one. Nonetheless, it still does not mean that what we obtained constitutes a good fit to the training data. Thus, we need to perform the *analysis of residuals*.

Interestingly, the average of residuals is always zero:

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i) = 0.$$

Therefore, if we want to summarise the residuals into a single number, we can use, for example, the root mean squared error instead:

$$\text{RMSE}(\mathbf{y}, \hat{\mathbf{y}}) = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}.$$

```
np.sqrt(np.mean(r**2))
## 6.948470091176111
```




Figure 9.9. Example residuals in a simple linear regression task.

Hopefully we can see that RMSE is a function of SSR that we have already sought to minimise.

Alternatively, we can compute the mean absolute error:

$$\text{MAE}(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|.$$

```
np.mean(np.abs(r))
## 5.207073583769202
```

MAE is nicely interpretable: it measures by how many kilograms we err *on average*. Not bad.

Exercise 9.10 Fit a regression line explaining weight as a function of the waist circumference and compute the corresponding RMSE and MAE. Are they better than when hip circumference is used as an explanatory variable?

Note Generally, fitting simple (involving one independent variable) linear models can only make sense for highly linearly correlated variables. Interestingly, if $\mathbf{y}$ and $\mathbf{x}$ are both standardised, and r is their Pearson's coefficient, then the least squares solution is given by $y = rx$.

To verify whether a fitted model is not extremely wrong (e.g., when we fit a linear model to data that clearly follows a different functional relationship), a plot of residuals against the fitted values can be of help; see Figure 9.10. Ideally, the points are

expected to be aligned totally at random therein, without any dependence structure (homoscedasticity).

```
plt.plot(y_pred, r, "o", alpha=0.1)
plt.axhline(0, ls="--", color="red") # horizontal line at y=0
plt.xlabel("fitted values")
plt.ylabel("residuals")
plt.show()
```




Figure 9.10. Residuals vs fitted values for the linear model explaining weight as a function of hip circumference. The variance of residuals slightly increases as $\hat{y}_i$ increases. This is not ideal, but it could be much worse than this.

Exercise 9.11 Compare¹¹ the RMSE and MAE for the k -nearest neighbour regression curves depicted in the left side of Figure 9.7. Also, draw the residuals vs fitted plot.

For linear models fitted using the least squares method, we have:

$$\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2 = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2.$$

In other words, the variance of the dependent variable (left) can be decomposed into the sum of the variance of the predictions and the averaged squared residuals. Multiplying it by n , we have that the *total* sum of squares is equal to the *explained* sum of

¹¹ In k -nearest neighbour regression, we are not aiming to minimise anything in particular. If the model is performing well with respect to some metrics such as RMSE or MAE, we can consider ourselves lucky. Nevertheless, some asymptotic results guarantee the optimality of the outcomes generated for large sample sizes (e.g., consistency); see, e.g., [24].

squares plus the *residual* sum of squares:

$$\text{TSS} = \text{ESS} + \text{RSS}.$$

We yearn for ESS to be as close to TSS as possible. Equivalently, it would be jolly nice to have RSS equal to 0.

The *coefficient of determination* (unadjusted R-Squared, sometimes referred to as simply the *score*) is a popular normalised, unitless measure that is easier to interpret than raw ESS or RSS when we have no domain-specific knowledge of the modelled problem. It is given by:

$$R^2(\mathbf{y}, \hat{\mathbf{y}}) = \frac{\text{ESS}}{\text{TSS}} = 1 - \frac{\text{RSS}}{\text{TSS}} = 1 - \frac{s_r^2}{s_y^2}.$$

```
1 - np.var(y_train-y_pred)/np.var(y_train)
## 0.8959634726270759
```

The coefficient of determination in the current context¹² is thus the proportion of variance of the dependent variable explained by the independent variables in the model. The closer it is to 1, the better. A dummy model that always returns the mean of $\mathbf{y}$ gives R-squared of 0.

In our case, $R^2 \simeq 0.9$ is high, which indicates a rather good fit.

Note (*) There are certain statistical results that can be relied upon provided that the residuals are independent random variables with expectation zero and the same variance (e.g., the Gauss–Markov theorem). Further, if they are normally distributed, then we have several hypothesis tests available (e.g., for the significance of coefficients). This is why in various textbooks such assumptions are additionally verified. But we do not go that far in this introductory course.

9.2.5 Multiple regression (*)

As another example, let's fit a model involving two independent variables: arm and hip circumference:

```
X_train = np.insert(body[:, [4, 5]], 2, 1, axis=1) # append a column of 1s
res = scipy.linalg.lstsq(X_train, y_train)
c = res[0]
np.round(c, 2)
## array([ 1.3 ,  0.9 , -63.38])
```

¹² For a model that is *not* generated via least squares, the coefficient of determination can also be negative, particularly when the fit is extremely bad. Also, note that this measure is dataset-dependent. Therefore, it ought not to be used for comparing models explaining different dependent variables.

We fitted the plane:

$$\text{weight} = 1.3 \text{ arm circumference} + 0.9 \text{ hip circumference} - 63.38.$$

We skip the visualisation part for we do not expect it to result in a readable plot: these are multidimensional data. The coefficient of determination is:

```
y_pred = c @ X_train.T
r = y_train - y_pred
1-np.var(r)/np.var(y_train)
## 0.9243996585518783
```

Root mean squared error:

```
np.sqrt(np.mean(r**2))
## 5.923223870044695
```

Mean absolute error:

```
np.mean(np.abs(r))
## 4.431548244333892
```

It is a slightly better model than the previous one. We can predict the participants' weights more accurately, at the cost of an increased model's complexity.

9.2.6 Variable transformation and linearisable models (**)

We are not restricted merely to linear functions of the input variables. By applying arbitrary transformations upon the columns of the design matrix, we can cover many diverse scenarios.

For instance, a polynomial model involving two variables:

$$g(v_1, v_2) = \beta_0 + \beta_1 v_1 + \beta_2 v_1^2 + \beta_3 v_1 v_2 + \beta_4 v_2 + \beta_5 v_2^2,$$

can be obtained by substituting $x_1 = 1$, $x_2 = v_1$, $x_3 = v_1^2$, $x_4 = v_1 v_2$, $x_5 = v_2$, $x_6 = v_2^2$, and then fitting a linear model involving six variables:

$$f(x_1, x_2, \dots, x_6) = c_1 x_1 + c_2 x_2 + \dots + c_6 x_6.$$

The design matrix is made of rubber, it can handle almost anything. If we have a linear model, but with respect to transformed data, the algorithm does not care. This is the beauty of the underlying mathematics; see also [12].

A creative modeller can also turn models such as $u = ce^{av}$ into $y = ax + b$ by replacing $y = \log u$, $x = v$, and $b = \log c$. There are numerous possibilities based on the properties of the log and exp functions listed in [Section 5.2](#). We call them *linearisable models*.

As an example, let's model the life expectancy at birth in different countries as a function of their GDP per capita (PPP).

We will consider four different models:

1. $y = c_1 + c_2x$ (linear),
2. $y = c_1 + c_2x + c_3x^2$ (quadratic),
3. $y = c_1 + c_2x + c_3x^2 + c_4x^3$ (cubic),
4. $y = c_1 + c_2 \log x$ (logarithmic).

Here are the helper functions that create the model matrices:

```
def make_model_matrix1(x):
    return x.reshape(-1, 1)**[0, 1]

def make_model_matrix2(x):
    return x.reshape(-1, 1)**[0, 1, 2]

def make_model_matrix3(x):
    return x.reshape(-1, 1)**[0, 1, 2, 3]

def make_model_matrix4(x):
    return (np.log(x)).reshape(-1, 1)**[0, 1]

make_model_matrix1.__name__ = "linear model"
make_model_matrix2.__name__ = "quadratic model"
make_model_matrix3.__name__ = "cubic model"
make_model_matrix4.__name__ = "logarithmic model"

model_matrix_makers = [
    make_model_matrix1,
    make_model_matrix2,
    make_model_matrix3,
    make_model_matrix4
]
x_original = world[:, 0]
Xs_train = [ make_model_matrix(x_original)
              for make_model_matrix in model_matrix_makers ]
```

Fitting the models:

```
y_train = world[:, 1]
cs = [ scipy.linalg.lstsq(X_train, y_train)[0]
       for X_train in Xs_train ]
```

Their coefficients of determination are equal to:

```
for i in range(len(Xs_train)):
    R2 = 1 - np.var(y_train - cs[i] @ Xs_train[i].T)/np.var(y_train)
    print(f"{model_matrix_makers[i].__name__:20} R2={R2:.3f}")
## linear model           R2=0.431
## quadratic model        R2=0.567
```

(continues on next page)

(continued from previous page)

```
## cubic model          R2=0.607
## logarithmic model    R2=0.651
```

The logarithmic model is thus the best (out of the models we considered). The four models are depicted in Figure 9.11.

```
plt.plot(x_original, y_train, "o", alpha=0.1)
_x = np.linspace(x_original.min(), x_original.max(), 101).reshape(-1, 1)
for i in range(len(model_matrix_makers)):
    _y = cs[i] @ model_matrix_makers[i](_x).T
    plt.plot(_x, _y, label=model_matrix_makers[i].__name__)
plt.legend()
plt.xlabel("per capita GDP PPP")
plt.ylabel("life expectancy (years)")
plt.show()
```



Figure 9.11. Different models for life expectancy vs GDP.

Exercise 9.12 Draw box plots and histograms of residuals for each model as well as the scatter plots of residuals vs fitted values.

9.2.7 Descriptive vs predictive power (**)

We approximated the life vs GDP relationship using a few different functions. Nevertheless, we see that the foregoing quadratic and cubic models possibly do not make much sense, semantically speaking. Sure, as far as individual points in the training set are concerned, they do fit the data more closely than the linear model. After all, they have smaller mean squared errors (again: at these given points). Looking at the way

they behave, one does not need a university degree in economics/social policy to conclude that they are not the best *description* of how the reality behaves (on average).

Important Naturally, a model's goodness of fit to observed data tends to improve as the model's complexity increases. The Razor principle (by William of Ockham et al.) advises that if some phenomenon can be explained in many different ways, the simplest explanation should be chosen (*do not multiply entities* [here: introduce independent variables] *without necessity*).

In particular, the more independent variables we have in the model, the greater the R^2 coefficient will be. We can try correcting for this phenomenon by considering the *adjusted* R^2 :

$$\bar{R}^2(\mathbf{y}, \hat{\mathbf{y}}) = 1 - (1 - R^2(\mathbf{y}, \hat{\mathbf{y}})) \frac{n-1}{n-m-1},$$

which, to some extent, penalises more complex models.

Note (**) Model quality measures adjusted for the number of model parameters, m , can also be useful in automated variable selection. For example, the Akaike Information Criterion is a popular measure given by:

$$\text{AIC} = 2m + n \log(\text{SSR}) - n \log n.$$

Furthermore, the Bayes Information Criterion is defined via:

$$\text{BIC} = m \log n + n \log(\text{SSR}) - n \log n.$$

Unfortunately, they are both dependent on the scale of $\mathbf{y}$.

We should also be concerned with quantifying a model's *predictive* power, i.e., how well does it generalise to data points that we do not have now (or pretend we do not have) but might face in the future. As we observe the modelled reality only at a few different points, the question is how the model performs when filling the gaps between the dots it connects.

In particular, we must be careful when *extrapolating* the data, i.e., making predictions outside of its usual domain. For example, the linear model predicts the following life expectancy for an imaginary country with \$500 000 per capita GDP:

```
cs[0] @ model_matrix_makers[0](np.array([500000])).T
## array([164.3593753])
```

and the quadratic one gives:

```
cs[1] @ model_matrix_makers[1](np.array([500000])).T
## array([-364.10630779])
```

Nonsense.

Example 9.13 Consider a theoretical illustration, where a true model of some reality is $y = 5 + 3x^3$.

```
def true_model(x):
    return 5 + 3*(x**3)
```

Still, for some reason we are only able to gather a small ($n = 25$) sample from this model. What is even worse, it is subject to some measurement error:

```
np.random.seed(42)
x = np.random.rand(25) # random xs on [0, 1]
y = true_model(x) + 0.2*np.random.randn(len(x)) # true_model(x) + noise
```

The least-squares fitting of $y = c_1 + c_2x^3$ to the above gives:

```
X03 = x.reshape(-1, 1)**[0, 3]
c03 = scipy.linalg.lstsq(X03, y)[0]
ssr03 = np.sum((y-c03 @ X03.T)**2)
np.round(c03, 2)
## array([5.01, 3.13])
```

which is not too far, but still somewhat¹³ distant from the true coefficients, 5 and 3.

We can also fit a more flexible cubic polynomial, $y = c_1 + c_2x + c_3x^2 + c_4x_3$:

```
X0123 = x.reshape(-1, 1)**[0, 1, 2, 3]
c0123 = scipy.linalg.lstsq(X0123, y)[0]
ssr0123 = np.sum((y-c0123 @ X0123.T)**2)
np.round(c0123, 2)
## array([4.89, 0.32, 0.57, 2.23])
```

In terms of the SSR, this more complex model of course explains the training data more accurately:

```
ssr03, ssr0123
## (1.0612111154029558, 0.9619488226837537)
```

Yet, it is farther away from the truth (which, whilst performing the fitting task based only on given x and y , is unknown). We may thus say that the first model generalises better on yet-to-be-observed data; see Figure 9.12 for an illustration.

¹³ For large n , we expect to pinpoint the true coefficients exactly. In our scenario (independent, normally distributed errors with the expectation of 0), the least squares method is the maximum likelihood estimator of the model parameters. As a consequence, it is consistent.

```

_x = np.linspace(0, 1, 101)
plt.plot(x, y, "o")
plt.plot(_x, true_model(_x), "--", label="true model")
plt.plot(_x, c0123 @ (_x.reshape(-1, 1)**[0, 1, 2, 3]).T,
        label="fitted model y=x**[0, 1, 2, 3]")
plt.plot(_x, c03 @ (_x.reshape(-1, 1)**[0, 3]).T,
        label="fitted model y=x**[0, 3]")
plt.legend()
plt.show()

```




Figure 9.12. The true (theoretical) model vs some guesstimates (fitted based on noisy data). More degrees of freedom is not always better.

Example 9.14 (**) We defined the sum of squared residuals (and its function, the root mean squared error) by means of the averaged deviation from the reference values. They are subject to error themselves, though. Even though they are our best-shot approximation of the truth, they should be taken with a degree of scepticism.

In the previous example, given the true (reference) model f defined over the domain D (in our case, $f(x) = 5 + 3x^3$ and $D = [0, 1]$) and an empirically fitted model $\hat{f}$, we can compute the square root of the integrated squared error over the whole D :

$$\text{RMSE}(f, \hat{f}) = \sqrt{\int_D (f(x) - \hat{f}(x))^2 dx}.$$

For polynomials and other simple functions, RMSE can be computed analytically. More generally, we can approximate it numerically by sampling the above at sufficiently many points and applying the trapezoidal rule (e.g., [77]). As this can be an educative programming exercise, let's consider a range of polynomial models of different degrees.

```

cs, rmse_train, rmse_test = [], [], [] # result containers
ps = np.arange(1, 10) # polynomial degrees
for p in ps: # for each polynomial degree:
    c = scipy.linalg.lstsq(x.reshape(-1, 1)**np.arange(p+1), y)[0] # fit
    cs.append(c)

    y_pred = c @ (x.reshape(-1, 1)**np.arange(p+1)).T # predictions
    rmse_train.append(np.sqrt(np.mean((y-y_pred)**2))) # RMSE

    _x = np.linspace(0, 1, 101) # many _xs
    _y = c @ (_x.reshape(-1, 1)**np.arange(p+1)).T # f(_x)
    _r = (true_model(_x) - _y)**2 # residuals
    rmse_test.append(np.sqrt(0.5*np.sum(
        np.diff(_x)*(_r[1:]+_r[:-1]) # trapezoidal rule for integration
    )))

plt.plot(ps, rmse_train, label="RMSE (training set)")
plt.plot(ps, rmse_test, label="RMSE (theoretical)")
plt.legend()
plt.yscale("log")
plt.xlabel("model complexity (polynomial degree)")
plt.show()

```




Figure 9.13. Small RMSE on training data does not necessarily imply good generalisation abilities.

Figure 9.13 shows that a model's ability to make correct generalisations onto unseen data improves as the complexity increases, at least initially. However, then it becomes worse. It is a typical behaviour. In fact, the model with the smallest RMSE on the training set, overfits to the input sample, see Figure 9.14.


```
plt.plot(x, y, "o")
plt.plot(_x, true_model(_x), "--", label="true model")
for i in [0, 1, 8]:
    plt.plot(_x, cs[i] @ (_x.reshape(-1, 1)**np.arange(ps[i]+1)).T,
             label=f"fitted degree-{ps[i]} polynomial")
plt.legend()
plt.show()
```




Figure 9.14. Under- and overfitting to training data.

Important When evaluating a model's quality in terms of predictive power on unseen data, we should go beyond inspecting its behaviour merely on the points from the training sample. As the *truth* is usually not known (if it were, we would not need any guessing), a common approach in case where we have a dataset of a considerable size is to divide it (randomly; see [Section 10.5.4](#)) into two parts:

- training sample (say, 60%) – used to fit a model,
- test sample (the remaining 40%) – used to assess its quality (e.g., by means of RMSE).

This might *emulate* an environment where some new data arrives later, see [Section 12.3.3](#) for more details.

Furthermore, if model selection is required, we may apply a training/validation/test split (say, 60/20/20%; see [Section 12.3.4](#)). Here, many models are constructed on the training set, the validation set is used to compute the metrics and choose the best model, and then the test set gives the final model's valuation to assure its usefulness/uselessness (because we do not want it to overfit to the test set).

Overall, models must never be blindly trusted. Common sense must always be applied. The fact that we fitted something using a sophisticated procedure on a dataset that was hard to obtain does not justify its use. Mediocre models must be discarded, and we should move on, regardless of how much time/resources we have invested whilst developing them. Too many bad models go into production and make our daily lives harder. We need to end this madness.

9.2.8 Fitting regression models with `scikit-learn` (*)

`scikit-learn`¹⁴ (`sklearn`; [75]) is a huge Python package built on top of `numpy`, `scipy`, and `matplotlib`. It has a consistent API and implements or provides wrappers for many regression, classification, clustering, and dimensionality reduction algorithms (amongst others).

Important `scikit-learn` is very convenient. Nevertheless, it allows us to fit models even if we do not understand the mathematics behind them. This is dangerous: it is like driving a sports car without the necessary skills while wearing a blindfold. Advanced students and practitioners will appreciate it, but when used by beginners, it needs to be handled with care. We should not mistake something's being easily accessible with its being safe to use. Remember that if we are given a procedure for which we are unable to provide its definition/mathematical properties/explain its idealised version in pseudocode, we are expected to refrain from using it (see Rule#7).

Due to this, we shall only present a quick demo of `scikit-learn`'s API. We will do that by fitting a multiple linear regression model again for the weight as a function of the arm and the hip circumference:

```
X_train = body[:, [4, 5]]
y_train = body[:, 0]
```

In `scikit-learn`, once we construct an object representing the model to be fitted, the `fit` method determines the optimal parameters.

```
import sklearn.linear_model
lm = sklearn.linear_model.LinearRegression(fit_intercept=True)
lm.fit(X_train, y_train)
lm.intercept_, lm.coef_
## (-63.383425410947524, array([1.30457807, 0.8986582 ]))
```

We, of course, obtained the same solution as with `scipy.linalg.lstsq`.

Computing the predicted values can be done via the `predict` method. For example, we can calculate the coefficient of determination:

¹⁴ <https://scikit-learn.org/stable/index.html>

```
y_pred = lm.predict(X_train)
import sklearn.metrics
sklearn.metrics.r2_score(y_train, y_pred)
## 0.9243996585518783
```

This function is convenient, but can we really recall the formula for the score and what it really measures?

9.2.9 Ill-conditioned model matrices (**)

Our approach to regression analysis relies on solving an optimisation problem (the method least squares). Nevertheless, sometimes the “optimal” solution that the algorithm returns might have nothing to do with the *true* minimum. And this is despite the fact that we have the theoretical results stating that the solution is unique¹⁵ (the objective is convex). The problem stems from our using the computer’s finite-precision floating point arithmetic; compare [Section 5.5.6](#).

Let’s fit a degree-4 polynomial to the life expectancy vs per capita GDP dataset.

```
x_original = world[:, 0]
X_train = (x_original.reshape(-1, 1))**[0, 1, 2, 3, 4]
y_train = world[:, 1]
cs = dict()
```

We store the estimated model coefficients in a dictionary because many methods will follow next. First, **scipy**:

```
res = scipy.linalg.lstsq(X_train, y_train)
cs["scipy_X"] = res[0]
cs["scipy_X"]
## array([ 2.33103950e-16,  6.42872371e-12,  1.34162021e-07,
##          -2.33980973e-12,  1.03490968e-17])
```

If we drew the fitted polynomial now (see [Figure 9.15](#)), we would see that the fit is unbelievably bad. The result returned by **scipy.linalg.lstsq** is now not at all optimal. All coefficients are approximately equal to 0.

It turns out that the fitting problem is extremely *ill-conditioned* (and it is not the algorithm’s fault): GDPs range from very small to very large ones. Furthermore, taking them to the fourth power breeds numbers of ever greater range. Finding the least squares solution involves some form of matrix inverse (not necessarily directly) and our model matrix may be close to singular (one that is not invertible).

As a measure of the model matrix’s ill-conditioning, we often use the *condition num-*

¹⁵ There are methods in statistical learning where there might be multiple local minima – this is even more difficult; see [Section 12.4.4](#).

ber, denoted $\kappa(\mathbf{X}^T)$. It is the ratio of the largest to the smallest *singular values*¹⁶ of $\mathbf{X}^T$, which are returned by the `scipy.linalg.lstsq` method itself:

```
s = res[3]      # singular values of X_train.T
s
## array([5.63097211e+20, 7.90771769e+14, 4.48366565e+09, 6.77575417e+04,
##        5.76116463e+00])
```

Note that they are already sorted nonincreasingly. The condition number $\kappa(\mathbf{X}^T)$ is equal to:

```
s[0] / s[-1] # condition number (largest/smallest singular value)
## 9.774017018467431e+19
```

As a rule of thumb, if the condition number is 10^k , we are losing k digits of numerical precision when performing the underlying computations. As the foregoing number is exceptionally large, we are thus currently faced with a very ill-conditioned problem. If the values in $\mathbf{X}$ or $\mathbf{y}$ are perturbed even slightly, we might expect significant changes in the computed regression coefficients.

Note (**) The least squares regression problem can be solved by means of the singular value decomposition of the model matrix, see [Section 9.3.4](#). Let $\mathbf{USQ}$ be the SVD of $\mathbf{X}^T$. Then $\mathbf{c} = \mathbf{US}^{-1}\mathbf{Qy}$, with $\mathbf{S}^{-1} = \text{diag}(1/s_{1,1}, \dots, 1/s_{m,m})$. As $s_{1,1} \geq \dots \geq s_{m,m}$ gives the singular values of $\mathbf{X}^T$, the aforementioned condition number can simply be computed as $s_{1,1}/s_{m,m}$.

Let's verify the method used by `scikit-learn`. As it fits the intercept separately, we expect it to be slightly better-behaving; still, it is merely a wrapper around `scipy.linalg.lstsq`, but with a different API.

```
import sklearn.linear_model
lm = sklearn.linear_model.LinearRegression(fit_intercept=True)
lm.fit(X_train[:, 1:], y_train)
cs["sklearn"] = np.r_[lm.intercept_, lm.coef_]
cs["sklearn"]
## array([ 6.92257708e+01,  5.05752755e-13,  1.38835643e-08,
##        -2.18869346e-13,  9.09347772e-19])
```

Here is the condition number of the underlying model matrix:

```
lm.singular_[0] / lm.singular_[-1]
## 1.402603229842849e+16
```

¹⁶ (**) Being themselves the square roots of eigenvalues of $\mathbf{X}^T\mathbf{X}$. Equivalently, $\kappa(\mathbf{X}^T) = \|(\mathbf{X}^T)^{-1}\| \|\mathbf{X}^T\|$ with respect to the spectral norm. Seriously, we really need to get good grasp of linear algebra to become successful data scientists.

The condition number is also enormous. Still, **scikit-learn** did not warn us about this being the case (insert frowning face emoji here). Had we trusted the solution returned by it, we would end up with conclusions from our data analysis built on sand. As we said in [Section 9.2.8](#), the package designers assumed that the users know what they are doing. This is okay, we are all adults here, although some of us are still learning.

Overall, if the model matrix is close to singular, the computation of its inverse is prone to enormous numerical errors. One way of dealing with this is to remove highly correlated variables (the multicollinearity problem). Interestingly, standardisation can *sometimes* make the fitting more numerically stable.

Let $\mathbf{Z}$ be a standardised version of the model matrix $\mathbf{X}$ with the intercept part (the column of 1s) not included, i.e., with $\mathbf{z}_{:,j} = (\mathbf{x}_{:,j} - \bar{x}_j)/s_j$ where $\bar{x}_j$ and s_j denotes the arithmetic mean and the standard deviation of the j -th column in $\mathbf{X}$. If $(d_1, \dots, d_{m-1})$ is the least squares solution for $\mathbf{Z}$, then the least squares solution to the underlying original regression problem is:

$$\mathbf{c} = \left(\bar{y} - \sum_{j=1}^{m-1} \frac{d_j}{s_j} \bar{x}_j, \frac{d_1}{s_1}, \frac{d_2}{s_2}, \dots, \frac{d_{m-1}}{s_{m-1}} \right),$$

with the first term corresponding to the intercept.

Let's test this approach with `scipy.linalg.lstsq`:

```
means = np.mean(X_train[:, 1:], axis=0)
stds = np.std(X_train[:, 1:], axis=0)
Z_train = (X_train[:, 1:] - means) / stds
resZ = scipy.linalg.lstsq(Z_train, y_train)
c_scipyZ = resZ[0] / stds
cs["scipy_Z"] = np.r_[np.mean(y_train) - (c_scipyZ @ means.T), c_scipyZ]
cs["scipy_Z"]
## array([ 6.35946784e+01,  1.04541932e-03, -2.41992445e-08,
##         2.39133533e-13, -8.13307828e-19])
```

The condition number is:

```
s = resZ[3]
s[0] / s[-1]
## 139.4279225737234
```

It is still far from perfect (we would prefer a value close to 1) but nevertheless it is a significant improvement over what we had before.

[Figure 9.15](#) depicts the three fitted models, each claiming to be *the* solution to the original regression problem. Note that, luckily, we know that in our case the logarithmic model is better than the polynomial one.

```

plt.plot(x_original, y_train, "o", alpha=0.1)
_x = np.linspace(x_original.min(), x_original.max(), 101).reshape(-1, 1)
_X = _x**[0, 1, 2, 3, 4]
for lab, c in cs.items():
    ssr = np.sum((y_train - c @ X_train.T)**2)
    plt.plot(_x, c @ _X.T, label=f"{lab:10} SSR={ssr:.2f}")
plt.legend()
plt.ylim(20, 120)
plt.xlabel("per capita GDP PPP")
plt.ylabel("life expectancy (years)")
plt.show()

```




Figure 9.15. Ill-conditioned model matrix can give a very wrong model.

Important Always check the model matrix's condition number.

Exercise 9.15 Check the condition numbers of all the models fitted so far in this chapter via the least squares method.

To be strict, if we read a paper in, say, social or medical sciences (amongst others) where the researchers fit a regression model but do not provide the model matrix's condition number, it is worthwhile to doubt the conclusions they make.

On a final note, we might wonder why the standardisation is not done automatically by the least squares solver. As usual with most numerical methods, there is no one-fits-all solution: e.g., when there are columns of extremely small variance or there are outliers in data. This is why we need to study all the topics deeply: to be able to respond flexibly to many different scenarios ourselves.

9.3 Finding interesting combinations of variables (*)

9.3.1 Dot products, angles, collinearity, and orthogonality (*)

Let's note that the dot product (Section 8.3) has a nice geometrical interpretation:

$$\mathbf{x} \cdot \mathbf{y} = \|\mathbf{x}\| \|\mathbf{y}\| \cos \alpha,$$

where α is the angle between two given vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$. In plain English, it is the product of the magnitudes of the two vectors and the cosine of the angle between them.

We can retrieve the cosine part by computing the dot product of the *normalised* vectors, i.e., such that their magnitudes are equal to 1:

$$\cos \alpha = \frac{\mathbf{x}}{\|\mathbf{x}\|} \cdot \frac{\mathbf{y}}{\|\mathbf{y}\|}.$$

For example, consider two vectors in $\mathbb{R}^2$, $\mathbf{u} = (1/2, 0)$ and $\mathbf{v} = (\sqrt{2}/2, \sqrt{2}/2)$, which are depicted in Figure 9.16.

```
u = np.array([0.5, 0])
v = np.array([np.sqrt(2)/2, np.sqrt(2)/2])
```

Their dot product is equal to:

```
np.sum(u*v)
## 0.3535533905932738
```

The dot product of their normalised versions, i.e., the cosine of the angle between them is:

```
u_norm = u/np.sqrt(np.sum(u*u))
v_norm = v/np.sqrt(np.sum(v*v)) # BTW: this vector is already normalised
np.sum(u_norm*v_norm)
## 0.7071067811865476
```

The angle itself can be determined by referring to the inverse of the cosine function, i.e., arccosine.

```
np.arccos(np.sum(u_norm*v_norm)) * 180/np.pi
## 45.0
```

Notice that we converted the angle from radians to degrees.

Important If two vectors are collinear (*codirectional*, one is a scaled version of another, angle 0), then $\cos 0 = 1$. If they point in opposite directions ($\pm\pi = \pm 180^\circ$ angle), then




Figure 9.16. Example vectors and the angle between them.

$\cos \pm\pi = -1$. For vectors that are *orthogonal* (perpendicular, $\pm\frac{\pi}{2} = \pm 90^\circ$ angle), we get $\cos \pm\frac{\pi}{2} = 0$.

Note (**) The standard deviation s of a vector $\mathbf{x} \in \mathbb{R}^n$ that has already been centred (whose components' mean is 0) is a scaled version of its magnitude, i.e., $s = \|\mathbf{x}\|/\sqrt{n}$. Looking at the definition of the Pearson linear correlation coefficient (Section 9.1.1), we see that it is the dot product of the standardised versions of two vectors $\mathbf{x}$ and $\mathbf{y}$ divided by the number of elements therein. If the vectors are centred, we can rewrite the formula equivalently as $r(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}}{\|\mathbf{x}\|} \cdot \frac{\mathbf{y}}{\|\mathbf{y}\|}$ and thus $r(\mathbf{x}, \mathbf{y}) = \cos \alpha$. It is not easy to imagine vectors in high-dimensional spaces, but from this observation we can at least imply the fact that r is bounded between -1 and 1. In this context, being not linearly correlated corresponds to the vectors' orthogonality.

9.3.2 Geometric transformations of points (*)

For certain square matrices of size $m \times m$, matrix multiplication can be thought of as an application of the corresponding geometrical transformation of points in $\mathbb{R}^m$

Let $\mathbf{X}$ be a matrix of shape $n \times m$, which we treat as representing the coordinates of n points in an m -dimensional space. For instance, if we are given a diagonal matrix:

$$\mathbf{S} = \text{diag}(s_1, s_2, \dots, s_m) = \begin{bmatrix} s_1 & 0 & \dots & 0 \\ 0 & s_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & s_m \end{bmatrix},$$

then $\mathbf{XS}$ represents *scaling* (stretching) with respect to the individual axes of the coordinate system because:

$$\mathbf{XS} = \begin{bmatrix} s_1 x_{1,1} & s_2 x_{1,2} & \cdots & s_m x_{1,m} \\ s_1 x_{2,1} & s_2 x_{2,2} & \cdots & s_m x_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ s_1 x_{n-1,1} & s_2 x_{n-1,2} & \cdots & s_m x_{n-1,m} \\ s_1 x_{n,1} & s_2 x_{n,2} & \cdots & s_m x_{n,m} \end{bmatrix}.$$

The above can be expressed in **numpy** without referring to the matrix multiplication. A notation like `X * np.array([s1, s2, ..., sm]).reshape(1, -1)` will suffice (elementwise multiplication and proper shape broadcasting).

Furthermore, let $\mathbf{Q}$ is an *orthonormal*¹⁷ matrix, i.e., a square matrix whose columns and rows are unit vectors (normalised), all orthogonal to each other:

- $\|\mathbf{q}_{i,\cdot}\| = 1$ for all i ,
- $\mathbf{q}_{i,\cdot} \cdot \mathbf{q}_{k,\cdot} = 0$ for all i, k ,
- $\|\mathbf{q}_{\cdot,j}\| = 1$ for all j ,
- $\mathbf{q}_{\cdot,j} \cdot \mathbf{q}_{\cdot,k} = 0$ for all j, k .

In such a case, $\mathbf{XQ}$ represents a combination of rotations and reflections.

Important By definition, a matrix $\mathbf{Q}$ is *orthonormal* if and only if $\mathbf{Q}^T \mathbf{Q} = \mathbf{Q} \mathbf{Q}^T = \mathbf{I}$. It is due to the $\cos \pm \frac{\pi}{2} = 0$ interpretation of the dot products of normalised orthogonal vectors.

In particular, the matrix representing the rotation in $\mathbb{R}^2$ about the origin $(0, 0)$ by the counterclockwise angle α :

$$\mathbf{R}(\alpha) = \begin{bmatrix} \cos \alpha & \sin \alpha \\ -\sin \alpha & \cos \alpha \end{bmatrix},$$

is orthonormal (which can be easily verified using the basic trigonometric equalities). Furthermore:

$$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix},$$

represent the two reflections, one against the x- and the other against the y-axis, respectively. Both are orthonormal matrices too.

Consider a dataset $\mathbf{X}'$ in $\mathbb{R}^2$:

¹⁷ Orthonormal matrices are sometimes simply referred to as orthogonal ones.

```
np.random.seed(12345)
Xp = np.random.randn(10000, 2) * 0.25
```

and its scaled, rotated, and translated (shifted) version:

$$\mathbf{X} = \mathbf{X}' \begin{bmatrix} 2 & 0 \\ 0 & 0.5 \end{bmatrix} \begin{bmatrix} \cos \frac{\pi}{6} & \sin \frac{\pi}{6} \\ -\sin \frac{\pi}{6} & \cos \frac{\pi}{6} \end{bmatrix} + \begin{bmatrix} 3 & 2 \end{bmatrix}.$$

```
t = np.array([3, 2])
S = np.diag([2, 0.5])
S
## array([[2. , 0. ],
##        [0. , 0.5]])
alpha = np.pi/6
Q = np.array([
    [ np.cos(alpha), np.sin(alpha)],
    [-np.sin(alpha), np.cos(alpha)]
])
Q
## array([[ 0.8660254,  0.5      ],
##        [-0.5      ,  0.8660254]])
X = Xp @ S @ Q + t
```




Figure 9.17. A dataset and its scaled, rotated, and shifted version.

We can consider $\mathbf{X} = \mathbf{X}'\mathbf{S}\mathbf{Q} + \mathbf{t}$ a version of $\mathbf{X}'$ in a new coordinate system (basis), see [Figure 9.17](#). Each column in the transformed matrix is a shifted linear combination of

the columns in the original matrix:

$$\mathbf{x}_{:,j} = t_j + \sum_{k=1}^m (s_{k,k} q_{k,j}) \mathbf{x}'_{:,k}.$$

The computing of such linear combinations of columns is not rare during a dataset's preprocessing step, especially if they are on the same scale or are unitless. As a matter of fact, the standardisation itself is a form of scaling and translation.

Exercise 9.16 Assume that we have a dataset with two columns representing the number of apples and the number of oranges in clients' baskets. What orthonormal and scaling transforms should be applied to obtain a matrix that gives the total number of fruits and surplus apples (e.g., to convert a row (4, 7) to (11, -3))?

9.3.3 Matrix inverse (*)

The *inverse* of a square matrix $\mathbf{A}$ (if it exists) is denoted by $\mathbf{A}^{-1}$. It is the matrix fulfilling the identity:

$$\mathbf{A}^{-1} \mathbf{A} = \mathbf{A} \mathbf{A}^{-1} = \mathbf{I}.$$

Noting that the identity matrix $\mathbf{I}$ is the neutral element of the matrix multiplication, the above is thus the analogue of the inverse of a scalar: something like $3 \cdot 3^{-1} = 3 \cdot \frac{1}{3} = \frac{1}{3} \cdot 3 = 1$.

Important For any invertible matrices of admissible shapes, it might be shown that the following noteworthy properties hold:

- $(\mathbf{A}^{-1})^T = (\mathbf{A}^T)^{-1}$,
 - $(\mathbf{AB})^{-1} = \mathbf{B}^{-1} \mathbf{A}^{-1}$,
 - a matrix equality $\mathbf{A} = \mathbf{BC}$ holds if and only if $\mathbf{AC}^{-1} = \mathbf{BCC}^{-1} = \mathbf{B}$; this is also equivalent to $\mathbf{B}^{-1} \mathbf{A} = \mathbf{B}^{-1} \mathbf{BC} = \mathbf{C}$.
-

Matrix inverse to identify the inverses of geometrical transformations. Knowing that $\mathbf{X} = \mathbf{X}' \mathbf{S} \mathbf{Q} + \mathbf{t}$, we can recreate the original matrix by applying:

$$\mathbf{X}' = (\mathbf{X} - \mathbf{t})(\mathbf{S} \mathbf{Q})^{-1} = (\mathbf{X} - \mathbf{t}) \mathbf{Q}^{-1} \mathbf{S}^{-1}.$$

It is worth knowing that if $\mathbf{S} = \text{diag}(s_1, s_2, \dots, s_m)$ is a diagonal matrix, then its inverse is $\mathbf{S}^{-1} = \text{diag}(1/s_1, 1/s_2, \dots, 1/s_m)$, which we can denote by $(1/\mathbf{S})$. In addition, the inverse of an orthonormal matrix $\mathbf{Q}$ is always equal to its transpose, $\mathbf{Q}^{-1} = \mathbf{Q}^T$. Luckily, we will not be inverting other matrices in this introductory course.

As a consequence:

$$\mathbf{X}' = (\mathbf{X} - \mathbf{t}) \mathbf{Q}^T (1/\mathbf{S}).$$

Let's verify this numerically (testing equality up to some inherent round-off error):

```
np.allclose(Xp, (X-t) @ Q.T @ np.diag(1/np.diag(S)))
## True
```

9.3.4 Singular value decomposition (*)

It turns out that given any real $n \times m$ matrix $\mathbf{X}$ with $n \geq m$, we can find an interesting scaling and orthonormal transform that, when applied on a dataset whose columns are already normalised, yields exactly $\mathbf{X}$.

Namely, the singular value decomposition (SVD in the so-called compact form) is a factorisation:

$$\mathbf{X} = \mathbf{U}\mathbf{S}\mathbf{Q},$$

where:

- $\mathbf{U}$ is an $n \times m$ semi-orthonormal matrix (its columns are orthonormal vectors; we have $\mathbf{U}^T \mathbf{U} = \mathbf{I}$),
- $\mathbf{S}$ is an $m \times m$ diagonal matrix such that $s_{1,1} \geq s_{2,2} \geq \dots \geq s_{m,m} \geq 0$,
- $\mathbf{Q}$ is an $m \times m$ orthonormal matrix.

Important In data analysis, we usually apply the SVD on matrices that have already been centred (so that their column means are all 0).

For example:

```
import scipy.linalg
n = X.shape[0]
X_centred = X - np.mean(X, axis=0)
U, s, Q = scipy.linalg.svd(X_centred, full_matrices=False)
```

And now:

```
U[:6, :] # preview the first few rows
## array([[ -0.00195072,  0.00474569],
##        [ -0.00510625, -0.00563582],
##        [ 0.01986719,  0.01419324],
##        [ 0.00104386,  0.00281853],
##        [ 0.00783406,  0.01255288],
##        [ 0.01025205, -0.0128136 ]])
```

The norms of all the columns in $\mathbf{U}$ are all equal to 1 (and hence standard deviations are $1/\sqrt{n}$). Consequently, they are on the same scale:

```
np.std(U, axis=0), 1/np.sqrt(n) # compare
## (array([0.01, 0.01]), 0.01)
```

What is more, they are orthogonal: their dot products are all equal to 0. Regarding what we said about Pearson's linear correlation coefficient and its relation to dot products of normalised vectors, we imply that the columns in $\mathbf{U}$ are not linearly correlated. In some sense, they form *independent* dimensions.

Now, we have $\mathbf{S} = \text{diag}(s_1, \dots, s_m)$, with the elements on the diagonal being:

```
s
## array([49.72180455, 12.5126241 ])
```

The elements on the main diagonal of $\mathbf{S}$ are used to scale the corresponding columns in $\mathbf{U}$. The fact that they are ordered decreasingly means that the first column in $\mathbf{US}$ has the greatest standard deviation, the second column has the second greatest variability, and so forth.

```
S = np.diag(s)
US = U @ S
np.std(US, axis=0) # equal to s/np.sqrt(n)
## array([0.49721805, 0.12512624])
```

Multiplying $\mathbf{US}$ by $\mathbf{Q}$ simply rotates and/or reflects the dataset. This brings $\mathbf{US}$ to a new coordinate system where, by construction, the dataset projected onto the direction determined by the first row in $\mathbf{Q}$, i.e., $\mathbf{q}_1$, has the largest variance, projection onto $\mathbf{q}_2$ has the second largest variance, and so on.

```
Q
## array([[ 0.86781968,  0.49687926],
##        [-0.49687926,  0.86781968]])
```

This is why we refer to the rows in $\mathbf{Q}$ as *principal directions* (or *components*). Their scaled versions (proportional to the standard deviations along them) are depicted in [Figure 9.18](#). Note that we have more or less recreated the steps needed to construct $\mathbf{X}$ from $\mathbf{X}'$ (by the way we generated $\mathbf{X}'$, we expect it to have linearly uncorrelated columns; yet, $\mathbf{X}'$ and $\mathbf{U}$ have different column variances).

```
plt.plot(X_centred[:, 0], X_centred[:, 1], "o", alpha=0.1)
plt.arrow(
    0, 0, Q[0, 0]*s[0]/np.sqrt(n), Q[0, 1]*s[0]/np.sqrt(n), width=0.02,
    facecolor="red", edgecolor="white", length_includes_head=True, zorder=2)
plt.arrow(
    0, 0, Q[1, 0]*s[1]/np.sqrt(n), Q[1, 1]*s[1]/np.sqrt(n), width=0.02,
    facecolor="red", edgecolor="white", length_includes_head=True, zorder=2)
plt.show()
```

9.3.5 Dimensionality reduction with SVD (*)

Consider an example three-dimensional dataset:




Figure 9.18. Principal directions of an example dataset (scaled so that they are proportional to the standard deviations along them).

```
chainlink = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/clustering/fcps_chainlink.csv")
```

Section 7.4 said that the plotting is always done on a two-dimensional surface (be it the computer screen or book page). We can look at the dataset only from one *angle* at a time.

In particular, a scatter plot matrix only depicts the dataset from the perspective of the axes of the Cartesian coordinate system (standard basis); see Figure 9.19 (we used a function we defined in Section 7.4.3).

```
pairplot(chainlink, ["axis1", "axis2", "axis3"]) # our function
plt.show()
```

These viewpoints by no means must reveal the true geometric structure of the dataset. However, we know that we can rotate the virtual camera and find some more *interesting* angle. It turns out that our dataset represents two nonintersecting rings, hopefully visible Figure 9.20.

```
fig = plt.figure()
ax = fig.add_subplot(1, 3, 1, projection="3d", facecolor="#ffffff00")
ax.scatter(chainlink[:, 0], chainlink[:, 1], chainlink[:, 2])
ax.view_init(elev=45, azim=45, vertical_axis="z")
ax = fig.add_subplot(1, 3, 2, projection="3d", facecolor="#ffffff00")
ax.scatter(chainlink[:, 0], chainlink[:, 1], chainlink[:, 2])
ax.view_init(elev=37, azim=0, vertical_axis="z")
```

(continues on next page)




Figure 9.19. Views from the perspective of the main axes.

(continued from previous page)

```
ax = fig.add_subplot(1, 3, 3, projection="3d", facecolor="#ffffff00")
ax.scatter(chainlink[:, 0], chainlink[:, 1], chainlink[:, 2])
ax.view_init(elev=10, azim=150, vertical_axis="z")
plt.show()
```

It turns out that we may find a noteworthy viewpoint using the SVD. Namely, we can perform the decomposition of a centred dataset which we denote by $\mathbf{X}$:

$$\mathbf{X} = \mathbf{U}\mathbf{S}\mathbf{Q}.$$

```
import scipy.linalg
X_centered = chainlink - np.mean(chainlink, axis=0)
U, s, Q = scipy.linalg.svd(X_centered, full_matrices=False)
```




Figure 9.20. Different views of the same dataset.

Then, considering its rotated/reflected version:

$$\mathbf{P} = \mathbf{X}\mathbf{Q}^{-1} = \mathbf{U}\mathbf{S},$$

we know that its first column has the highest variance, the second column has the second highest variability, and so on. It might indeed be worth looking at that dataset from that *most informative* perspective.

Figure 9.21 gives the scatter plot for $\mathbf{p}_{\cdot,1}$ and $\mathbf{p}_{\cdot,2}$. Maybe it does not reveal the true geometric structure of the dataset (no single two-dimensional projection can do that), but at least it is better than the initial ones (from the pairs plot).

```
P2 = U[:, :2] @ np.diag(s[:2]) # the same as (U@np.diag(s))[:, :2]
plt.plot(P2[:, 0], P2[:, 1], "o")
plt.axis("equal")
plt.show()
```

What we just did is a kind of *dimensionality reduction*. We found a viewpoint (in the form of an orthonormal matrix, being a mixture of rotations and reflections) on $\mathbf{X}$ such that its orthonormal projection onto the first two axes of the Cartesian coordinate system is the most informative¹⁸ (in terms of having the highest variance along these axes).

¹⁸ (**) The Eckart–Young–Mirsky theorem states that $\mathbf{U}_{:,k}\mathbf{S}_{k,k:k}\mathbf{Q}_{:,k}$, (where “ k ” denotes “the first k rows or columns”) is the best rank- k approximation of $\mathbf{X}$ with respect to both the Frobenius and spectral norms.




Figure 9.21. The view from the two principal axes.

9.3.6 Principal component analysis (*)

Principal component analysis (PCA) is a fancy name for the entire process involving our brainstorming upon what happens along the projections onto the most variable dimensions. It can be used not only for data visualisation and deduplication, but also for feature engineering (as it creates new columns that are linear combinations of existing ones).

Consider a few chosen countrywise 2016 [Sustainable Society Indices](#)¹⁹:

```
ssi = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/ssi_2016_indicators.csv",
    comment="#") # more on data frames in the next chapter
X = np.array(ssi.iloc[:, [3, 5, 13, 15, 19] ]) # select columns, make matrix
n = X.shape[0]
X[:6, :] # preview
## array([[ 9.32      ,  8.13333333,  8.386      ,  8.5757      ,  5.46249573],
##        [ 8.74      ,  7.71666667,  7.346      ,  6.8426      ,  6.2929302 ],
##        [ 5.11      ,  4.31666667,  8.788      ,  9.2035      ,  3.91062849],
##        [ 9.61      ,  7.93333333,  5.97      ,  5.5232      ,  7.75361284],
##        [ 8.95      ,  7.81666667,  8.032      ,  8.2639      ,  4.42350654],
##        [10.      ,  8.65      ,  1.      ,  1.      ,  9.66401848]])
```

Each index is on the scale from 0 to 10. These are, in this order:

1. Safe Sanitation,
2. Healthy Life,

¹⁹ <https://ssi.wi.th-koeln.de/>

3. Energy Use,
4. Greenhouse Gases,
5. Gross Domestic Product.

Above we displayed the data corresponding to six countries:

```
countries = list(ssi.iloc[:, 0]) # select the 1st column from the data frame
countries[:6] # preview
## ['Albania', 'Algeria', 'Angola', 'Argentina', 'Armenia', 'Australia']
```

This is a five-dimensional dataset. We cannot easily visualise it. Observing that the pairs plot does not reveal too much is left as an exercise. Let's thus perform the SVD decomposition of a standardised version of this dataset, $\mathbf{Z}$ (recall that the centring is necessary, at the very least).

```
Z = (X - np.mean(X, axis=0))/np.std(X, axis=0)
U, s, Q = scipy.linalg.svd(Z, full_matrices=False)
```

The standard deviations of the data projected onto the consecutive principal components (columns in $\mathbf{U}$) are:

```
s/np.sqrt(n)
## array([2.02953531, 0.7529221 , 0.3943008 , 0.31897889, 0.23848286])
```

It is customary to check the ratios of the cumulative variances explained by the consecutive principal components, which is a normalised measure of their importances. We can compute them by calling:

```
np.cumsum(s**2)/np.sum(s**2)
## array([0.82380272, 0.93718105, 0.96827568, 0.98862519, 1.          ])
```

As, in some sense, the variability within the first two components covers c. 94% of the variability of the whole dataset, we can restrict ourselves only to a two-dimensional projection of this dataset. The rows in $\mathbf{Q}$ define the *loadings*, which give the coefficients defining the linear combinations of the rows in $\mathbf{Z}$ that correspond to the principal components.

Let's try to find their interpretation.

```
np.round(Q[0, :], 2) # loadings - the first principal axis
## array([-0.43, -0.43, 0.44, 0.45, -0.47])
```

The first row in $\mathbf{Q}$ consists of similar values, but with different signs. We can consider them a scaled version of the average Energy Use (column 3), Greenhouse Gases (4), and MINUS Safe Sanitation (1), MINUS Healthy Life (2), MINUS Gross Domestic Product (5). We could call this a measure of a country's overall eco-unfriendliness(?) because countries with low Healthy Life and high Greenhouse Gasses will score highly on this scale.

```
np.round(Q[1, :], 2) # loadings - the second principal axis
## array([ 0.52,  0.5 ,  0.52,  0.45, -0.02])
```

The second row in **Q** defines a scaled version of the average of Safe Sanitation (1), Healthy Life (2), Energy Use (3), and Greenhouse Gases (4), almost completely ignoring the GDP (5). Can we call it a measure of industrialisation? Something like this. But this naming is just for fun²⁰.

Figure 9.22 is a scatter plot of the countries projected onto the said two principal directions. For readability, we only display a few chosen labels. This is merely a projection/approximation, but it might be an interesting one for some decision makers.

```
P2 = U[:, :2] @ np.diag(s[:2]) # == Y @ Q[:, :2].T
plt.plot(P2[:, 0], P2[:, 1], "o", alpha=0.1)
which = [ # hand-crafted/artisan
    141, 117, 69, 123, 35, 80, 93, 45, 15, 2, 60, 56, 14,
    104, 122, 8, 134, 128, 0, 94, 114, 50, 34, 41, 33, 77,
    64, 67, 152, 135, 148, 99, 149, 126, 111, 57, 20, 63
]
for i in which:
    plt.text(P2[i, 0], P2[i, 1], countries[i], ha="center")
plt.axis("equal")
plt.xlabel("1st principal component (eco-unfriendliness?)")
plt.ylabel("2nd principal component (industrialisation?)")
plt.show()
```

Exercise 9.17 Perform a principal component analysis of the body dataset. Project the points onto a two-dimensional plane.

9.4 Further reading

Other approaches to regression via linear models include ridge and lasso, the latter having the nice property of automatically getting rid of noninformative variables from the model. Furthermore, instead of minimising squared residuals, we can also consider, e.g., least absolute deviation.

There are many other approaches to dimensionality reduction, also nonlinear ones, including kernel PCA, feature agglomeration via hierarchical clustering, autoencoders, t-SNE, UMAP, etc.

A popular introductory text in statistical learning is [50]. We recommend [2, 10, 11, 24, 26] for more advanced students. Computing-orientated students could benefit from checking out [68].

²⁰ Nonetheless, someone might take these results seriously and scribble a research thesis about it. Mathematics, unlike the brains of ordinary mortals, does not need our imperfect interpretations/fairy tales to function properly. We need more maths in our lives.




Figure 9.22. An example principal component analysis of countries.

9.5 Exercises

Exercise 9.18 Why correlation is not causation?

Exercise 9.19 What does the linear correlation of 0.9 mean? What about the rank correlation of 0.9? And the linear correlation of 0.0?

Exercise 9.20 How is Spearman's coefficient related to Pearson's one?

Exercise 9.21 State the optimisation problem behind the least squares fitting of linear models.

Exercise 9.22 What are the different ways of the numerical summarising of residuals?

Exercise 9.23 Why is it important for the residuals to be homoscedastic?

Exercise 9.24 Is a more complex model always better?

Exercise 9.25 Why must extrapolation be handled with care?

Exercise 9.26 Why did we say that novice users should refrain from using *scikit-learn*?

Exercise 9.27 What is the condition number of a model matrix and why is it worthwhile to always check it?

Exercise 9.28 What is the geometrical interpretation of the dot product of two normalised vectors?

Exercise 9.29 How can we verify if two vectors are orthonormal? What is an orthonormal projection? What is the inverse of an orthonormal matrix?

Exercise 9.30 *What is the inverse of a diagonal matrix?*

Exercise 9.31 *Characterise the general properties of the three matrices obtained by performing the singular value decomposition of a given matrix of shape $n \times m$.*

Exercise 9.32 *How can we obtain the first principal component of a given centred matrix?*

Exercise 9.33 *How can we compute the ratios of the variances explained by the consecutive principal components?*

Part IV

Heterogeneous data

Introducing data frames

`numpy` arrays are an extremely versatile tool for performing data analysis activities and other numerical computations of various kinds. Even though it is theoretically possible otherwise, in practice, we only store elements of the same type there: most often numbers.

`pandas`¹ [66] is amongst over *one hundred thousand*² open-source packages and repositories that use `numpy` to provide additional data wrangling functionality. It was originally written by Wes McKinney but was heavily inspired by the `data.frame`³ objects in S and R as well as tables in relational (think: SQL) databases and spreadsheets.

`pandas` defines a few classes, of which the most important are:

- `DataFrame` – for representing tabular data (matrix-like) with columns of possibly different types, in particular a mix of numerical and categorical variables,
- `Series` – vector-like objects for storing individual columns,
- `Index` and its derivatives – vector-like (usually) objects for labelling individual rows and columns in `DataFrames` and items in `Series` objects,
- `SeriesGroupBy` and `DataFrameGroupBy` – which model observations grouped by a categorical variable or a combination of factors ([Chapter 12](#)),

together with many methods for:

- transforming/aggregating/processing data, also in groups determined by categorical variables or products thereof,
- reshaping (e.g., from wide to long format) and joining datasets,
- importing/exporting data from/to various sources and formats, e.g., CSV and HDF5 files or relational databases,
- handling missing data,

all of which we introduce in this part.

It is customary to import the `pandas` package under the following alias:

¹ <https://pandas.pydata.org/>

² <https://libraries.io/pypi/numpy>

³ Data frames were first introduced in the 1991 version of the S language [15].

```
import pandas as pd
```

Important Let's repeat: **pandas** is built on top of **numpy** and most objects therein can be processed by **numpy** functions as well. Many other functions, e.g., in **scikit-learn**, accept both **DataFrame** and **ndarray** objects, but often convert the former to the latter internally to enable data processing using fast, lower-level C/C++/Fortran routines.

What we have learnt so far⁴ still applies. But there is more, hence this part.

10.1 Creating data frames

Data frames can be created, amongst others, using the **DataFrame** class constructor, which can be fed, for example, with a **numpy** matrix:

```
np.random.seed(123)
pd.DataFrame(
    np.random.rand(4, 3),
    columns=["a", "b", "c"]
)
##           a           b           c
## 0  0.696469  0.286139  0.226851
## 1  0.551315  0.719469  0.423106
## 2  0.980764  0.684830  0.480932
## 3  0.392118  0.343178  0.729050
```

Notice that rows and columns are labelled (and how readable that is).

A dictionary of vector-like objects of equal lengths is another common data source:

```
np.random.seed(123)
df = pd.DataFrame(dict(
    a = np.round(np.random.rand(5), 2),
    b = [1, 2.5, np.nan, 4, np.nan],
    c = [True, True, False, False, True],
    d = ["A", "B", "C", None, "E"],
    e = ["spam", "spam", "bacon", "spam", "eggs"],
    f = np.array([
        "2021-01-01", "2022-02-02", "2023-03-03", "2024-04-04", "2025-05-05"
    ], dtype="datetime64[D]"),
    g = [
        ["spam"], ["bacon", "spam"], None, ["eggs", "bacon", "spam"], ["ham"]
    ]
))
```

(continues on next page)

⁴ If, by any chance, some overenthusiastic readers decided to start this superb book at this chapter, it is now the time to go back to the [Preface](#) and learn everything in the right order. See you later.

(continued from previous page)

```
],
))
df
##      a      b      c      d      e      f      g
## 0  0.70  1.0   True    A  spam 2021-01-01  [spam]
## 1  0.29  2.5   True    B  spam 2022-02-02  [bacon, spam]
## 2  0.23  NaN  False    C  bacon 2023-03-03      None
## 3  0.55  4.0  False   None spam 2024-04-04  [eggs, bacon, spam]
## 4  0.72  NaN   True    E  eggs 2025-05-05  [ham]
```

The above illustrates the possibility of having columns of different types.

Exercise 10.1 Check out the `pandas.DataFrame.from_dict` and `pandas.DataFrame.from_records` functions in the [documentation](#)⁵. Use them to create example data frames.

Further, data frames can be read from files in different formats, for instance, CSV:

```
body = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
                  "teaching-data/master/marek/nhanes_adult_female_bmx_2020.csv",
                  comment="#")
body.head() # display the first few rows (five by default)
##      BMXWT  BMXHT  BMXARML  BMXLEG  BMXARMC  BMXHIP  BMXWAIST
## 0    97.1   160.2    34.7    40.8    35.8   126.1    117.9
## 1    91.1   152.7    33.5    33.0    38.5   125.5    103.1
## 2    73.0   161.2    37.4    38.0    31.8   106.2    92.0
## 3    61.7   157.4    38.0    34.7    29.0   101.0    90.5
## 4    55.4   154.6    34.6    34.0    28.3    92.5    73.2
```

Reading from URLs and local files is also supported; compare [Section 13.6.1](#).

Exercise 10.2 Check out other `pandas.read_*` functions in the `pandas` documentation, e.g., for importing spreadsheets, Apache Parquet and HDF5 files, scraping tables from HTML documents, or reading data from relational databases. We will discuss some of them in more detail later.

Exercise 10.3 (*) Large files that do not fit into computer's memory (but not too large) can still be read with `pandas.read_csv`. Check out the meaning of the `usecols`, `dtype`, `skiprows`, and `nrows` arguments. On a side note, sampling is mentioned in [Section 10.5.4](#) and chunking in [Section 13.2](#).

10.1.1 Data frames are matrix-like

Data frames are modelled through `numpy` matrices. We can thus already feel quite at home with them.

For example, a data frame, it is easy to fetch its number of rows and columns:

⁵ <https://pandas.pydata.org/docs>

```
df.shape
## (5, 7)
```

or the type of each column:

```
df.dtypes # returns a Series object; see below
## a      float64
## b      float64
## c      bool
## d      object
## e      object
## f      datetime64[s]
## g      object
## dtype: object
```

Recall that `numpy` arrays are equipped with the `dtype` slot.

10.1.2 Series

There is a separate class for storing individual data frame columns: it is called `Series`.

```
s = df.loc[:, "a"] # extract the `a` column; alternatively: df.a
s
## 0      0.70
## 1      0.29
## 2      0.23
## 3      0.55
## 4      0.72
## Name: a, dtype: float64
```

Data frames with one column are printed out slightly differently. We get the column name at the top, but do not have the `dtype` information at the bottom.

```
s.to_frame() # or: pd.DataFrame(s)
##      a
## 0  0.70
## 1  0.29
## 2  0.23
## 3  0.55
## 4  0.72
```

Indexing will be discussed later.

Important It is crucial to know when we are dealing with a `Series` and when with a `DataFrame` object as each of them defines a slightly different set of methods.

We will now be relying on object-orientated syntax (compare [Section 2.2.3](#)) much more frequently than before.

Example 10.4 *By calling:*

```
s.mean()
## 0.498000000000000005
```

we refer to **pandas.Series.mean** (which returns a scalar), whereas:

```
df.mean(numeric_only=True)
## a    0.498
## b    2.500
## c    0.600
## dtype: float64
```

uses **pandas.DataFrame.mean** (which yields a Series).

Look up these two methods in the **pandas** manual. Note that their argument list is slightly different.

Objects of the class **Series** are vector-like:

```
s.shape
## (5,)
s.dtype
## dtype('float64')
```

They are wrappers around **numpy** arrays.

```
s.values
## array([0.7 , 0.29, 0.23, 0.55, 0.72])
```

Most importantly, **numpy** functions can be called directly on them:

```
np.mean(s)
## 0.498000000000000005
```

As a consequence, what we covered in the part of this book that dealt with vector processing still holds for data frame columns (but there will be more).

Series can also be *named*.

```
s.name
## 'a'
```

This is convenient, especially when we convert them to a data frame as the name sets the label of the newly created column:

```
s.rename("spam").to_frame()
##      spam
## 0    0.70
## 1    0.29
## 2    0.23
```

(continues on next page)

(continued from previous page)

```
## 3 0.55
## 4 0.72
```

10.1.3 Index

Another important class is called `Index`⁶. We use it for storing element or axes labels.

The index (lowercase) *slot* of a data frame stores an object of the class `Index` (or one of its derivatives) that gives the row names:

```
df.index # row labels
## RangeIndex(start=0, stop=5, step=1)
```

It represents a sequence (0, 1, 2, 3, 4). Furthermore, the column slot gives:

```
df.columns # column labels
## Index(['a', 'b', 'c', 'd', 'e', 'f', 'g'], dtype='object')
```

Also, we can label the individual elements in `Series` objects:

```
s.index
## RangeIndex(start=0, stop=5, step=1)
```

The `set_index` method can be applied to make a data frame column act as a sequence of row labels:

```
df2 = df.set_index("e")
df2
```

| | a | b | c | d | f | g |
|-------|------|-----|-------|-----------------|---------------------|---------------|
| e | | | | | | |
| spam | 0.70 | 1.0 | True | A 2021-01-01 | | [spam] |
| spam | 0.29 | 2.5 | True | B 2022-02-02 | | [bacon, spam] |
| bacon | 0.23 | NaN | False | C 2023-03-03 | | None |
| spam | 0.55 | 4.0 | False | None 2024-04-04 | [eggs, bacon, spam] | |
| eggs | 0.72 | NaN | True | E 2025-05-05 | | [ham] |

This `Index` object is named:

```
df2.index.name
## 'e'
```

We can also rename the axes on the fly:

```
df2.rename_axis(index="ROWS", columns="COLS")
## COLS      a      b      c      d      f      g
```

(continues on next page)

⁶ The name `Index` is confusing not only because it clashes with the *index* operator (square brackets), but also the concept of an *index* in relational databases. In `pandas`, we can have nonunique row names.

(continued from previous page)

```
## ROWS
## spam    0.70  1.0   True    A 2021-01-01           [spam]
## spam    0.29  2.5   True    B 2022-02-02           [bacon, spam]
## bacon   0.23  NaN  False   C 2023-03-03           None
## spam    0.55  4.0  False   None 2024-04-04 [eggs, bacon, spam]
## eggs    0.72  NaN   True    E 2025-05-05           [ham]
```

Having a named index slot is handy when converting a vector of row labels back to a standalone column:

```
df2.rename_axis(index="NEW_COLUMN").reset_index()
##  NEW_COLUMN    a    b    c    d    f    g
## 0      spam    0.70  1.0   True   A 2021-01-01    [spam]
## 1      spam    0.29  2.5   True   B 2022-02-02    [bacon, spam]
## 2     bacon    0.23  NaN  False   C 2023-03-03      None
## 3      spam    0.55  4.0  False  None 2024-04-04 [eggs, bacon, spam]
## 4     eggs    0.72  NaN   True   E 2025-05-05    [ham]
```

But we can also do it this way:

```
df2.reset_index(names="Spanish Inquisition")
## Spanish Inquisition    a    b    ...    d    f    g
## 0      spam    0.70  1.0    ...   A 2021-01-01    [spam]
## 1      spam    0.29  2.5    ...   B 2022-02-02    [bacon, spam]
## 2     bacon    0.23  NaN    ...   C 2023-03-03      None
## 3      spam    0.55  4.0    ...  None 2024-04-04 [eggs, bacon, spam]
## 4     eggs    0.72  NaN    ...   E 2025-05-05    [ham]
##
## [5 rows x 7 columns]
```

There is also an option to drop the current index whatsoever and to replace it with the default label sequence, i.e., 0, 1, 2, ...:

```
df2.reset_index(drop=True)
##      a    b    c    d    f    g
## 0  0.70  1.0   True   A 2021-01-01    [spam]
## 1  0.29  2.5   True   B 2022-02-02    [bacon, spam]
## 2  0.23  NaN  False   C 2023-03-03      None
## 3  0.55  4.0  False  None 2024-04-04 [eggs, bacon, spam]
## 4  0.72  NaN   True   E 2025-05-05    [ham]
```

Take note of the fact that `reset_index`, and many other methods that we have used so far, do not modify the data frame in place.

Important We will soon get used to calling `reset_index(drop=True)` frequently: sometimes more than once in a single series of commands.

Exercise 10.5 Use the `pandas.DataFrame.rename` method to change the name of the *a* column in *df* to *spam*.

Also, a *hierarchical index* – one that is comprised of more than one level – is possible. For example, here is a sorted (see Section 10.6.1) version of *df* with a new index based on two columns at the same time:

```
df.sort_values("e", ascending=False).set_index(["e", "c"])
##           a      b      d      f      g
## e      c
## spam  True  0.70  1.0    A  2021-01-01    [spam]
##       True  0.29  2.5    B  2022-02-02    [bacon, spam]
##       False 0.55  4.0   None 2024-04-04  [eggs, bacon, spam]
## eggs  True  0.72  NaN    E  2025-05-05    [ham]
## bacon False 0.23  NaN    C  2023-03-03    None
```

For the sake of readability, the consecutive repeated *spams* were not printed.

Example 10.6 Hierarchical indexes might arise after aggregating data in groups. For example:

```
nhanes = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_p_demo_bmx_2020.csv",
    comment="#").rename({
        "BMXBMI": "bmival",
        "RIAGENDR": "gender",
        "DMDBORN4": "usborn"
    }, axis=1)
```

In Chapter 12, we will get used to writing:

```
res = nhanes.groupby(["gender", "usborn"])["bmival"].mean()
res # BMI by gender and US born-ness
## gender  usborn
## 1      1      25.734110
##      2      27.405251
## 2      1      27.120261
##      2      27.579448
##      77      28.725000
##      99      32.600000
## Name: bmival, dtype: float64
```

This returned a *Series* object with a hierarchical index. If we do not fancy it, `reset_index` comes to our rescue:

```
res.reset_index()
##   gender  usborn  bmival
## 0      1      1  25.734110
## 1      1      2  27.405251
## 2      2      1  27.120261
## 3      2      2  27.579448
```

(continues on next page)

(continued from previous page)

```
## 4      2      77  28.725000
## 5      2      99  32.600000
```

10.2 Aggregating data frames

Here is another toy data frame:

```
np.random.seed(123)
df = pd.DataFrame(dict(
    u = np.round(np.random.rand(5), 2),
    v = np.round(np.random.randn(5), 2),
    w = ["spam", "bacon", "spam", "eggs", "sausage"]
), index=["a", "b", "c", "d", "e"])
df
##      u      v      w
## a  0.70  0.32  spam
## b  0.29 -0.05  bacon
## c  0.23 -0.20  spam
## d  0.55  1.98  eggs
## e  0.72 -1.62  sausage
```

All **numpy** functions can be applied on individual columns, i.e., objects of the type `Series`, because they are vector-like.

```
u = df.loc[:, "u"] # extract the `u` column (gives a Series; see below)
np.quantile(u, [0, 0.5, 1])
## array([0.23, 0.55, 0.72])
```

Most **numpy** functions also work if they are fed with data frames, but we will need to extract the numeric columns manually.

```
uv = df.loc[:, ["u", "v"]] # select two columns (a DataFrame; see below)
np.quantile(uv, [0, 0.5, 1], axis=0)
## array([[ 0.23, -1.62],
##        [ 0.55, -0.05],
##        [ 0.72,  1.98]])
```

Sometimes the results will automatically be coerced to a `Series` object with the `index` slot set appropriately:

```
np.mean(uv, axis=0)
## u      0.498
## v      0.086
## dtype: float64
```

For convenience, many operations are also available as methods for the `Series` and `DataFrame` classes, e.g., `mean`, `median`, `min`, `max`, `quantile`, `var`, `std`, and `skew`.

```
df.mean(numeric_only=True)
## u    0.498
## v    0.086
## dtype: float64
df.quantile([0, 0.5, 1], numeric_only=True)
##          u          v
## 0.0    0.23 -1.62
## 0.5    0.55 -0.05
## 1.0    0.72  1.98
```

Also note the `describe` method, which returns a few statistics at the same time.

```
df.describe()
##          u          v
## count  5.000000  5.000000
## mean    0.498000  0.086000
## std     0.227969  1.289643
## min     0.230000 -1.620000
## 25%     0.290000 -0.200000
## 50%     0.550000 -0.050000
## 75%     0.700000  0.320000
## max     0.720000  1.980000
```

Exercise 10.7 Check out the `pandas.DataFrame.agg` method that can apply all aggregates given by a list of functions. Compose a call equivalent to `df.describe()`.

Note (*) Let's stress that above we see the corrected for bias (but still only asymptotically unbiased) version of standard deviation, given by $\sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$; compare [Section 5.1](#). In `pandas`, `std` methods assume `ddof=1` by default, whereas we recall that `numpy` uses `ddof=0`.

```
np.round([u.std(), np.std(u), np.std(np.array(u)), u.std(ddof=0)], 3)
## array([0.228, 0.204, 0.204, 0.204])
```

This is an unfortunate inconsistency between the two packages, but please do not blame the messenger.

10.3 Transforming data frames

By applying the already well-known vectorised mathematical functions from **numpy**, we can transform each data cell and return an object of the same type as the input one.

```
np.exp(df.loc[:, "u"])
## a    2.013753
## b    1.336427
## c    1.258600
## d    1.733253
## e    2.054433
## Name: u, dtype: float64
np.exp(df.loc[:, ["u", "v"]])
##              u              v
## a    2.013753    1.377128
## b    1.336427    0.951229
## c    1.258600    0.818731
## d    1.733253    7.242743
## e    2.054433    0.197899
```

When applying the binary arithmetic, relational, and logical operators on an object of the class **Series** and a scalar or a **numpy** vector, the operations are performed elementwisely. For instance, here is a standardised version of the **u** column:

```
u = df.loc[:, "u"]
(u - np.mean(u)) / np.std(u)
## a    0.990672
## b   -1.020098
## c   -1.314357
## d    0.255025
## e    1.088759
## Name: u, dtype: float64
```

Arithmetic operators act on the elements with corresponding labels. For two objects having identical index slots (this is the most common scenario), this is the same as elementwise vectorisation. For instance:

```
df.loc[:, "u"] > df.loc[:, "v"] # here: elementwise comparison
## a    True
## b    True
## c    True
## d   False
## e    True
## dtype: bool
```

For transforming many numerical columns at once, it is worthwhile either to convert them to a numeric matrix explicitly and then use the basic **numpy** functions:

```
uv = np.array(df.loc[:, ["u", "v"]])
uv2 = (uv-np.mean(uv, axis=0))/np.std(uv, axis=0)
uv2
## array([[ 0.99067229,  0.20286225],
##        [-1.0200982 , -0.11790285],
##        [-1.3143573 , -0.24794275],
##        [ 0.25502455,  1.64197052],
##        [ 1.08875866, -1.47898717]])
```

or to use the `pandas.DataFrame.apply` method which invokes a given function on each column separately:

```
uv2 = df.loc[:, ["u", "v"]].apply(lambda x: (x-np.mean(x))/np.std(x))
uv2
##           u           v
## a  0.990672  0.202862
## b -1.020098 -0.117903
## c -1.314357 -0.247943
## d  0.255025  1.641971
## e  1.088759 -1.478987
```

Anticipating what we cover in the next section, in both cases, we can write `df.loc[:, ["u", "v"]] = uv2` to replace the old content. Also, new columns can be added based on the transformed versions of the existing ones. For instance:

```
df.loc[:, "uv_squared"] = (df.loc[:, "u"] * df.loc[:, "v"])**2
df
##           u           v           w  uv_squared
## a  0.70  0.32      spam      0.050176
## b  0.29 -0.05     bacon      0.000210
## c  0.23 -0.20      spam      0.002116
## d  0.55  1.98     eggs      1.185921
## e  0.72 -1.62  sausage      1.360489
```

Example 10.8 (*) *Arithmetic operations on objects with different index slots are vectorised labelwisely:*

```
x = pd.Series([1, 10, 100, 1000, 10000], index=["a", "b", "a", "a", "c"])
x
## a      1
## b     10
## a    100
## a   1000
## c  10000
## dtype: int64
y = pd.Series([1, 2, 3, 4, 5], index=["b", "b", "a", "d", "c"])
y
## b      1
```

(continues on next page)

(continued from previous page)

```
## b    2
## a    3
## d    4
## c    5
## dtype: int64
```

And now:

```
x * y
## a      3.0
## a     300.0
## a    3000.0
## b      10.0
## b     20.0
## c    50000.0
## d         NaN
## dtype: float64
```

Here, each element in the first *Series* named *a* was multiplied by each (there was only one) element labelled *a* in the second *Series*. For *d*, there were no matches, hence the result's being marked as missing; compare [Chapter 15](#). Thus, it behaves like a full outer join-type operation; see [Section 10.6.3](#).

The above is different from elementwise vectorisation in *numpy*:

```
np.array(x) * np.array(y)
## array([ 1, 20, 300, 4000, 50000])
```

Labelwise vectorisation can be useful in certain contexts. However, we need to be aware of this (yet another) incompatibility between the two packages.

10.4 Indexing Series objects

Recall that each *DataFrame* and *Series* object is equipped with a slot called *index*, which is an object of the class *Index* (or subclass thereof), giving the row and element labels, respectively. It turns out that we may apply the *index* operator, `[...]`, to subset these objects not only through the *indexers* known from the *numpy* part (e.g., numerical ones, i.e., by position) but also ones that pinpoint the items via their labels. We have thus quite a lot of index-ing to discuss.

For illustration, we will be playing with two objects of the type *Series*:

```
np.random.seed(123)
b = pd.Series(np.round(np.random.rand(10), 2))
b.index = np.random.permutation(np.arange(10))
```

(continues on next page)

(continued from previous page)

```

b
## 2    0.70
## 1    0.29
## 8    0.23
## 7    0.55
## 9    0.72
## 4    0.42
## 5    0.98
## 6    0.68
## 3    0.48
## 0    0.39
## dtype: float64

```

and:

```

c = b.copy()
c.index = list("abcdefghij")
c
## a    0.70
## b    0.29
## c    0.23
## d    0.55
## e    0.72
## f    0.42
## g    0.98
## h    0.68
## i    0.48
## j    0.39
## dtype: float64

```

They consist of the same values, in the same order, but have different labels (index slots). In particular, *b*'s labels are integers that *do not* match the physical element positions (where 0 would denote the first element, etc.).

Important For *numpy* vectors, we had four different indexing schemes: via a scalar (extracts an element at a given position), a slice, an integer vector, and a logical vector. Series objects are *additionally* labelled. Therefore, they can also be accessed through the contents of the index slot.

10.4.1 Do not use [...] directly (in the current version of pandas)

Applying the index operator, [...], directly on Series is currently not a wise idea:

```

b[0] # do not use it
## 0.39
b[ [0] ] # do not use it

```

(continues on next page)

(continued from previous page)

```
## 0    0.39
## dtype: float64
```

both do not select the first item, but the item labelled 0. However, the undermentioned two calls fall back to position-based indexing.

```
b[:1] # do not use it: it will change in the future!
## 2    0.7
## dtype: float64
c[0] # there is no label `0` (do not use it: it will change in the future!)
## 0.7
```

Confusing? Well, with some self-discipline, the solution is easy:

Important In the current version of **pandas**, we recommend abstaining from applying `[...]` directly on `Series` and `DataFrame` objects. We should explicitly refer to the `loc[...]` and `iloc[...]` accessors for the label- and position-based filtering, respectively.

In the future, direct call to `[...]` on `Series` will use label-based indexing (as we will see below, `b[:5]` will hence not mean “select the first five rows”). At the same time, `[...]` on `DataFrame` currently serves as a label-based selector of columns only.

10.4.2 `loc[...]`

`Series.loc[...]` implements label-based indexing.

```
b.loc[0]
## 0.39
```

This returned the element labelled 0. On the other hand, `c.loc[0]` will raise a `KeyError` because `c` consists of string labels only. But in this case, we can write:

```
c.loc["j"]
## 0.39
```

Next, we can use lists of labels to select a *subset*.

```
b.loc[ [0, 1, 0] ]
## 0    0.39
## 1    0.29
## 0    0.39
## dtype: float64
c.loc[ ["j", "b", "j"] ]
## j    0.39
## b    0.29
```

(continues on next page)

(continued from previous page)

```
## j    0.39
## dtype: float64
```

The result is always of the type `Series`.

Slicing behaves differently as the range is *inclusive* (sic!⁷) at both sides:

```
b.loc[1:7]
## 1    0.29
## 8    0.23
## 7    0.55
## dtype: float64
b.loc[0:4:-1]
## 0    0.39
## 3    0.48
## 6    0.68
## 5    0.98
## 4    0.42
## dtype: float64
c.loc["d":"g"]
## d    0.55
## e    0.72
## f    0.42
## g    0.98
## dtype: float64
```

They return all elements between the two indicated labels.

Note Be careful that if there are repeated labels, then we will be returning *all* (sic!⁸) the matching items:

```
d = pd.Series([1, 2, 3, 4], index=["a", "b", "a", "c"])
d.loc["a"]
## a    1
## a    3
## dtype: int64
```

The result is not a scalar but a `Series` object.

10.4.3 `iloc[...]`

Here are some examples of position-based indexing with the `iloc[...]` accessor. It is worth stressing that, fortunately, its behaviour is consistent with its `numpy` coun-

⁷ Inconsistency; but makes sense when selecting column ranges.

⁸ Inconsistency; but makes sense for hierarchical indexes with repeated.

terpart, i.e., the ordinary square brackets applied on objects of the class `ndarray`. For example:

```
b.iloc[0] # the same: c.iloc[0]
## 0.7
```

returns the first element.

```
b.iloc[1:7] # the same: b.iloc[1:7]
## 1    0.29
## 8    0.23
## 7    0.55
## 9    0.72
## 4    0.42
## 5    0.98
## dtype: float64
```

returns the second, third, ..., seventh element (not including `b.iloc[7]`, i.e., the eighth one).

10.4.4 Logical indexing

Indexing using a logical vector-like object is also available. For this purpose, we will usually be using `loc[...]` with either a logical `Series` object of identical index slot as the subsetted object, or a Boolean `numpy` vector.

```
b.loc[(b > 0.4) & (b < 0.6)]
## 7    0.55
## 4    0.42
## 3    0.48
## dtype: float64
```

For `iloc[...]`, the indexer must be unlabelled, e.g., be an ordinary `numpy` vector.

10.5 Indexing data frames

10.5.1 `loc[...]` and `iloc[...]`

For data frames, `iloc` and `loc` can be applied too. Now, however, they require *two* arguments: a row and a column selector. For example:

```
np.random.seed(123)
df = pd.DataFrame(dict(
    u = np.round(np.random.rand(5), 2),
    v = np.round(np.random.randn(5), 2),
    w = ["spam", "bacon", "spam", "eggs", "sausage"],
```

(continues on next page)

(continued from previous page)

```
x = [True, False, True, False, True]
))
```

And now:

```
df.loc[ df.loc[:, "u"] > 0.5, "u":"w" ]
##           u           v           w
## 0  0.70  0.32      spam
## 3  0.55  1.98      eggs
## 4  0.72 -1.62  sausage
```

It selected the rows where the values in the `u` column are greater than 0.5 and then returns all columns between `u` and `w` (inclusive!).

Furthermore:

```
df.iloc[:3, :].loc[:, ["u", "w"]]
##           u           w
## 0  0.70      spam
## 1  0.29      bacon
## 2  0.23      spam
```

It fetched the first three rows (by position; `iloc` is necessary) and then selects two indicated columns.

Compare this to:

```
df.loc[:3, ["u", "w"]] # df[:3, ["u", "w"]] does not even work; please don't
##           u           w
## 0  0.70      spam
## 1  0.29      bacon
## 2  0.23      spam
## 3  0.55      eggs
```

which has four (!) rows.

Important Getting a scrambled numeric index that does not match the physical positions is not rare: for instance, in the context of data frame sorting ([Section 10.6.1](#)):

```
df2 = df.sort_values("v")
df2
##           u           v           w           x
## 4  0.72 -1.62  sausage      True
## 2  0.23 -0.20      spam      True
## 1  0.29 -0.05      bacon     False
## 0  0.70  0.32      spam      True
## 3  0.55  1.98      eggs     False
```

Note how different are the following results:

```
df2.loc[:3, :] # up to label 3, inclusive
##      u      v      w      x
## 4  0.72 -1.62  sausage  True
## 2  0.23 -0.20    spam  True
## 1  0.29 -0.05   bacon False
## 0  0.70  0.32    spam  True
## 3  0.55  1.98    eggs False
df2.iloc[:3, :] # always: the first three
##      u      v      w      x
## 4  0.72 -1.62  sausage  True
## 2  0.23 -0.20    spam  True
## 1  0.29 -0.05   bacon False
```

Important We can frequently write `df.u` as a shorter version of `df.loc[:, "u"]`. This improves the readability in contexts such as:

```
df.loc[(df.u >= 0.5) & (df.u <= 0.7), ["u", "w"]]
##      u      w
## 0  0.70  spam
## 3  0.55  eggs
```

This accessor is, sadly, not universal. We can verify this by considering a data frame with a column named, e.g., `mean`: it clashes with the built-in method. As a workaround, we should either use `loc[...]` or rename the column, for instance, like `Mean` or `MEAN`.

Exercise 10.9 Use `pandas.DataFrame.drop` to select all columns except `v` in `df`.

Exercise 10.10 Use `pandas.Series.isin` (amongst others) to select all rows with `spam` and `bacon` on the `df`'s menu.

Exercise 10.11 In the `tips`⁹ dataset, select data on male customers where the total bills were in the `[10, 20]` interval. Also, select Saturday and Sunday records where the tips were greater than \$5.

10.5.2 Adding rows and columns

`loc[...]` can also be used to add new columns to an existing data frame:

```
df.loc[:, "y"] = df.loc[:, "u"]**2 # or df.loc[:, "y"] = df.u**2
df
##      u      v      w      x      y
## 0  0.70  0.32    spam  True  0.4900
## 1  0.29 -0.05   bacon False  0.0841
## 2  0.23 -0.20    spam  True  0.0529
```

(continues on next page)

⁹ <https://github.com/gagolews/teaching-data/raw/master/other/tips.csv>

(continued from previous page)

```
## 3  0.55  1.98    eggs  False  0.3025
## 4  0.72 -1.62  sausage   True  0.5184
```

Important Notation like “`df.new_column = ...`” does not work. As we said, only `loc` and `iloc` are universal. For other accessors, this is not necessarily the case.

Exercise 10.12 Use `pandas.DataFrame.insert` to add a new column not necessarily at the end of `df`.

Exercise 10.13 Use `pandas.DataFrame.assign` to add a new column and replace an existing one with another, not necessarily of the same dtype.

Exercise 10.14 Use `pandas.DataFrame.concat` to add a few more rows to `df`.

10.5.3 Modifying items

In the current version of `pandas`, modifying particular elements gives a warning:

```
df.loc[:, "u"].iloc[0] = 7 # the same as df.u.iloc[0] = 7
## SettingWithCopyWarning:
## A value is trying to be set on a copy of a slice from a DataFrame
df.loc[:, "u"].iloc[0] # testing
## 7.0
```

To remedy this, it is best to create a copy of a column, modify it, and then to replace the old contents with the new ones.

```
u = df.loc[:, "u"].copy()
u.iloc[0] = 42 # or a whole for loop to process them all, or whatever
df.loc[:, "u"] = u
df.loc[:, "u"].iloc[0] # testing
## 42.0
```

10.5.4 Pseudorandom sampling and splitting

As a simple application of what we have covered so far, let's consider some ways to sample several rows from an existing data frame. We may need them, e.g., when our datasets are too large to fit into memory or make data analysis too slow to run. In the most rudimentary scenarios, we can use the `pandas.DataFrame.sample` method, e.g., to:

- select five rows, without replacement,
- select 20% rows, with replacement,
- rearrange all the rows.

For example:

```
body = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
                  "teaching-data/master/marek/nhanes_adult_female_bmx_2020.csv",
                  comment="#")
body.sample(5, random_state=123) # 5 rows without replacement
```

| | BMXWT | BMXHT | BMXARML | BMXLEG | BMXARMC | BMXHIP | BMXWAIST |
|---------|-------|-------|---------|--------|---------|--------|----------|
| ## 4214 | 58.4 | 156.2 | 35.2 | 34.7 | 27.2 | 99.5 | 77.5 |
| ## 3361 | 73.7 | 161.0 | 36.5 | 34.5 | 29.0 | 107.6 | 98.2 |
| ## 3759 | 61.4 | 164.6 | 37.5 | 40.4 | 26.9 | 93.5 | 84.4 |
| ## 3733 | 120.4 | 158.8 | 33.5 | 34.6 | 40.5 | 147.2 | 129.3 |
| ## 1121 | 123.5 | 157.5 | 35.5 | 29.0 | 50.5 | 143.0 | 136.4 |

Notice the `random_state` argument which controls the seed of the pseudorandom number generator: this way, we get reproducible results. Alternatively, we could call `numpy.random.seed`.

Exercise 10.15 Show how the three aforementioned scenarios can be implemented manually using `iloc[...]` and `numpy.random.permutation` or `numpy.random.choice`.

Exercise 10.16 Can `pandas.read_csv` be used to read only a random sample of rows from a CSV file?

In machine learning practice, we are used to training and evaluating machine learning models on different (mutually disjoint) subsets of the whole data frame.

For instance, [Section 12.3.3](#) mentions that we may be interested in performing the so-called *training/test split* (partitioning), where 80% (or 60% or 70%) of the randomly selected rows would constitute the first new data frame and the remaining 20% (or 40% or 30%, respectively) would go to the second one.

Given a data frame like:

```
df = body.head(10) # this is just an example
df
```

| | BMXWT | BMXHT | BMXARML | BMXLEG | BMXARMC | BMXHIP | BMXWAIST |
|------|-------|-------|---------|--------|---------|--------|----------|
| ## 0 | 97.1 | 160.2 | 34.7 | 40.8 | 35.8 | 126.1 | 117.9 |
| ## 1 | 91.1 | 152.7 | 33.5 | 33.0 | 38.5 | 125.5 | 103.1 |
| ## 2 | 73.0 | 161.2 | 37.4 | 38.0 | 31.8 | 106.2 | 92.0 |
| ## 3 | 61.7 | 157.4 | 38.0 | 34.7 | 29.0 | 101.0 | 90.5 |
| ## 4 | 55.4 | 154.6 | 34.6 | 34.0 | 28.3 | 92.5 | 73.2 |
| ## 5 | 62.0 | 144.7 | 32.5 | 34.2 | 29.8 | 106.7 | 84.8 |
| ## 6 | 66.2 | 166.5 | 37.5 | 37.6 | 32.0 | 96.3 | 95.7 |
| ## 7 | 75.9 | 154.5 | 35.4 | 37.6 | 32.7 | 107.7 | 98.7 |
| ## 8 | 77.2 | 159.2 | 38.5 | 40.5 | 35.7 | 102.0 | 97.5 |
| ## 9 | 91.6 | 174.5 | 36.1 | 45.9 | 35.2 | 121.3 | 100.3 |

one way to perform the aforementioned split is to generate a random permutation of the set of row indexes:

```
np.random.seed(123) # reproducibility matters
idx = np.random.permutation(df.shape[0])
```

(continues on next page)

(continued from previous page)

```
idx
## array([4, 0, 7, 5, 8, 3, 1, 6, 9, 2])
```

Then, we pick the first 80% for the training set:

```
k = int(df.shape[0]*0.8)
df.iloc[idx[:k], :]
```

| ## | BMXWT | BMXHT | BMXARML | BMXLEG | BMXARMC | BMXHIP | BMXWAIST |
|------|-------|-------|---------|--------|---------|--------|----------|
| ## 4 | 55.4 | 154.6 | 34.6 | 34.0 | 28.3 | 92.5 | 73.2 |
| ## 0 | 97.1 | 160.2 | 34.7 | 40.8 | 35.8 | 126.1 | 117.9 |
| ## 7 | 75.9 | 154.5 | 35.4 | 37.6 | 32.7 | 107.7 | 98.7 |
| ## 5 | 62.0 | 144.7 | 32.5 | 34.2 | 29.8 | 106.7 | 84.8 |
| ## 8 | 77.2 | 159.2 | 38.5 | 40.5 | 35.7 | 102.0 | 97.5 |
| ## 3 | 61.7 | 157.4 | 38.0 | 34.7 | 29.0 | 101.0 | 90.5 |
| ## 1 | 91.1 | 152.7 | 33.5 | 33.0 | 38.5 | 125.5 | 103.1 |
| ## 6 | 66.2 | 166.5 | 37.5 | 37.6 | 32.0 | 96.3 | 95.7 |

The remaining ones produce the test set:

```
df.iloc[idx[k:], :]
```

| ## | BMXWT | BMXHT | BMXARML | BMXLEG | BMXARMC | BMXHIP | BMXWAIST |
|------|-------|-------|---------|--------|---------|--------|----------|
| ## 9 | 91.6 | 174.5 | 36.1 | 45.9 | 35.2 | 121.3 | 100.3 |
| ## 2 | 73.0 | 161.2 | 37.4 | 38.0 | 31.8 | 106.2 | 92.0 |

Exercise 10.17 In the *wine_quality_all*¹⁰ dataset, leave out all but the white wines. Partition the resulting data frame randomly into three data frames: *wines_train* (60% of the rows), *wines_validate* (another 20% of the rows), and *wines_test* (the remaining 20%).

Exercise 10.18 Compose a function *kfold* which takes a data frame *df* and an integer *k* > 1 as arguments. Return a list of data frames resulting in stemming from randomly partitioning *df* into *k* disjoint chunks of equal (or almost equal if that is not possible) sizes.

10.5.5 Hierarchical indexes (*)

Consider a data frame with a hierarchical index:

```
np.random.seed(123)
df = pd.DataFrame(dict(
    year = np.repeat([2023, 2024, 2025], 4),
    quarter = np.tile(["Q1", "Q2", "Q3", "Q4"], 3),
    data = np.round(np.random.rand(12), 2)
)).set_index(["year", "quarter"])
df
```

| ## | | data |
|----|--------------|------|
| ## | year quarter | |
| ## | 2023 Q1 | 0.70 |

(continues on next page)

¹⁰ https://github.com/gagolews/teaching-data/raw/master/other/wine_quality_all.csv

(continued from previous page)

```
##      Q2      0.29
##      Q3      0.23
##      Q4      0.55
## 2024 Q1      0.72
##      Q2      0.42
##      Q3      0.98
##      Q4      0.68
## 2025 Q1      0.48
##      Q2      0.39
##      Q3      0.34
##      Q4      0.73
```

The index has both levels named, but this is purely for aesthetic reasons.

Indexing using `loc[...]` by default relates to the first level of the hierarchy:

```
df.loc[2023, :]
##      data
## quarter
## Q1      0.70
## Q2      0.29
## Q3      0.23
## Q4      0.55
```

Note that we selected *all* rows corresponding to a given label and dropped (!) this level of the hierarchy.

Another example:

```
df.loc[ [2023, 2025], : ]
##      data
## year quarter
## 2023 Q1      0.70
##      Q2      0.29
##      Q3      0.23
##      Q4      0.55
## 2025 Q1      0.48
##      Q2      0.39
##      Q3      0.34
##      Q4      0.73
```

To access deeper levels, we can use tuples as indexers:

```
df.loc[ (2023, "Q1"), : ]
## data      0.7
## Name: (2023, Q1), dtype: float64
df.loc[ [(2023, "Q1"), (2024, "Q3")], : ]
##      data
## year quarter
```

(continues on next page)

(continued from previous page)

```
## 2023 Q1      0.70
## 2024 Q3      0.98
```

In certain scenarios, though, it will probably be much easier to subset a hierarchical index by using `reset_index` and `set_index` creatively (together with `loc[...]` and `pandas.Series.isin`, etc.).

Let's stress again that the ``:`` operator can only be used *directly* within the square brackets. Nonetheless, we can always use the `slice` constructor to create a slice in any context:

```
df.loc[ (slice(None), ["Q1", "Q3"]), : ] # :, ["Q1", "Q3"]
##          data
## year quarter
## 2023 Q1      0.70
##      Q3      0.23
## 2024 Q1      0.72
##      Q3      0.98
## 2025 Q1      0.48
##      Q3      0.34
df.loc[ (slice(None, None, -1), slice("Q2", "Q3")), : ] # :-1, "Q2":"Q3"
##          data
## year quarter
## 2025 Q3      0.34
##      Q2      0.39
## 2024 Q3      0.98
##      Q2      0.42
## 2023 Q3      0.23
##      Q2      0.29
```

10.6 Further operations on data frames

One of the many roles of data frames is to represent tables of values for their nice presentation, e.g., in reports from data analysis or research papers. Here are some functions that can aid in their formatting.

10.6.1 Sorting

Consider another example dataset: the yearly (for 2018) average [air quality data](https://discover.data.vic.gov.au/dataset/epa-air-watch-all-sites-air-quality-hourly-averages-yearly)¹¹ in the Australian state of Victoria.

¹¹ <https://discover.data.vic.gov.au/dataset/epa-air-watch-all-sites-air-quality-hourly-averages-yearly>


```

air = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
                  "teaching-data/master/marek/air_quality_2018_means.csv",
                  comment="#")
air = (
    air.
    loc[air.param_id.isin(["BPM2.5", "NO2"]), :].
    reset_index(drop=True)
)

```

We chose two air quality parameters using `pandas.Series.isin`, which determines whether each element in a Series is enlisted in a given sequence. We could also have used `set_index` and `loc[...]` for that.

Notice that the preceding code spans many lines. We needed to enclose it in round brackets to avoid a syntax error. Alternatively, we could have used backslashes at the end of each line.

Anyway, here is the data frame:

```

air
##          sp_name param_id      value
## 0    Alphington  BPM2.5    7.848758
## 1    Alphington   NO2    9.558120
## 2    Altona North   NO2    9.467912
## 3    Churchill    BPM2.5    6.391230
## 4    Dandenong     NO2    9.800705
## 5    Footscray     BPM2.5    7.640948
## 6    Footscray     NO2    10.274531
## 7    Geelong South BPM2.5    6.502762
## 8    Geelong South   NO2    5.681722
## 9    Melbourne CBD BPM2.5    8.072998
## 10   Moe           BPM2.5    6.427079
## 11   Morwell East   BPM2.5    6.784596
## 12   Morwell South  BPM2.5    6.512849
## 13   Morwell South   NO2    5.124430
## 14   Traralgon     BPM2.5    8.024735
## 15   Traralgon     NO2    5.776333

```

`sort_values` is a convenient means to order the rows with respect to one criterion, be it numeric or categorical.

```

air.sort_values("value", ascending=False)
##          sp_name param_id      value
## 6    Footscray     NO2    10.274531
## 4    Dandenong     NO2    9.800705
## 1    Alphington     NO2    9.558120
## 2    Altona North     NO2    9.467912
## 9    Melbourne CBD  BPM2.5    8.072998
## 14   Traralgon     BPM2.5    8.024735
## 0    Alphington    BPM2.5    7.848758

```

(continues on next page)

(continued from previous page)

```
## 5      Footscray    BPM2.5    7.640948
## 11     Morwell East    BPM2.5    6.784596
## 12     Morwell South    BPM2.5    6.512849
## 7      Geelong South    BPM2.5    6.502762
## 10           Moe      BPM2.5    6.427079
## 3      Churchill    BPM2.5    6.391230
## 15     Traralgon      NO2      5.776333
## 8      Geelong South    NO2      5.681722
## 13     Morwell South    NO2      5.124430
```

It is also possible to take into account more sorting criteria:

```
air.sort_values(["param_id", "value"], ascending=[True, False])
##      sp_name param_id    value
## 9  Melbourne CBD    BPM2.5    8.072998
## 14   Traralgon    BPM2.5    8.024735
## 0   Alphington    BPM2.5    7.848758
## 5      Footscray    BPM2.5    7.640948
## 11     Morwell East    BPM2.5    6.784596
## 12     Morwell South    BPM2.5    6.512849
## 7      Geelong South    BPM2.5    6.502762
## 10           Moe      BPM2.5    6.427079
## 3      Churchill    BPM2.5    6.391230
## 6      Footscray      NO2    10.274531
## 4      Dandenong      NO2    9.800705
## 1      Alphington      NO2    9.558120
## 2      Altona North      NO2    9.467912
## 15     Traralgon      NO2    5.776333
## 8      Geelong South      NO2    5.681722
## 13     Morwell South      NO2    5.124430
```

Here, in each group of identical parameters, we get a decreasing order with respect to the value.

Exercise 10.19 Compare the ordering with respect to `param_id` and `value` vs `value` and then `param_id`.

Note (*) Lamentably, `DataFrame.sort_values` by default does not use a stable algorithm. If a data frame is sorted with respect to one criterion, and then we reorder it with respect to another one, tied observations are not guaranteed to be listed in the original order:

```
(pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
  "teaching-data/master/marek/air_quality_2018_means.csv",
  comment="#")
  .sort_values("sp_name")
  .sort_values("param_id")
  .set_index("param_id"))
```

(continues on next page)

(continued from previous page)

```
.loc[["BPM2.5", "NO2"], :]  
.reset_index()  
##   param_id      sp_name      value  
## 0   BPM2.5  Melbourne CBD  8.072998  
## 1   BPM2.5           Moe    6.427079  
## 2   BPM2.5   Footscray    7.640948  
## 3   BPM2.5  Morwell East    6.784596  
## 4   BPM2.5   Churchill    6.391230  
## 5   BPM2.5  Morwell South    6.512849  
## 6   BPM2.5   Traralgon    8.024735  
## 7   BPM2.5   Alphington    7.848758  
## 8   BPM2.5  Geelong South    6.502762  
## 9       NO2  Morwell South    5.124430  
## 10      NO2   Traralgon    5.776333  
## 11      NO2  Geelong South    5.681722  
## 12      NO2  Altona North    9.467912  
## 13      NO2   Alphington    9.558120  
## 14      NO2   Dandenong    9.800705  
## 15      NO2   Footscray   10.274531
```

We lost the ordering based on station names in the two subgroups. To switch to a mergesort-like method (timsort), we should pass `kind="stable"`.

```
(pd.read_csv("https://raw.githubusercontent.com/gagolews/" +  
"teaching-data/master/marek/air_quality_2018_means.csv",  
comment="#")  
.sort_values("sp_name")  
.sort_values("param_id", kind="stable") # !  
.set_index("param_id")  
.loc[["BPM2.5", "NO2"], :]  
.reset_index()  
##   param_id      sp_name      value  
## 0   BPM2.5   Alphington    7.848758  
## 1   BPM2.5   Churchill    6.391230  
## 2   BPM2.5   Footscray    7.640948  
## 3   BPM2.5  Geelong South    6.502762  
## 4   BPM2.5  Melbourne CBD    8.072998  
## 5   BPM2.5           Moe    6.427079  
## 6   BPM2.5  Morwell East    6.784596  
## 7   BPM2.5  Morwell South    6.512849  
## 8   BPM2.5   Traralgon    8.024735  
## 9       NO2   Alphington    9.558120  
## 10      NO2  Altona North    9.467912  
## 11      NO2   Dandenong    9.800705  
## 12      NO2   Footscray   10.274531  
## 13      NO2  Geelong South    5.681722  
## 14      NO2  Morwell South    5.124430  
## 15      NO2   Traralgon    5.776333
```

Exercise 10.20 (*) Perform identical reorderings but using only `loc[...]`, `iloc[...]`, and `numpy.argsort`.

10.6.2 Stacking and unstacking (long/tall and wide forms)

Let's discuss some further ways to transform data frames that benefit from, make sense thanks to, or are possible because they can have columns of various types.

The `air` dataset is in the *long* (*tall*) format. All measurements are *stacked* one after/below another. Such a form is convenient for data storage, especially where there are only a few recorded values but many possible combinations of levels (sparse data).

The long format might not be optimal in all data processing activities, though; compare [101]. In the part of this book on matrix processing, it was much more natural for us to have a single *observation* in each row (e.g., data for each measurement station).

We can *unstack* the `air` data frame easily:

```
air_wide = air.set_index(["sp_name", "param_id"]).unstack().loc[:, "value"]
air_wide
```

| ## param_id | BPM2.5 | NO2 |
|------------------|----------|-----------|
| ## sp_name | | |
| ## Alphington | 7.848758 | 9.558120 |
| ## Altona North | NaN | 9.467912 |
| ## Churchill | 6.391230 | NaN |
| ## Dandenong | NaN | 9.800705 |
| ## Footscray | 7.640948 | 10.274531 |
| ## Geelong South | 6.502762 | 5.681722 |
| ## Melbourne CBD | 8.072998 | NaN |
| ## Moe | 6.427079 | NaN |
| ## Morwell East | 6.784596 | NaN |
| ## Morwell South | 6.512849 | 5.124430 |
| ## Traralgon | 8.024735 | 5.776333 |

This is the so-called *wide* format.

A missing value is denoted by NaN (not-a-number); see [Section 15.1](#) for more details. Interestingly, we obtained a hierarchical index in the columns (sic!) slot. Hence, to drop the last level of the hierarchy, we had to add the `loc[...]` part. Also notice that the index and columns slots are named.

The other way around, we can use the `stack` method:

```
air_wide.T.rename_axis(index="location", columns="param").\
    stack().rename("value").reset_index()
```

| ## | location | param | value |
|------|----------|---------------|----------|
| ## 0 | BPM2.5 | Alphington | 7.848758 |
| ## 1 | BPM2.5 | Churchill | 6.391230 |
| ## 2 | BPM2.5 | Footscray | 7.640948 |
| ## 3 | BPM2.5 | Geelong South | 6.502762 |

(continues on next page)

(continued from previous page)

```
## 4  BPM2.5  Melbourne CBD  8.072998
## 5  BPM2.5           Moe    6.427079
## 6  BPM2.5  Morwell East  6.784596
## 7  BPM2.5  Morwell South 6.512849
## 8  BPM2.5    Traralgon  8.024735
## 9      NO2    Alphington 9.558120
## 10     NO2    Altona North 9.467912
## 11     NO2    Dandenong  9.800705
## 12     NO2    Footscray 10.274531
## 13     NO2    Geelong South 5.681722
## 14     NO2    Morwell South 5.124430
## 15     NO2    Traralgon  5.776333
```

We used the data frame transpose (T) to get a location-major order (less boring an outcome in this context). Missing values are gone now. We do not need them anymore. Nevertheless, passing `dropna=False` would help us identify the combinations of location and param for which the readings are not provided.

10.6.3 Joining (merging)

In database design, it is common to normalise the datasets. We do this to avoid the duplication of information and pathologies stemming from them (e.g., [21]).

Example 10.21 *The air quality parameters are separately described in another data frame:*

```
param = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
                    "teaching-data/master/marek/air_quality_2018_param.csv",
                    comment="#")
param = param.rename(dict(param_std_unit_of_measure="unit"), axis=1)
param
##   param_id          param_name  unit  param_short_name
## 0      API  Airborne particle index  none  Visibility Reduction
## 1  BPM2.5  BAM  Particles < 2.5 micron  ug/m3             PM2.5
## 2      CO          Carbon Monoxide  ppm                  CO
## 3  HPM10          Hivol PM10      ug/m3                  NaN
## 4      NO2          Nitrogen Dioxide  ppb             NO2
## 5      O3              Ozone      ppb              O3
## 6  PM10      TEOM Particles <10micron  ug/m3             PM10
## 7  PPM2.5          Partisol PM2.5  ug/m3             NaN
## 8      SO2          Sulfur Dioxide  ppb             SO2
```

We could have stored them alongside the `air` data frame, but that would be a waste of space. Also, if we wanted to modify some datum (note, e.g., the annoying double space in `param_name` for `BPM2.5`), we would have to update all the relevant records.

Instead, we can always match the records in both data frames that have the same `param_ids`, and join (merge) these datasets only when we really need this.

Let's discuss the possible join operations by studying two toy datasets:

```
A = pd.DataFrame({
    "x": ["a0", "a1", "a2", "a3"],
    "y": ["b0", "b1", "b2", "b3"]
})
A
```

| | x | y |
|---|----|----|
| 0 | a0 | b0 |
| 1 | a1 | b1 |
| 2 | a2 | b2 |
| 3 | a3 | b3 |

and:

```
B = pd.DataFrame({
    "x": ["a0", "a2", "a2", "a4"],
    "z": ["c0", "c1", "c2", "c3"]
})
B
```

| | x | z |
|---|----|----|
| 0 | a0 | c0 |
| 1 | a2 | c1 |
| 2 | a2 | c2 |
| 3 | a4 | c3 |

They both have one column somewhat *in common*, *x*.

The *inner (natural) join* returns the records that have a match in both datasets:

```
pd.merge(A, B, on="x")
```

| | x | y | z |
|---|----|----|----|
| 0 | a0 | b0 | c0 |
| 1 | a2 | b2 | c1 |
| 2 | a2 | b2 | c2 |

The *left join* of *A* with *B* guarantees to return all the records from *A*, even those which are not matched by anything in *B*.

```
pd.merge(A, B, how="left", on="x")
```

| | x | y | z |
|---|----|----|-----|
| 0 | a0 | b0 | c0 |
| 1 | a1 | b1 | NaN |
| 2 | a2 | b2 | c1 |
| 3 | a2 | b2 | c2 |
| 4 | a3 | b3 | NaN |

The *right join* of *A* with *B* is the same as the left join of *B* with *A*:

```
pd.merge(A, B, how="right", on="x")
```

(continues on next page)

(continued from previous page)

```
##      x      y      z
## 0  a0    b0    c0
## 1  a2    b2    c1
## 2  a2    b2    c2
## 3  a4   NaN    c3
```

Finally, the *full outer join* is the set-theoretic union of the left and the right join:

```
pd.merge(A, B, how="outer", on="x")
##      x      y      z
## 0  a0    b0    c0
## 1  a1    b1   NaN
## 2  a2    b2    c1
## 3  a2    b2    c2
## 4  a3    b3   NaN
## 5  a4   NaN    c3
```

Exercise 10.22 Join *air_quality_2018_value*¹² with *air_quality_2018_point*¹³, and then with *air_quality_2018_param*¹⁴.

Exercise 10.23 Normalise *air_quality_2018*¹⁵ so that you get the three separate data frames mentioned in the previous exercise (*value*, *point*, and *param*).

Exercise 10.24 (*) In the National Health and Nutrition Examination Survey, each participant is uniquely identified by their sequence number (SEQN). This token is mentioned in numerous datasets, including:

- *demographic variables*¹⁶,
- *body measures*¹⁷,
- *audiometry*¹⁸,
- and *many more*¹⁹.

Join a few chosen datasets that you find interesting.

10.6.4 Set-theoretic operations and removing duplicates

Here are two not at all disjoint sets of imaginary persons:

¹² https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018_value.csv.gz

¹³ https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018_point.csv

¹⁴ https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018_param.csv

¹⁵ https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018.csv.gz

¹⁶ https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/P_DEMO.htm

¹⁷ https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/P_BMX.htm

¹⁸ https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/AUX_J.htm

¹⁹ <https://wwwn.cdc.gov/Nchs/Nhanes/continuousnhanes/default.aspx?BeginYear=2017>

```
A = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
               "teaching-data/master/marek/some_birth_dates1.csv",
               comment="#")
```

```
A
##           Name  BirthDate
## 0  Paitoon Ornwimol  26.06.1958
## 1    Antónia Lata  20.05.1935
## 2  Bertoldo Mallozzi  17.08.1972
## 3   Nedeljko Bukv  19.12.1921
## 4   Micha Kitchen  17.09.1930
## 5  Mefodiy Shachar  01.10.1914
## 6    Paul Meckler  29.09.1968
## 7  Katarzyna Lasko  20.10.1971
## 8     Åge Trelstad  07.03.1935
## 9  Duchanee Panomyaong  19.06.1952
```

and:

```
B = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
               "teaching-data/master/marek/some_birth_dates2.csv",
               comment="#")
```

```
B
##           Name  BirthDate
## 0  Hushang Naigamwala  25.08.1991
## 1      Zhen Wei  16.11.1975
## 2   Micha Kitchen  17.09.1930
## 3   Jodoc Alwin  16.11.1969
## 4    Igor Mazał  14.05.2004
## 5  Katarzyna Lasko  20.10.1971
## 6  Duchanee Panomyaong  19.06.1952
## 7  Mefodiy Shachar  01.10.1914
## 8    Paul Meckler  29.09.1968
## 9    Noe Tae-Woong  11.07.1970
## 10   Åge Trelstad  07.03.1935
```

In both datasets, there is a single categorical column whose elements uniquely identify each record (i.e., Name). In the language of relational databases, we would call it the *primary key*. In such a case, implementing the set-theoretic operations is relatively easy, as we can refer to the `pandas.Series.isin` method.

First, $A \cap B$ (intersection), includes only the rows that are *both* in *A* and in *B*:

```
A.loc[A.Name.isin(B.Name), :]
```

```
##           Name  BirthDate
## 4   Micha Kitchen  17.09.1930
## 5  Mefodiy Shachar  01.10.1914
## 6    Paul Meckler  29.09.1968
## 7  Katarzyna Lasko  20.10.1971
## 8     Åge Trelstad  07.03.1935
## 9  Duchanee Panomyaong  19.06.1952
```


Second, $A \setminus B$ (difference), gives all the rows that are in A *but not* in B :

```
A.loc[~A.Name.isin(B.Name), :]
```

| | Name | BirthDate |
|------|-------------------|------------|
| ## 0 | Paitoon Ornwimol | 26.06.1958 |
| ## 1 | Antónia Lata | 20.05.1935 |
| ## 2 | Bertoldo Mallozzi | 17.08.1972 |
| ## 3 | Nedeljko Bukv | 19.12.1921 |

Third, $A \cup B$ (union), returns the rows that exist in A or are in B :

```
pd.concat((A, B.loc[~B.Name.isin(A.Name), :]))
```

| | Name | BirthDate |
|------|---------------------|------------|
| ## 0 | Paitoon Ornwimol | 26.06.1958 |
| ## 1 | Antónia Lata | 20.05.1935 |
| ## 2 | Bertoldo Mallozzi | 17.08.1972 |
| ## 3 | Nedeljko Bukv | 19.12.1921 |
| ## 4 | Micha Kitchen | 17.09.1930 |
| ## 5 | Mefodiy Shachar | 01.10.1914 |
| ## 6 | Paul Meckler | 29.09.1968 |
| ## 7 | Katarzyna Lasko | 20.10.1971 |
| ## 8 | Åge Trelstad | 07.03.1935 |
| ## 9 | Duchanee Panomyaong | 19.06.1952 |
| ## 0 | Hushang Naigamwala | 25.08.1991 |
| ## 1 | Zhen Wei | 16.11.1975 |
| ## 3 | Jodoc Alwin | 16.11.1969 |
| ## 4 | Igor Mazał | 14.05.2004 |
| ## 9 | Noe Tae-Woong | 11.07.1970 |

There are no duplicate rows in any of the above outputs.

Exercise 10.25 Determine $(A \cup B) \setminus (A \cap B) = (A \setminus B) \cup (B \setminus A)$ (symmetric difference).

Exercise 10.26 (*) Determine the union, intersection, and difference of the *wine_sample1*²⁰ and *wine_sample2*²¹ datasets, where there is no column uniquely identifying the observations. Hint: consider using `pandas.concat` and `pandas.DataFrame.duplicated` or `pandas.DataFrame.drop_duplicates`.

10.6.5 ...and (too) many more

Looking at the list of methods for the `DataFrame` and `Series` classes in the `pandas` package's [documentation](#)²², we can see that they are abundant. Together with the object-orientated syntax, we will often find ourselves appreciating the high readability of even complex operation chains such as `data.drop_duplicates().groupby(["year", "month"]).mean().reset_index()`; see [Chapter 12](#).

Nevertheless, the methods are probably too plentiful to our taste. Their developers

²⁰ https://github.com/gagolews/teaching-data/raw/master/other/wine_sample1.csv

²¹ https://github.com/gagolews/teaching-data/raw/master/other/wine_sample2.csv

²² <https://pandas.pydata.org/pandas-docs/stable/reference/index.html>

were overgenerous. They wanted to include a list of all the possible *verbs* related to data analysis, even if they can be trivially expressed by a composition of 2-3 simpler operations from `numpy` or `scipy` instead.

As strong advocates of minimalism, we would rather save ourselves from being overloaded with too much new information. This is why our focus in this book is on developing the most *transferable*²³ skills. Our approach is also slightly more hygienic. We do not want the reader to develop a hopeless mindset, the habit of looking everything up on the internet when faced with even the simplest kinds of problems. We have brains for a reason.

10.7 Exercises

Exercise 10.27 *How are data frames different from matrices?*

Exercise 10.28 *What are the use cases of the `name` slot in `Series` and `Index` objects?*

Exercise 10.29 *What is the purpose of `set_index` and `reset_index`?*

Exercise 10.30 *Why learning `numpy` is crucial when someone wants to become a proficient user of `pandas`?*

Exercise 10.31 *What is the difference between `iloc[...]` and `loc[...]`?*

Exercise 10.32 *Why applying the index operator `[...]` directly on a `Series` or `DataFrame` object is discouraged?*

Exercise 10.33 *What is the difference between `index`, `Index`, and `columns`?*

Exercise 10.34 *How can we compute the arithmetic mean and median of all the numeric columns in a data frame, using a single line of code?*

Exercise 10.35 *What is a training/test split and how to perform it using `numpy` and `pandas`?*

Exercise 10.36 *What is the difference between stacking and unstacking? Which one yields a wide (as opposed to long) format?*

Exercise 10.37 *Name different data frame join (merge) operations and explain how they work.*

Exercise 10.38 *How does sorting with respect to more than one criterion work?*

Exercise 10.39 *Name the set-theoretic operations on data frames.*

²³ This is also in line with the observation that Python with `pandas` is not the only environment where we can work with data frames; e.g., base R and Julia with `DataFrame.jl` support that too.

Handling categorical data

So far, we have been mostly dealing with *quantitative* (numeric) data, on which we were able to apply various mathematical operations, such as computing the arithmetic mean or taking the square thereof. Naturally, not every transformation must always make sense in every context (e.g., multiplying temperatures – what does it mean when we say that it is twice as hot today as compared to yesterday?), but still, the possibilities were plenty.

Qualitative data (also known as categorical data, factors, or enumerated types) such as eye colour, blood type, or a flag whether a patient is ill, on the other hand, take a small number of unique values. They support an extremely limited set of admissible operations. Namely, we can only determine whether two entities are equal or not.

In datasets involving many features (Chapter 12), categorical variables are often used for observation *grouping* (e.g., so that we can compute the best and average time for marathoners in each age category or draw box plots for finish times of men and women separately). Also, they may serve as target variables in statistical classification tasks (e.g., so that we can determine if an email is “spam” or “not spam”).

11.1 Representing and generating categorical data

Common ways to represent a categorical variable with l distinct levels $\{L_1, L_2, \dots, L_l\}$ is by storing it as:

- a vector of strings,
- a vector of integers between 0 (inclusive) and l (exclusive¹).

These two are easily interchangeable.

For $l = 2$ (binary data), another convenient representation is by means of logical vectors. This can be extended to a so-called one-hot encoded representation using a logical vector of length l .

¹ This coincides with the possible indexes into an array of length l . In some other languages, e.g., R, we would use integers between 1 and l (inclusive). Nevertheless, a dataset creator is free to encode the labels however they want. For example, DMBORN4 in NHANES has: 1 (born in 50 US states or Washington, DC), 2 (others), 77 (refused to answer), and 99 (do not know).

Consider the data on the original whereabouts of the top 16 marathoners (the 37th Warsaw Marathon dataset):

```
marathon = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/37_pzu_warsaw_marathon_simplified.csv",
    comment="#")
cntrs = np.array(marathon.country, dtype="str")
cntrs16 = cntrs[:16]
cntrs16
## array(['KE', 'KE', 'KE', 'ET', 'KE', 'KE', 'ET', 'MA', 'PL', 'PL', 'IL',
##        'PL', 'KE', 'KE', 'PL', 'PL'], dtype='<U2')
```

These are two-letter ISO 3166 country codes encoded as strings (notice the `dtype="str"` argument).

Calling `pandas.unique` determines the set of distinct categories:

```
cat_cntrs16 = pd.unique(cntrs16)
cat_cntrs16
## array(['KE', 'ET', 'MA', 'PL', 'IL'], dtype='<U2')
```

Hence, `cntrs16` is a categorical vector of length $n = 16$ (`len(cntrs16)`) with data assuming one of $l = 5$ different levels (`len(cat_cntrs16)`).

Note We could have also used `numpy.unique` (Section 5.5.3) but it would sort the distinct values lexicographically. In other words, they would *not* be listed in the order of appearance.

11.1.1 Encoding and decoding factors

To *encode* a label vector using a set of consecutive nonnegative integers, we can call `pandas.factorize`:

```
codes_cntrs16, cat_cntrs16 = pd.factorize(cntrs16) # sort=False
cat_cntrs16
## array(['KE', 'ET', 'MA', 'PL', 'IL'], dtype='<U2')
codes_cntrs16
## array([0, 0, 0, 1, 0, 0, 1, 2, 3, 3, 4, 3, 0, 0, 3, 3])
```

The code sequence 0, 0, 0, 1, ... corresponds to the first, the first, the first, the second, ... level in `cat_cntrs16`, i.e., Kenya, Kenya, Kenya, Ethiopia,

Important When we represent categorical data using numeric codes, it is possible to introduce non-occurring levels. Such information can be useful, e.g., we could explicitly indicate that there were no runners from Australia in the top 16.

Even though we can *represent* categorical variables using a set of integers, it does

not mean that they become instances of a quantitative type. Arithmetic operations thereon do not really make sense.

The values between 0 (inclusive) and 5 (exclusive) can be used to index a given array of length $l = 5$. As a consequence, to *decode* our factor, we can call:

```
cat_cntrs16[codes_cntrs16]
## array(['KE', 'KE', 'KE', 'ET', 'KE', 'KE', 'ET', 'MA', 'PL', 'PL', 'IL',
##        'PL', 'KE', 'KE', 'PL', 'PL'], dtype='<U2')
```

We can use any other set of labels now:

```
np.array(["Kenya", "Ethiopia", "Morocco", "Poland", "Israel"])[codes_cntrs16]
## array(['Kenya', 'Kenya', 'Kenya', 'Ethiopia', 'Kenya', 'Kenya',
##        'Ethiopia', 'Morocco', 'Poland', 'Poland', 'Israel', 'Poland',
##        'Kenya', 'Kenya', 'Poland', 'Poland'], dtype='<U8')
```

It is an instance of the *relabelling* of a categorical variable.

Exercise 11.1 (**) *Here is a way of recoding a variable, i.e., changing the order of labels and permuting the numeric codes:*

```
new_codes = np.array([3, 0, 2, 4, 1]) # an example permutation of labels
new_cat_cntrs16 = cat_cntrs16[new_codes]
new_cat_cntrs16
## array(['PL', 'KE', 'MA', 'IL', 'ET'], dtype='<U2')
```

Then we make use of the fact that **numpy.argsort** applied on a vector representing a permutation, determines its very inverse:

```
new_codes_cntrs16 = np.argsort(new_codes)[codes_cntrs16]
new_codes_cntrs16
## array([1, 1, 1, 4, 1, 1, 4, 2, 0, 0, 3, 0, 1, 1, 0, 0])
```

Verification:

```
np.all(cntrs16 == new_cat_cntrs16[new_codes_cntrs16])
## True
```

Exercise 11.2 (**) *Determine the set of unique values in `cntrs16` in the order of appearance (and not sorted lexicographically), but without using `pandas.unique` nor `pandas.factorize`. Then, encode `cntrs16` using this level set.*

Hint: check out the `return_index` argument to `numpy.unique` and `numpy.searchsorted`.

Furthermore, **pandas** includes² a special dtype for storing categorical data. Namely, we can write:

² https://pandas.pydata.org/pandas-docs/stable/user_guide/categorical.html

```
cntrs16_series = pd.Series(cntrs16, dtype="category")
```

or, equivalently:

```
cntrs16_series = pd.Series(cntrs16).astype("category")
```

These two yield a Series object displayed as if it was represented using string labels:

```
cntrs16_series.head() # preview
## 0    KE
## 1    KE
## 2    KE
## 3    ET
## 4    KE
## dtype: category
## Categories (5, object): ['ET', 'IL', 'KE', 'MA', 'PL']
```

Instead, it is encoded using the aforementioned numeric representation:

```
np.array(cntrs16_series.cat.codes)
## array([2, 2, 2, 0, 2, 2, 0, 3, 4, 4, 1, 4, 2, 2, 4, 4], dtype=int8)
cntrs16_series.cat.categories
## Index(['ET', 'IL', 'KE', 'MA', 'PL'], dtype='object')
```

This time the labels are sorted lexicographically.

Most often we will be storing categorical data in data frames as ordinary strings, unless a relabelling on the fly is required:

```
(marathon.iloc[:16, :].country.astype("category")
 .cat.reorder_categories(
     ["KE", "IL", "MA", "ET", "PL"]
 )
 .cat.rename_categories(
     ["Kenya", "Israel", "Morocco", "Ethiopia", "Poland"]
 ).astype("str")
).head()
## 0    Kenya
## 1    Kenya
## 2    Kenya
## 3    Ethiopia
## 4    Kenya
## Name: country, dtype: object
```

11.1.2 Binary data as logical and probability vectors

Binary data is a special case of the qualitative setting, where we only have $l = 2$ categories. For example:

- o, e.g., healthy/fail/off/non-spam/absent(...),

- 1, e.g., ill/success/on/spam/present/...).

Usually, the *interesting* or *noteworthy* category is denoted by 1.

Important When converting logical to numeric, `False` becomes 0 and `True` becomes 1. Conversely, 0 is converted to `False` and anything else (including -0.326) to `True`.

Hence, instead of working with vectors of 0s and 1s, we might equivalently be playing with logical arrays. For example:

```
np.array([True, False, True, True, False]).astype(int)
## array([1, 0, 1, 1, 0])
```

The other way around:

```
np.array([-2, -0.326, -0.000001, 0.0, 0.1, 1, 7643]).astype(bool)
## array([ True,  True,  True, False,  True,  True,  True])
```

or, equivalently:

```
np.array([-2, -0.326, -0.000001, 0.0, 0.1, 1, 7643]) != 0
## array([ True,  True,  True, False,  True,  True,  True])
```

Important It is not rare to work with vectors of probabilities, where the i -th element therein, say $p[i]$, denotes the likelihood of an observation's belonging to the class 1. Consequently, the probability of being a member of the class 0 is $1-p[i]$. In the case where we would rather work with *crisp* classes, we can simply apply the conversion ($p \geq 0.5$) to get a logical vector.

Exercise 11.3 Given a numeric vector x , create a vector of the same length as x whose i -th element is equal to "yes" if $x[i]$ is in the unit interval and to "no" otherwise. Use `numpy.where`, which can act as a vectorised version of the *if* statement.

11.1.3 One-hot encoding (*)

Let x be a vector of n integer labels in $\{0, \dots, l-1\}$. Its *one-hot*-encoded version is a 0/1 (or, equivalently, logical) matrix $\mathbf{R}$ of shape $n \times l$ such that $r_{i,j} = 1$ if and only if $x_i = j$.

For example, if $x = (0, 1, 2, 1)$ and $l = 4$, then:

$$\mathbf{R} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}.$$

One can easily verify that each row consists of one and only one 1 (the number of 1s

per one row is 1). Such a representation is adequate when solving a multiclass classification problem by means of l binary classifiers. For example, if *spam*, *bacon*, and *hot dogs* are on the menu, then *spam* is encoded as $(1, 0, 0)$, i.e., yeah-spam, nah-bacon, and nah-hot dog. We can build three binary classifiers, each narrowly specialising in sniffing one particular type of food.

Example 11.4 Write a function to one-hot encode a given categorical vector represented using character strings.

Example 11.5 Compose a function to decode a one-hot-encoded matrix.

Example 11.6 (*) We can also work with matrices like $\mathbf{P} \in [0, 1]^{n \times l}$, where $p_{i,j}$ denotes the probability of the i -th object's belonging to the j -th class. Given an example matrix of this kind, verify that in each row the probabilities sum to 1 (up to a small numeric error). Then, decode such a matrix by choosing the greatest element in each row.

11.1.4 Binning numeric data (revisited)

Numerical data can be converted to categorical via binning (quantisation). Even though this causes information loss, it may open some new possibilities. In fact, we needed binning to draw all the histograms in [Chapter 4](#). Also, reporting observation counts in each bin instead of raw data enables us to include them in printed reports (in the form of tables).

Note We are strong proponents of openness and transparency. Thus, we always encourage all entities (governments, universities, non-profits, corporations, etc.) to share raw, unabridged versions of their datasets under the terms of some open data license. This is to enable public scrutiny and to get the most out of the possibilities they can bring for benefit of the community.

Of course, sometimes the sharing of unprocessed information can violate the privacy of the subjects. In such a case, it might be worthwhile to communicate them in a binned form.

Note Rounding is a kind of binning. In particular, `numpy.round` rounds to the nearest tenths, hundredths, ..., as well as tens, hundreds, It is useful if data are naturally imprecise, and we do not want to give the impression that it is otherwise. Nonetheless, rounding can easily introduce tied observations, which are problematic on their own; see [Section 5.5.3](#).

Consider the 16 best marathon finish times (in minutes):

```
mins = np.array(marathon.mins)
mins16 = mins[:16]
mins16
```

(continues on next page)

(continued from previous page)

```
## array([129.32, 130.75, 130.97, 134.17, 134.68, 135.97, 139.88, 143.2 ,
##        145.22, 145.92, 146.83, 147.8 , 149.65, 149.88, 152.65, 152.88])
```

`numpy.searchsorted` can determine the interval where each value in `mins` falls.

```
bins = [130, 140, 150]
codes_mins16 = np.searchsorted(bins, mins16)
codes_mins16
## array([0, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 3, 3])
```

By default, the intervals are of the form $(a, b]$ (not including a , including b). The code 0 corresponds to values less than the first bin edge, whereas the code 3 represent items greater than or equal to the last boundary.

`pandas.cut` gives us another interface to the same binning method. It returns a vector-like object with `dtype="category"`, with very readable labels generated automatically (and ordered; see [Section 11.4.7](#)):

```
cut_mins16 = pd.Series(pd.cut(mins16, [-np.inf, 130, 140, 150, np.inf]))
cut_mins16.iloc[ [0, 1, 6, 7, 13, 14, 15] ].astype("str") # preview
## 0      (-inf, 130.0]
## 1      (130.0, 140.0]
## 6      (130.0, 140.0]
## 7      (140.0, 150.0]
## 13     (140.0, 150.0]
## 14     (150.0, inf]
## 15     (150.0, inf]
## dtype: object
cut_mins16.cat.categories.astype("str")
## Index(['(-inf, 130.0]', '(130.0, 140.0]', '(140.0, 150.0]',
##        '(150.0, inf]'],
##        dtype='object')
```

Example 11.7 (*) We can create a set of the corresponding categories manually, for example, as follows:

```
bins2 = np.r_[-np.inf, bins, np.inf]
np.array(
    [f"({bins2[i]}, {bins2[i+1]})" for i in range(len(bins2)-1)]
)
## array(['(-inf, 130.0]', '(130.0, 140.0]', '(140.0, 150.0]',
##        '(150.0, inf]'], dtype='<U14')
```

Exercise 11.8 (*) Check out the `numpy.histogram_bin_edges` function which tries to determine some informative interval boundaries based on a few simple heuristics. Recall that `numpy.linspace` and `numpy.geomspace` can be used for generating equidistant bin edges on linear and logarithmic scales, respectively.

11.1.5 Generating pseudorandom labels

`numpy.random.choice` creates a pseudorandom sample with categories picked with given probabilities:

```
np.random.seed(123)
np.random.choice(
    a=["spam", "bacon", "eggs", "tempeh"],
    p=[ 0.7,    0.1,   0.15,   0.05],
    replace=True,
    size=16
)
## array(['spam', 'spam', 'spam', 'spam', 'bacon', 'spam', 'tempeh', 'spam',
##        'spam', 'spam', 'spam', 'bacon', 'spam', 'spam', 'spam', 'bacon'],
##        dtype='<U6')
```

If we generate a sufficiently long vector, we will expect "spam" to occur approximately 70% of the time, and "tempeh" to be drawn in 5% of the cases, etc.

11.2 Frequency distributions

11.2.1 Counting

`pandas.Series.value_counts` creates a frequency table in the form of a Series object equipped with a readable index (element labels):

```
pd.Series(cnts16).value_counts() # sort=True, ascending=False
## KE    7
## PL    5
## ET    2
## MA    1
## IL    1
## Name: count, dtype: int64
```

By default, data are ordered with respect to the counts, decreasingly.

If we *already* have an array of integer codes between 0 and $l - 1$, `numpy.bincount` will return the number of times each code appears therein.

```
counts_cnts16 = np.bincount(codes_cnts16)
counts_cnts16
## array([7, 2, 1, 5, 1])
```

A vector of counts can easily be turned into a vector of proportions (fractions) or percentages (if we multiply them by 100):

```
counts_cntrs16 / np.sum(counts_cntrs16) * 100
## array([43.75, 12.5 ,  6.25, 31.25,  6.25])
```

Almost 31.25% of the top runners were from Poland (this marathon is held in Warsaw after all...).

Exercise 11.9 Using `numpy.argsort`, sort `counts_cntrs16` increasingly together with the corresponding items in `cat_cntrs16`.

11.2.2 Two-way contingency tables: Factor combinations

Some datasets may bear many categorical columns, each having possibly different levels. Let's now consider all the runners in the marathon dataset:

```
marathon.loc[:, "age"] = marathon.category.str.slice(1) # first two chars
marathon.loc[marathon.age >= "60", "age"] = "60+" # not enough 70+ y olds
marathon = marathon.loc[:, ["sex", "age", "country"]]
marathon.head()
##   sex age country
## 0  M  20      KE
## 1  M  20      KE
## 2  M  20      KE
## 3  M  20      ET
## 4  M  30      KE
```

The three columns are: sex, age (in 10-year brackets), and country. We can, of course, analyse the data distribution in each column individually, but this we leave as an exercise. Instead, we note that some interesting patterns might also arise when we study the *combinations* of levels of different variables.

Here are the levels of the sex and age variables:

```
pd.unique(marathon.sex)
## array(['M', 'F'], dtype=object)
pd.unique(marathon.age)
## array(['20', '30', '50', '40', '60+'], dtype=object)
```

We have $2 \cdot 5 = 10$ combinations thereof. We can use `pandas.DataFrame.value_counts` to determine the number of observations at each two-dimensional level:

```
counts2 = marathon.loc[:, ["sex", "age"]].value_counts()
counts2
## sex age
## M   30  2200
##     40  1708
##     20   879
##     50   541
## F   30   449
```

(continues on next page)

(continued from previous page)

```
##      40      262
##      20      240
## M    60+      170
## F     50       43
##      60+       19
## Name: count, dtype: int64
```

These can be converted to a two-way *contingency table*, which is a matrix that gives the number of occurrences of each pair of values from the two factors:

```
V = counts2.unstack(fill_value=0)
V
## age    20    30    40    50    60+
## sex
## F     240   449   262   43    19
## M     879  2200  1708  541   170
```

For example, there were 19 women aged at least 60 amongst the marathoners. Jolly good.

The *marginal* (one-dimensional) frequency distributions can be recreated by computing the rowwise and columnwise sums of V:

```
np.sum(V, axis=1)
## sex
## F    1013
## M    5498
## dtype: int64
np.sum(V, axis=0)
## age
## 20    1119
## 30    2649
## 40    1970
## 50     584
## 60+    189
## dtype: int64
```

11.2.3 Combinations of even more factors

`pandas.DataFrame.value_counts` can also be used with a combination of more than two categorical variables:

```
counts3 = (marathon
    .loc[
        marathon.country.isin(["PL", "UA", "SK"]),
        ["country", "sex", "age"]
    ]
    .value_counts())
```

(continues on next page)

(continued from previous page)

```

.rename("count")
.reset_index()
)
counts3
##    country sex  age  count
## 0      PL   M   30  2081
## 1      PL   M   40  1593
## 2      PL   M   20   824
## 3      PL   M   50   475
## 4      PL   F   30   422
## 5      PL   F   40   248
## 6      PL   F   20   222
## 7      PL   M  60+   134
## 8      PL   F   50    26
## 9      PL   F  60+     8
## 10     UA   M   30     8
## 11     UA   M   20     8
## 12     UA   M   50     3
## 13     UA   F   30     2
## 14     UA   M   40     2
## 15     SK   M  60+     1
## 16     SK   F   50     1
## 17     SK   M   40     1

```

The display is in the *long* format (compare [Section 10.6.2](#)) because we cannot nicely print a three-dimensional array. Yet, we can always partially unstack the dataset, for aesthetic reasons:

```

counts3.set_index(["country", "sex", "age"]).unstack()
##              count
## age              20      30      40      50      60+
## country sex
## PL      F    222.0   422.0   248.0   26.0    8.0
##          M    824.0  2081.0  1593.0  475.0   134.0
## SK      F      NaN     NaN     NaN     1.0    NaN
##          M      NaN     NaN     1.0    NaN     1.0
## UA      F      NaN     2.0     NaN     NaN    NaN
##          M     8.0     8.0     2.0     3.0    NaN

```

Let's again appreciate how versatile the concept of data frames is. Not only can we represent data to be investigated (one row per observation, variables possibly of mixed types) but also we can store the results of such analyses (neatly formatted tables).

11.3 Visualising factors

Methods for visualising categorical data are by no means fascinating (unless we use them as grouping variables in more complex datasets, but this is a topic that we cover in [Chapter 12](#)).

11.3.1 Bar plots

Example data:

```
x = (marathon.age.astype("category")
     .cat.reorder_categories(["20", "30", "40", "50", "60+"])
     .value_counts(sort=False)
)
x
## age
## 20      1119
## 30      2649
## 40      1970
## 50        584
## 60+       189
## Name: count, dtype: int64
```

Bar plots are self-explanatory and hence will do the trick most of the time; see [Figure 11.1](#).

```
ind = np.arange(len(x)) # 0, 1, 2, 3, 4
plt.bar(ind, height=x, color="lightgray", edgecolor="black", alpha=0.8)
plt.xticks(ind, x.index)
plt.show()
```

The `ind` vector gives the x-coordinates of the bars; here: consecutive integers. By calling `matplotlib.pyplot.xticks` we assign them readable labels.

Exercise 11.10 Draw a bar plot for the five most prevalent foreign (i.e., excluding Polish) marathoners' original whereabouts. Add a bar that represents "all other" countries. Depict percentages instead of counts, so that the total bar height is 100%. Assign a different colour to each bar.

A bar plot is a versatile tool for visualising the counts also in the two-variable case; see [Figure 11.2](#). Let's now use a pleasant wrapper around `matplotlib.pyplot.bar` offered by the statistical data visualisation package called [seaborn](#)³ [99] (written by Michael Waskom).

³ <https://seaborn.pydata.org/>




Figure 11.1. Runners' age categories.

```
import seaborn as sns
v = (marathon.loc[:, ["sex", "age"]].value_counts(sort=False)
     .rename("count").reset_index()
    )
sns.barplot(x="age", hue="sex", y="count", data=v)
plt.show()
```




Figure 11.2. Number of runners by age category and sex.

Note It is customary to call a single function from `seaborn` and then perform a series of additional calls to `matplotlib` to tweak the display details. We should remember that the former uses the latter to achieve its goals, not vice versa. `seaborn` is particularly convenient for plotting grouped data.

Exercise 11.11 (*) Draw a similar chart using `matplotlib.pyplot.bar`.

Exercise 11.12 (**) Create a stacked bar plot similar to the one in Figure 11.3, where we have horizontal bars for data that have been normalised so that, for each sex, their sum is 100%.




Figure 11.3. An example stacked bar plot: age distribution for different sexes amongst all the runners.

11.3.2 Political marketing and statistics

Even such a simple chart as bar plot can be manipulated. In the second round of the 2025 Polish presidential elections, Karol Nawrocki won against Rafał Trzaskowski. The results might have been presented by some right-wing outlets in a way depicted in Figure 11.4.

```
plt.subplot(xlim=(0, 3), ylim=(49, 51), ylabel="%")
plt.bar([1, 2], height=[50.89, 49.11], width=0.25,
        color="lightgray", edgecolor="black", alpha=0.8)
plt.xticks([1, 2], ["Nawrocki", "Trzaskowski"])
plt.show()
```

On the other hand, a media conglomerate from the opposite side of the political spectrum could have reported them like in Figure 11.5.




Figure 11.4. Flawless victory

```
plt.bar([1, 2], height=[50.89, 49.11], width=0.25,
        color="lightgray", edgecolor="black", alpha=0.8)
plt.xticks([1, 2], ["Nawrocki", "Trzaskowski"])
plt.ylabel("%")
plt.xlim(0, 3)
plt.ylim(0, 250)
plt.yticks([0, 100])
plt.show()
```

Important We must always read the axis tick marks. Also, when drawing bar plots, we must never trick the reader for this is unethical; compare Rule#9. More issues in statistical deception are explored, e.g., in [98].

11.3.3 .

We are definitely not going to discuss the infamous pie charts because their use in data analysis has been widely criticised for a long time (it is difficult to judge the ratios of areas of their slices). Do not draw them. Ever. Good morning.

11.3.4 Pareto charts (*)

It is often the case that most instances of something's happening (usually 70–90%) are due to only a few causes (10–30%). This is known as the *Pareto principle* (with 80% vs 20% being an often cited rule of thumb).




Figure 11.5. So close

Example 11.13 *Chapter 6 modelled the US cities' population dataset using the Pareto distribution (the very same Pareto, but a different, yet related mathematical object). We discovered that only c. 14% of the settlements (those with 10 000 or more inhabitants) are home to as much as 84% of the population. Hence, we may say that this data domain follows the Pareto rule.*

Here is a [dataset](#)⁴ fabricated by the Clinical Excellence Commission in New South Wales, Australia, listing the most frequent causes of medication errors:

```
cat_med = np.array([
    "Unauthorised drug", "Wrong IV rate", "Wrong patient", "Dose missed",
    "Underdose", "Wrong calculation", "Wrong route", "Wrong drug",
    "Wrong time", "Technique error", "Duplicated drugs", "Overdose"
])
counts_med = np.array([1, 4, 53, 92, 7, 16, 27, 76, 83, 3, 9, 59])
np.sum(counts_med) # total number of medication errors
## 430
```

Let's display the dataset ordered with respect to the counts, decreasingly:

```
med = pd.Series(counts_med, index=cat_med).sort_values(ascending=False)
med
## Dose missed          92
## Wrong time           83
## Wrong drug           76
## Overdose             59
## Wrong patient        53
```

(continues on next page)

⁴ <https://www.cec.health.nsw.gov.au/CEC-Academy/quality-improvement-tools/pareto-charts>

(continued from previous page)

```
## Wrong route      27
## Wrong calculation 16
## Duplicated drugs  9
## Underdose        7
## Wrong IV rate     4
## Technique error   3
## Unauthorised drug 1
## dtype: int64
```

Pareto charts may aid in visualising the datasets where the Pareto principle is likely to hold, at least approximately. They include bar plots with some extras:

- bars are listed in decreasing order,
- the cumulative percentage curve is added.

The plotting of the Pareto chart is a little tricky because it involves using two different Y axes (as usual, fine-tuning the figure and studying the manual of the `matplotlib` package is left as an exercise.)

```
x = np.arange(len(med)) # 0, 1, 2, ...
p = 100.0*med/np.sum(med) # percentages

fig, ax1 = plt.subplots()
ax1.set_xticks(x-0.5, med.index, rotation=60)
ax1.set_ylabel("%")
ax1.bar(x, height=p, color="lightgray", edgecolor="black")

ax2 = ax1.twinx() # creates a new coordinate system with a shared x-axis
ax2.plot(x, np.cumsum(p), "ro-")
ax2.grid(visible=False)
ax2.set_ylabel("cumulative %")

fig.tight_layout()
plt.show()
```

Figure 11.6 shows that the first five causes (less than 40%) correspond to c. 85% of the medication errors. More precisely, the cumulative probabilities are:

```
med.cumsum()/np.sum(med)
## Dose missed      0.213953
## Wrong time       0.406977
## Wrong drug       0.583721
## Overdose         0.720930
## Wrong patient    0.844186
## Wrong route      0.906977
## Wrong calculation 0.944186
## Duplicated drugs  0.965116
## Underdose        0.981395
```

(continues on next page)




Figure 11.6. An example Pareto chart: the most frequent causes for medication errors.

(continued from previous page)

```
## Wrong IV rate      0.990698
## Technique error    0.997674
## Unauthorised drug  1.000000
## dtype: float64
```

Note that there is an explicit assumption here that a single error is only due to a single cause. Also, we presume that each medication error has a similar degree of severity.

Policymakers and quality controllers often rely on similar simplifications. They most probably are going to be addressing only the top causes. If we ever wondered why some processes (mal)function the way they do, foregoing is a hint. Inventing something more effective yet so simple at the same time requires much more effort.

It would be also nice to report the number of cases where no mistakes are made and the cases where errors are insignificant. Healthcare workers are doing a wonderful job for our communities, especially in the public system. Why add to their stress?

11.3.5 Heat maps

Two-way contingency tables can be depicted by means of a heat map, where each count influences the corresponding cell's colour intensity; see Figure 11.7.

```
V = marathon.loc[:, ["sex", "age"]].value_counts().unstack(fill_value=0)
sns.heatmap(V, annot=True, fmt="d", cmap=plt.colormaps.get_cmap("copper"))
plt.show()
```

As an exercise, draw a similar heat map using `matplotlib.pyplot.imshow`.




Figure 11.7. A heat map for the marathoners' sex and age category.

11.4 Aggregating and comparing factors

11.4.1 Mode

The only operation on categorical data on which we can rely is counting.

```
counts = marathon.country.value_counts()
counts.head()
## country
## PL    6033
## GB     71
## DE     38
## FR     33
## SE     30
## Name: count, dtype: int64
```

Therefore, as far as qualitative data aggregation is concerned, what we are left with is the *mode*, i.e., the most frequently occurring value.

```
counts.index[0] # counts are sorted
## 'PL'
```

Important A mode might be ambiguous.

It turns out that amongst the fastest 22 runners (a nicely round number), there is a tie between Kenya and Poland – both meet our definition of a mode:

```
counts = marathon.country.iloc[:22].value_counts()
counts
## country
## KE      7
## PL      7
## ET      3
## IL      3
## MA      1
## MD      1
## Name: count, dtype: int64
```

To avoid any bias, it is always best to report all the potential mode candidates:

```
counts.loc[counts == counts.max()].index
## Index(['KE', 'PL'], dtype='object', name='country')
```

If one value is required, though, we can pick one at random (calling `numpy.random.choice`).

11.4.2 Binary data as logical vectors

Recall that we are used to representing binary data as logical vectors or, equivalently, vectors of 0s and 1s.

Perhaps the most useful arithmetic operation on logical vectors is the sum. For example:

```
np.sum(marathon.country == "PL")
## 6033
```

This gave the number of elements that are equal to "PL" because the sum of 0s and 1s is equal to the number of 1s in the sequence. Note that `(country == "PL")` is a logical vector that represents a binary categorical variable with levels: not-Poland (False) and Poland (True).

If we divide the preceding result by the length of the vector, we will get the proportion:

```
np.mean(marathon.country == "PL")
## 0.9265857779142989
```

Roughly 93% of the runners were from Poland. As this is greater than 0.5, "PL" is definitely the mode.

Exercise 11.14 What is the meaning of `numpy.all`, `numpy.any`, `numpy.min`, `numpy.max`, `numpy.cumsum`, and `numpy.cumprod` applied on logical vectors?

Note (**) Having the 0/1 (or zero/non-zero) vs False/True correspondence permits us to perform some logical operations using integer arithmetic. In mathematics, 0 is the annihilator of multiplication and the neutral element of addition, whereas 1 is the neutral element of multiplication. In particular, assuming that p and q are logical values and a and b are numeric ones, we have, what follows:

- $p+q \neq 0$ means that at least one value is True and $p+q == 0$ if and only if both are False;
- more generally, $p+q == 2$ if both elements are True, $p+q == 1$ if only one is True (we call it exclusive-or, XOR), and $p+q == 0$ if both are False;
- $p*q \neq 0$ means that both values are True and $p*q == 0$ holds whenever at least one is False;
- $1-p$ corresponds to the negation of p (changes 1 to 0 and 0 to 1);
- $p*a + (1-p)*b$ is equal to a if p is True and equal to b otherwise.

11.4.3 Pearson chi-squared test (*)

The Kolmogorov–Smirnov test described in [Section 6.2.3](#) verifies whether a given sample differs significantly from a hypothesised *continuous*⁵ distribution, i.e., it works for *numeric* data.

For binned/categorical data, we can use a classical and easy-to-understand test developed by Karl Pearson in 1900. It is supposed to judge whether the differences between the observed proportions $\hat{p}_1, \dots, \hat{p}_l$ and the theoretical ones $p_1, \dots, p_l$ are significantly large or not:

$$\begin{cases} H_0 : \hat{p}_i = p_i \text{ for all } i = 1, \dots, l & \text{(null hypothesis)} \\ H_1 : \hat{p}_i \neq p_i \text{ for some } i = 1, \dots, l & \text{(alternative hypothesis)} \end{cases}$$

Having such a test is beneficial, e.g., when the data we have at hand are based on small surveys that are supposed to serve as estimates of what might be happening in a larger population.

Going back to our political example from [Section 11.3.2](#), it turns out that one of the pre-election polls indicated that $c = 516$ out of $n = 1017$ people would vote for the first candidate. We have $\hat{p}_1 = 50.74\%$ (Nawrocki) and $\hat{p}_2 = 49.26\%$ (Trzaskowski). If we would like to test whether the observed proportions are significantly different from each other, we could test them against the theoretical distribution $p_1 = 50\%$ and $p_2 = 50\%$, which states that there is a tie between the competitors (up to a sampling error).

⁵ (*) There exists a discrete version of the Kolmogorov–Smirnov test, but it is defined in a different way than in [Section 6.2.3](#); compare [4, 18].

A natural test statistic is based on the relative squared differences:

$$\hat{T} = n \sum_{i=1}^l \frac{(\hat{p}_i - p_i)^2}{p_i}.$$

```
c, n = 516, 1017
p_observed = np.array([c, n-c]) / n
p_expected = np.array([0.5, 0.5])
T = n * np.sum( (p_observed-p_expected)**2 / p_expected )
T
## 0.2212389380530986
```

Similarly to the continuous case in [Section 6.2.3](#), we reject the null hypothesis, if:

$$\hat{T} \geq K.$$

The critical value K is computed based on the fact that, if the null hypothesis is true, $\hat{T}$ follows the χ^2 (chi-squared, hence the name of the test) distribution with $l - 1$ degrees of freedom, see `scipy.stats.chi2`.

We thus need to query the theoretical quantile function to determine the test statistic that is not exceeded in 99.9% of the trials (under the null hypothesis):

```
alpha = 0.001 # significance level
scipy.stats.chi2.ppf(1-alpha, len(p_observed)-1)
## 10.827566170662733
```

As $\hat{T} < K$ (because $0.22 < 10.83$), we cannot deem the two proportions significantly different. In other words, this poll did not indicate (at the significance level 0.1%) any of the candidates as a clear winner.

Exercise 11.15 Assuming $n = 1017$, determine the smallest c , i.e., the number of respondents claiming they would vote for Nawrocki, that leads to the rejection of the null hypothesis.

11.4.4 Two-sample Pearson chi-squared test (*)

Let's consider the data depicted in [Figure 11.3](#) and test whether the runners' age distributions differ significantly between men and women.

```
V = marathon.loc[:, ["sex", "age"]].value_counts().unstack(fill_value=0)
c1, c2 = np.array(V) # the first row, the second row
c1 # women
## array([240, 449, 262, 43, 19])
c2 # men
## array([ 879, 2200, 1708, 541, 170])
```

There are $l = 5$ age categories. First, denote the total number of observations in both groups by n' and n'' .


```
n1 = c1.sum()
n2 = c2.sum()
n1, n2
## (1013, 5498)
```

The observed proportions in the first group (females), denoted by $\hat{p}'_1, \dots, \hat{p}'_l$, are, respectively:

```
p1 = c1/n1
p1
## array([0.23692004, 0.44323791, 0.25863771, 0.04244817, 0.01875617])
```

Here are the proportions in the second group (males), $\hat{p}''_1, \dots, \hat{p}''_l$:

```
p2 = c2/n2
p2
## array([0.15987632, 0.40014551, 0.31065842, 0.09839942, 0.03092033])
```

We would like to verify whether the corresponding proportions are equal (up to some sampling error):

$$\begin{cases} H_0 : \hat{p}'_i = \hat{p}''_i \text{ for all } i = 1, \dots, l & (\text{null hypothesis}) \\ H_1 : \hat{p}'_i \neq \hat{p}''_i \text{ for some } i = 1, \dots, l & (\text{alternative hypothesis}) \end{cases}$$

In other words, we want to determine whether the categorical data in the two groups come from the same discrete probability distribution.

Taking the estimated expected proportions:

$$\bar{p}_i = \frac{n'_i \hat{p}'_i + n''_i \hat{p}''_i}{n' + n''},$$

for all $i = 1, \dots, l$, the test statistic this time is equal to:

$$\hat{T} = n' \sum_{i=1}^l \frac{(\hat{p}'_i - \bar{p}_i)^2}{\bar{p}_i} + n'' \sum_{i=1}^l \frac{(\hat{p}''_i - \bar{p}_i)^2}{\bar{p}_i},$$

which is a variation on the one-sample theme presented in the previous subsection.

```
pp = (n1*p1+n2*p2)/(n1+n2)
T = n1 * np.sum( (p1-pp)**2 / pp ) + n2 * np.sum( (p2-pp)**2 / pp )
T
## 75.31373854741857
```

If the null hypothesis is true, the test statistic *approximately* follows the χ^2 distribution with $l - 1$ degrees of freedom⁶. The critical value K is equal to:

⁶ Notice that [77] in Section 14.3 recommends l degrees of freedom, but we do not agree with this rather informal reasoning. Also, simple Monte Carlo simulations suggest that $l - 1$ is a better candidate.

```
alpha = 0.001 # significance level
scipy.stats.chi2.ppf(1-alpha, len(p1)-1)
## 18.46682695290317
```

As $\hat{T} \geq K$ (because $75.31 \geq 18.47$), we reject the null hypothesis. And so, the runners' age distribution differs across sexes (at significance level 0.1%).

11.4.5 Measuring association (*)

Let's consider the Australian Bureau of Statistics [National Health Survey 2018⁷](https://www.abs.gov.au/statistics/health/health-conditions-and-risks/national-health-survey-first-results/2017-18) data on the prevalence of certain medical conditions as a function of age. Here is the extracted contingency table:

```
l = [
    "Arthritis", "Asthma", "Back problems", "Cancer (malignant neoplasms)",
    "Chronic obstructive pulmonary disease", "Diabetes mellitus",
    "Heart, stroke and vascular disease", "Kidney disease",
    "Mental and behavioural conditions", "Osteoporosis"],
    ["15-44", "45-64", "65+"]
]
C = 1000*np.array([
    [ 360.2,    1489.0,    1772.2],
    [1069.7,     741.9,    433.7],
    [1469.6,    1513.3,    955.3],
    [  28.1,     162.7,    237.5],
    [ 103.8,     207.0,    251.9],
    [ 135.4,     427.3,    607.7],
    [  94.0,     344.4,    716.0],
    [  29.6,      67.7,    123.3],
    [2218.9,    1390.6,    725.0],
    [  36.1,     312.3,    564.7],
]).astype(int)
pd.DataFrame(C, index=l[0], columns=l[1])
```

| | 15-44 | 45-64 | 65+ |
|---------------------------------------|---------|---------|---------|
| Arthritis | 360000 | 1489000 | 1772000 |
| Asthma | 1069000 | 741000 | 433000 |
| Back problems | 1469000 | 1513000 | 955000 |
| Cancer (malignant neoplasms) | 28000 | 162000 | 237000 |
| Chronic obstructive pulmonary disease | 103000 | 207000 | 251000 |
| Diabetes mellitus | 135000 | 427000 | 607000 |
| Heart, stroke and vascular disease | 94000 | 344000 | 716000 |
| Kidney disease | 29000 | 67000 | 123000 |
| Mental and behavioural conditions | 2218000 | 1390000 | 725000 |
| Osteoporosis | 36000 | 312000 | 564000 |

Cramér's V is one of a few ways to measure the degree of association between two

⁷ <https://www.abs.gov.au/statistics/health/health-conditions-and-risks/national-health-survey-first-results/2017-18>

categorical variables. It is equal to 0 (the lowest possible value) if the two variables are independent (there is no association between them) and 1 (the highest possible value) if they are tied.

Given a two-way contingency table C with n rows and m columns and assuming that:

$$T = \sum_{i=1}^n \sum_{j=1}^m \frac{(c_{i,j} - e_{i,j})^2}{e_{i,j}},$$

where:

$$e_{i,j} = \frac{\left(\sum_{k=1}^m c_{i,k}\right) \left(\sum_{k=1}^n c_{k,j}\right)}{\sum_{i=1}^n \sum_{j=1}^m c_{i,j}},$$

the Cramér coefficient is given by:

$$V = \sqrt{\frac{T}{\min\{n-1, m-1\} \sum_{i=1}^n \sum_{j=1}^m c_{i,j}}}.$$

Here, $c_{i,j}$ gives the actually observed counts and $e_{i,j}$ denotes the number that we would expect to see if the two variables were really independent.

```
scipy.stats.contingency.association(C)
## 0.316237999724298
```

Hence, there might be a small association between age and the prevalence of certain conditions. In other words, some conditions might be more prevalent in some age groups than others.

Exercise 11.16 Compute the Cramér V using only *numpy* functions.

Example 11.17 (**) We can easily verify the hypothesis whether V does not differ significantly from 0, i.e., whether the variables are independent. Looking at T , we see that this is essentially the test statistic in Pearson's chi-squared goodness-of-fit test.

```
E = C.sum(axis=1).reshape(-1, 1) * C.sum(axis=0).reshape(1, -1) / C.sum()
T = np.sum((C-E)**2 / E)
T
## 3715440.465191512
```

If the data are really independent, T follows the chi-squared distribution $n + m - 1$. As a consequence, the critical value K is equal to:

```
alpha = 0.001 # significance level
scipy.stats.chi2.ppf(1-alpha, C.shape[0] + C.shape[1] - 1)
## 32.90949040736021
```

As T is much greater than K , we conclude (at significance level 0.1%) that the health conditions are not independent of age.

Exercise 11.18 (*) Take a look at Table 19: Comorbidity of selected chronic conditions in the *National Health Survey 2018*⁸, where we clearly see that many disorders co-occur. Visualise them on some heat maps and bar plots (including data grouped by sex and age).

11.4.6 Binned numeric data

Generally, modes are meaningless for continuous data, where repeated values are – at least theoretically – highly unlikely. It might make sense to compute them on binned data, though.

Looking at a histogram, e.g., in Figure 4.2, the mode is the interval corresponding to the highest bar (hopefully assuming there is only one). If we would like to obtain a single number, we can choose for example the middle of this interval as the mode.

For numeric data, the mode will heavily depend on the coarseness and type of binning; compare Figure 4.4 and Figure 6.8. Thus, the question “what is the most popular income” is overall a difficult one to answer.

Exercise 11.19 Compute some informative modes for the *uk_income_simulated_2020*⁹ dataset. Play around with different numbers of bins on linear and logarithmic scales and see how they affect the mode.

11.4.7 Ordinal data (*)

The case where the categories can be linearly ordered, is called *ordinal data*. For instance, Australian university grades are: F (fail) < P (pass) < C (credit) < D (distinction) < HD (high distinction), some questionnaires use Likert-type scales such as “strongly disagree” < “disagree” < “neutral” < “agree” < “strongly agree”, etc.

With a linear ordering we can go beyond the mode. Due to the existence of order statistics and observation ranks, we can also easily define sample quantiles. Nevertheless, the standard methods for resolving ties will not work: we need to be careful.

For example, the median of a sample of student grades (P, P, C, D, HD) is C, but (P, P, C, D, HD, HD) is either C or D - we can choose one at random or just report that the solution is ambiguous (C+? D-? C/D?).

Another option, of course, is to treat ordinal data as numbers (e.g., F = 0, P = 1, ..., HD = 4). In the latter example, the median would be equal to 2.5.

There are some cases, though, where the conversion of labels to consecutive integers is far from optimal. For instance, where it gives the impression that the “distance” between different levels is always equal (linear).

Exercise 11.20 (**) The *grades_results*¹⁰ dataset represents the grades (F, P, C, D, HD) of

⁸ <https://www.abs.gov.au/statistics/health/health-conditions-and-risks/national-health-survey-first-results/2017-18>

⁹ https://github.com/gagolews/teaching-data/raw/master/marek/uk_income_simulated_2020.txt

¹⁰ https://github.com/gagolews/teaching-data/raw/master/marek/grades_results.txt

100 students attending an imaginary course in an Australian university. You can load it in the form of an ordered categorical *Series* by calling:

```
grades = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/grades_results.txt", dtype="str")
grades = pd.Series(pd.Categorical(grades,
    categories=["F", "P", "C", "D", "HD"], ordered=True))
grades.value_counts() # note the order of labels
## F      30
## P      29
## C      23
## HD     22
## D      19
## Name: count, dtype: int64
```

How would you determine the average grade represented as a number between 0 and 100, taking into account that for a P you need at least 50%, C is given for $\geq 60\%$, D for $\geq 70\%$, and HD for only (!) 80% of the points. Come up with a pessimistic, optimistic, and best-shot estimate, and then compare your result to the true corresponding scores listed in the *grades_scores*¹¹ dataset.

11.5 Exercises

Exercise 11.21 Does it make sense to compute the arithmetic mean of a categorical variable?

Exercise 11.22 Name the basic use cases for categorical data.

Exercise 11.23 (*) What is a Pareto chart?

Exercise 11.24 How can we deal with the case of the mode being nonunique?

Exercise 11.25 What is the meaning of the sum and mean for binary data (logical vectors)?

Exercise 11.26 What is the meaning of `numpy.mean((x > 0) & (x < 1))`, where `x` is a numeric vector?

Exercise 11.27 List some ways to visualise multidimensional categorical data (combinations of two or more factors).

Exercise 11.28 (*) State the null hypotheses verified by the one- and two-sample chi-squared goodness-of-fit tests.

Exercise 11.29 (*) How is Cramér's V defined and what values does it take?

¹¹ https://github.com/gagolews/teaching-data/raw/master/marek/grades_scores.txt

Processing data in groups

Consider another subset of the US Centres for Disease Control and Prevention National Health and Nutrition Examination Survey, this time carrying some body measures (`P_BMX`¹) together with demographics (`P_DEMO`²).

```
nhanes = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_p_demo_bmx_2020.csv",
    comment="#")
nhanes = (
    nhanes
    .loc[
        (nhanes.DMDBORN4 <= 2) & (nhanes.RIDAGEYR >= 18),
        ["RIDAGEYR", "BMXWT", "BMXHT", "BMXBMI", "RIAGENDR", "DMDBORN4"]
    ] # age >= 18 and only US and non-US born
    .rename({
        "RIDAGEYR": "age",
        "BMXWT": "weight",
        "BMXHT": "height",
        "BMXBMI": "bmival",
        "RIAGENDR": "gender",
        "DMDBORN4": "usborn"
    }, axis=1) # rename columns
    .dropna() # remove missing values
    .reset_index(drop=True)
)
```

We consider only the adult (at least 18 years old) participants, whose country of birth (the US or not) is well-defined. Let's recode the `usborn` and `gender` variables (for readability), and introduce the BMI categories:

```
nhanes = nhanes.assign(
    usborn=( # recode usborn
        nhanes.usborn.astype("category")
        .cat.rename_categories(["yes", "no"]).astype("str")
    ),
    gender=( # recode gender
        nhanes.gender.astype("category")
        .cat.rename_categories(["male", "female"]).astype("str")
    )
)
```

(continues on next page)

¹ https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/P_BMX.htm

² https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/P_DEMO.htm

(continued from previous page)

```

),
bmicat=pd.cut(                                     # new column
    nhanes.bmival,
    bins= [0, 18.5, 25, 30, np.inf],
    labels=["underweight", "normal", "overweight", "obese"]
)
)

```

Here is a preview of this data frame:

```

nhanes.head()
##   age  weight  height  bmival  gender  usborn  bmicat
##  0   29   97.1   160.2   37.8  female    no    obese
##  1   49   98.8   182.3   29.7   male     yes  overweight
##  2   36   74.3   184.2   21.9   male     yes   normal
##  3   68  103.7   185.3   30.2   male     yes   obese
##  4   76   83.3   177.1   26.6   male     yes  overweight

```

We have a mix of categorical (gender, US born-ness, BMI category) and numerical (age, weight, height, BMI) variables. Unless we had encoded qualitative variables as integers, this would not be possible with plain matrices, at least not with a single one.

In this section, we will treat the categorical columns as grouping variables. This way, we will be able to e.g., summarise or visualise the data *in each group* separately. Suffice it to say that it is likely that data distributions vary across different factor levels. It is much like having many data frames stored in one object, e.g., the heights of women and men separately.

nhanes is thus an example of heterogeneous data at their best.

12.1 Basic methods

DataFrame and Series objects are equipped with the **groupby** methods, which assist in performing a wide range of operations in data groups defined by one or more data frame columns (compare [100]). They return objects of the classes DataFrameGroupBy and SeriesGroupby:

```

type(nhanes.groupby("gender"))
## <class 'pandas.core.groupby.generic.DataFrameGroupBy'>
type(nhanes.groupby("gender").height) # or (...)[ "height" ]
## <class 'pandas.core.groupby.generic.SeriesGroupBy'>

```

Important When browsing the list of available attributes in the **pandas** manual, it is

worth knowing that `DataFrameGroupBy` and `SeriesGroupBy` are separate types. Still, they have many methods and slots in common.

Exercise 12.1 Skim through the [documentation](#)³ of the said classes.

For example, the `pandas.DataFrameGroupBy.size` method determines the number of observations in each group:

```
nhanes.groupby("gender").size()
## gender
## female    4514
## male      4271
## dtype: int64
```

It returns an object of the type `Series`. We can also perform the grouping with respect to a combination of levels in two qualitative columns:

```
nhanes.groupby(["gender", "bmicat"], observed=False).size()
## gender bmicat
## female underweight    93
##         normal       1161
##         overweight    1245
##         obese        2015
## male   underweight    65
##         normal       1074
##         overweight    1513
##         obese        1619
## dtype: int64
```

This yields a `Series` with a hierarchical index (Section 10.1.3). Nevertheless, we can always call `reset_index` to convert it to standalone columns:

```
nhanes.groupby(
    ["gender", "bmicat"], observed=True
).size().rename("counts").reset_index()
##   gender    bmicat  counts
## 0  female underweight    93
## 1  female   normal   1161
## 2  female overweight   1245
## 3  female   obese    2015
## 4   male underweight    65
## 5   male   normal   1074
## 6   male overweight   1513
## 7   male   obese    1619
```

Take note of the `rename` part. It gave us some readable column names.

Furthermore, it is possible to group rows in a data frame using a list of any `Series`

³ <https://pandas.pydata.org/pandas-docs/stable/reference/groupby.html>

objects, i.e., not just column names in a given data frame; see [Section 16.2.3](#) for an example.

Exercise 12.2 (*) *Note the difference between `pandas.GroupBy.count` and `pandas.GroupBy.size` methods (by reading their documentation).*

12.1.1 Aggregating data in groups

The `DataFrameGroupBy` and `SeriesGroupBy` classes are equipped with several well-known aggregation functions. For example:

```
nhanes.groupby("gender").mean(numeric_only=True).reset_index()
##   gender      age      weight      height      bmvial
## 0  female  48.956580  78.351839  160.089189  30.489189
## 1   male  49.653477  88.589932  173.759541  29.243620
```

The arithmetic mean was computed only on numeric columns⁴. Alternatively, we can apply the aggregate only on specific columns:

```
nhanes.groupby("gender")[["weight", "height"]].mean().reset_index()
##   gender      weight      height
## 0  female  78.351839  160.089189
## 1   male  88.589932  173.759541
```

Another example:

```
nhanes.groupby(["gender", "bmicat"], observed=False)["height"].\
    mean().reset_index()
##   gender      bmicat      height
## 0  female  underweight  161.976344
## 1  female    normal    160.149182
## 2  female  overweight  159.573012
## 3  female     obese    160.286452
## 4   male  underweight  174.073846
## 5   male    normal    173.443855
## 6   male  overweight  173.051685
## 7   male     obese    174.617851
```

Further, the most common aggregates that we described in [Section 5.1](#) can be generated by calling the `describe` method:

```
nhanes.groupby("gender").height.describe().T
## gender      female      male
## count  4514.000000  4271.000000
```

(continues on next page)

⁴ (*) In this example, we called `pandas.GroupBy.mean`. Note that it has slightly different functionality from `pandas.DataFrame.mean` and `pandas.Series.mean`, which all needed to be implemented separately so that we can use them in complex operation chains. Still, they all call the underlying `numpy.mean` function. Object-orientated programming has its pros (more expressive syntax) and cons (sometimes more redundancy in the API design).

(continued from previous page)

```
## mean      160.089189    173.759541
## std        7.035483      7.702224
## min       131.100000    144.600000
## 25%       155.300000    168.500000
## 50%       160.000000    173.800000
## 75%       164.800000    178.900000
## max       189.300000    199.600000
```

We have applied the transpose (T) to get a more readable (“tall”) result.

If different aggregates are needed, we can call **aggregate** to apply a custom list of functions:

```
(nhanes.
  groupby("gender")["height", "weight"].
  aggregate(["mean", len, lambda x: (np.max(x)+np.min(x))/2]).
  reset_index()
)
##   gender      height      len <lambda_0>      weight      len <lambda_0>
##   female  160.089189  4514      160.2  78.351839  4514      143.45
##   male    173.759541  4271      172.1  88.589932  4271      139.70
```

Interestingly, the result’s columns slot uses a hierarchical index.

Note The column names in the output object are generated by reading the applied functions’ `__name__` slots; see, e.g., `print(np.mean.__name__)`.

```
mr = lambda x: (np.max(x)+np.min(x))/2
mr.__name__ = "midrange"
(nhanes.
  loc[:, ["gender", "height", "weight"]].
  groupby("gender").
  aggregate(["mean", mr]).
  reset_index()
)
##   gender      height      weight
##   female  160.089189  160.2  78.351839  143.45
##   male    173.759541  172.1  88.589932  139.70
```

12.1.2 Transforming data in groups

We can easily transform individual columns relative to different data groups by means of the **transform** method for GroupBy objects.

```
def std0(x, axis=None):
    return np.std(x, axis=axis, ddof=0)
std0.__name__ = "std0"

def standardise(x):
    return (x - np.mean(x, axis=0)) / std0(x, axis=0)

nhanes.loc[:, "height_std"] = (
    nhanes.
    loc[:, ["height", "gender"]].
    groupby("gender").
    transform(standardise)
)

nhanes.head()
```

| | age | weight | height | bmival | gender | usborn | bmocat | height_std |
|------|-----|--------|--------|--------|--------|--------|------------|------------|
| ## 0 | 29 | 97.1 | 160.2 | 37.8 | female | no | obese | 0.015752 |
| ## 1 | 49 | 98.8 | 182.3 | 29.7 | male | yes | overweight | 1.108960 |
| ## 2 | 36 | 74.3 | 184.2 | 21.9 | male | yes | normal | 1.355671 |
| ## 3 | 68 | 103.7 | 185.3 | 30.2 | male | yes | obese | 1.498504 |
| ## 4 | 76 | 83.3 | 177.1 | 26.6 | male | yes | overweight | 0.433751 |

The new column gives the *relative* z-scores: a woman with a relative z-score of 0 has height of 160.1 cm, whereas a man with the same z-score has height of 173.8 cm.

We can check that the means and standard deviations in both groups are equal to 0 and 1:

```
(nhanes.
    loc[:, ["gender", "height", "height_std"]].
    groupby("gender").
    aggregate(["mean", std0])
)
```

| | height | height_std |
|-----------|------------|------------|
| | mean | std0 |
| ## gender | | |
| ## female | 160.089189 | 7.034703 |
| ## male | 173.759541 | 7.701323 |

Note that we needed to use a custom function for computing the standard deviation with `ddof=0`. This is likely a bug in `pandas` that `nhanes.groupby("gender").aggregate([np.std])` somewhat passes `ddof=1` to `numpy.std`,

Exercise 12.3 Create a data frame comprised of the five tallest men and the five tallest women.

12.1.3 Manual splitting into subgroups (*)

It turns out that `GroupBy` objects and their derivatives are *iterable*; compare [Section 3.4](#). As a consequence, the grouped data frames and series can be easily processed manually in case where the built-in methods are insufficient (i.e., not so rarely).

Let's consider a small sample of our data frame.

```
grouped = (nhanes.head()
           .loc[:, ["gender", "weight", "height"]].groupby("gender")
           )
list(grouped)
## [('female',   gender weight height
##  0  female    97.1  160.2), ('male',   gender weight height
##  1  male     98.8  182.3
##  2  male     74.3  184.2
##  3  male    103.7  185.3
##  4  male     83.3  177.1)]
```

The way the output is formatted is imperfect, so we need to contemplate it for a tick. We see that when iterating through a GroupBy object, we get access to pairs giving all the levels of the grouping variable and the subsets of the input data frame corresponding to these categories.

Here is a simple example where we make use of the earlier fact:

```
for level, df in grouped:
    # level is a string label
    # df is a data frame - we can do whatever we want
    print(f"There are {df.shape[0]} subject(s) with gender='{level}'")
## There are 1 subject(s) with gender='female'.
## There are 4 subject(s) with gender='male'.
```

We see that splitting followed by manual processing of the chunks in a loop is tedious in the case where we would merely like to compute some simple aggregates. These scenarios are extremely common. No wonder why the **pandas** developers introduced a convenient interface in the form of the **pandas.DataFrame.groupby** and **pandas.Series.groupby** methods and the **DataFrameGroupBy** and **SeriesGroupby** classes. Still, for more ambitious tasks, the low-level way to perform the splitting will come in handy.

Exercise 12.4 (**) Using the manual splitting and **matplotlib.pyplot.boxplot**, draw a box-and-whisker plot of heights grouped by BMI category (four boxes side by side).

Exercise 12.5 (**) Using the manual splitting, compute the relative z-scores of the height column separately for each BMI category.

12.2 Plotting data in groups

The **seaborn** package is particularly convenient for plotting grouped data – it is highly interoperable with **pandas**.

12.2.1 Series of box plots

Figure 12.1 depicts a box plot with four boxes side by side:

```
sns.boxplot(x="bmival", y="gender", hue="usborn",
            data=nhanes, palette="Paired")
plt.show()
```




Figure 12.1. The distribution of BMIs for different genders and countries of birth.

Let's contemplate for a while how easy it is now to compare the BMI distributions in different groups. Here, we have two grouping variables, as specified by the `y` and `hue` arguments.

Exercise 12.6 Create a similar series of violin plots.

Exercise 12.7 (*) Add the average BMIs in each group to the above box plot using `matplotlib.pyplot.plot`. Check `ylim` to determine the range on the y-axis.

12.2.2 Series of bar plots

On the other hand, Figure 12.2 shows a bar plot representing a contingency table. It was obtained in a different way from that used in Chapter 11:

```
sns.barplot(
    y="counts", x="gender", hue="bmocat", palette="Paired",
    data=(
        nhanes.
        groupby(["gender", "bmocat"], observed=False).
        size().
        rename("counts").
    )
)
```

(continues on next page)

(continued from previous page)

```

    reset_index()
)
)
plt.show()

```




Figure 12.2. Number of persons for each gender and BMI category.

Exercise 12.8 Draw a similar bar plot where the bar heights sum to 100% for each gender.

Exercise 12.9 Using the two-sample chi-squared test, verify whether the BMI category distributions for men and women differ significantly from each other.

12.2.3 Semitransparent histograms

Figure 12.3 illustrates that playing with semitransparent objects can make comparisons more intuitive (the `alpha` argument).

```

sns.histplot(data=nhanes, x="weight", hue="usborn", alpha=0.33,
             element="step", stat="density", common_norm=False)
plt.show()

```

By passing `common_norm=False`, we scaled each histogram separately, so that it represents a density function (area under each curve is 1). It is the behaviour we desire when the samples are of different lengths.

12.2.4 Scatter plots with group information

Scatter plots for grouped data can display category information using points of different colours and/or styles, compare Figure 12.4.




Figure 12.3. The weight distribution of the US-born participants has a higher mean and variance.

```
sns.scatterplot(x="height", y="weight", hue="gender", style="gender",
                data=nhanes, alpha=0.2, markers=["o", "v"])
plt.show()
```




Figure 12.4. Weight vs height grouped by gender.

12.2.5 Grid (trellis) plots

Grid plot (also known as trellis, panel, conditioning, or lattice plots) are a way to visualise data separately for each factor level. All the plots share the same coordinate ranges which makes them easily comparable. For instance, [Figure 12.5](#) depicts a series of histograms of weights grouped by a combination of two categorical variables.

```
grid = sns.FacetGrid(nhanes, col="gender", row="usborn")
grid = grid.map(sns.histplot, "weight", stat="density", color="lightgray")
plt.show()
```




Figure 12.5. Distribution of weights for different genders and countries of birth.

Exercise 12.10 Pass `hue="bmicat"` additionally to `seaborn.FacetGrid`.

Important Grid plots can feature any kind of data visualisation we have discussed so far (e.g., histograms, bar plots, scatter plots).

Exercise 12.11 Draw a trellis plot with scatter plots of weight vs height grouped by BMI category and gender.

12.2.6 Kolmogorov–Smirnov test for comparing ECDFs (*)

Figure 12.6 compares the empirical cumulative distribution functions of the weight distributions for US and non-US born participants.

```
for usborn, weight in nhanes.groupby("usborn").weight:
    sns.ecdfplot(data=weight, legend=False, label=usborn)
plt.legend(title="usborn")
plt.show()
```




Figure 12.6. Empirical cumulative distribution functions of weight distributions for different birthplaces.

We have used manual splitting of the weight column into subgroups and then plotted the two ECDFs separately because a call to `seaborn.ecdfplot(data=nhanes, x="weight", hue="usborn")` does not honour our wish to use alternating line styles (most likely due to a bug).

A two-sample Kolmogorov–Smirnov test can be used to check whether two ECDFs $\hat{F}'_n$ (e.g., the weight of the US-born participants) and $\hat{F}''_m$ (e.g., the weight of non-US-born

persons) are significantly different from each other:

$$\begin{cases} H_0 : \hat{F}'_n = \hat{F}''_n & (\text{null hypothesis}) \\ H_1 : \hat{F}'_n \neq \hat{F}''_n & (\text{two-sided alternative}) \end{cases}$$

The test statistic will be a variation of the one-sample setting discussed in [Section 6.2.3](#). Namely, let:

$$\hat{D}_{n,m} = \sup_{t \in \mathbb{R}} |\hat{F}'_n(t) - \hat{F}''_m(t)|.$$

Computing it is slightly trickier than in the previous case⁵. Luckily, an appropriate procedure is available in `scipy.stats`:

```
x12 = nhanes.set_index("usborn").weight
x1 = x12.loc["yes"] # the first sample
x2 = x12.loc["no"] # the second sample
Dnm = scipy.stats.ks_2samp(x1, x2)[0]
Dnm
## 0.22068075889911914
```

Assuming significance level $\alpha = 0.001$, the critical value is approximately (for larger n and m) equal to:

$$K_{n,m} = \sqrt{-\frac{\log(\alpha/2)(n+m)}{2nm}}.$$

```
alpha = 0.001
np.sqrt(-np.log(alpha/2) * (len(x1)+len(x2)) / (2*len(x1)*len(x2)))
## 0.04607410479813944
```

As usual, we reject the null hypothesis when $\hat{D}_{n,m} \geq K_{n,m}$, which is exactly the case here (at significance level 0.1%). In other words, weights of US- and non-US-born participants differ significantly.

Important Frequentist hypothesis testing only takes into account the deviation between distributions that is explainable due to sampling effects (the assumed randomness of the data generation process). For large sample sizes, even very small deviations⁶ will be deemed *statistically significant*, but it does not mean that we consider them as *practically significant*.

For instance, we might discover that a very costly, environmentally unfriendly, and generally inconvenient for everyone upgrade leads to a process' improvement: we reject the null hypothesis stating that two distributions are equal. Nevertheless, a careful

⁵ Remember that this is an introductory course, and we are still being very generous here. We encourage the readers to upskill themselves (later, of course) not only in mathematics, but also in programming (e.g., algorithms and data structures).

⁶ Including those that are merely due to round-off errors.

inspection told us that the gains are roughly 0.5%. In such a case, it is worthwhile to apply good old common sense and refrain from implementing it.

Exercise 12.12 Compare between the ECDFs of weights of men and women who are between 18 and 25 years old. Determine whether they are significantly different.

Important Some statistical textbooks and many research papers in the social sciences (amongst many others) employ the significance level of $\alpha = 5\%$, which is often criticised as too high⁷. Many stakeholders aggressively push towards constant improvements in terms of inventing bigger, better, faster, more efficient things. In this context, larger α generates more *sensational* discoveries: it considers smaller differences as already significant. This all adds to what we call the reproducibility crisis in the empirical sciences.

We, on the other hand, claim that it is better to err on the side of being cautious. This, in the long run, is more sustainable.

12.2.7 Comparing quantiles (*)

Plotting quantiles in two samples against each other can also give us some further (informal) insight with regard to the possible distributional differences. Figure 12.7 depicts an example Q-Q plot (see also the one-sample version in Section 6.2.2), where we see that the distributions have similar shapes (points more or less lie on a straight line), but they are shifted and/or scaled (if they were, they would be on the identity line).

```
x = nhanes.weight.loc[nhanes.usborn == "yes"]
y = nhanes.weight.loc[nhanes.usborn == "no"]
xd = np.sort(x)
yd = np.sort(y)
if len(xd) > len(yd): # interpolate between quantiles in a longer sample
    xd = np.quantile(xd, np.arange(1, len(yd)+1)/(len(yd)+1))
else:
    yd = np.quantile(yd, np.arange(1, len(xd)+1)/(len(xd)+1))
plt.plot(xd, yd, "o")
plt.axline((xd[len(xd)//2], xd[len(xd)//2]), slope=1,
           linestyle=":", color="gray") # identity line
plt.xlabel(f"Sample quantiles (weight; usborn=yes)")
plt.ylabel(f"Sample quantiles (weight; usborn=no)")
plt.show()
```

Notice that we interpolated between the quantiles in a larger sample to match the length of the shorter vector.

⁷ For similar reasons, we do not introduce the notion of p-values. Most practitioners tend to misunderstand them anyway.




Figure 12.7. A two-sample Q-Q plot.

12.3 Classification tasks (*)

Consider a small sample of white, rather sweet wines from a much larger [wine quality](#)⁸ dataset.

```
wine_train = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/other/sweetwhitewine_train2.csv",
    comment="#")
wine_train.head()
##      alcohol      sugar  bad
## 0  10.625271  10.340159   0
## 1   9.066111  18.593274   1
## 2  10.806395   6.206685   0
## 3  13.432876   2.739529   0
## 4   9.578162   3.053025   0
```

We are given each wine's alcohol and residual sugar content, as well as a binary categorical variable stating whether a group of sommeliers deem a given beverage rather bad (1) or not (0). [Figure 12.8](#) reveals that subpar wines are fairly low in... alcohol and, to some extent, sugar.

```
sns.scatterplot(x="alcohol", y="sugar", data=wine_train,
    hue="bad", style="bad", markers=["o", "v"], alpha=0.5)
```

(continues on next page)

⁸ <http://archive.ics.uci.edu/ml/datasets/Wine+Quality>

(continued from previous page)

```
plt.xlabel("alcohol")
plt.ylabel("sugar")
plt.legend(title="bad")
plt.show()
```




Figure 12.8. Scatter plot for sugar vs alcohol content for white, rather sweet wines, and whether they are considered bad (1) or drinkable (0) by some experts.

Someone answer the door! We have a delivery of a few new wine bottles. Interestingly, their alcohol and sugar contents have been given on their respective labels.

```
wine_test = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/other/sweetwhitewine_test2.csv",
    comment="#").iloc[:, :-1]
wine_test.head()
##      alcohol      sugar
## 0    9.315523  10.041971
## 1   12.909232   6.814249
## 2    9.051020  12.818683
## 3    9.567601  11.091827
## 4    9.494031  12.053790
```

We would like to determine which of the wines from the test set might be not-bad without asking an expert for their opinion. In other words, we would like to exercise a *classification* task (see, e.g., [10, 50]). More formally:

Important Assume we are given a set of training points $\mathbf{X} \in \mathbb{R}^{n \times m}$ and the corresponding reference outputs $\mathbf{y} \in \{L_1, L_2, \dots, L_l\}^n$ in the form of a categorical vari-

able with l distinct levels. The aim of a *classification* algorithm is to predict what the outputs for each point from a possibly different dataset $\mathbf{X}' \in \mathbb{R}^{n' \times m}$, i.e., $\hat{\mathbf{y}}' \in \{L_1, L_2, \dots, L_l\}^{n'}$, might be.

In other words, we are asked to fill the gaps in a categorical variable. Recall that in a regression problem (Section 9.2), the reference outputs were numerical.

Exercise 12.13 Which of the following are instances of classification problems and which are regression tasks?

- Detect email spam.
- Predict a market stock price (good luck with that).
- Assess credit risk.
- Detect tumour tissues in medical images.
- Predict the time-to-recovery of cancer patients.
- Recognise smiling faces on photographs (kind of creepy).
- Detect unattended luggage in airport security camera footage.

What kind of data should you gather to tackle them?

12.3.1 K-nearest neighbour classification (*)

One of the simplest approaches to classification is based on the information about a test point's nearest neighbours living in the training sample; compare Section 8.4.4.

Fix $k \geq 1$. Namely, to classify some $\mathbf{x}' \in \mathbb{R}^m$:

1. Find the indexes $N_k(\mathbf{x}') = \{i_1, \dots, i_k\}$ of the k points from $\mathbf{X}$ closest to $\mathbf{x}'$, i.e., ones that fulfil for all $j \notin \{i_1, \dots, i_k\}$:

$$\|\mathbf{x}_{i_1} - \mathbf{x}'\| \leq \dots \leq \|\mathbf{x}_{i_k} - \mathbf{x}'\| \leq \|\mathbf{x}_{j^*} - \mathbf{x}'\|.$$

2. Classify $\mathbf{x}'$ as $\hat{y}' = \text{mode}(y_{i_1}, \dots, y_{i_k})$, i.e., assign it the label that most frequently occurs amongst its k nearest neighbours. If a mode is nonunique, resolve the ties at random.

It is thus a similar algorithm to k -nearest neighbour regression (Section 9.2.1). We only replaced the *quantitative* mean with the *qualitative* mode.

This is a variation on the theme: if you don't know what to do in a given situation, try to mimic what the majority of people around you are doing or saying. For instance, if you don't know what to think about a particular wine, discover that amongst the five most similar ones (in terms of alcohol and sugar content) three are said to be awful. Now you can claim that you don't like it because it's not sweet enough. Thanks to this, others will take you for a very refined wine taster.

Let's apply a 5-nearest neighbour classifier on the standardised version of the dataset.

As we are about to use a technique based on pairwise distances, it would be best if the variables were on the same scale. Thus, we first compute the z-scores for the training set:

```
X_train = np.array(wine_train.loc[:, ["alcohol", "sugar"]])
means = np.mean(X_train, axis=0)
sds = np.std(X_train, axis=0)
Z_train = (X_train-means)/sds
```

Then, we determine the z-scores for the test set:

```
Z_test = (np.array(wine_test.loc[:, ["alcohol", "sugar"]]) - means) / sds
```

Let's stress that we referred to the aggregates computed for the training set. This is a representative example of a situation where we cannot simply use a built-in method from **pandas**. Instead, we apply what we have learnt about **numpy**.

To make the predictions, we will use the following function:

```
def knn_class(X_test, X_train, y_train, k):
    nnis = scipy.spatial.KDTree(X_train).query(X_test, k)[1]
    nnls = y_train[nnis] # same as: y_train[nnis.reshape(-1)].reshape(-1, k)
    return scipy.stats.mode(nnls.reshape(-1, k), axis=1, keepdims=False)[0]
```

First, we fetched the indexes of each test point's nearest neighbours (amongst the points in the training set). Then, we read their corresponding labels; they are stored in a matrix with k columns. Finally, we computed the modes in each row. As a consequence, we have each point in the test set classified.

And now:

```
k = 5
y_train = np.array(wine_train.bad)
y_pred = knn_class(Z_test, Z_train, y_train, k)
y_pred[:10] # preview
## array([1, 0, 0, 1, 1, 0, 1, 0, 0, 1])
```

Note `scipy.stats.mode` does not resolve possible ties at random: e.g., the mode of $(1, 1, 1, 2, 2, 2)$ is always 1. Nevertheless, in our case, k is odd and the number of possible classes is $l = 2$, so the mode is always unique.

Figure 12.9 shows how nearest neighbour classification categorises different regions of a section of the two-dimensional plane. The greater the k , the smoother the decision boundaries. Naturally, in regions corresponding to few training points, we do not expect the classification accuracy to be acceptable⁹.

⁹ (*) As an exercise, we could author a fixed-radius classifier; compare Section 8.4.4. In sparsely populated regions, the decision might be “unknown”.


```

x1 = np.linspace(Z_train[:, 0].min(), Z_train[:, 0].max(), 100)
x2 = np.linspace(Z_train[:, 1].min(), Z_train[:, 1].max(), 100)
xg1, xg2 = np.meshgrid(x1, x2)
Xg12 = np.column_stack((xg1.reshape(-1), xg2.reshape(-1)))
ks = [5, 25]
for i in range(len(ks)):
    plt.subplot(1, len(ks), i+1)
    yg12 = knn_class(Xg12, Z_train, y_train, ks[i])
    plt.scatter(Z_train[y_train == 0, 0], Z_train[y_train == 0, 1],
                c="black", marker="o", alpha=0.5)
    plt.scatter(Z_train[y_train == 1, 0], Z_train[y_train == 1, 1],
                c="#DF536B", marker="v", alpha=0.5)
    plt.contourf(x1, x2, yg12.reshape(len(x2), len(x1)),
                 cmap="gist_heat", alpha=0.5)
    plt.title(f"$k={ks[i]}$", fontdict=dict(fontsize=10))
    plt.xlabel("alcohol")
    if i == 0: plt.ylabel("sugar")
plt.show()

```




Figure 12.9. k -nearest neighbour classification of a whole, dense, two-dimensional grid of points for different k .

Example 12.14 (*) *The same with the `scikit-learn` package:*

```

import sklearn.neighbors
knn = sklearn.neighbors.KNeighborsClassifier(k)
knn.fit(Z_train, y_train)
y_pred2 = knn.predict(Z_test)

```

We can verify that the results match by calling:

```
np.all(y_pred2 == y_pred)
## True
```

12.3.2 Assessing prediction quality (*)

It is time to reveal the truth: our test wines, it turns out, have already been assessed by some experts.

```
y_test = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
                    "teaching-data/master/other/sweetwhitewine_test2.csv",
                    comment="#")
y_test = np.array(y_test.bad)
y_test[:10] # preview
## array([1, 0, 0, 1, 0, 0, 1, 0, 1, 1])
```

The *accuracy* score is the most straightforward measure of the similarity between these true labels (denoted $\mathbf{y}' = (y'_1, \dots, y'_{n'})$) and the ones predicted by the classifier (denoted $\hat{\mathbf{y}}' = (\hat{y}'_1, \dots, \hat{y}'_{n'})$). It is defined as a ratio between the correctly classified instances and all the instances:

$$\text{Accuracy}(\mathbf{y}', \hat{\mathbf{y}}') = \frac{\sum_{i=1}^{n'} \mathbf{1}(y'_i = \hat{y}'_i)}{n'},$$

where the *indicator function* $\mathbf{1}(y'_i = \hat{y}'_i) = 1$ if and only if $y'_i = \hat{y}'_i$ and 0 otherwise. Computing it for our test sample gives:

```
np.mean(y_test == y_pred)
## 0.706
```

Thus, 71% of the wines were correctly classified with regard to their true quality. Before we get too enthusiastic, let's note that our dataset is slightly *imbalanced* in terms of the distribution of label counts:

```
pd.Series(y_test).value_counts() # contingency table
## 0    330
## 1    170
## Name: count, dtype: int64
```

It turns out that the majority of the wines (330 out of 500) in our sample are *truly* delicious. Notice that a dummy classifier which labels *all* the wines as great would have accuracy of 66%. Our *k*-nearest neighbour approach to wine quality assessment is not that usable after all.

It is therefore always beneficial to analyse the corresponding *confusion matrix*, which is a two-way contingency table summarising the correct decisions and errors we make.

```
C = pd.DataFrame(
    dict(y_pred=y_pred, y_test=y_test)
).value_counts().unstack(fill_value=0)
C
## y_test      0      1
## y_pred
## 0          272    89
## 1           58    81
```

In the binary classification case ($l = 2$) such as this one, its entries are usually referred to as (see also the table below):

- TN – the number of cases where the true $y'_i = 0$ and the predicted $\hat{y}'_i = 0$ (true negative),
- TP – the number of instances such that the true $y'_i = 1$ and the predicted $\hat{y}'_i = 1$ (true positive),
- FN – how many times the true $y'_i = 1$ but the predicted $\hat{y}'_i = 0$ (false negative),
- FP – how many times the true $y'_i = 0$ but the predicted $\hat{y}'_i = 1$ (false positive).

The terms *positive* and *negative* refer to the output predicted by a classifier, i.e., they indicate whether some $\hat{y}'_i$ is equal to 1 and 0, respectively.

Table 12.1. The different cases of true vs predicted labels in a binary classification task ($l = 2$)

| | $\hat{y}'_i = 1$ | $\hat{y}'_i = 0$ |
|------------|-------------------------------|--------------------------------|
| $y'_i = 1$ | True Positive | False Negative (Type II error) |
| $y'_i = 0$ | False Positive (Type I error) | True Negative |

Ideally, the number of false positives and false negatives should be as low as possible. The accuracy score only takes the raw number of true negatives (TN) and true positives (TP) into account:

$$\text{Accuracy}(\mathbf{y}', \hat{\mathbf{y}}') = \frac{\text{TN} + \text{TP}}{\text{TN} + \text{TP} + \text{FN} + \text{FP}}.$$

Consequently, it might not be a valid metric in imbalanced classification problems.

There are, fortunately, some more meaningful measures in the case where class 1 is less prevalent and where mispredicting it is considered more hazardous than making an inaccurate prediction with respect to class 0. After all, most will agree that it is better to be surprised by a vino mislabelled as bad, than be disappointed with a highly recommended product where we have already built some expectations around it. Further, getting a virus infection not recognised where we are genuinely sick can be more

dangerous for the people around us than being asked to stay at home with nothing but a headache.

Precision answers the question: If the classifier outputs 1, what is the probability that this is indeed true?

$$\text{Precision}(\mathbf{y}', \hat{\mathbf{y}}') = \frac{\text{TP}}{\text{TP} + \text{FP}} = \frac{\sum_{i=1}^{n'} y'_i \hat{y}'_i}{\sum_{i=1}^{n'} \hat{y}'_i}.$$

```
C = np.array(C) # convert to matrix
C[1,1]/(C[1,1]+C[1,0]) # precision
## 0.5827338129496403
np.sum(y_test*y_pred)/np.sum(y_pred) # equivalently
## 0.5827338129496403
```

When a classifier labels a vino as bad, in 58% of cases it is veritably undrinkable.

Recall (sensitivity, hit rate, or true positive rate) addresses the question: If the true class is 1, what is the probability that the classifier will detect it?

$$\text{Recall}(\mathbf{y}', \hat{\mathbf{y}}') = \frac{\text{TP}}{\text{TP} + \text{FN}} = \frac{\sum_{i=1}^{n'} y'_i \hat{y}'_i}{\sum_{i=1}^{n'} y'_i}.$$

```
C[1,1]/(C[1,1]+C[0,1]) # recall
## 0.4764705882352941
np.sum(y_test*y_pred)/np.sum(y_test) # equivalently
## 0.4764705882352941
```

Only 48% of the really bad wines will be filtered out by the classifier.

The *F measure* (or F_1 measure), is the harmonic¹⁰ mean of precision and recall in the case where we would rather have them aggregated into a single number:

$$F(\mathbf{y}', \hat{\mathbf{y}}') = \frac{1}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = \left(\frac{1}{2} \left(\text{Precision}^{-1} + \text{Recall}^{-1} \right) \right)^{-1} = \frac{\text{TP}}{\text{TP} + \frac{\text{FP} + \text{FN}}{2}}.$$

```
C[1,1]/(C[1,1]+0.5*C[0,1]+0.5*C[1,0]) # F
## 0.5242718446601942
```

Overall, we can conclude that our classifier is rather weak.

Exercise 12.15 Would you use precision or recall in the following settings:

¹⁰ (*) For any vector of nonnegative values, its minimum $\leq$ its harmonic mean $\leq$ its arithmetic mean.

- *medical diagnosis*,
- *medical screening*,
- *suggestions of potential matches in a dating app*,
- *plagiarism detection*,
- *wine recommendation?*

12.3.3 Splitting into training and test sets (*)

The training set was used as a source of knowledge about our problem domain. The k -nearest neighbour classifier is technically *model-free*. As a consequence, to generate a new prediction, we need to be able to query all the points in the database every time.

Nonetheless, most statistical/machine learning algorithms, by construction, generalise the patterns discovered in the dataset in the form of mathematical functions (often-times, very complicated ones), that are fitted by minimising some error metric. Linear regression analysis by means of the least squares approximation uses exactly this kind of approach. Logistic regression for a binary response variable would be a conceptually similar classifier, but it is beyond our introductory course.

Either way, we used a separate *test set* to verify the quality of our classifier on so-far *unobserved* data, i.e., its *predictive* capabilities. We do not want our model to fit to the training data too closely. This could lead to its being completely useless when filling the gaps between the points it was exposed to. This is like being a student who can only repeat what the teacher says, and when faced with a slightly different real-world problem, they panic and say complete gibberish.

In the preceding example, the training and test sets were created by yours truly. Normally, however, the data scientists split a single data frame into two parts themselves; see [Section 10.5.4](#). This way, they can *mimic* the situation where some *test* observations become available after the learning phase is complete.

12.3.4 Validating many models (parameter selection) (**)

In statistical modelling, there often are many *hyperparameters* that need to be tweaked. For example:

- which independent variables should be used for model building,
- what is the best way to preprocess them; e.g., which of them should be standardised,
- if an algorithm has some tunable parameters, what is their best combination; for instance, which k should we use in the k -nearest neighbours search.

At initial stages of data analysis, we usually tune them up by trial and error. Later, but this is already beyond the scope of this introductory course, we are used to exploring all the possible combinations thereof (exhaustive grid search) or making use of some local search-based heuristics (e.g., greedy optimisers such as hill climbing).

These always involve verifying the performance of *many* different classifiers, for example, 1-, 3-, 9, and 15-nearest neighbours-based ones. For each of them, we need to compute separate quality metrics, e.g., the F measures. Then, we promote the classifier which enjoys the highest score. Unfortunately, if we do it recklessly, this can lead to *overfitting*, this time with respect to the test set. The obtained metrics might be too optimistic and can poorly reflect the real performance of the solution on future data.

Assuming that our dataset carries a decent number of observations, to overcome this problem, we can perform a random *training/validation/test split*:

- *training sample* (e.g., 60% of randomly chosen rows) – for model construction,
- *validation sample* (e.g., 20%) – used to tune the hyperparameters of many classifiers and to choose the best one,
- *test (hold-out) sample* (e.g., the remaining 20%) – used to assess the goodness of fit of the best classifier.

This common sense-based approach is not limited to classification. We can validate different regression models in the same way.

Important We would like to obtain a valid estimate of a classifier's performance on previously unobserved data. For this reason, the test (hold-out) sample must neither be used in the training nor the validation phase.

Exercise 12.16 Determine the best parameter setting for the *k*-nearest neighbour classification of the *color* variable based on standardised versions of some physicochemical features (chosen columns) of wines in the *wine_quality_all*¹¹ dataset. Create a 60/20/20% dataset split. For each $k = 1, 3, 5, 7, 9$, compute the corresponding F measure on the validation test. Evaluate the quality of the best classifier on the test set.

Note (*) Instead of a training/validation/test split, we can use various *cross-validation* techniques, especially on smaller datasets. For instance, in a *5-fold cross-validation*, we split the original training set randomly into five disjoint parts: *A, B, C, D, E* (more or less of the same size). We use each combination of four chunks as training sets and the remaining part as the validation set, for which we generate the predictions and then compute, say, the F measure:

| training set | validation set | F measure |
|--------------------------|----------------|-----------|
| $B \cup C \cup D \cup E$ | <i>A</i> | F_A |
| $A \cup C \cup D \cup E$ | <i>B</i> | F_B |
| $A \cup B \cup D \cup E$ | <i>C</i> | F_C |
| $A \cup B \cup C \cup E$ | <i>D</i> | F_D |
| $A \cup B \cup C \cup D$ | <i>E</i> | F_E |

¹¹ https://github.com/gagolews/teaching-data/raw/master/other/wine_quality_all.csv

In the end, we can determine the average F measure, $(F_A + F_B + F_C + F_D + F_E)/5$, as a basis for assessing different classifiers' quality.

Once the best classifier is chosen, we can use the whole training sample to fit the final model and then consider the separate test sample to assess its quality.

Furthermore, for highly imbalanced labels, some form of *stratified sampling* might be necessary. Such problems are typically explored in more advanced courses in statistical learning.

Exercise 12.17 (**) Redo [Exercise 12.16](#), but this time maximise the F measure obtained by a 5-fold cross-validation.

12.4 Clustering tasks (*)

So far, we have been implicitly assuming that either each dataset comes from a single homogeneous distribution, or we have a categorical variable that naturally defines the groups that we can split the dataset into. Nevertheless, it might be the case that we are given a sample coming from a distribution mixture, where some subsets behave differently, but a grouping variable has not been provided at all (e.g., we have height and weight data but no information about the subjects' sexes).

Clustering methods (also known as segmentation or quantisation; see, e.g., [2, 103]) partition a dataset into groups based only on the spatial structure of the points' relative densities. In the undermentioned *k*-means method, the cluster structure is determined based on the points' proximity to *k* carefully chosen group centroids; compare [Section 8.4.2](#).

12.4.1 K-means method (*)

Fix $k \geq 2$. In the *k*-means method¹², we seek *k* pivot points, $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^m$, such that the sum of squared distances between the input points in $\mathbf{X} \in \mathbb{R}^{n \times m}$ and their closest pivots is minimised:

$$\text{minimise } \sum_{i=1}^n \min \{ \|\mathbf{x}_{i,\cdot} - \mathbf{c}_1\|^2, \|\mathbf{x}_{i,\cdot} - \mathbf{c}_2\|^2, \dots, \|\mathbf{x}_{i,\cdot} - \mathbf{c}_k\|^2 \} \quad \text{w.r.t. } \mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k.$$

Let's introduce the *label vector* $\mathbf{l}$ such that:

$$l_i = \arg \min_j \|\mathbf{x}_{i,\cdot} - \mathbf{c}_j\|^2,$$

¹² We do not have to denote the number of clusters by *k*. We could be speaking about the 2-means, 3-means, *l*-means, or *ü*-means method too. Nevertheless, some mainstream practitioners consider *k*-means as a kind of a brand name, let's thus refrain from adding to their confusion. Another widely known algorithm is called fuzzy (weighted) *c*-means [8].

i.e., it is the index of the pivot closest to $\mathbf{x}_{i,\cdot}$.

We will consider all the points $\mathbf{x}_{i,\cdot}$ with i such that $l_i = j$ as belonging to the same, j -th, *cluster* (point group). This way $\mathbf{l}$ defines a *partition* of the original dataset into k nonempty, mutually disjoint subsets.

Now, the aforementioned optimisation task can be equivalently rewritten as:

$$\text{minimise } \sum_{i=1}^n \|\mathbf{x}_{i,\cdot} - \mathbf{c}_{l_i}\|^2 \quad \text{w.r.t. } \mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k.$$

We refer to the objective function as the (total) *within-cluster sum of squares* (WCSS). This version looks easier, but it is only some false impression: l_i s depend on $\mathbf{c}_j$ s. They vary *together*. We have just made them less explicit.

Given a fixed label vector $\mathbf{l}$ representing a partition, $\mathbf{c}_j$ is the centroid (Section 8.4.2) of the points assigned thereto:

$$\mathbf{c}_j = \frac{1}{n_j} \sum_{i:l_i=j} \mathbf{x}_{i,\cdot},$$

where $n_j = |\{i : l_i = j\}|$ gives the number of i s such that $l_i = j$, i.e., the size of the j -th cluster.

Here is an example dataset (see below for a scatter plot):

```
X = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/blobs1.txt", delimiter=",")
```

We can call `scipy.cluster.vq.kmeans2` to find $k = 2$ clusters:

```
import scipy.cluster.vq
C, l = scipy.cluster.vq.kmeans2(X, 2)
```

The discovered cluster centres are stored in a matrix with k rows and m columns, i.e., the j -th row gives $\mathbf{c}_j$.

```
C
## array([[ 0.99622971,  1.052801  ],
##        [-0.90041365, -1.08411794]])
```

The label vector is:

```
l
## array([1, 1, 1, ..., 0, 0, 0], dtype=int32)
```

As usual in Python, indexing starts at 0. So for $k = 2$ we only obtain the labels 0 and 1.

Figure 12.10 depicts the two clusters together with the cluster centroids. We use `l` as a colour selector in `my_colours[l]` (this is a clever instance of the integer vector-based

indexing). It seems that we correctly discovered the very natural partitioning of this dataset into two clusters.

```
plt.scatter(X[:, 0], X[:, 1], c=np.array(["black", "#DF536B"])[l])
plt.plot(C[:, 0], C[:, 1], "yX")
plt.axis("equal")
plt.show()
```




Figure 12.10. Two clusters discovered by the k -means method. Cluster centroids are marked separately.

Here are the cluster sizes:

```
np.bincount(l) # or, e.g., pd.Series(l).value_counts()
## array([1017, 1039])
```

The label vector `l` can be added as a new column in the dataset. Here is a preview:

```
Xl = pd.DataFrame(dict(X1=X[:, 0], X2=X[:, 1], l=l))
Xl.sample(5, random_state=42) # some randomly chosen rows
##          X1          X2  l
## 184 -0.973736 -0.417269  1
## 1724  1.432034  1.392533  0
## 251 -2.407422 -0.302862  1
## 1121  2.158669 -0.000564  0
## 1486  2.060772  2.672565  0
```

We can now enjoy all the techniques for processing data in groups that we have discussed so far. In particular, computing the columnwise means gives nothing else than the above cluster centroids:

```

Xl.groupby("l").mean()
##           X1           X2
## l
## 0  0.996230  1.052801
## 1 -0.900414 -1.084118

```

The label vector `l` can be recreated by computing the distances between all the points and the centroids and then picking the indexes of the closest pivots:

```

l_test = np.argmin(scipy.spatial.distance.cdist(X, C), axis=1)
np.all(l_test == l) # verify they are identical
## True

```

Important By construction¹³, the *k*-means method can only detect clusters of convex shapes (such as Gaussian blobs).

Exercise 12.18 Perform the clustering of the *wut_isolation*¹⁴ dataset and notice how non-sensical, geometrically speaking, the returned clusters are.

Exercise 12.19 Determine a clustering of the *wut_twosplashes*¹⁵ dataset and display the results on a scatter plot. Compare them with those obtained on the standardised version of the dataset. Recall what we said about the Euclidean distance and its perception being disturbed when a plot's aspect ratio is not 1:1.

Note (*) An even simpler classifier than the *k*-nearest neighbours one builds upon the concept of the nearest centroids. Namely, it first determines the centroids (componentwise arithmetic means) of the points in each class. Then, a new point (from the test set) is assigned to the class whose centroid is the closest thereto. The implementation of such a classifier is left as a rather straightforward exercise for the reader. As an application, we recommend using it to extrapolate the results generated by the *k*-means method (for different *ks*) to previously unobserved data, e.g., all points on a dense equidistant grid.

12.4.2 Solving *k*-means is hard (*)

Alas, the *k*-means method – the identification of label vectors/cluster centres that minimise the total within-cluster sum of squares – relies on solving a computationally hard combinatorial optimisation problem (e.g., [61]). In other words, the search for the *truly* (i.e., globally) optimal solution takes, for larger *n* and *k*, an impractically long time.

As a consequence, we must rely on some approximate algorithms which all have one

¹³ (*) And its relation to Voronoi diagrams.

¹⁴ https://github.com/gagolews/teaching-data/raw/master/clustering/wut_isolation.csv

¹⁵ https://github.com/gagolews/teaching-data/raw/master/clustering/wut_twosplashes.csv

drawback in common. Namely, whatever they return can be *suboptimal*. Hence, they can constitute a possibly meaningless solution.

The documentation of `scipy.cluster.vq.kmeans2` is, of course, honest about it. It states that the method *attempts to minimise the Euclidean distance between observations and centroids*. Further, `sklearn.cluster.KMeans`, which implements a similar algorithm, mentions that the procedure is *very fast [...], but it falls in local minima. That is why it can be useful to restart it several times*.

To understand what it all means, it will be very educational to study this issue in more detail. This is because the discussed approach to clustering is not the only hard problem in data science (selecting an optimal set of independent variables with respect to AIC or BIC in linear regression is another example).

12.4.3 Lloyd algorithm (*)

Technically, there is no such thing as *the k-means algorithm*. There are many procedures, based on numerous different heuristics, that attempt to solve the *k-means problem*. Unfortunately, neither of them is perfect for it is not possible.

Perhaps the most widely known and easiest to understand method is traditionally attributed to Lloyd [64]. It is based on the fixed-point iteration. For a given $\mathbf{X} \in \mathbb{R}^{n \times m}$ and $k \geq 2$:

1. Pick initial cluster centres $\mathbf{c}_1, \dots, \mathbf{c}_k$ by sampling k points from the input dataset without replacement.
2. For each point in the dataset, $\mathbf{x}_{i,\cdot}$, determine the index of its closest centre l_i :

$$l_i = \arg \min_j \|\mathbf{x}_{i,\cdot} - \mathbf{c}_j\|^2.$$

3. Compute the centroids of the clusters defined by the label vector $\mathbf{l}$, i.e., for every $j = 1, 2, \dots, k$:

$$\mathbf{c}_j = \frac{1}{n_j} \sum_{i:l_i=j} \mathbf{x}_{i,\cdot},$$

where $n_j = |\{i : l_i = j\}|$ gives the size of the j -th cluster.

4. If the objective function (total within-cluster sum of squares) has not changed significantly since the last iteration (say, the absolute value of the difference between the last and the current loss is less than 10^{-9}), then stop and return the current $\mathbf{c}_1, \dots, \mathbf{c}_k$ as the result. Otherwise, go to Step 2.

Exercise 12.20 (*) *Implement the Lloyd algorithm in the form of a function `kmeans(X, C)`, where X is the data matrix ($n \times m$) and where the rows in C , being a $k \times m$ matrix, give the initial cluster centres.*

12.4.4 Local minima (*)

The way the foregoing algorithm is constructed implies what follows.

Important Lloyd's method guarantees that the centres $c_1, \dots, c_k$ it returns cannot be significantly improved any further by repeating Steps 2 and 3 of the algorithm. Still, it does not necessarily mean that they yield the *globally* optimal (the best possible) WCSS. We might as well get stuck in a *local* minimum, where there is no better positioning thereof in the *neighbourhoods* of the current cluster centres; compare Figure 12.11. Yet, had we looked beyond them, we could have found a superior solution.




Figure 12.11. An example function (of only one variable; our problem is much higher-dimensional) with many local minima. How can we be sure there is no better minimum outside of the depicted interval?

A variant of the Lloyd method is given in `scipy.cluster.vq.kmeans2`, where the initial cluster centres are picked at random. Let's test its behaviour by analysing three chosen categories from the 2016 [Sustainable Society Indices](https://ssii.wi.th-koeln.de/)¹⁶ dataset.

```
ssi = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/ssi_2016_categories.csv",
    comment="#")
X = ssi.set_index("Country").loc[:,
    ["PersonalDevelopmentAndHealth", "WellBalancedSociety", "Economy"]
].rename({
    "PersonalDevelopmentAndHealth": "Health",
```

(continues on next page)

¹⁶ <https://ssii.wi.th-koeln.de/>

(continued from previous page)

```

        "WellBalancedSociety": "Balance",
        "Economy": "Economy"
    }, axis=1) # rename columns
n = X.shape[0]
X.loc[["Australia", "Germany", "Poland", "United States"], :] # preview
##
## Country
## Australia      8.590927  6.105539  7.593052
## Germany        8.629024  8.036620  5.575906
## Poland         8.265950  7.331700  5.989513
## United States  8.357395  5.069076  3.756943

```

It is a three-dimensional dataset, where each point (row) corresponds to a different country. Let's find a partition into $k = 3$ clusters.

```

k = 3
np.random.seed(123) # reproducibility matters
C1, l1 = scipy.cluster.vq.kmeans2(X, k)
C1
## array([[7.99945084, 6.50033648, 4.36537659],
##        [7.6370645 , 4.54396676, 6.89893746],
##        [6.24317074, 3.17968018, 3.60779268]])

```

The objective function (total within-cluster sum of squares) at the returned cluster centres is equal to:

```

import scipy.spatial.distance
def get_wcss(X, C):
    D = scipy.spatial.distance.cdist(X, C)**2
    return np.sum(np.min(D, axis=1))

get_wcss(X, C1)
## 446.5221283436733

```

Is it acceptable or not necessarily? We are unable to tell. What we can do, however, is to run the algorithm again, this time from a different starting point:

```

np.random.seed(1234) # different seed - different initial centres
C2, l2 = scipy.cluster.vq.kmeans2(X, k)
C2
## array([[7.80779013, 5.19409177, 6.97790733],
##        [6.31794579, 3.12048584, 3.84519706],
##        [7.92606993, 6.35691349, 3.91202972]])
get_wcss(X, C2)
## 437.51120966832775

```

It is a better solution (we are lucky; it might as well have been worse). But is it the best possible? Again, we cannot tell, alone in the dark.

Does a potential suboptimality affect the way the data points are grouped? It is indeed the case here. Let's look at the contingency table for the two label vectors:

```
pd.DataFrame(dict(l1=l1, l2=l2)).value_counts().unstack(fill_value=0)
## l2    0    1    2
## l1
## 0      8    0   43
## 1     39    6    0
## 2      0   57    1
```

Important Clusters are essentially unordered. The label vector (1, 1, 2, 2, 1, 3) represents the same clustering as the label vectors (3, 3, 2, 2, 3, 1) and (2, 2, 3, 3, 2, 1).

By looking at the contingency table, we see that clusters 0, 1, and 2 in *l1* correspond, respectively, to clusters 2, 0, and 1 in *l2* (via a kind of majority voting). We can relabel the elements in *l1* to get a more readable result:

```
l1p = np.array([2, 0, 1])[l1]
pd.DataFrame(dict(l1p=l1p, l2=l2)).value_counts().unstack(fill_value=0)
## l2    0    1    2
## l1p
## 0     39    6    0
## 1      0   57    1
## 2      8    0   43
```

It is an improvement. It turns out that 8+6+1 countries are categorised differently. We definitely want to avoid a diplomatic crisis stemming from our not knowing that the algorithm might return suboptimal solutions.

Exercise 12.21 (*) *Determine which countries are affected.*

12.4.5 Random restarts (*)

There will never be any guarantees, but we can increase the probability of generating a satisfactory solution by simply restarting the method multiple times from many randomly chosen points and picking the best¹⁷ solution (the one with the smallest WCSS) identified as the result.

Let's perform 1000 such *restarts*:

```
wcss, Cs = [], []
for i in range(1000):
```

(continues on next page)

¹⁷ If we have many different heuristics, each aiming to approximate a solution to the *k*-means problem, from the practical point of view it does not really matter which one returns the best solution – they are merely our tools to achieve a higher goal. Ideally, we could run all of them many times and get the result that corresponds to the smallest WCSS. It is crucial to *do our best* to find the optimal set of cluster centres – the more approaches we test, the better the chance of success.

(continued from previous page)

```
C, l = scipy.cluster.vq.kmeans2(X, k, seed=i)
Cs.append(C)
wcss.append(get_wcss(X, C))
```

The best of the local minima (no guarantee that it is the global one, again) is:

```
np.min(wcss)
## 437.51120966832775
```

It corresponds to the cluster centres:

```
Cs[np.argmin(wcss)]
## array([[7.80779013, 5.19409177, 6.97790733],
##        [7.92606993, 6.35691349, 3.91202972],
##        [6.31794579, 3.12048584, 3.84519706]])
```

They are the same as C2 above (up to a permutation of labels). We were lucky¹⁸, after all.

It is very educational to look at the distribution of the objective function at the identified local minima to see that, proportionally, in the case of this dataset, it is not rare to converge to a really bad solution; see [Figure 12.12](#).

```
plt.hist(wcss, bins=100)
plt.show()
```

Also, [Figure 12.13](#) depicts all the cluster centres to which the algorithm converged. We see that we should not be trusting the results generated by a single run of a heuristic solver to the k -means problem.

Example 12.22 (*) *The `scikit-learn` package has an algorithm that is similar to the Lloyd's one. The method is equipped with the `n_init` parameter (which defaults to 10) which automatically applies the aforementioned restarting.*

```
import sklearn.cluster
np.random.seed(123)
km = sklearn.cluster.KMeans(k, n_init=10)
km.fit(X)
## KMeans(n_clusters=3, n_init=10)
km.inertia_ # WCSS - not optimal!
## 437.54671889589287
```

Still, there are no guarantees: the solution is suboptimal too. As an exercise, pass `n_init=100`, `n_init=1000`, and `n_init=10000` and determine the returned WCSS.

Note It is theoretically possible that a developer from the `scikit-learn` team, when

¹⁸ Mind who is the benevolent dictator of the pseudorandom number generator's seed.




Figure 12.12. Within-cluster sum of squares at the results returned by different runs of the k -means algorithm. Sometimes we might be very unlucky.

they see the preceding result, will make a tweak in the algorithm so that after an update to the package, the returned minimum will be better. This cannot be deemed a bug fix, though, as there are no bugs here. Improving the behaviour of the method in this example will lead to its degradation in others. There is no free lunch in optimisation.

Note Some datasets are more well-behaving than others. The k -means method is *overall* quite usable, but we must always be cautious.

We recommend performing at least 100 random restarts. Also, if a report from data analysis does not say anything about the number of tries performed, we are advised to assume that the results are gibberish¹⁹. People will complain about our being a pain, but we know better; compare Rule#9.

Exercise 12.23 Run the k -means method, $k = 8$, on the *shipu_unbalance*²⁰ dataset from many random sets of cluster centres. Note the value of the total within-cluster sum of squares. Also, plot the cluster centres discovered. Do they make sense? Compare these to the case where you

¹⁹ For instance, R's `stats::kmeans` automatically uses `nstart=1`. It is not rare, unfortunately, that data analysts only stick with the default arguments.

²⁰ https://github.com/gagolews/teaching-data/raw/master/clustering/shipu_unbalance.csv




Figure 12.13. Traces of different cluster centres our k -means algorithm converged to. Some are definitely not optimal, and therefore the method must be restarted a few times to increase the likelihood of pinpointing the true solution.

start the method from the following cluster centres which are close to the global minimum.

$$\mathbf{C} = \begin{bmatrix} -15 & 5 \\ -12 & 10 \\ -10 & 5 \\ 15 & 0 \\ 15 & 10 \\ 20 & 5 \\ 25 & 0 \\ 25 & 10 \end{bmatrix}.$$

12.5 Further reading

An overall noteworthy introduction to classification is [50] and [10]. Nevertheless, as we said earlier, we recommend going through a solid course in matrix algebra and mathematical statistics first, e.g., [22, 42] and [23, 40, 41]. For advanced theoretical (probabilistic, information-theoretic) results, see, e.g., [11, 24].

Hierarchical clustering algorithms (see, e.g., [34, 69]) are also worthwhile as they do not require asking for a fixed number of clusters. Furthermore, density-based algorithms (DBSCAN and its variants) [14, 27, 62] utilise the notion of fixed-radius search that we introduced in Section 8.4.4.

There are a few ways that aim to assess the quality of clustering results, but their meaningfulness is somewhat limited; see [38] for discussion.

12.6 Exercises

Exercise 12.24 Name the data type of the objects that the `DataFrame.groupby` method returns.

Exercise 12.25 What is the relationship between the `GroupBy`, `DataFrameGroupBy`, and `SeriesGroupBy` classes?

Exercise 12.26 What are relative z-scores and how can we compute them?

Exercise 12.27 Why and when the accuracy score might not be the best way to quantify a classifier's performance?

Exercise 12.28 What is the difference between recall and precision, both in terms of how they are defined and where they are the most useful?

Exercise 12.29 Explain how the *k*-nearest neighbour classification and regression algorithms work. Why do we say that they are model-free?

Exercise 12.30 In the context of *k*-nearest neighbour classification, why it might be important to resolve the potential ties at random when computing the mode of the neighbours' labels?

Exercise 12.31 What is the purpose of a training/test and a training/validation/test set split?

Exercise 12.32 Give the formula for the total within-cluster sum of squares.

Exercise 12.33 Are there any cluster shapes that cannot be detected by the *k*-means method?

Exercise 12.34 Why do we say that solving the *k*-means problem is hard?

Exercise 12.35 Why restarting Lloyd's algorithm many times is necessary? Why are reports from data analysis that do not mention the number of restarts not trustworthy?

Accessing databases

pandas is convenient for working with data that fit into memory and which can be stored in individual CSV files. Still, larger information banks in a shared environment will often be made available to us via relational (structured) databases such as PostgreSQL or MariaDB, or a wide range of commercial products.

Most commonly, we use SQL (Structured Query Language) to define the data chunks¹ we want to analyse. Then, we fetch them from the database driver in the form of a **pandas** data frame. This enables us to perform the operations we are already familiar with, e.g., various transformations or visualisations.

Below we make a quick introduction to the basics of SQL using **SQLite**², which is a lightweight, flat-file, and server-less open-source database management system. Overall, SQLite is a sensible choice for data of even hundreds or thousands of gigabytes in size that fit on a single computer's disk. This is more than enough for playing with our data science projects or prototyping more complex solutions.

Important In this chapter, we will learn that the syntax of SQL is very readable: it is modelled after the natural (English) language. The purpose of this introduction is not to compose own queries nor to design new databanks. The latter is covered by a separate course on database systems; see, e.g., [17, 21].

13.1 Example database

In this chapter, we will be working with a simplified data dump of the Q&A site **Travel Stack Exchange**³, which we **downloaded**⁴ on 2017-10-31. It consists of five separate data frames.

¹ Technically, there are ways to use **pandas** with data that do not fit into memory. However, SQL is usually a more versatile choice. If we have too much data, we can always fetch their random samples (this is what statistics is for) or pre-aggregate the information on the server side. This should be sufficient for most intermediate-level users.

² <https://sqlite.org/>

³ <https://travel.stackexchange.com/>

⁴ <https://archive.org/details/stackexchange>

First, Tags gives, amongst others, topic categories (TagName) and how many questions mention them (Count):

```
Tags = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/travel_stackexchange_com_2017/Tags.csv.gz",
    comment="#")
Tags.head(3)
```

| | Count | ExcerptPostId | Id | TagName | WikiPostId |
|------|-------|---------------|----|-----------|------------|
| ## 0 | 104 | 2138.0 | 1 | cruising | 2137.0 |
| ## 1 | 43 | 357.0 | 2 | caribbean | 356.0 |
| ## 2 | 43 | 319.0 | 4 | vacations | 318.0 |

Second, Users provides information on the registered users.

```
Users = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/travel_stackexchange_com_2017/Users.csv.gz",
    comment="#")
Users.head(3)
```

| | AccountId | Age | CreationDate | ... | Reputation | UpVotes | Views |
|------|-----------|------|-------------------------|-----|------------|---------|-------|
| ## 0 | -1.0 | NaN | 2011-06-21T15:16:44.253 | ... | 1.0 | 2472.0 | 0.0 |
| ## 1 | 2.0 | 40.0 | 2011-06-21T20:10:03.720 | ... | 101.0 | 1.0 | 31.0 |
| ## 2 | 7598.0 | 32.0 | 2011-06-21T20:11:02.490 | ... | 101.0 | 1.0 | 14.0 |

[3 rows x 11 columns]

Third, Badges recalls all rewards handed to the users (UserId) for their engaging in various praiseworthy activities:

```
Badges = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/travel_stackexchange_com_2017/Badges.csv.gz",
    comment="#")
Badges.head(3)
```

| | Class | Date | Id | Name | TagBased | UserId |
|------|-------|-------------------------|----|----------------|----------|--------|
| ## 0 | 3 | 2011-06-21T20:16:48.910 | 1 | Autobiographer | False | 2 |
| ## 1 | 3 | 2011-06-21T20:16:48.910 | 2 | Autobiographer | False | 3 |
| ## 2 | 3 | 2011-06-21T20:16:48.910 | 3 | Autobiographer | False | 4 |

Fourth, Posts lists all the questions and answers (the latter do not have ParentId set to NaN).

```
Posts = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/travel_stackexchange_com_2017/Posts.csv.gz",
    comment="#")
Posts.head(3)
```

| | AcceptedAnswerId | ... | ViewCount |
|------|------------------|-----|-----------|
| ## 0 | 393.0 | ... | 419.0 |
| ## 1 | NaN | ... | 1399.0 |
| ## 2 | NaN | ... | NaN |

[3 rows x 17 columns]

Fifth, `Votes` list all the up-votes (`VoteTypeId` equal to 2) and down-votes (`VoteTypeId` of 3) to all the posts.

```
Votes = pd.read_csv("https://raw.githubusercontent.com/gagolews/ +
                    "teaching-data/master/travel_stackexchange_com_2017/Votes.csv.gz",
                    comment="#")
```

```
Votes.head(3)
```

| ## | BountyAmount | CreationDate | Id | PostId | UserId | VoteTypeId |
|------|--------------|-------------------------|----|--------|--------|------------|
| ## 0 | NaN | 2011-06-21T00:00:00.000 | 1 | 1 | NaN | 2 |
| ## 1 | NaN | 2011-06-21T00:00:00.000 | 2 | 1 | NaN | 2 |
| ## 2 | NaN | 2011-06-21T00:00:00.000 | 3 | 2 | NaN | 2 |

Exercise 13.1 See the [README⁵](#) file for a detailed description of each column. Note that rows are uniquely defined by their respective `Ids`. They are relations between the data frames, e.g., `Users`. `Id` vs `Badges`. `UserId`, `Posts`. `Id` vs `Votes`. `PostId`, etc. Moreover, for privacy reasons, some `UserIds` might be missing. In such a case, they are encoded with a not-a-number; compare [Chapter 15](#).

13.2 Exporting data to a database

Let's establish a connection with a new SQLite database. In our case, this will be an ordinary file stored on the computer's disk:

```
import tempfile, os.path
dbfile = os.path.join(tempfile.mkdtemp(), "travel.db")
print(dbfile)
## /tmp/tmp6fih9lv/travel.db
```

It defines the file path (compare [Section 13.6.1](#)) where the database is going to be stored. We use a randomly generated filename inside the local file system's (we are on Linux) temporary directory, `/tmp`. This is just a pleasant exercise, and we will not be using this database afterwards. The reader might prefer setting a filename relative to the current working directory (as given by `os.getcwd`), e.g., `dbfile = "travel.db"`.

We can now connect to the database:

```
import sqlite3
conn = sqlite3.connect(dbfile)
```

The database might now be queried: we can add new tables, insert new rows, and retrieve records.

⁵ https://github.com/gagolews/teaching-data/raw/master/travel_stackexchange_com_2017/README.md

Important In the end, we must not forget about the call to `conn.close()`.

Our data are already in the form of **pandas** data frames. Therefore, exporting them to the database is straightforward. We only need to make a series of calls to the **pandas.DataFrame.to_sql** method.

```
Tags.to_sql("Tags", conn, index=False)
Users.to_sql("Users", conn, index=False)
Badges.to_sql("Badges", conn, index=False)
Posts.to_sql("Posts", conn, index=False)
Votes.to_sql("Votes", conn, index=False)
```

Note (*) It is possible to export data that do not fit into memory by reading them in chunks of considerable, but not too large, sizes. In particular **pandas.read_csv** has the `nrows` argument that lets us read several rows from a file connection; see [Section 13.6.4](#). Then, **pandas.DataFrame.to_sql(..., if_exists="append")** can be used to append new rows to an existing table.

Exporting data can be done without **pandas** as well, e.g., when they are to be fetched from XML or JSON files (compare [Section 13.5](#)) and processed manually, row by row. Intermediate-level SQL users can call **conn.execute("CREATE TABLE t...")**, followed by **conn.executemany("INSERT INTO t VALUES(?, ?, ?)", l)**, and then **conn.commit()**. This will create a new table (here: named `t`) populated by a list of records (e.g., in the form of tuples or **numpy** vectors). For more details, see the [manual](#)⁶ of the **sqlite3** package.

13.3 Exercises on SQL vs pandas

We can use **pandas** to fetch the results of any SQL query in the form of a data frame. For example:

```
pd.read_sql_query("""
    SELECT * FROM Tags LIMIT 3
""", conn)
```

| ## | Count | ExcerptPostId | Id | TagName | WikiPostId |
|------|-------|---------------|----|-----------|------------|
| ## 0 | 104 | 2138.0 | 1 | cruising | 2137.0 |
| ## 1 | 43 | 357.0 | 2 | caribbean | 356.0 |
| ## 2 | 43 | 319.0 | 4 | vacations | 318.0 |

⁶ <https://docs.python.org/3/library/sqlite3.html>

This query selected all columns (`SELECT *`) and the first three rows (`LIMIT 3`) from the `Tags` table.

Exercise 13.2 For the afore- and all the undermentioned SQL queries, write the equivalent Python code that generates the same result using **pandas** functions and methods. In each case, there might be more than one equally fine solution. In case of any doubt about the meaning of the queries, please refer to the [SQLite documentation](https://sqlite.org/lang.html)⁷. Example solutions are provided at the end of this section.

Example 13.3 For a reference query:

```
res1a = pd.read_sql_query("""
    SELECT * FROM Tags LIMIT 3
""", conn)
```

The equivalent **pandas** implementation might look like:

```
res1b = Tags.head(3)
```

To verify that the results are equal, we can call:

```
pd.testing.assert_frame_equal(res1a, res1b) # no error == OK
```

No error message means that the test is passed. The cordial thing about the `assert_frame_equal` function is that it ignores small round-off errors introduced by arithmetic operations.

Nonetheless, the results generated by **pandas** might be the same up to the reordering of rows. In such a case, before calling `pandas.testing.assert_frame_equal`, we can invoke `DataFrame.sort_values` on both data frames to sort them with respect to 1 or 2 chosen columns.

13.3.1 Filtering

Exercise 13.4 From `Tags`, select two columns `TagName` and `Count` and rows for which `TagName` is equal to one of the three choices provided.

```
res2a = pd.read_sql_query("""
    SELECT TagName, Count
    FROM Tags
    WHERE TagName IN ('poland', 'australia', 'china')
""", conn)
res2a
##      TagName  Count
## 0      china    443
## 1  australia    411
## 2      poland    139
```

Hint: use `pandas.Series.isin`.

⁷ <https://sqlite.org/lang.html>

Exercise 13.5 Select a set of columns from *Posts* whose rows fulfil a given set of conditions.

```
res3a = pd.read_sql_query("""
    SELECT Title, Score, ViewCount, FavoriteCount
    FROM Posts
    WHERE PostTypeId=1 AND
           ViewCount>=10000 AND
           FavoriteCount BETWEEN 35 AND 100
""", conn)
res3a
```

| ## | | Title | ... | FavoriteCount |
|------|---|-------|-----|---------------|
| ## 0 | When traveling to a country with a different c... | ... | ... | 35.0 |
| ## 1 | How can I do a "broad" search for flights? | ... | ... | 49.0 |
| ## 2 | Tactics to avoid getting harassed by corrupt p... | ... | ... | 42.0 |
| ## 3 | Flight tickets: buy two weeks before even duri... | ... | ... | 36.0 |
| ## 4 | OK we're all adults here, so really, how on ea... | ... | ... | 79.0 |
| ## 5 | How to intentionally get denied entry to the U... | ... | ... | 53.0 |
| ## 6 | How do you know if Americans genuinely/literal... | ... | ... | 79.0 |
| ## 7 | OK, we are all adults here, so what is a bidet... | ... | ... | 38.0 |
| ## 8 | How to cope with too slow Wi-Fi at hotel? | ... | ... | 41.0 |
| ## | | | | |
| ## | [9 rows x 4 columns] | | | |

13.3.2 Ordering

Exercise 13.6 Select the *Title* and *Score* columns from *Posts* where *ParentId* is missing (i.e., the post is, in fact, a question) and *Title* is well-defined. Then, sort the results by the *Score* column, decreasingly (descending order). Finally, return only the first five rows (e.g., top five scoring questions).

```
res4a = pd.read_sql_query("""
    SELECT Title, Score
    FROM Posts
    WHERE ParentId IS NULL AND Title IS NOT NULL
    ORDER BY Score DESC
    LIMIT 5
""", conn)
res4a
```

| ## | | Title | Score |
|------|---|-------|-------|
| ## 0 | OK we're all adults here, so really, how on ea... | | 306 |
| ## 1 | How do you know if Americans genuinely/literal... | | 254 |
| ## 2 | How to intentionally get denied entry to the U... | | 219 |
| ## 3 | Why are airline passengers asked to lift up wi... | | 210 |
| ## 4 | Why prohibit engine braking? | | 178 |

Hint: use `pandas.DataFrame.sort_values` and `numpy.isnan` or `pandas.isnull`.

13.3.3 Removing duplicates

Exercise 13.7 Get all unique badge names for the user with *Id*=23.

```
res5a = pd.read_sql_query("""
    SELECT DISTINCT Name
    FROM Badges
    WHERE UserId=23
""", conn)
res5a
##              Name
## 0      Supporter
## 1       Student
## 2       Teacher
## 3       Scholar
## 4          Beta
## 5   Nice Question
## 6        Editor
## 7   Nice Answer
## 8     Yearling
## 9  Popular Question
## 10    Taxonomist
## 11  Notable Question
```

Hint: use `pandas.DataFrame.drop_duplicates`.

Exercise 13.8 For each badge handed to the user with *Id*=23, extract the award year store it in a new column named *Year*. Then, select only the unique pairs (*Name*, *Year*).

```
res6a = pd.read_sql_query("""
    SELECT DISTINCT
        Name,
        CAST(strftime('%Y', Date) AS FLOAT) AS Year
    FROM Badges
    WHERE UserId=23
""", conn)
res6a
##              Name      Year
## 0      Supporter  2011.0
## 1       Student  2011.0
## 2       Teacher  2011.0
## 3       Scholar  2011.0
## 4          Beta  2011.0
## 5   Nice Question  2011.0
## 6        Editor  2012.0
## 7   Nice Answer  2012.0
## 8     Yearling  2012.0
## 9   Nice Question  2012.0
## 10  Nice Question  2013.0
## 11     Yearling  2013.0
## 12  Popular Question  2014.0
```

(continues on next page)

(continued from previous page)

```
## 13      Yearling    2014.0
## 14      Taxonomist  2014.0
## 15      Notable Question  2015.0
## 16      Nice Question  2017.0
```

Hint: use `Badges.Date.astype("datetime64[ns]").dt.strftime("%Y")` followed by `astype("float")`; see [Chapter 16](#).

13.3.4 Grouping and aggregating

Exercise 13.9 Count how many badges of each type were won by the user with `Id=23`. Also, for each badge type, compute the minimal, average, and maximal receiving year. Return only the top four badges (with respect to the counts).

```
res7a = pd.read_sql_query("""
SELECT
    Name,
    COUNT(*) AS Count,
    MIN(CAST(strftime('%Y', Date) AS FLOAT)) AS MinYear,
    AVG(CAST(strftime('%Y', Date) AS FLOAT)) AS MeanYear,
    MAX(CAST(strftime('%Y', Date) AS FLOAT)) AS MaxYear
FROM Badges
WHERE UserId=23
GROUP BY Name
ORDER BY Count DESC
LIMIT 4
""", conn)
res7a
```

| ## | Name | Count | MinYear | MeanYear | MaxYear |
|------|------------------|-------|---------|----------|---------|
| ## 0 | Nice Question | 4 | 2011.0 | 2013.25 | 2017.0 |
| ## 1 | Yearling | 3 | 2012.0 | 2013.00 | 2014.0 |
| ## 2 | Popular Question | 3 | 2014.0 | 2014.00 | 2014.0 |
| ## 3 | Notable Question | 2 | 2015.0 | 2015.00 | 2015.0 |

Exercise 13.10 Count how many unique combinations of pairs (`Name`, `Year`) for the badges won by the user with `Id=23` are there. Then, return only the rows having `Count` greater than 1 and order the results by `Count` decreasingly. In other words, list the badges received more than once in any given year.

```
res8a = pd.read_sql_query("""
SELECT
    Name,
    CAST(strftime('%Y', Date) AS FLOAT) AS Year,
    COUNT(*) AS Count
FROM Badges
WHERE UserId=23
GROUP BY Name, Year
```

(continues on next page)

(continued from previous page)

```

HAVING Count > 1
ORDER BY Count DESC
""" , conn)
res8a
##           Name      Year  Count
## 0 Popular Question  2014.0      3
## 1 Notable Question  2015.0      2

```

Note that *WHERE* is performed before *GROUP BY*, and *HAVING* is applied thereafter.

13.3.5 Joining

Exercise 13.11 Join (merge) *Tags*, *Posts*, and *Users* for all posts with *OwnerUserId* not equal to -1 (i.e., the tags which were created by “alive” users). Return the top six records with respect to *Tags.Count*.

```

res9a = pd.read_sql_query("""
SELECT Tags.TagName, Tags.Count, Posts.OwnerUserId,
       Users.Age, Users.Location, Users.DisplayName
FROM Tags
JOIN Posts ON Posts.Id=Tags.WikiPostId
JOIN Users ON Users.AccountId=Posts.OwnerUserId
WHERE OwnerUserId != -1
ORDER BY Tags.Count DESC, Tags.TagName ASC
LIMIT 6
""", conn)
res9a
##           TagName  Count  ...           Location  DisplayName
## 0         canada    802  ...    Mumbai, India         hitec
## 1         europe    681  ... Philadelphia, PA         Adam Tuttle
## 2 visa-refusals    554  ...   New York, NY Benjamin Pollack
## 3      australia    411  ...    Mumbai, India         hitec
## 4             eu    204  ... Philadelphia, PA         Adam Tuttle
## 5 new-york-city    204  ...    Mumbai, India         hitec
##
## [6 rows x 6 columns]

```

Exercise 13.12 First, create an auxiliary (temporary) table named *UpVotesTab*, where we store the information about the number of up-votes (*VoteTypeId*=2) that each post has received. Then, join (merge) this table with *Posts* and fetch some details about the five questions (*PostTypeId*=1) with the most up-votes.

```

res10a = pd.read_sql_query("""
SELECT UpVotesTab.*, Posts.Title FROM
(
    SELECT PostId, COUNT(*) AS UpVotes
    FROM Votes
    WHERE VoteTypeId=2

```

(continues on next page)

(continued from previous page)

```

        GROUP BY PostId
    ) AS UpVotesTab
    JOIN Posts ON UpVotesTab.PostId=Posts.Id
    WHERE Posts.PostTypeId=1
    ORDER BY UpVotesTab.UpVotes DESC LIMIT 5
""" , conn)
res10a
##      PostId  UpVotes                                     Title
## 0      3080      307  OK we're all adults here, so really, how on ea...
## 1     38177       254  How do you know if Americans genuinely/literal...
## 2     24540       221  How to intentionally get denied entry to the U...
## 3     20207       211  Why are airline passengers asked to lift up wi...
## 4     96447       178                                     Why prohibit engine braking?

```

13.3.6 Solutions to exercises

In this section, we provide example solutions to the above exercises.

Example 13.13 To obtain a result equivalent to *res2a*, we need simple data filtering only:

```

res2b = (
    Tags.
    loc[
        Tags.TagName.isin(["poland", "australia", "china"]),
        ["TagName", "Count"]
    ].
    reset_index(drop=True)
)

```

Let's verify whether the two data frames are identical:

```
pd.testing.assert_frame_equal(res2a, res2b) # no error == OK
```

Example 13.14 To generate *res3a* with *pandas* only, we need some more complex filtering with *loc[...]*:

```

res3b = (
    Posts.
    loc[
        (Posts.PostTypeId == 1) & (Posts.ViewCount >= 10000) &
        (Posts.FavoriteCount >= 35) & (Posts.FavoriteCount <= 100),
        ["Title", "Score", "ViewCount", "FavoriteCount"]
    ].
    reset_index(drop=True)
)
pd.testing.assert_frame_equal(res3a, res3b) # no error == OK

```

Example 13.15 For *res4a*, some filtering and sorting is all we need:

```

res4b = (
    Posts.
    loc[
        Posts.ParentId.isna() & (~Posts.Title.isna()),
        ["Title", "Score"]
    ].
    sort_values("Score", ascending=False).
    head(5).
    reset_index(drop=True)
)
pd.testing.assert_frame_equal(res4a, res4b) # no error == OK

```

Example 13.16 The key to `res5a` is the `pandas.DataFrame.drop_duplicates` method:

```

res5b = (
    Badges.
    loc[Badges.UserId == 23, ["Name"]].
    drop_duplicates().
    reset_index(drop=True)
)
pd.testing.assert_frame_equal(res5a, res5b) # no error == OK

```

Example 13.17 For `res6a`, we first need to add a new column to the copy of `Badges`:

```

Badges2 = Badges.copy() # otherwise we would destroy the original object
Badges2.loc[:, "Year"] = (
    Badges2.Date.astype("datetime64[ns]").dt.strftime("%Y").astype("float")
)

```

Then, we apply some filtering and the removal of duplicated rows:

```

res6b = (
    Badges2.
    loc[Badges2.UserId == 23, ["Name", "Year"]].
    drop_duplicates().
    reset_index(drop=True)
)
pd.testing.assert_frame_equal(res6a, res6b) # no error == OK

```

Example 13.18 For `res7a`, we can use `pandas.DataFrameGroupBy.aggregate`:

```

Badges2 = Badges.copy()
Badges2.loc[:, "Year"] = (
    Badges2.Date.astype("datetime64[ns]").dt.strftime("%Y").astype("float")
)
res7b = (
    Badges2.
    loc[Badges2.UserId == 23, ["Name", "Year"]].
    groupby("Name")["Year"].

```

(continues on next page)

(continued from previous page)

```

aggregate([len, "min", "mean", "max"]).
sort_values("len", ascending=False).
head(4).
reset_index()
)
res7b.columns = ["Name", "Count", "MinYear", "MeanYear", "MaxYear"]

```

Had we not converted *Year* to *float*, we would obtain a meaningless average year, without any warning.

Unfortunately, the rows in *res7a* and *res7b* are ordered differently. For testing, we need to re-order them in the same way:

```

pd.testing.assert_frame_equal(
    res7a.sort_values(["Name", "Count"]).reset_index(drop=True),
    res7b.sort_values(["Name", "Count"]).reset_index(drop=True)
) # no error == OK

```

Example 13.19 For *res8a*, we first count the number of values in each group:

```

Badges2 = Badges.copy()
Badges2.loc[:, "Year"] = (
    Badges2.Date.astype("datetime64[ns]").dt.strftime("%Y").astype("float")
)
res8b = (
    Badges2.
loc[ Badges2.UserId == 23, ["Name", "Year"] ].
groupby(["Name", "Year"]).
size().
rename("Count").
reset_index()
)

```

The *HAVING* part is performed after *WHERE* and *GROUP BY*.

```

res8b = (
    res8b.
loc[ res8b.Count > 1, : ].
sort_values("Count", ascending=False).
reset_index(drop=True)
)
pd.testing.assert_frame_equal(res8a, res8b) # no error == OK

```

Example 13.20 To obtain a result equivalent to *res9a*, we need to merge *Posts* with *Tags*, and then merge the result with *Users*:

```

res9b = pd.merge(Posts, Tags, left_on="Id", right_on="WikiPostId")
res9b = pd.merge(Users, res9b, left_on="AccountId", right_on="OwnerUserId")

```

Then, some filtering and sorting will do the trick:

```

res9b = (
    res9b.
    loc[
        (res9b.OwnerUserId != -1) & (~res9b.OwnerUserId.isna()),
        ["TagName", "Count", "OwnerUserId", "Age", "Location", "DisplayName"]
    ].
    sort_values(["Count", "TagName"], ascending=[False, True]).
    head(6).
    reset_index(drop=True)
)

```

In SQL, “not equals to -1” implies *IS NOT NULL*.

```
pd.testing.assert_frame_equal(res9a, res9b) # no error == OK
```

Example 13.21 To obtain a result equivalent to *res10a*, we first need to create an auxiliary data frame that corresponds to the subquery.

```

UpVotesTab = (
    Votes.
    loc[Votes.VoteTypeId==2, :].
    groupby("PostId").
    size().
    rename("UpVotes").
    reset_index()
)

```

And now:

```

res10b = pd.merge(UpVotesTab, Posts, left_on="PostId", right_on="Id")
res10b = (
    res10b.
    loc[res10b.PostTypeId==1, ["PostId", "UpVotes", "Title"]].
    sort_values("UpVotes", ascending=False).
    head(5).
    reset_index(drop=True)
)
pd.testing.assert_frame_equal(res10a, res10b) # no error == OK

```

13.4 Closing the database connection

We said we should not forget about:

```
conn.close()
```

This gives some sense of closure. Such a relief.

13.5 Common data serialisation formats for the Web

CSV files are an all-round way to exchange *tabular* data between different programming and data analysis environments.

For unstructured or non-tabularly-structured data, XML and JSON (and its superset, YAML) are the common formats of choice, especially for communicating with different Web APIs.

To ensure we can fetch data in these formats, we recommend solving the undermentioned exercises. Sadly, often this will require some tedious labour, neither art nor science; see also [95] and [20].

Exercise 13.22 Consider the Web API for *accessing*⁸ the on-street parking bay sensor data in Melbourne, VIC, Australia. Using the `json` package, convert the *data*⁹ in the JSON format to a data frame.

Exercise 13.23 Australian Radiation Protection and Nuclear Safety Agency *publishes*¹⁰ UV data for different Aussie cities. Using the `xml` package, convert *this XML dataset*¹¹ to a data frame.

Exercise 13.24 (*) Check out the English Wikipedia article with a *list of 20th-century classical composers*¹². Using `pandas.read_html`, convert the Climate Data table included therein to a data frame.

Exercise 13.25 (*) Using the `lxml` package, author a function that converts each bullet list featured in a given Wikipedia article (e.g., *this one*¹³), to a list of strings.

Exercise 13.26 (**) Import an archived version of a *Stack Exchange*¹⁴ site that you find interesting and store it in an SQLite database. You can find the relevant data dumps *here*¹⁵.

Exercise 13.27 (**) *Download*¹⁶ and then import an archived version of one of the wikis hosted by the *Wikimedia Foundation*¹⁷ (e.g., the whole English Wikipedia) so that it can be stored in an SQLite database.

⁸ <https://data.melbourne.vic.gov.au/explore/dataset/on-street-parking-bay-sensors/api>

⁹ <https://data.melbourne.vic.gov.au/api/explore/v2.1/catalog/datasets/on-street-parking-bay-sensors/exports/json>

¹⁰ <https://www.arpansa.gov.au/our-services/monitoring/ultraviolet-radiation-monitoring/ultraviolet-radiation-data-information>

¹¹ <https://uvdata.arpansa.gov.au/xml/uvvalues.xml>

¹² https://en.wikipedia.org/wiki/List_of_20th-century_classical_composers

¹³ https://en.wikipedia.org/wiki/Category:Fr%C3%A9d%C3%A9ric_Chopin

¹⁴ <https://stackexchange.com/>

¹⁵ <https://archive.org/details/stackexchange>

¹⁶ https://meta.wikimedia.org/wiki/Data_dumps

¹⁷ <https://wikimediafoundation.org/>

13.6 Working with many files

For the mass-processing of many files, it is worth knowing of some functions for dealing with file paths, searching for files, etc. Usually, we will be looking up ways to complete specific tasks at hand, e.g., how to read data from a ZIP-like archive, on the internet. After all, contrary to the basic operations of vectors, matrices, and data frames, they are not amongst the actions that we perform frequently.

Good development practices related to data storage are described in [49].

13.6.1 File paths

UNIX-like operating systems, including GNU/Linux and m**OS, use slashes, “/”, as path separators, e.g., “/home/marek/file.csv”. Win**s, however, uses backslashes, “\”, which have a special meaning in character strings (escape sequences; see [Section 2.1.3](#)). Therefore, they should be input as, e.g., “c:\\users\\marek\\file.csv”. Alternatively, we can use *raw* strings, where the backslash is treated literally, e.g., r“c:\users\marek\file.csv”.

When constructing file paths programmatically, it is thus best to rely on `os.path.join`, which takes care of the system-specific nuances.

```
import os.path
os.path.join("~", "Desktop", "file.csv") # we are on GNU/Linux
## '~/Desktop/file.csv'
```

The tilde, “~”, denotes the current user’s *home* directory.

For storing auxiliary data, we can use the system’s temporary directory. See the `tempfile` module for functions that generate appropriate file paths therein. For instance, a subdirectory inside the temporary directory can be created via a call to `tempfile.mkdtemp`.

Important We will frequently be referring to file paths relative to the working directory of the currently executed Python session (e.g., from which IPython/Jupyter notebook server was started); see `os.getcwd`.

All non-absolute file names (ones that do not start with “~”, “/”, “c:\\”, and the like), for example, “filename.csv” or `os.path.join("subdir", "filename.csv")` are always *relative* to the current working directory.

For instance, if the working directory is “/home/marek/projects/python”, then “filename.csv” refers to “/home/marek/projects/python/filename.csv”.

Also, “..” denotes the current working directory’s parent directory. Thus, “../filename2.csv” resolves to “/home/marek/projects/filename2.csv”.

Exercise 13.28 Print the current working directory by calling `os.getcwd`. Next, download the file `air_quality_2018_param`¹⁸ and save it in the current Python session's working directory (e.g., in your web browser, right-click on the web page's canvas and select Save Page As...). Load with `pandas.read_csv` by passing `"air_quality_2018_param.csv"` as the input path.

Exercise 13.29 (*) Download the aforementioned file programmatically (if you have not done so yet) using the `requests` module.

13.6.2 File search

`glob.glob` and `os.listdir` generate a list of files in a given directory (and possibly all its subdirectories).

`os.path.isdir` and `os.path.isfile` determine the type of a given object in the file system.

Exercise 13.30 Write a function that computes the total size of all the files in a given directory and all its subdirectories.

13.6.3 Exception handling

Accessing resources on the disk or the internet can lead to errors, for example, when the file is not found. The `try...except` statement can be used if we want to be able to react to any of the envisaged errors

```
try:
    # statements to execute
    x = pd.read_csv("file_not_found.csv")
    print(x.head()) # not executed if the above raises an error
except OSError:
    # if an exception occurs, we can handle it here
    print("File has not been found")
## File has not been found
```

For more details, refer to the [documentation](#)¹⁹.

13.6.4 File connections (*)

Basic ways of opening and reading from/writing to file connections are described in the [Python documentation](#)²⁰. Section 14.3.5 shows an example where we create a Markdown file manually.

They may be useful for processing large files chunk by chunk. In particular, `pandas.read_csv` accepts a file handler (see `open` and `numpy.lib.npyio.DataSource`). Then, by passing the `nrows` argument we can request a specific number of records to be read at the current position.

¹⁸ https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018_param.csv

¹⁹ <https://docs.python.org/3/tutorial/errors.html>

²⁰ <https://docs.python.org/3/tutorial/inputoutput.html>

With `pandas.to_csv` we can also append portions of data frames to a file.

13.7 Further reading

`pandas` is not the only package that brings data frames to the Python world; check out `polars`. Furthermore, `Dask` and `modin` can be helpful with data that do not fit into memory (and when we do not want to rely on sampling or chunking).

13.8 Exercises

Exercise 13.31 Find an example of an XML and JSON file. Which one is more human-readable? Do they differ in terms of capabilities?

Exercise 13.32 What is wrong with constructing file paths like `"~" + "\\ " + "filename.csv"`?

Exercise 13.33 What are the benefits of using a SQL database management system in data science activities?

Exercise 13.34 (*) How can we populate a database with gigabytes of data read from many CSV files?

Part V

Other data types

In [35], it is noted that effective processing of character strings is needed at various stages of data analysis pipelines: from data cleansing and preparation, through information extraction, to report generation; compare, e.g., [95] and [20]. *Pattern searching, string collation and sorting, normalisation, transliteration, and formatting are ubiquitous in text mining, natural language processing, and bioinformatics.* Means for the handling of string data should be included in each statistician's or data scientist's repertoire to complement their numerical computing and data wrangling skills.

In this chapter, we discuss the handiest string operations in base Python, together with their vectorised versions in **numpy** and **pandas**. We also mention some more advanced features of the Unicode ICU library.

14.1 Basic string operations

Recall from [Section 2.1.3](#) that the `str` class represents individual character strings:

```
x = "spam"
type(x)
## <class 'str'>
```

There are a few binary operators overloaded for strings, e.g., ``+`` stands for string concatenation:

```
x + " and eggs"
## 'spam and eggs'
```

``*`` duplicates a given string:

```
x * 3
## 'spamspamspam'
```

[Chapter 3](#) noted that `str` is a sequential type. As a consequence, we can extract individual code points and create substrings using the index operator:

```
x[-1] # last letter
## 'm'
```

Strings are immutable, but parts thereof can always be reused in conjunction with the concatenation operator:

```
x[:2] + "ecial"
## 'special'
```

14.1.1 Unicode as the universal encoding

It is worth knowing that all strings in Python (from version 3.0) use Unicode¹, which is a universal encoding capable of representing c. 150 000 characters covering letters and numbers in contemporary and historic alphabets/scripts, mathematical, political, phonetic, and other symbols, emojis, etc.

Note Despite the wide support for Unicode, sometimes our own or other readers' display (e.g., web browsers when viewing an HTML version of the output report) might not be able to *render* all code points properly, e.g., due to missing fonts. Still, we can rest assured that they are processed correctly if string functions are applied thereon.

14.1.2 Normalising strings

Dirty text data are a pain, especially if similar (semantically) tokens are encoded in many different ways. For the sake of string matching, we might want, e.g., the German "groß", "GROSS", and " gross " to compare all equal.

`str.strip` removes whitespaces (spaces, tabs, newline characters) at both ends of strings (see also `str.lstrip` and `str.rstrip` for their nonsymmetric versions).

`str.lower` and `str.upper` change letter case. For caseless comparison/matching, `str.casefold` might be a slightly better option as it unfolds many more code point sequences:

```
"Groß".lower(), "Groß".upper(), "Groß".casefold()
## ('groß', 'GROSS', 'gross')
```

Note (*) More advanced string *transliteration*² can be performed by means of the **ICU**³ (International Components for Unicode) library. Its Python bindings are provided by the **PyICU** package. Unfortunately, the package is not easily available on Win***s.

For instance, converting all code points to ASCII (English) might be necessary when

¹ (*) More precisely, Python strings are UTF-8-encoded. Most web pages and APIs are nowadays served in UTF-8. But we can still occasionally encounter files encoded in ISO-8859-1 (Western Europe), Windows-1250 (Eastern Europe), Windows-1251 (Cyrillic), GB18030 and Big5 (Chinese), EUC-KR (Korean), Shift-JIS and EUC-JP (Japanese), amongst others. They can be converted using the `str.decode` method.

² <https://unicode-org.github.io/icu/userguide/transforms/general>

³ <https://icu.unicode.org/>

identifiers are expected to miss some diacritics that would normally be included (as in "Gągolewski" vs "Gagolewski"):

```
import icu # PyICU package
(icu.Transliterator
 .createInstance("Lower; Any-Latin; Latin-ASCII")
 .transliterate(
     "Χαίρετε! Groß gżegżółka – © La Niña – köszönöm – Gągolewski"
 )
)
## 'chairete! gross gżegżółka - (C) la nina - koszonom - gagolewski'
```

Converting between different **Unicode Normalisation Forms**⁴ (also available in the **unicodedata** package and via **pandas.Series.str.normalize**) might be used for the removal of some formatting nuances:

```
icu.Transliterator.createInstance("NFKD; NFC").transliterate("%qr²")
## '¹/₄qr²'
```

14.1.3 Substring searching and replacing

Determining if a string has a particular fixed substring can be done in several ways.

For instance, the **in** operator verifies whether a particular substring occurs at least once:

```
food = "bacon, spam, spam, srapatapam, eggs, and spam"
"spam" in food
## True
```

The **str.count** method determines the number of occurrences of a substring:

```
food.count("spam")
## 3
```

To locate the first pattern appearance, we call **str.index**:

```
food.index("spam")
## 7
```

str.replace substitutes matching substrings with new content:

```
food.replace("spam", "veggies")
## 'bacon, veggies, veggies, srapatapam, eggs, and veggies'
```

Exercise 14.1 Read the manual of the following methods: **str.startswith**, **str.endswith**, **str.find**, **str.rfind**, **str.rindex**, **str.removeprefix**, and **str.removesuffix**.

⁴ <https://www.unicode.org/faq/normalization.html>

The splitting of long strings at specific fixed delimiters can be done via:

```
food.split(", ")
## ['bacon', 'spam', 'spam', 'srapatapam', 'eggs', 'and spam']
```

See also `str.partition`. The `str.join` method implements the inverse operation:

```
", ".join(["spam", "bacon", "eggs", "spam"])
## 'spam, bacon, eggs, spam'
```

Moreover, [Section 14.4](#) will discuss pattern matching with regular expressions. They can be useful in, amongst others, extracting more abstract data chunks (numbers, URLs, email addresses, IDs) from strings.

14.1.4 Locale-aware services in ICU (*)

Recall that relational operators such as `<` and `>=` perform the lexicographic comparing of strings (like in a dictionary or an encyclopedia):

```
"spam" > "egg"
## True
```

We have: `"a" < "aa" < "aaaaaaaaaaaa" < "ab" < "aba" < "abb" < "b" < "ba" < "baaaaaa" < "bb" < "Spanish Inquisition"`.

The lexicographic ordering (character-by-character, from left to right) is not necessarily appropriate for strings with numerals:

```
"a9" < "a123" # 1 is smaller than 9
## False
```

Additionally, it only takes into account the numeric codes (see [Section 14.4.3](#)) corresponding to each Unicode character. Consequently, it does not work well with non-English alphabets:

```
"MIELONECZKA" < "MIELONECZKI"
## False
```

In Polish, *A with ogonek* (Ą) is expected to sort after A and before B, let alone I. However, their corresponding numeric codes in the Unicode table are: 260 (Ą), 65 (A), 66 (B), and 73 (I). The resulting ordering is thus incorrect, as far as natural language processing is concerned.

It is best to perform string collation using the services provided by ICU. Here is an example of German phone book-like collation where "ö" is treated the same as "oe":

```
c = icu.Collator.createInstance(icu.Locale("de_DE@collation=phonebook"))
c.setStrength(0) # ignore case and some diacritics
c.compare("Löwe", "loewe")
## 0
```

A result of 0 means that the strings are deemed equal.

In some languages, contractions occur, e.g., in Slovak and Czech, two code points "ch" are treated as a single entity and are sorted after "h":

```
icu.Collator.createInstance(icu.Locale("sk_SK")).compare("chladný", "hladný")
## 1
```

This means that we have "chladný" > "hladný" (the first argument is greater than the second one). Compare the above to something similar in Polish:

```
icu.Collator.createInstance(icu.Locale("pl_PL")).compare("chłodny", "hardy")
## -1
```

That is, "chłodny" < "hardy" (the first argument is less than the second one).

Also, with ICU, numeric collation is possible:

```
c = icu.Collator.createInstance()
c.setAttribute(
    icu.UCollAttribute.NUMERIC_COLLATION,
    icu.UCollAttributeValue.ON
)
c.compare("a9", "a123")
## -1
```

Which is the correct result: "a9" is less than "a123" (compare the above to the example where we used the ordinary `<`).

14.1.5 String operations in pandas

String sequences in `pandas.Series` are by default⁵ using the broadest possible object data type:

```
pd.Series(["spam", "bacon", "spam"])
## 0    spam
## 1    bacon
## 2    spam
## dtype: object
```

This allows for missing values encoding by means of the `None` object (which is of the type `None`, not `str`); compare [Section 15.1](#).

Vectorised versions of base string operations are available via the `pandas.Series.str` accessor. We thus have `pandas.Series.str.strip`, `pandas.Series.str.split`, `pandas.Series.str.find`, and so forth. For instance:

⁵ https://pandas.pydata.org/pandas-docs/stable/user_guide/text.html

```
x = pd.Series(["spam", "bacon", None, "buckwheat", "spam"])
x.str.upper()
## 0      SPAM
## 1     BACON
## 2     None
## 3  BUCKWHEAT
## 4      SPAM
## dtype: object
```

But there is more. For example, a function to compute the length of each string:

```
x.str.len()
## 0    4.0
## 1    5.0
## 2    NaN
## 3    9.0
## 4    4.0
## dtype: float64
```

Vectorised concatenation of strings can be performed using the overloaded `+` operator:

```
x + " and spam"
## 0    spam and spam
## 1    bacon and spam
## 2             NaN
## 3    buckwheat and spam
## 4    spam and spam
## dtype: object
```

To concatenate all items into a single string, we call:

```
x.str.cat(sep="; ")
## 'spam; bacon; buckwheat; spam'
```

Conversion to numeric:

```
pd.Series(["1.3", "-7", None, "3523"]).astype(float)
## 0    1.3
## 1   -7.0
## 2    NaN
## 3  3523.0
## dtype: float64
```

Select substrings:

```
x.str.slice(2, -1) # like x.iloc[i][2:-1] for all i
## 0    a
## 1    co
```

(continues on next page)

(continued from previous page)

```
## 2      None
## 3    ckwhea
## 4        a
## dtype: object
```

Replace substrings:

```
x.str.slice_replace(0, 2, "tofu") # like x.iloc[i][2:-1] = "tofu"
## 0      tofuam
## 1    tofucon
## 2      None
## 3  tofuckwheat
## 4      tofuam
## dtype: object
```

Exercise 14.2 Consider the *nasaweather_glaciers*⁶ data frame. All glaciers are assigned 11/12-character unique identifiers as defined by the WGMS convention that forms the glacier ID number by combining the following five elements:

1. 2-character political unit (the first two letters of the ID),
2. 1-digit continent code (the third letter),
3. 4-character drainage code (the next four),
4. 2-digit free position code (the next two),
5. 2- or 3-digit local glacier code (the remaining ones).

Extract the five chunks and store them as independent columns in the data frame.

14.1.6 String operations in numpy (*)

There is a huge overlap between the **numpy** and **pandas** capabilities for string handling, with the latter being more powerful. After all, **numpy** is a workhorse for *numerical* computing. Still, some readers might find what follows useful.

As mentioned in our introduction to **numpy** vectors, objects of the type `ndarray` can store not only numeric and logical data, but also character strings. For example:

```
x = np.array(["spam", "bacon", "egg"])
x
## array(['spam', 'bacon', 'egg'], dtype='<U5')
```

Here, the data type “<U5” means that we deal with Unicode strings of length no greater than five. Unfortunately, replacing elements with too long a content will spawn truncated strings:

⁶ https://github.com/gagolews/teaching-data/raw/master/other/nasaweather_glaciers.csv

```
x[2] = "buckwheat"
x
## array(['spam', 'bacon', 'buckw'], dtype='<U5')
```

To remedy this, we first need to recast the vector manually:

```
x = x.astype("<U10")
x[2] = "buckwheat"
x
## array(['spam', 'bacon', 'buckwheat'], dtype='<U10')
```

Conversion from/to numeric is also possible:

```
np.array(["1.3", "-7", "3523"]).astype(float)
## array([ 1.300e+00, -7.000e+00,  3.523e+03])
np.array([1, 3.14, -5153]).astype(str)
## array(['1.0', '3.14', '-5153.0'], dtype='<U32')
```

The `numpy.char7` module includes several vectorised versions of string routines, most of which we have already discussed. For example:

```
x = np.array([
    "spam", "spam, bacon, and spam",
    "spam, eggs, bacon, spam, spam, and spam"
])
np.char.split(x, ", ")
## array([list(['spam']), list(['spam', 'bacon', 'and spam']),
##        list(['spam', 'eggs', 'bacon', 'spam', 'spam', 'and spam'])],
##        dtype=object)
np.char.count(x, "spam")
## array([1, 2, 4])
```

Vectorised operations that we would normally perform through the binary operators (i.e., `+`, `*`, `<`, etc.) are available through standalone functions:

```
np.char.add(["spam", "bacon"], " and spam")
## array(['spam and spam', 'bacon and spam'], dtype='<U14')
np.char.equal(["spam", "bacon", "spam"], "spam")
## array([ True, False,  True])
```

The function that returns the length of each string is also noteworthy:

```
np.char.str_len(x)
## array([ 4, 21, 39])
```

⁷ <https://numpy.org/doc/stable/reference/routines.char.html>

14.2 Working with string lists

`pandas` nicely supports lists of strings of varying lengths. For instance:

```
x = pd.Series([
    "spam",
    "spam, bacon, spam",
    "potatoes",
    None,
    "spam, eggs, bacon, spam, spam"
])
xs = x.str.split(", ", regex=False)
xs
## 0          [spam]
## 1    [spam, bacon, spam]
## 2          [potatoes]
## 3                None
## 4    [spam, eggs, bacon, spam, spam]
## dtype: object
```

And now, e.g., looking at the last element:

```
xs.iloc[-1]
## ['spam', 'eggs', 'bacon', 'spam', 'spam']
```

reveals that it is indeed a list of strings.

There are a few vectorised operations that enable us to work with such variable length lists, such as concatenating all strings:

```
xs.str.join("; ")
## 0          spam
## 1    spam; bacon; spam
## 2          potatoes
## 3                None
## 4    spam; eggs; bacon; spam; spam
## dtype: object
```

selecting, say, the first string in each list:

```
xs.str.get(0)
## 0      spam
## 1      spam
## 2    potatoes
## 3        None
## 4      spam
## dtype: object
```

or slicing:

```
xs.str.slice(0, -1) # like xs.iloc[i][0:-1] for all i
## 0                []
## 1                [spam, bacon]
## 2                []
## 3                None
## 4                [spam, eggs, bacon, spam]
## dtype: object
```

Exercise 14.3 (*) Using `pandas.merge`, join the `countries`⁸, `world_factbook_2020`⁹, and `ssi_2016_dimensions`¹⁰ datasets based on the country names. Note that some manual data cleansing will be necessary beforehand.

Exercise 14.4 (**) Given a `Series` object `xs` that includes lists of strings, convert it to a `0/1` representation.

1. Determine the list of all unique strings; let's call it `xu`.
2. Create a dataframe `x` with `xs.shape[0]` rows and `len(xu)` columns such that `x.iloc[i, j]` is equal to 1 if `xu[j]` is amongst `xs.loc[i]` and equal to 0 otherwise. Set the column names to `xs`.
3. Given `x` (and only `x`: neither `xs` nor `xu`), perform the inverse operation.

For example, for the above `xs` object, `x` should look like:

```
##   bacon  eggs  potatoes  spam
## 0      0     0         0     1
## 1      1     0         0     1
## 2      0     0         1     0
## 3      0     0         0     0
## 4      1     1         0     1
```

14.3 Formatted outputs for reproducible report generation

Some good development practices related to reproducible report generation are discussed in [86, 104, 105]. Note that the paradigm of literate programming was introduced by D. Knuth in [58].

Reports from data analysis can be prepared, e.g., in Jupyter Notebooks or by writing directly to Markdown files which we can later compile to PDF or HTML. Below we briefly discuss how to output nicely formatted objects programmatically.

⁸ <https://github.com/gagolews/teaching-data/raw/master/other/countries.csv>

⁹ https://github.com/gagolews/teaching-data/raw/master/marek/world_factbook_2020.csv

¹⁰ https://github.com/gagolews/teaching-data/raw/master/marek/ssi_2016_dimensions.csv

14.3.1 Formatting strings

Inclusion of textual representation of data stored in existing objects can easily be done using f-strings (formatted string literals; see [Section 2.1.3](#)) of the type `f"...{expression}..."`. For instance:

```
pi = 3.14159265358979323846
f"n = {pi:.2f}"
## 'n = 3.14'
```

creates a string showing the value of the variable `pi` formatted as a float rounded to two places after the decimal separator.

Note (**) Similar functionality can be achieved using the `str.format` method:

```
"n = {:.2f}".format(pi)
## 'n = 3.14'
```

as well as the `%` operator overloaded for strings, which uses `sprintf`-like value placeholders known to some readers from other programming languages (such as C):

```
"n = %.2f" % pi
## 'n = 3.14'
```

14.3.2 str and repr

The `str` and `repr` functions can create string representations of many objects:

```
x = np.array([1, 2, 3])
str(x)
## '[1 2 3]'
repr(x)
## 'array([1, 2, 3])'
```

The former is more human-readable, and the latter is slightly more technical. Note that `repr` often returns an output that can be interpreted as executable Python code with no or few adjustments. Nonetheless, `pandas` objects are amongst the many exceptions to this rule.

14.3.3 Aligning strings

`str.center`, `str.ljust`, `str.rjust` can be used to centre-, left-, or right-align a string so that it is of at least given width. This might make the display thereof more aesthetic. Very long strings, possibly containing whole text paragraphs can be dealt with using the `wrap` and `shorten` functions from the `textwrap` package.

14.3.4 Direct Markdown output in Jupyter

Further, with IPython/Jupyter, we can output strings that will be directly interpreted as Markdown-formatted:

```
import IPython.display
x = 2+2
out = f"*Result*:  $2^2=2\cdot 2={x}$ ." # LaTeX math
IPython.display.Markdown(out)
```

Result: $2^2 = 2 \cdot 2 = 4$.

Recall from [Section 1.2.5](#) that Markdown is a very flexible markup¹¹ language that allows us to define itemised and numbered lists, mathematical formulae, tables, images, etc.

On a side note, data frames can be nicely prepared for display in a report using `pandas.DataFrame.to_markdown`.

14.3.5 Manual Markdown file output (*)

We can also generate Markdown code programmatically in the form of standalone `.md` files:

```
import tempfile, os.path
filename = os.path.join(tempfile.mkdtemp(), "test-report.md")
f = open(filename, "w") # open for writing (overwrite if exists)
f.write("**Yummy Foods** include, but are not limited to:\n\n")
x = ["spam", "bacon", "eggs", "spam"]
for e in x:
    f.write(f"* {e}\n")
f.write("\nAnd now for something *completely* different:\n\n")
f.write("Rank | Food\n")
f.write("-----|-----\n")
for i in range(len(x)):
    f.write(f"{i+1:4} | {x[i][::-1]:10}\n")
f.close()
```

Here is the resulting raw Markdown source file:

```
with open(filename, "r") as f: # will call f.close() automatically
    out = f.read()
print(out)
## **Yummy Foods** include, but are not limited to:
##
## * spam
```

(continues on next page)

¹¹ (*) Markdown is amongst many markup languages. Other learn-worthy ones include HTML (for the Web) and LaTeX (especially for the beautiful typesetting of maths, print-ready articles, and books, e.g., PDF; see [72] for a comprehensive introduction).

(continued from previous page)

```
## * bacon
## * eggs
## * spam
##
## And now for something *completely* different:
##
## Rank | Food
## -----|-----
##      1 | maps
##      2 | nocab
##      3 | sgge
##      4 | maps
```

We can convert it to other formats, including HTML, PDF, EPUB, ODT, and even presentations by running¹² the **pandoc**¹³ tool. We may also embed it directly inside an IPython/Jupyter notebook:

```
IPython.display.Markdown(out)
```

Yummy Foods include, but are not limited to:

- spam
- bacon
- eggs
- spam

And now for something *completely* different:

| Rank | Food |
|------|-------|
| 1 | maps |
| 2 | nocab |
| 3 | sgge |
| 4 | maps |

Note Figures created in **matplotlib** can be exported to PNG, SVG, or PDF files using the **matplotlib.pyplot.savefig** function. We can include them manually in a Markdown document using the `![description](filename)` syntax.

Note (*) IPython/Jupyter Notebooks can be converted to different formats using the

¹² External programs can be executed using `subprocess.run`.

¹³ <https://pandoc.org/>

`jupyter-nbconvert`¹⁴ command line tool. `jupyter-text`¹⁵ can create notebooks from ordinary text files. Literate programming with mixed R and Python is possible with the R packages `knitr`¹⁶ and `reticulate`¹⁷. See [76] for an overview of many more options.

14.4 Regular expressions (*)

This section contains large excerpts from yours truly's other work [35].

Regular expressions (regexes) provide concise grammar for defining systematic patterns which can be sought in character strings. Examples of such patterns include: specific fixed substrings, emojis of any kind, standalone sequences of lower-case Latin letters ("words"), substrings that can be interpreted as real numbers (with or without fractional parts, also in scientific notation), telephone numbers, email addresses, or URLs.

Theoretically, the concept of regular pattern matching dates to the so-called regular languages and finite state automata [57]; see also [79] and [52]. Regexes, in the form as we know it today, were already present in one of the pre-UNIX implementations of the command-line text editor `qed` [80] (the predecessor of the well-known `sed`).

Most programming languages and text editors (including `Kate`¹⁸ and `VSCodium`¹⁹) support regex-based pattern finding and replacing. This is why regular expressions should be amongst the instruments at every data scientist's disposal.

14.4.1 Regex matching with `re` (*)

In Python, the `re` module implements a regular expression matching engine. It accepts patterns that follow similar syntax to the one available in the Perl language.

Before we proceed with a detailed discussion on how to read and write regular expressions, let's first review some of the methods for identifying the matching substrings. Below we use the `r"\bni+\b"` regex as an example. It catches "n" followed by at least one "i" that begins and ends at a *word boundary*. In other words, we seek "ni", "nii", "niii", etc. which may be considered standalone words.

In particular, `re.findall` extracts all non-overlapping matches to a given regex:

```
import re
x = "We're the knights who say ni! niiii! ni! ni!!!!!!!"
```

(continues on next page)

¹⁴ <https://pypi.org/project/nbconvert>

¹⁵ <https://jupyter-text.readthedocs.io/en/latest>

¹⁶ <https://yihui.org/knitr>

¹⁷ <https://rstudio.github.io/reticulate>

¹⁸ <https://kate-editor.org/>

¹⁹ <https://vscodium.com/>

(continued from previous page)

```
re.findall(r"\bni+\b", x)
## ['ni', 'niiiii', 'ni', 'niiiiiiii']
```

The order of arguments is (*look for what, where*), not vice versa.

Important We used the `r"..."` prefix to input a string so that `"\b"` is not treated as an escape sequence which denotes the backspace character. Otherwise, the foregoing would have to be written as `"\\bni+\\b"`.

If we had not insisted on matching at the word boundaries (i.e., if we used the simple `"ni+"` regex instead), we would also match the `"ni"` in `"knights"`.

The `re.search` function returns an object of the class `re.Match` that enables us to get some more information about the first match:

```
r = re.search(r"\bni+\b", x)
r.start(), r.end(), r.group()
## (26, 28, 'ni')
```

It includes the start and the end position (index) as well as the match itself. If the regex contains *capture groups* (more details follow), we can also pinpoint the matches thereto.

Moreover, `re.finditer` returns an iterable object that includes the same details, but now about all the matches:

```
rs = re.finditer(r"\bni+\b", x)
for r in rs:
    print((r.start(), r.end(), r.group()))
## (26, 28, 'ni')
## (30, 36, 'niiiii')
## (38, 40, 'ni')
## (42, 52, 'niiiiiiii')
```

`re.split` divides a string into chunks separated by matches to a given regex:

```
re.split(r"!s+", x)
## ["We're the knights who say ni", 'niiiii', 'ni', 'niiiiiii!']
```

The `"!s+"` regex matches the exclamation mark followed by one or more whitespace characters.

Using `re.sub`, each match can be replaced with a given string:

```
re.sub(r"\bni+\b", "nu", x)
## "We're the knights who say nu! nu! nu! nu!"
```

Note (**) More flexible replacement strings can be generated by passing a custom function as the second argument:

```
re.sub(r"\bni+\b", lambda m: "n" + "u"*(m.end()-m.start()-1), x)
## "We're the knights who say nu! nuuuuu! nu! nuuuuuuuu!"
```

14.4.2 Regex matching with pandas (*)

The `pandas.Series.str` accessor also defines a number of vectorised functions that utilise the `re` package's matcher.

Example Series object:

```
x = pd.Series(["ni!", "niii, ni, nii!", None, "spam, bacon", "nii, ni!"])
x
## 0          ni!
## 1    niii, ni, nii!
## 2          None
## 3    spam, bacon
## 4      nii, ni!
## dtype: object
```

Here are the most notable functions:

```
x.str.contains(r"\bni+\b")
## 0      True
## 1      True
## 2      None
## 3     False
## 4      True
## dtype: object
x.str.count(r"\bni+\b")
## 0      1.0
## 1      3.0
## 2      NaN
## 3      0.0
## 4      2.0
## dtype: float64
x.str.replace(r"\bni+\b", "nu", regex=True)
## 0      nu!
## 1    nu, nu, nu!
## 2      None
## 3    spam, bacon
## 4      nu, nu!
```

(continues on next page)

(continued from previous page)

```

## dtype: object
x.str.findall(r"\bni+\b")
## 0      [ni]
## 1    [niiii, ni, nii]
## 2      None
## 3      []
## 4    [nii, ni]
## dtype: object
x.str.split(r",\s+") # a comma, one or more whitespaces
## 0      [ni!]
## 1    [niiii, ni, nii!]
## 2      None
## 3    [spam, bacon]
## 4    [nii, ni!]
## dtype: object

```

In the two last cases, we get lists of strings as results.

Also, later we will mention `pandas.Series.str.extract` and `pandas.Series.str.extractall` which work with regexes that include capture groups.

Note (*) If we intend to seek matches to the same pattern in many different strings without the use of `pandas`, it might be faster to precompile a regex first, and then use the `re.Pattern.findall` method instead or `re.findall`:

```

p = re.compile(r"\bni+\b") # returns an object of the class `re.Pattern`
p.findall("We're the Spanish Inquisition ni! ni! niiii! nininiiiiiii!")
## ['ni', 'ni', 'niiii']

```

14.4.3 Matching individual characters (*)

In the coming subsections, we review the most essential elements of the regex syntax as we did in [35]. One general introduction to regexes is [31]. The `re` module flavour is summarised in the official [manual](https://docs.python.org/3/library/re.html)²⁰, see also [60].

We begin by discussing different ways to define character sets. In this part, determining the length of all matching substrings will be straightforward.

Important The following characters have special meaning to the regex engine: “.”, “\”, “[”, “(”, “)”, “[“, “]”, “{“, “}”, “^”, “\$”, “*”, “+”, and “?”.

Any regular expression that contains none of the preceding characters behaves like a fixed pattern:

²⁰ <https://docs.python.org/3/library/re.html>

```
re.findall("spam", "spam, eggs, spam, bacon, sausage, and spam")
## ['spam', 'spam', 'spam']
```

There are three occurrences of a pattern that is comprised of four code points, “s” followed by “p”, then by “a”, and ending with “m”.

If we want to include a special character as part of a regular expression so that it is treated literally, we will need to escape it with a backslash, “\”.

```
re.findall(r"\.", "spam...")
## ['.', '.', '.']
```

Matching anything (almost) (*)

The (unescaped) dot, “.”, matches any code point except the newline.

```
x = "Spam, ham,\njam, SPAM, eggs, and spam"
re.findall("..am", x, re.IGNORECASE)
## ['Spam', ' ham', 'SPAM', 'spam']
```

It extracted non-overlapping substrings of length four that end with “am”, case-insensitively.

The dot’s insensitivity to the newline character is motivated by the need to maintain compatibility with tools such as **grep** (when searching within text files in a line-by-line manner). This behaviour can be altered by setting the DOTALL flag.

```
re.findall("..am", x, re.DOTALL|re.IGNORECASE) # '|' is the bitwise OR
## ['Spam', ' ham', '\njam', 'SPAM', 'spam']
```

Defining character sets (*)

Sets of characters can be introduced by enumerating their members within a pair of square brackets. For instance, “[abc]” denotes the set {*a*, *b*, *c*} – such a regular expression matches one (and only one) symbol from this set. Moreover, in:

```
re.findall("[hj]am", x)
## ['ham', 'jam']
```

the “[hj]am” regex matches: “h” or “j”, followed by “a”, followed by “m”. In other words, “ham” and “jam” are the only two strings that are matched by this pattern (unless matching is done case-insensitively).

Important The following characters, if used within square brackets, may be treated not literally: “\”, “[“, “]”, “^”, “_”, “&”, “~”, and “|”.

To include them as-is in a character set, the backslash-escape must be used. For example, “[\\[\\]\\]” matches a backslash or a square bracket.

Complementing sets (*)

Including “^” (the caret) after the opening square bracket denotes a set’s complement. Hence, “[^abc]” matches any code point except “a”, “b”, and “c”. Here is an example where we seek any substring that consists of four non-spaces:

```
x = "Nobody expects the Spanish Inquisition!"
re.findall("[^ ][^ ][^ ][^ ]", x)
## ['Nobo', 'expe', 'Span', 'Inqu', 'isit', 'ion!']
```

Defining code point ranges (*)

Each Unicode character can be referenced by its unique [numeric code](#)²¹. For instance, “a” is assigned code U+0061 and “z” is mapped to U+007A. In the pre-Unicode era (mostly with regard to the ASCII codes, ≤ U+007F, representing English letters, decimal digits, as well as some punctuation and control characters), we were used to relying on specific code ranges. For example, “[a-z]” denotes the set comprised of all characters with codes between U+0061 and U+007A, i.e., lowercase letters of the English (Latin) alphabet.

```
re.findall("[0-9A-Za-z]", "Gągolewski")
## ['G', 'g', 'o', 'l', 'e', 'w', 's', 'k', 'i']
```

This pattern denotes the union of three code ranges: ASCII upper- and lowercase letters and digits. Nowadays, in the processing of text in natural languages, this notation should be avoided. Note the missing “ą” (Polish “a” with ogonek) in the result.

Using predefined character sets (*)

Consider a string:

```
x = "aqbß&AAB 12  ,.;'! \t-+=\n[]{}\"","
```

Some glyphs are not available in the PDF version of this book because we did not install the required fonts, e.g., the Arabic digit 4 or left and right arrows. However, they are well-defined at the program level.

Noteworthy Unicode-aware code point classes include the *word* characters:

```
re.findall(r"\w", x)
## ['a', 'q', 'b', 'ß', 'β', 'Æ', 'A', 'A', 'B', ' ', '1', '2', ' ', ' ', ' ']
```

decimal digits:

²¹ <https://www.unicode.org/charts>

```
re.findall(r"\d", x)
## ['1', '2', ' ', ' ']
```

and whitespaces:

```
re.findall(r"\s", x)
## [' ', '\t', '\n']
```

Moreover, e.g., “\w” is equivalent to “[^\\w]”, i.e., denotes the set’s complement.

14.4.4 Alternating and grouping subexpressions (*)

Alternation operator (*)

The alternation operator, “|” (the pipe or bar), matches either its left or its right branch. For instance:

```
x = "spam, egg, ham, jam, algae, and an amalgam of spam, all al dente"
re.findall("spam|ham", x)
## ['spam', 'ham', 'spam']
```

Grouping subexpressions (*)

The “|” operator has very low precedence (otherwise, we would match “spamam” or “spaham” above instead). If we want to introduce an alternative of *subexpressions*, we need to group them using the “(?:...)” syntax. For instance, “(?:sp|h)am” matches either “spam” or “ham”.

Notice that the bare use of the round brackets, “(...)” (i.e., without the “(?:...)” part, has the side-effect of creating new capturing groups; see below for more details.

Also, matching is always done left-to-right, on the first-come, first-served (greedy) basis. Consequently, if the left branch is a subset of the right one, the latter will never be matched. In particular, “(?:al|alga|algae)” can only match “al”. To fix this, we can write “(?:algae|alga|al)”.

Non-grouping parentheses (*)

Some parenthesised subexpressions – those in which the opening bracket is followed by the question mark – have a distinct meaning. In particular, “(?:#...)” denotes a free-format comment that is ignored by the regex parser:

```
re.findall(
    "(?# match 'sp' or 'h')(?:sp|h)(?# and 'am')am(?:# or match 'egg')egg",
    x
)
## ['spam', 'egg', 'ham', 'spam']
```

This is just horrible. Luckily, constructing more sophisticated regexes by concatenating subfragments thereof is more readable:

```
re.findall(
    "(?:sp|h)" + # match either 'sp' or 'h'
    "am" +      # followed by 'am'
    "|" +       # ... or ...
    "egg",      # just match 'egg'
    x
)
## ['spam', 'egg', 'ham', 'spam']
```

What is more, e.g., “(?*i*)” enables the case-insensitive mode.

```
re.findall("(?i)spam", "Spam spam SPAMITY spam")
## ['Spam', 'spam', 'SPAM', 'spam']
```

14.4.5 Quantifiers (*)

More often than not, a *variable* number of instances of the same subexpression needs to be captured. Sometimes we want to make its presence optional. These can be achieved by means of the following quantifiers:

- “?” matches 0 or 1 time;
- “*” matches 0 or more times;
- “+” matches 1 or more times;
- “{*n*,*m*}” matches between *n* and *m* times;
- “{*n*,}” matches at least *n* times;
- “{*n*}” matches exactly *n* times.

These operators are applied onto the directly preceding atoms. For example, “*ni*+” captures “*ni*”, “*nii*”, “*niii*”, etc., but neither “*n*” alone nor “*ninini*” altogether.

By default, the quantifiers are greedy – they match the repeated subexpression as many times as possible. The “?” suffix (forming quantifiers such as “??”, “*?”, “+?”, and so forth) tries with as few occurrences as possible (to obtain a match still).

Greedy:

```
x = "sp(AM)(maps)(SP)am"
re.findall(r"\(.+\)", x)
## ['(AM)(maps)(SP)']
```

Lazy:

```
re.findall(r"\(.+?\)", x)
## ['(AM)', '(maps)', '(SP)']
```

Greedy (but clever):

```
re.findall(r"\([^)]+\)", x)
## ['(AM)', '(maps)', '(SP)']
```

The first regex is greedy: it matches an opening bracket, then as many characters as possible (including “)”) that are followed by a closing bracket. The two other patterns terminate as soon as the first closing bracket is found.

More examples:

```
x = "spamamamnomnomnomammmmmmmmm"
re.findall("sp(?:am|nom)+", x)
## ['spamamamnomnomnomam']
re.findall("sp(?:am|nom)+?", x)
## ['spam']
```

And:

```
re.findall("sp(?:am|nom)+?m*", x)
## ['spam']
re.findall("sp(?:am|nom)+?m+", x)
## ['spamamamnomnomnomammmmmmmmm']
```

Let’s stress that the quantifier is applied to the subexpression that stands directly before it. Grouping parentheses can be used in case they are needed.

```
x = "12, 34.5, 678.901234, 37...629, ..."
re.findall(r"\d+(\?:\.\d+)", x)
## ['34.5', '678.901234']
```

matches digits, a dot, and another series of digits.

```
re.findall(r"\d+(\?:\.\d+)?", x)
## ['12', '34.5', '678.901234', '37', '629']
```

finds digits which are possibly (but not necessarily) followed by a dot and a digit sequence.

Exercise 14.5 Write a regex that extracts all *#hashtags* from a string `#omg #SoEasy`.

14.4.6 Capture groups and references thereto (**)

Round-bracketed subexpressions (without the “?:” prefix) form the so-called *capture groups* that can be extracted separately or be referred to in other parts of the same regex.

Extracting capture group matches (**)

The preceding statement can be nicely verified by calling `re.findall`:

```
x = "name='Sir Launcelot', quest='Seek Grail', favcolour='blue'"
re.findall(r"(\w+)='(?:)'", x)
## [('name', 'Sir Launcelot'), ('quest', 'Seek Grail'), ('favcolour', 'blue')]
```

It returned the matches to the individual capture groups, not the whole matching sub-strings.

re.find and **re.finditer** can pinpoint each component:

```
r = re.search(r"(\w+)='(?:)'", x)
print("whole (0):", (r.start(), r.end(), r.group()))
print("      1 :", (r.start(1), r.end(1), r.group(1)))
print("      2 :", (r.start(2), r.end(2), r.group(2)))
## whole (0): (0, 20, "name='Sir Launcelot'")
##          1 : (0, 4, 'name')
##          2 : (6, 19, 'Sir Launcelot')
```

Here its vectorised version using **pandas**, returning the first match:

```
y = pd.Series([
    "name='Sir Launcelot'",
    "quest='Seek Grail'",
    "favcolour='blue', favcolour='yel.. Aaargh!'"
])
y.str.extract(r"(\w+)='(?:)'")
##          0          1
## 0      name  Sir Launcelot
## 1      quest   Seek Grail
## 2  favcolour      blue
```

We see that the findings are conveniently presented in the data frame form. The first column gives the matches to the first capture group. All matches can be extracted too:

```
y.str.extractall(r"(\w+)='(?:)'")
##          0          1
## match
## 0 0      name  Sir Launcelot
## 1 0      quest   Seek Grail
## 2 0  favcolour      blue
## 1 1  favcolour  yel.. Aaargh!
```

Recall that if we just need the grouping part of “(...)”, i.e., without the capturing feature, “(?:...)” can be applied.

Also, named capture groups defined like “(?P<name>...)” are supported.

```
y.str.extract("(?:\\w+)='(?P<value>.+?)'")
##          value
```

(continues on next page)

(continued from previous page)

```
## 0  Sir Launcelot
## 1    Seek Grail
## 2         blue
```

Replacing with capture group matches (**)

When using `re.sub` and `pandas.Series.str.replace`, matches to particular capture groups can be recalled in replacement strings. The match in its entirety is denoted by “\g<0>”, then “\g<1>” stores whatever was caught by the first capture group, and “\g<2>” is the match to the second capture group, etc.

```
re.sub(r"(\w+)='(.*?)'", r"\g<2> is a \g<1>", x)
## 'Sir Launcelot is a name, Seek Grail is a quest, blue is a favcolour'
```

Named capture groups can be referred to too:

```
re.sub(r"(?P<key>\w+)='(?P<value>.*?)'",
      r"\g<value> is a \g<key>", x)
## 'Sir Launcelot is a name, Seek Grail is a quest, blue is a favcolour'
```

Back-referencing (**)

Matches to capture groups can also be part of the regexes themselves. In such a context, e.g., “\1” denotes whatever has been consumed by the first capture group.

In general, parsing HTML code with regexes is not recommended, unless it is well-structured (which might be the case if it is generated programmatically; but we can always use the `lxml` package). Despite this, let’s consider the following examples:

```
x = "<p><em>spam</em></p><code>eggs</code>"
re.findall(r"<[a-z]+>.*?</[a-z]+>", x)
## ['<p><em>spam</em>', '<code>eggs</code>']
```

It did not match the correct closing HTML tag. But we can make this happen by writing:

```
re.findall(r"(<[a-z]+>).*?</\2>", x)
## [('<p><em>spam</em></p>', 'p'), ('<code>eggs</code>', 'code')]
```

This regex guarantees that the match will include all characters between the opening “<tag>” and the corresponding (not: any) closing “</tag>”.

Named capture groups can be referenced using the “(?P=name)” syntax:

```
re.findall(r"(<(?P<tagname>[a-z]+>).*?</(?P=tagname)>)", x)
## [('<p><em>spam</em></p>', 'p'), ('<code>eggs</code>', 'code')]
```

The angle brackets are part of the token.

14.4.7 Anchoring (*)

Lastly, let's mention the ways to match a pattern at a given abstract position within a string.

Matching at the beginning or end of a string (*)

“^” and “\$” match, respectively, start and end of the string (or each line within a string, if the `re.MULTILINE` flag is set).

```
x = pd.Series(["spam egg", "bacon spam", "spam", "egg spam bacon", "milk"])
rs = ["spam", "^spam", "spam$", "spam$|^spam", "^spam$"] # regexes to test
```

The five regular expressions match "spam", respectively, anywhere within the string, at the beginning, at the end, at the beginning or end, and in strings that are equal to the pattern itself. We can check this by calling:

```
pd.concat([x.str.contains(r) for r in rs], axis=1, keys=rs)
##      spam  ^spam  spam$  spam$|^spam  ^spam$
## 0    True    True   False           True   False
## 1    True   False    True           True   False
## 2    True    True    True           True    True
## 3    True   False   False          False   False
## 4   False   False   False          False   False
```

Exercise 14.6 *Compose a regex that does the same job as `str.strip`.*

Matching at word boundaries (*)

What is more, “\b” matches at a “word boundary”, e.g., near spaces, punctuation marks, or at the start/end of a string (i.e., wherever there is a transition between a word, “\w”, and a non-word character, “\W”, or vice versa).

In the following example, we match all stand-alone numbers (this regular expression is imperfect, though):

```
re.findall(r"[+]?[b\d+(?:\.\d+)?[b]", "+12, 34.5, -5.3243")
## ['+12', '34.5', '-5.3243']
```

Looking behind and ahead (**)

There is a way to guarantee that a pattern occurrence begins or ends with a match to a subexpression: “(?<=...)” denotes the *look-behind*, whereas “...(?=...)” designates a *look-ahead*.

```
x = "I like spam, spam, eggs, and spam."
re.findall(r"\b\w+\b(?:[,.]")", x)
## ['spam', 'spam', 'eggs', 'spam']
```

This regex captured words that end with a comma or a dot

Moreover, “(?<!\. . .) . . .” and “. . . (?!\. . .)” are their *negated* versions (negative look-behind/ahead).

```
re.findall(r"\b\w+\b(?![,\.])", x)
## ['I', 'like', 'and']
```

This time, we matched the words that end with neither a comma nor a dot.

14.5 Exercises

Exercise 14.7 List some ways to normalise character strings.

Exercise 14.8 (**) What are the challenges of processing non-English text?

Exercise 14.9 What are the problems with the “[A-Za-z]” and “[A-z]” character sets?

Exercise 14.10 Name the two ways to turn on case-insensitive regex matching.

Exercise 14.11 What is a word boundary?

Exercise 14.12 What is the difference between the “^” and “\$” anchors?

Exercise 14.13 When would we prefer using “[0-9]” instead of “[d]”?

Exercise 14.14 What is the difference between the “?”, “??”, “*”, “*?”, “+”, and “+?” quantifiers?

Exercise 14.15 Does “. ” match all the characters?

Exercise 14.16 What are named capture groups and how can we refer to the matches thereto in `re.sub`?

Exercise 14.17 Write a regex that extracts all standalone numbers accepted by Python, including 12.123, -53, +1e-9, -1.2423e10, 4. and .2.

Exercise 14.18 Author a regex that matches all email addresses.

Exercise 14.19 Indite a regex that matches all URLs starting with `http://` or `https://`.

Exercise 14.20 Cleanse the `warsaw_weather`²² dataset so that it contains analysable numeric data.

²² https://github.com/gagolews/teaching-data/raw/master/marek/warsaw_weather.csv

Missing, censored, and questionable data

Up to now, we have been mostly assuming that observations are of decent quality, i.e., trustworthy. It would be nice if that was always the case, but it is not.

In this chapter, we briefly address the most rudimentary methods for dealing with *suspicious* observations: outliers, missing, censored, imprecise, and incorrect data.

15.1 Missing data

Consider an excerpt from National Health and Nutrition Examination Survey that we played with in [Chapter 12](#):

```
nhanes = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/nhanes_p_demo_bmx_2020.csv",
    comment="#")
nhanes.loc[:, ["BMXWT", "BMXHT", "RIDAGEYR", "BMIHEAD", "BMXHEAD"]].head()
```

| ## | BMXWT | BMXHT | RIDAGEYR | BMIHEAD | BMXHEAD |
|------|-------|-------|----------|---------|---------|
| ## 0 | NaN | NaN | 2 | NaN | NaN |
| ## 1 | 42.2 | 154.7 | 13 | NaN | NaN |
| ## 2 | 12.0 | 89.3 | 2 | NaN | NaN |
| ## 3 | 97.1 | 160.2 | 29 | NaN | NaN |
| ## 4 | 13.6 | NaN | 2 | NaN | NaN |

Some of the columns bear NaN (not-a-number) values. They are used here to encode *missing* (not available) data. Previously, we decided not to be bothered by them: a shy call to `dropna` resulted in their removal. But we are curious now.

The reasons behind why some items are missing might be numerous, in particular:

- a participant did not know the answer to a given question;
- someone refused to answer a given question;
- a person did not take part in the study anymore (attrition, death, etc.);
- an item was not applicable (e.g., number of minutes spent cycling weekly when someone answered they did not learn to ride a bike yet);
- a piece of information was not collected, e.g., due to the lack of funding or a failure of a piece of equipment.

15.1.1 Representing and detecting missing values

Sometimes missing values are specially encoded, especially in CSV files, e.g., with `-1`, `0`, `9999`, `numpy.inf`, `-numpy.inf`, or `None`, strings such as `"NA"`, `"N/A"`, `"Not Applicable"`, `"---"`. This is why we must always inspect our datasets carefully. To assure consistent representation, we can convert them to `NaN` (as in: `numpy.nan`) in numeric (floating-point) columns or to Python's `None` otherwise.

Vectorised functions such as `numpy.isnan` (or, more generally, `numpy.isfinite`) and `pandas.isnull` as well as `isna` methods for the `DataFrame` and `Series` classes verify whether an item is missing or not.

For instance, here are the counts and proportions of missing values in selected columns of `nhanes`:

```
nhanes.isna().apply([np.sum, np.mean]).T.nlargest(5, "sum") # top 5 only
##              sum      mean
## BMIHEAD      14300.0  1.000000
## BMIRECUM      14257.0  0.996993
## BMIHT         14129.0  0.988042
## BMXHEAD       13990.0  0.978322
## BMIHIP        13924.0  0.973706
```

Looking at the column descriptions on the data provider's [website](#)¹, for example, `BMIHEAD` stands for “Head Circumference Comment”, whereas `BMXHEAD` is “Head Circumference (cm)”, but these were only collected for infants.

Exercise 15.1 Read the column descriptions (refer to the comments in the CSV file for the relevant URLs) to identify the possible reasons for some of the records in `nhanes` being missing.

Exercise 15.2 Learn about the difference between the `pandas.DataFrameGroupBy.size` and `pandas.DataFrameGroupBy.count` methods.

15.1.2 Computing with missing values

Our using `NaN` to denote a missing piece of information is merely an ugly (but functional) hack². The original use case for not-a-number is to represent the results of incorrect operations, e.g., logarithms of negative numbers or subtracting two infinite entities. We thus need extra care when handling them.

Generally, arithmetic operations on missing values yield a result that is undefined as well:

```
np.nan + 2 # "don't know" + 2 == "don't know"
## nan
np.mean([1, np.nan, 2, 3])
## nan
```

¹ https://wwwn.cdc.gov/Nchs/Nhanes/2017-2018/P_BMX.htm

² (*) The R environment, on the other hand, supports missing values out-of-the-box.

There are versions of certain aggregation functions that ignore missing values whatsoever: `numpy.nanmean`, `numpy.nanmin`, `numpy.nanmax`, `numpy.nanpercentile`, `numpy.nanstd`, etc.

```
np.nanmean([1, np.nan, 2, 3])
## 2.0
```

Regrettably, running these aggregation functions directly on Series objects ignores missing entities by default. Compare an application of `numpy.mean` on a Series instance vs on a vector:

```
x = nhanes.head().loc[:, "BMXHT"] # some example Series, whatever
np.mean(x), np.mean(np.array(x))
## (134.73333333333332, nan)
```

This is unfortunate a behaviour as this way we might miss (sic!) the presence of missing values. Therefore, it is crucial to have the dataset carefully inspected in advance.

Also, NaN is of the floating-point type. As a consequence, it cannot be present in, amongst others, logical vectors.

```
x # preview
## 0      NaN
## 1    154.7
## 2     89.3
## 3    160.2
## 4      NaN
## Name: BMXHT, dtype: float64
y = (x > 100)
y
## 0    False
## 1     True
## 2    False
## 3     True
## 4    False
## Name: BMXHT, dtype: bool
```

Unfortunately, comparisons against missing values yield `False`, instead of the more semantically valid missing value. Hence, if we want to retain the missingness information (we do not know if a missing value is greater than 100), we need to do it manually:

```
y = y.astype("object") # required for numpy vectors, not for pandas Series
y[np.isnan(x)] = None
y
## 0    None
## 1     True
## 2    False
## 3     True
```

(continues on next page)

(continued from previous page)

```
## 4      None
## Name: BMXHT, dtype: object
```

Exercise 15.3 Read the *pandas* [documentation](#)³ about missing value handling.

15.1.3 Missing at random or not?

At a general level (from the mathematical modelling perspective), we may distinguish between a few missingness patterns [85]:

- *missing completely at random*: reasons are unrelated to data and probabilities of cases' being missing are all the same;
- *missing at random*: there are different probabilities of being missing within distinct groups (e.g., ethical data scientists might tend to refuse to answer specific questions);
- *missing not at random*: due to reasons unknown to us (e.g., data was collected at different times, there might be significant differences within the groups that we cannot easily identify, e.g., amongst participants with a background in mathematics where we did not ask about education or occupation).

It is important to try to determine the reason for missingness. This will usually imply the kinds of techniques that are suitable in specific cases.

15.1.4 Discarding missing values

We may try removing (discarding) rows or columns that carry at least one, some, or too many missing values. Nonetheless, such a scheme will obviously not work for small datasets, where each observation is precious⁴.

Also, we ought *not* to exercise data removal in situations where missingness is conditional (e.g., data only available for infants) or otherwise group-dependent (not completely at random). Otherwise, for example, it might result in an imbalanced dataset.

Exercise 15.4 With the *nhanes_p_demo_bmx_2020*⁵ dataset, perform what follows.

1. Remove all columns that are comprised of missing values only.
2. Remove all columns that are made of more than 20% missing values.
3. Remove all rows that only consist of missing values.
4. Remove all rows that bear at least one missing value.
5. Remove all columns that carry at least one missing value.

³ https://pandas.pydata.org/pandas-docs/stable/user_guide/missing_data.html

⁴ On the other hand, if we want to infer from small datasets, we should ask ourselves whether this is a good idea at all... It might be better to refrain from any data analysis than to come up with conclusions that are likely to be unjustified.

⁵ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_p_demo_bmx_2020.csv

Hint: `pandas.DataFrame.dropna` might be useful in the simplest cases, and `numpy.isnan` or `pandas.DataFrame.isna` with `loc[...]` or `iloc[...]` can be applied otherwise.

15.1.5 Mean imputation

When we cannot afford or it is inappropriate/inconvenient to proceed with the removal of missing observations or columns, we may try applying some missing value *imputation* techniques. Let's be clear, though: this is merely a replacement thereof by some *hopefully* adequate guesstimates.

Important In all kinds of reports from data analysis, we need to be explicit about the way we handle the missing values. Sometimes they might strongly affect the results.

Consider an example vector with missing values, comprised of heights of the adult participants of the NHANES study.

```
x = nhanes.loc[nhanes.loc[:, "RIDAGEYR"] >= 18, "BMXHT"]
```

The simplest approach is to replace each missing value with the corresponding column's mean. This does not change the overall average but decreases the variance.

```
xi = x.copy()
xi[np.isnan(xi)] = np.nanmean(xi)
```

Similarly, we could consider replacing missing values with the median, or – in the case of categorical data – the mode.

Furthermore, we expect heights to differ, on average, between sexes. Consequently, another imputation option is to replace the missing values with the corresponding within-group averages:

```
xg = x.copy()
g = nhanes.loc[nhanes.loc[:, "RIDAGEYR"] >= 18, "RIAGENDR"]
xg[np.isnan(xg) & (g == 1)] = np.nanmean(xg[g == 1]) # male
xg[np.isnan(xg) & (g == 2)] = np.nanmean(xg[g == 2]) # female
```

Unfortunately, whichever imputation method we choose, will artificially distort the data distribution and introduce some kind of bias; see Figure 15.1 for the histograms of `x`, `xi`, and `xg`. These effects can be obscured if we increase the histogram bins' widths, but they will still be present in the data. No surprise here: we added to the sample many identical values.

Exercise 15.5 With the `nhanes_p_demo_bmx_20206` dataset, perform what follows.

⁶ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_p_demo_bmx_2020.csv




Figure 15.1. The mean imputation distorts the data distribution.

1. For each numerical column, replace all missing values with the column averages.
2. For each categorical column, replace all missing values with the column modes.
3. For each numerical column, replace all missing values with the averages corresponding to a patient's sex (as given by the `RIAGENDR` column).

15.1.6 Imputation by classification and regression (*)

We can easily compose a missing value imputer based on averaging data from an observation's non-missing nearest neighbours; compare [Section 9.2.1](#) and [Section 12.3.1](#). This is an extension of the simple idea of finding the most *similar* observation (with respect to chosen criteria) to a given one and then borrowing non-missing measurements from it.

More generally, different regression or classification models can be built on non-missing data (training sample) and then the missing observations can be replaced by the values predicted by those models.

Note (**) Rubin (e.g., in [63]) suggests the use of a procedure called *multiple imputation* (see also [94]), where copies of the original datasets are created, missing values are imputed by sampling from some estimated distributions, the inference is made, and then the results are aggregated. An example implementation of such an algorithm is available in `sklearn.impute.IterativeImputer`.

15.2 Censored and interval data (*)

Censored data frequently appear in the context of reliability, risk analysis, and biostatistics, where the observed objects might *fail* (e.g., break down, die, withdraw; compare, e.g., [67]). Our introductory course cannot obviously cover everything. However, a beginner analyst needs to be at least aware of the existence of:

- *right-censored* data: we only know that the actual value is greater than the recorded one (e.g., we stopped the experiment on the reliability of light bulbs after 1000 hours, so those which still work will not have their time-of-failure precisely known);
- *left-censored* data: the true observation is less than the recorded one, e.g., we observe a component's failure, but we do not know for how long it has been in operation before the study has started.

In such cases, the recorded datum of, say, 1000, can essentially mean $[1000, \infty)$, $[0, 1000]$, or $(-\infty, 1000]$.

There might also be instances where we know that a value is in some interval $[a, b]$. There are numerical libraries that deal with *interval computations*, and some data analysis methods exist for dealing with such a scenario.

15.3 Incorrect data

Missing data can already be marked in a given sample. But we also might be willing to mark some existing values as missing, e.g., when they are incorrect. For example:

- for text data, misspelled words;
- for spatial data, GPS coordinates of places out of this world, nonexistent zip codes, or invalid addresses;
- for date-time data, misformatted date-time strings, incorrect dates such as “29 February 2011”, an event's start date being after the end date;
- for physical measurements, observations that do not meet specific constraints, e.g., negative ages, or heights of people over 300 centimetres;
- IDs of entities that simply do not exist (e.g., unregistered or deleted clients' accounts);

and so forth.

To be able to identify and handle incorrect data, we need specific knowledge of a particular domain. Optimally, data validation techniques should already be employed on the data collection stage. For instance, when a user submits an online form.

There can be many tools that can assist us with identifying erroneous observations, e.g., spell checkers such as [hunspell](https://hunspell.github.io/)⁷.

For smaller datasets, observations can also be inspected manually. In other cases, we might have to develop custom algorithms for detecting such bugs in data.

Exercise 15.6 *Given some data frame with numeric columns only, perform what follows.*

1. Check if all numeric values in each column are between 0 and 1000.
2. Check if all values in each column are unique.
3. Verify that all the rowwise sums add up to 1.0 (up to a small numeric error).
4. Check if the data frame consists of 0s and 1s only. Provided that this is the case, verify that for each row, if there is a 1 in some column, then all the columns to the right are filled with 1s too.

Many data validation methods can be reduced to operations on strings; see [Chapter 14](#). They may be as simple as writing a single regular expression or checking if a label is in a dictionary of possible values but also as difficult as writing your own parser for a custom context-sensitive grammar.

Exercise 15.7 *Once we import the data fetched from dirty sources, relevant information will have to be extracted from raw text, e.g., strings like "1" should be converted to floating-point numbers. In the sequel, we suggest several tasks that can aid in developing data validation skills involving some operations on text.*

Given an example data frame with text columns (manually invented, please be creative), perform what follows.

1. Remove trailing and leading whitespaces from each string.
2. Check if all strings can be interpreted as numbers, e.g., "23.43".
3. Verify if a date string in the YYYY-MM-DD format is correct.
4. Determine if a date-time string in the YYYY-MM-DD hh:mm:ss format is correct.
5. Check if all strings are of the form (+NN) NNN-NNN-NNN or (+NN) NNNN-NNN-NNN, where N denotes any digit (valid telephone numbers).
6. Inspect whether all strings are valid country names.
7. (*) Given a person's date of birth, sex, and Polish ID number [PESEL](#)⁸, check if that ID is correct.
8. (*) Determine if a string represents a correct International Bank Account Number (IBAN⁹) (note that IBANs have two check digits).
9. (*) Transliterate text to ASCII, e.g., "żółty ☺" to "zołty (C)".
10. (**) Using an external spell checker, determine if every string is a valid English word.

⁷ <https://hunspell.github.io/>

⁸ <https://en.wikipedia.org/wiki/PESEL>

⁹ https://en.wikipedia.org/wiki/International_Bank_Account_Number

11. (**) *Using an external spell checker, ascertain that every string is a valid English noun in the singular form.*
12. (**) *Resolve all abbreviations by means of a custom dictionary, e.g., "Kat." → "Katherine", "Gr." → "Grzegorz".*

15.4 Outliers

Another group of inspectionworthy observations consists of *outliers*. We can define them as the samples that reside in the areas of substantially lower density than their neighbours.

Outliers might be present due to an error, or their being otherwise anomalous, but they may also simply be interesting, original, or novel. After all, statistics does not give any meaning to data items; humans do.

What we do with outliers is a separate decision. We can get rid of them, correct them, replace them with a missing value (and then possibly impute), or analyse them separately. In particular, there is a separate subfield in statistics called extreme value theory that is interested in predicting the distribution of very large observations (e.g., for modelling floods, extreme rainfall, or temperatures); see, e.g., [6]. But this is a topic for a more advanced course; see, e.g., [53]. By then, let's stick with some simpler settings.

15.4.1 The 3/2 IQR rule for normally-distributed data

For unidimensional data (or individual columns in matrices and data frames), the first few smallest and largest observations should usually be inspected manually. For instance, it might happen that someone accidentally entered a patient's height in metres instead of centimetres: such cases are easily detectable. A data scientist is like a detective.

Let's recall the rule of thumb discussed in the section on box-and-whisker plots (Section 5.1.4). For data that are expected to come from a normal distribution, everything that does not fall into the interval $[Q_1 - 1.5\text{IQR}, Q_3 + 1.5\text{IQR}]$ can be considered suspicious. This definition is based on quartiles only, so it is not affected by potential outliers (they are robust aggregates; compare [53]). Plus, the magic constant 1.5 is nicely round and thus easy to memorise (an attractive feature). It is not too small and not too large; for the normal distribution $N(\mu, \sigma)$, the aforementioned interval corresponds to roughly $[\mu - 2.698\sigma, \mu + 2.698\sigma]$, and the probability of obtaining a value outside of it is c. 0.7%. In other words, for a sample of size 1000 that is *truly* normally distributed (not contaminated by anything), only seven observations will be flagged. It is not a problem to inspect them by hand.

Note (*) We can choose a different threshold. For instance, for the normal distribu-

tion $N(10, 1)$, even though the probability of observing a value greater than 15 is *theoretically* non-zero, it is smaller 0.000029%, so it is sensible to treat this observation as suspicious. On the other hand, we do not want to mark too many observations as outliers: inspecting them manually might be too labour-intensive.

Exercise 15.8 For each column in `nhanes_p_demo_bmx_2020`¹⁰, inspect a few smallest and largest observations and see if they make sense.

Exercise 15.9 Perform the foregoing separately for data in each group as defined by the `RIA-GENDR` column.

15.4.2 Unidimensional density estimation (*)

For skewed distributions such as the ones representing incomes, there might be nothing wrong, at least statistically speaking, with very large isolated observations.

For well-separated multimodal distributions on the real line, outliers may sometimes also fall in between the areas of high density.

Example 15.10 That neither box plots themselves, nor the 1.5IQR rule might not be ideal tools for multimodal data is exemplified in Figure 15.2. Here, we have a mixture of $N(10, 1)$ and $N(25, 1)$ samples and four potential outliers at 0, 15, 45, and 50.

```
x = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/blobs2.txt")
plt.subplot(1, 2, 1)
plt.boxplot(x, vert=False)
plt.yticks([])
plt.subplot(1, 2, 2)
plt.hist(x, bins=50, color="lightgray", edgecolor="black")
plt.ylabel("Count")

plt.show()
```

Fixed-radius search techniques discussed in Section 8.4 can be used for estimating the underlying probability density function. Given a data sample $\mathbf{x} = (x_1, \dots, x_n)$, consider¹¹:

$$\hat{f}_r(z) = \frac{1}{2rn} \sum_{i=1}^n |B_r(z)|,$$

where $|B_r(z)|$ denotes the number of observations from $\mathbf{x}$ whose distance to z is not greater than r , i.e., fall into the interval $[z - r, z + r]$.

```
n = len(x)
r = 1 # radius - feel free to play with different values
```

(continues on next page)

¹⁰ https://github.com/gagolews/teaching-data/raw/master/marek/nhanes_p_demo_bmx_2020.csv

¹¹ This is an instance of a kernel density estimator, with the simplest kernel: a rectangular one.




Figure 15.2. With box plots, we may fail to detect some outliers.

(continued from previous page)

```
import scipy.spatial
t = scipy.spatial.KDTree(x.reshape(-1, 1))
dx = pd.Series(t.query_ball_point(x.reshape(-1, 1), r)).str.len() / (2*r*n)
dx[:6] # preview
## 0    0.000250
## 1    0.116267
## 2    0.116766
## 3    0.166667
## 4    0.076098
## 5    0.156188
## dtype: float64
```

Then, points in the sample lying in low-density regions (i.e., all x_i such that $\hat{f}_r(x_i)$ is small) can be flagged for further inspection:

```
x[dx < 0.001]
## array([ 0.          , 13.57157922, 15.          , 45.          , 50.          ])
```

See Figure 15.3 for an illustration of $\hat{f}_r$. Of course, r must be chosen with care, just like the number of bins in a histogram.

```
z = np.linspace(np.min(x)-5, np.max(x)+5, 1001)
dz = pd.Series(t.query_ball_point(z.reshape(-1, 1), r)).str.len() / (2*r*n)
plt.plot(z, dz, label=f"density estimator ( $\hat{f}_r$ )")
plt.hist(x, bins=50, color="lightgray", edgecolor="black", density=True)
plt.ylabel("Density")
plt.show()
```




Figure 15.3. Density estimation based on fixed-radius search.

15.4.3 Multidimensional density estimation (*)

By far we should have become used to the fact that unidimensional data projections might lead to our losing too much information. Some values can seem perfectly fine when they are considered in isolation, but already plotting them in 2D reveals that the reality is more complex than that.

Figure 15.4 depicts the distributions of two natural projections of an example dataset:

```
X = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/blobs1.txt", delimiter=",")
plt.figure(figsize=(plt.rcParams["figure.figsize"][0], )*2) # width=height
plt.subplot(2, 2, 1)
plt.boxplot(X[:, 0], vert=False)
plt.yticks([1], ["X[:, 0]"])
plt.subplot(2, 2, 2)
plt.hist(X[:, 0], bins=20, color="lightgray", edgecolor="black")
plt.title("X[:, 0]")
plt.subplot(2, 2, 3)
plt.boxplot(X[:, 1], vert=False)
plt.yticks([1], ["X[:, 1]"])
plt.subplot(2, 2, 4)
plt.hist(X[:, 1], bins=20, color="lightgray", edgecolor="black")
plt.title("X[:, 1]")
plt.show()
```

There is nothing suspicious here. Or is there?

The scatter plot in Figure 15.5 reveals that the data consist of two well-separable blobs:




Figure 15.4. One-dimensional projections of the blobs1 dataset.

```
plt.plot(X[:, 0], X[:, 1], "o")
plt.axis("equal")
plt.show()
```

There are a few observations that we might mark as outliers. The truth is that yours truly injected eight junk points at the very end of the dataset. Ha.

```
X[-8:, :]
## array([[ -3. ,  3. ],
##        [ 3. ,  3. ],
##        [ 3. , -3. ],
##        [-3. , -3. ],
##        [-3.5,  3.5],
##        [-2.5,  2.5],
```

(continues on next page)




Figure 15.5. Scatter plot of the blobs1 dataset.

(continued from previous page)

```
##      [-2. ,  2. ],
##      [-1.5,  1.5]])
```

Handling multidimensional data requires slightly more sophisticated methods; see, e.g., [2]. A straightforward approach is to check if there are any points within an observation's radius of some assumed size $r > 0$. If that is not the case, we may consider it an outlier. This is a variation on the aforementioned unidimensional density estimation approach¹².

Example 15.11 Consider the following code chunk:

```
t = scipy.spatial.KDTree(X)
n = t.query_ball_point(X, 0.2) # r=0.2 (radius) - play with it yourself
c = np.array(pd.Series(n).str.len())
c[[0, 1, -2, -1]] # preview
## array([42, 30,  1,  1])
```

$c[i]$ gives the number of points within $X[i, :]$'s r -radius (with respect to the Euclidean distance), including the point itself. Consequently, $c[i]==1$ denotes a potential outlier; see Figure 15.6 for an illustration.

¹² (**) We can easily normalise the outputs to get a true 2D kernel density estimator, but multivariate statistics is beyond the scope of this course. In particular, that data might have fixed marginal distributions (projections onto 1D) but their multidimensional images might be very different is beautifully described by the copula theory [70].

```
plt.plot(X[c > 1, 0], X[c > 1, 1], "o", label="normal point")
plt.plot(X[c == 1, 0], X[c == 1, 1], "v", label="outlier")
plt.axis("equal")
plt.legend()
plt.show()
```




Figure 15.6. Outlier detection based on a fixed-radius search for the blobs1 dataset.

15.5 Exercises

Exercise 15.12 How can missing values be represented in *numpy* and *pandas*?

Exercise 15.13 Explain some basic strategies for dealing with missing values in numeric vectors.

Exercise 15.14 Why we ought to be very explicit about the way we handle missing and other suspicious data? Is it advisable to mark as missing (or remove completely) the observations that we dislike or otherwise deem inappropriate, controversial, dangerous, incompatible with our political views, etc.?

Exercise 15.15 Is replacing missing values with the sample arithmetic mean for income data (as in, e.g., the `uk_income_simulated_2020`¹³ dataset) a sensible strategy?

Exercise 15.16 What are the differences between data missing completely at random, missing at random, and missing not at random?

¹³ https://github.com/gagolews/teaching-data/raw/master/marek/uk_income_simulated_2020.txt

Exercise 15.17 *List some strategies for dealing with data that might contain outliers.*

So far, we have been using **numpy** and **pandas** mostly for storing:

- *independent* measurements, where each row gives, e.g., weight, height, ... records of a different subject; we often consider these a sample of a representative subset of one or more populations, each recorded at a particular point in time;
- data summaries to be reported in the form of tables or figures, e.g., frequency distributions giving counts for the corresponding categories or labels.

In this chapter, we will explore the most basic concepts related to the wrangling of *time series*, i.e., signals indexed by discrete time. Usually, a time series is a sequence of measurements sampled at equally spaced moments, e.g., a patient's heart rate probed every second, daily average currency exchange rates, or highest yearly temperatures recorded in some location.

16.1 Temporal ordering and line charts

Consider the midrange¹ daily temperatures in degrees Celsius at the Spokane International Airport (Spokane, WA, US) between 1889-08-01 (the first observation) and 2021-12-31 (the last observation).

```
temps = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/spokane_temperature.txt")
```

Let's preview the December 2021 data:

```
temps[-31:] # last 31 days
## array([ 11.9,  5.8,  0.6,  0.8, -1.9, -4.4, -1.9,  1.4, -1.9,
##        -1.4,  3.9,  1.9,  1.9, -0.8, -2.5, -3.6, -10. , -1.1,
##        -1.7, -5.3, -5.3, -0.3,  1.9, -0.6, -1.4, -5. , -9.4,
##        -12.8, -12.2, -11.4, -11.4])
```

Here are some data aggregates for the whole sample. First, the popular quantiles:

¹ Note that midrange, being the mean of the lowest and the highest observed temperature on a given day, is not a particularly good estimate of the average daily reading. This dataset is considered for illustrational purposes only.

```
np.quantile(temps, [0, 0.25, 0.5, 0.75, 1])
## array([-26.9,  2.2,  8.6, 16.4, 33.9])
```

Then, the arithmetic mean and standard deviation:

```
np.mean(temps), np.std(temps)
## (8.990273958441021, 9.16204388619955)
```

A graphical summary of the data distribution is depicted in Figure 16.1.

```
plt.violinplot(temps, vert=False, showextrema=False)
plt.boxplot(temps, vert=False)
plt.show()
```




Figure 16.1. Distribution of the midrange daily temperatures in Spokane in the period 1889–2021. Observations are treated as a bag of unrelated items (temperature on a “randomly chosen day” in a version of planet Earth where there is no climate change).

When computing data aggregates or plotting histograms, the order of elements does not matter. Contrary to the case of the *independent* measurements, vectors representing time series do not have to be treated simply as mixed bags of unrelated items.

Important In time series, for any given item x_i , its neighbouring elements x_{i-1} and x_{i+1} denote the recordings occurring directly before and after it. We can use this *temporal ordering* to model how *consecutive* measurements *depend* on each other, describe how they change over time, forecast future values, detect seasonal and long-time trends, and so forth.

Figure 16.2 depicts the data for 2021, plotted as a function of time. What we see is often referred to as a *line chart* (line graph): data points are connected by straight line segments. There are some visible seasonal variations, such as, well, obviously, that winter is colder than summer. There is also some natural variability on top of seasonal patterns typical for the Northern Hemisphere.

```
plt.plot(temps[-365:])
plt.xticks([0, 181, 364], ["2021-01-01", "2021-07-01", "2021-12-31"])
plt.show()
```




Figure 16.2. Line chart of midrange daily temperatures in Spokane for 2021.

16.2 Working with date-times and time-deltas

16.2.1 Representation: The UNIX epoch

`numpy.datetime64`² is a type to represent date-times. Usually, we will be creating dates from strings. For instance:

```
d = np.array([
    "1889-08-01", "1970-01-01", "1970-01-02", "2021-12-31", "today"
], dtype="datetime64[D]")
d
## array(['1889-08-01', '1970-01-01', '1970-01-02', '2021-12-31',
        '2026-01-20'], dtype='datetime64[D]')
```

² <https://numpy.org/doc/stable/reference/arrays.datetime.html>

Similarly with date-times:

```
dt = np.array(["1970-01-01T02:01:05", "now"], dtype="datetime64[s]")
dt
## array(['1970-01-01T02:01:05', '2026-01-20T09:29:53'],
##       dtype='datetime64[s]')
```

Important Internally, date-times are represented as the number of days or seconds (datetime64[D] or datetime64[s]) since the *UNIX Epoch*, 1970-01-01T00:00:00 in the UTC time zone.

Let's verify the preceding statement:

```
d.astype(int)
## array([-29372,         0,         1, 18992, 20473])
dt.astype(int)
## array([         7265, 1768901393])
```

When we think about it for a while, this is exactly what we expected.

Exercise 16.1 (*) *Compose a regular expression that extracts all dates in the YYYY-MM-DD format from a (possibly long) string and converts them to datetime64.*

16.2.2 Time differences

Computing date-time differences (time-deltas) is possible thanks to the numpy. timedelta64 objects:

```
d = np.timedelta64(1, "D") # minus 1 Day
## array(['1889-07-31', '1969-12-31', '1970-01-01', '2021-12-30',
##       '2026-01-19'], dtype='datetime64[D]')
dt + np.timedelta64(12, "h") # plus 12 hours
## array(['1970-01-01T14:01:05', '2026-01-20T21:29:53'],
##       dtype='datetime64[s]')
```

Also, numpy.arange (see also pandas.date_range) generates a sequence of equidistant date-times:

```
dates = np.arange("1889-08-01", "2022-01-01", dtype="datetime64[D]")
dates[:3] # preview
## array(['1889-08-01', '1889-08-02', '1889-08-03'], dtype='datetime64[D]')
dates[-3:] # preview
## array(['2021-12-29', '2021-12-30', '2021-12-31'], dtype='datetime64[D]')
```

16.2.3 Date-times in data frames

Dates and date-times can be emplaced in pandas data frames:

```
spokane = pd.DataFrame(dict(
    date=np.arange("1889-08-01", "2022-01-01", dtype="datetime64[D]"),
    temp=temps
))
spokane.head()
##           date  temp
## 0 1889-08-01  21.1
## 1 1889-08-02  20.8
## 2 1889-08-03  22.2
## 3 1889-08-04  21.7
## 4 1889-08-05  18.3
```

When we ask the date column to become the data frame's index (i.e., row labels), we will be able to select date ranges easily with `loc[...]` and string slices (refer to the manual of `pandas.DatetimeIndex` for more details).

```
spokane.set_index("date").loc["2021-12-25":, :].reset_index()
##           date  temp
## 0 2021-12-25  -1.4
## 1 2021-12-26  -5.0
## 2 2021-12-27  -9.4
## 3 2021-12-28 -12.8
## 4 2021-12-29 -12.2
## 5 2021-12-30 -11.4
## 6 2021-12-31 -11.4
```

Example 16.2 *Figure 16.3 depicts the temperatures in the last five years:*

```
x = spokane.set_index("date").loc["2017-01-01":, "temp"].reset_index()
plt.plot(x.date, x.temp)
plt.show()
```

The `pandas.to_datetime` function can also convert arbitrarily formatted date strings, e.g., "MM/DD/YYYY" or "DD.MM.YYYY" to Series of `datetime64s`.

```
dates = ["05.04.1991", "14.07.2022", "21.12.2042"]
dates = pd.Series(pd.to_datetime(dates, format="%d.%m.%Y"))
dates
## 0    1991-04-05
## 1    2022-07-14
## 2    2042-12-21
## dtype: datetime64[ns]
```

Exercise 16.3 *From the `birth_dates`³ dataset, select all people less than 18 years old (as of the current day).*

Several date-time functions and related properties can be referred to via the `pandas.Series.dt` accessor, which is similar to `pandas.Series.str` discussed in [Chapter 14](#).

³ https://github.com/gagolews/teaching-data/raw/master/marek/birth_dates.csv




Figure 16.3. Line chart of midrange daily temperatures in Spokane for 2017–2021.

For instance, converting date-time objects to strings following custom format specifiers can be performed with:

```
dates.dt.strftime("%d.%m.%Y")
## 0    05.04.1991
## 1    14.07.2022
## 2    21.12.2042
## dtype: object
```

We can also extract different date or time *fields*, such as date, time, year, month, day, dayofyear, hour, minute, second, etc. For example:

```
dates_ymd = pd.DataFrame(dict(
    year = dates.dt.year,
    month = dates.dt.month,
    day = dates.dt.day
))
dates_ymd
##   year  month  day
## 0  1991     4    5
## 1  2022     7   14
## 2  2042    12   21
```

The other way around, we should note that `pandas.to_datetime` can convert data frames with columns named year, month, day, etc., to date-time objects:

```
pd.to_datetime(dates_ymd)
## 0    1991-04-05
```

(continues on next page)

(continued from previous page)

```
## 1    2022-07-14
## 2    2042-12-21
## dtype: datetime64[ns]
```

Example 16.4 Let's extract the month and year parts of dates to compute the average monthly temperatures it the last 50 or so years:

```
x = spokane.set_index("date").loc["1970":, ].reset_index()
mean_monthly_temps = x.groupby([
    x.date.dt.year.rename("year"),
    x.date.dt.month.rename("month")
]).temp.mean().unstack()
mean_monthly_temps.head().round(1) # preview
```

| ## month | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|----------|------|------|-----|-----|------|------|------|------|------|-----|-----|------|
| ## year | | | | | | | | | | | | |
| ## 1970 | -3.4 | 2.3 | 2.8 | 5.3 | 12.7 | 19.0 | 22.5 | 21.2 | 12.3 | 7.2 | 2.2 | -2.4 |
| ## 1971 | -0.1 | 0.8 | 1.7 | 7.4 | 13.5 | 14.6 | 21.0 | 23.4 | 12.9 | 6.8 | 1.9 | -3.5 |
| ## 1972 | -5.2 | -0.7 | 5.2 | 5.6 | 13.8 | 16.6 | 20.0 | 21.7 | 13.0 | 8.4 | 3.5 | -3.7 |
| ## 1973 | -2.8 | 1.6 | 5.0 | 7.8 | 13.6 | 16.7 | 21.8 | 20.6 | 15.4 | 8.4 | 0.9 | 0.7 |
| ## 1974 | -4.4 | 1.8 | 3.6 | 8.0 | 10.1 | 18.9 | 19.9 | 20.1 | 15.8 | 8.9 | 2.4 | -0.8 |

Figure 16.4 depicts these data on a heat map. We rediscover the ultimate truth that winters are cold, whereas in the summertime the living is easy, what a wonderful world.

```
sns.heatmap(mean_monthly_temps)
plt.show()
```




Figure 16.4. Average monthly temperatures.

16.3 Basic operations

16.3.1 Iterated differences and cumulative sums revisited

Recall from [Section 5.5.1](#) the `numpy.diff` function and its almost-inverse, `numpy.cumsum`. The former can turn a time series into a vector of *relative changes* (deltas), $\Delta_i = x_{i+1} - x_i$.

```
x = temps[-7:] # last 7 days
x
## array([-1.4, -5. , -9.4, -12.8, -12.2, -11.4, -11.4])
```

The iterated differences (deltas) are:

```
d = np.diff(x)
d
## array([-3.6, -4.4, -3.4,  0.6,  0.8,  0. ])
```

For instance, between the second and the first day of the last week, the midrange temperature dropped by -3.6°C .

The other way around, here the cumulative sums of the deltas:

```
np.cumsum(d)
## array([-3.6, -8. , -11.4, -10.8, -10. , -10. ])
```

This turned deltas back to a shifted version of the original series. But we will need the first (root) observation therefrom to restore the dataset in full:

```
x[0] + np.append(0, np.cumsum(d))
## array([-1.4, -5. , -9.4, -12.8, -12.2, -11.4, -11.4])
```

Exercise 16.5 Consider the `euraud-20200101-20200630-no-na`⁴ dataset which lists daily EUR/AUD exchange rates in the first half of 2020 (remember COVID-19?), with missing observations removed. Using `numpy.diff`, compute the minimum, median, average, and maximum daily price changes. Also, draw a box and whisker plot for these deltas.

Example 16.6 (*) The exponential distribution family is sometimes used for the modelling of times between different events (deltas). It might be a sensible choice under the assumption that a system generates a constant number of events on average and that they occur independently of each other, e.g., for the times between requests to a cloud service during peak hours, wait times for the next pedestrian to appear at a crossing near the Southern Cross Station in Melbourne, or the amount of time it takes a bank teller to interact with a customer (there is a whole branch of applied mathematics called *queuing theory* that deals with this type of modelling).

⁴ <https://github.com/gagolews/teaching-data/raw/master/marek/euraud-20200101-20200630-no-na.txt>

An exponential family is identified by the scale parameter $s > 0$, being at the same time its expected value. The probability density function of $\text{Exp}(s)$ is given for $x \geq 0$ by:

$$f(x) = \frac{1}{s} e^{-x/s},$$

and $f(x) = 0$ otherwise. We need to be careful: some textbooks choose the parametrisation by $\lambda = 1/s$ instead of s . The **scipy** package also uses this convention.

Here is a pseudorandom sample where there are five events per minute on average:

```
np.random.seed(123)
l = 60/5 # 5 events per 60 seconds on average
d = scipy.stats.expon.rvs(size=1200, scale=l)
np.round(d[:8], 3) # preview
## array([14.307,  4.045,  3.087,  9.617, 15.253,  6.601, 47.412, 13.856])
```

This gave us the wait times between the events, in seconds.

A natural sample estimator of the scale parameter is:

```
np.mean(d)
## 11.839894504211724
```

The result is close to what we expected, i.e., $s = 12$ seconds between the events.

We can convert the current sample to date-times (starting at a fixed calendar date) as follows. Note that we will measure the deltas in milliseconds so that we do not loose precision; `datetime64` is based on integers, not floating-point numbers.

```
t0 = np.array("2022-01-01T00:00:00", dtype="datetime64[ms]")
d_ms = np.round(d*1000).astype(int) # in milliseconds
t = t0 + np.array(np.cumsum(d_ms), dtype="timedelta64[ms]")
t[:8] # preview
## array(['2022-01-01T00:00:14.307', '2022-01-01T00:00:18.352',
##        '2022-01-01T00:00:21.439', '2022-01-01T00:00:31.056',
##        '2022-01-01T00:00:46.309', '2022-01-01T00:00:52.910',
##        '2022-01-01T00:01:40.322', '2022-01-01T00:01:54.178'],
##        dtype='datetime64[ms]')
t[-2:] # preview
## array(['2022-01-01T03:56:45.312', '2022-01-01T03:56:47.890'],
##        dtype='datetime64[ms]')
```

As an exercise, let's apply binning and count how many events occur in each hour:

```
b = np.arange( # four 1-hour interval (five time points)
              "2022-01-01T00:00:00", "2022-01-01T05:00:00",
              1000*60*60, # number of milliseconds in 1 hour
              dtype="datetime64[ms]"
            )
np.histogram(t, bins=b)[0]
## array([305, 300, 274, 321])
```

We expect 5 events per second, i.e., 300 of them per hour. On a side note, from a course in statistics we know that for exponential inter-event times, the number of events per unit of time follows a Poisson distribution.

Exercise 16.7 (*) Consider the `wait_times`⁵ dataset that gives the times between some consecutive events, in seconds. Estimate the event rate per hour. Draw a histogram representing the number of events per hour.

Exercise 16.8 (*) Consider the `btcusd_ohlcv_2021_dates`⁶ dataset which gives the daily BTC/USD exchange rates in 2021:

```
btc = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/btcusd_ohlcv_2021_dates.csv",
    comment="#").loc[:, ["Date", "Close"]]
btc["Date"] = btc["Date"].astype("datetime64[s]")
btc.head(12)
##          Date      Close
## 0  2021-01-01  29374.152
## 1  2021-01-02  32127.268
## 2  2021-01-03  32782.023
## 3  2021-01-04  31971.914
## 4  2021-01-05  33992.430
## 5  2021-01-06  36824.363
## 6  2021-01-07  39371.043
## 7  2021-01-08  40797.609
## 8  2021-01-09  40254.547
## 9  2021-01-10  38356.441
## 10 2021-01-11  35566.656
## 11 2021-01-12  33922.961
```

Author a function that converts it to a lagged representation, being a convenient form for some machine learning algorithms.

1. Add the `Change` column that gives by how much the price changed since the previous day.
2. Add the `Dir` column indicating if the change was positive or negative.
3. Add the `Lag1`, ..., `Lag5` columns which give the Changes in the five preceding days.

The first few rows of the resulting data frame should look like this (assuming we do not want any missing values):

```
##          Date Close  Change Dir  Lag1  Lag2  Lag3  Lag4  Lag5
## 2021-01-07 39371  2546.68 inc  2831.93  2020.52 -810.11  654.76  2753.12
## 2021-01-08 40798  1426.57 inc  2546.68  2831.93  2020.52 -810.11  654.76
## 2021-01-09 40255  -543.06 dec  1426.57  2546.68  2831.93  2020.52 -810.11
## 2021-01-10 38356 -1898.11 dec  -543.06  1426.57  2546.68  2831.93  2020.52
## 2021-01-11 35567 -2789.78 dec -1898.11  -543.06  1426.57  2546.68  2831.93
## 2021-01-12 33923 -1643.69 dec -2789.78 -1898.11 -543.06  1426.57  2546.68
```

⁵ https://github.com/gagolews/teaching-data/raw/master/marek/wait_times.txt

⁶ https://github.com/gagolews/teaching-data/raw/master/marek/btcusd_ohlcv_2021_dates.csv

In the sixth row (representing 2021-01-12), *Lag1* corresponds to *Change* on 2021-01-11, *Lag2* gives the *Change* on 2021-01-10, and so forth.

To spice things up, make sure your code can generate any number (as defined by another parameter to the function) of lagged variables.

16.3.2 Smoothing with moving averages

With time series it makes sense to consider processing whole batches of consecutive points as there is a time dependence between them. In particular, we can consider computing different aggregates inside *rolling windows* of a particular size. For instance, the *k*-moving average of a given sequence $(x_1, x_2, \dots, x_n)$ is a vector $(y_1, y_2, \dots, y_{n-k+1})$ such that:

$$y_i = \frac{1}{k} (x_i + x_{i+1} + \dots + x_{i+k-1}) = \frac{1}{k} \sum_{j=1}^k x_{i+j-1},$$

i.e., the arithmetic mean of $k \leq n$ consecutive observations starting at x_i .

For example, here are the temperatures in the last seven days of December 2011:

```
x = spokane.set_index("date").iloc[-7:, :]
x
##                temp
## date
## 2021-12-25  -1.4
## 2021-12-26  -5.0
## 2021-12-27  -9.4
## 2021-12-28 -12.8
## 2021-12-29 -12.2
## 2021-12-30 -11.4
## 2021-12-31 -11.4
```

The 3-moving (rolling) average:

```
x.rolling(3, center=True).mean().round(2)
##                temp
## date
## 2021-12-25   NaN
## 2021-12-26  -5.27
## 2021-12-27  -9.07
## 2021-12-28 -11.47
## 2021-12-29 -12.13
## 2021-12-30 -11.67
## 2021-12-31   NaN
```

We get, in this order: the mean of the first three observations; the mean of the second, third, and fourth items; then the mean of the third, fourth, and fifth; and so forth. Notice that the observations were centred in such a way that we have the same number of

missing values at the start and end of the series. This way, we treat the first three-day moving average (the average of the temperatures on the first three days) as representative of the second day.

And now for something completely different; the 5-moving average:

```
x.rolling(5, center=True).mean().round(2)
##                temp
## date
## 2021-12-25      NaN
## 2021-12-26      NaN
## 2021-12-27    -8.16
## 2021-12-28   -10.16
## 2021-12-29   -11.44
## 2021-12-30      NaN
## 2021-12-31      NaN
```

Applying the moving average has the nice effect of *smoothing* out all kinds of broadly-conceived noise. To illustrate this, compare the temperature data for the last five years in [Figure 16.3](#) to their averaged versions in [Figure 16.5](#).

```
x = spokane.set_index("date").loc["2017-01-01":, "temp"]
x30 = x.rolling(30, center=True).mean()
x100 = x.rolling(100, center=True).mean()
plt.plot(x30, label="30-day moving average")
plt.plot(x100, "r--", label="100-day moving average")
plt.legend()
plt.show()
```

Exercise 16.9 (*) Other aggregation functions can be applied in rolling windows as well. Draw, in the same figure, the plots of the one-year moving minimums, medians, and maximums.

16.3.3 Detecting trends and seasonal patterns

Thanks to windowed aggregation, we can also detect general trends (when using longish windows). For instance, let's compute the ten-year moving averages for the last 50-odd years' worth of data:

```
x = spokane.set_index("date").loc["1970-01-01":, "temp"]
x10y = x.rolling(3653, center=True).mean()
```

Based on this, we can compute the detrended series:

```
xd = x - x10y
```

Seasonal patterns can be revealed by smoothening out the detrended version of the data, e.g., using a one-year moving average:




Figure 16.5. Line chart of 30- and 100-moving averages of the midrange daily temperatures in Spokane for 2017–2021.

```
xd1y = xd.rolling(365, center=True).mean()
```

Figure 16.6 illustrates this.

```
plt.plot(x10y, label="trend")
plt.plot(xd1y, "r--", label="seasonal pattern")
plt.legend()
plt.show()
```

Also, if we know the length of the seasonal pattern (in our case, 365-ish days), we can draw a seasonal plot, where we have a separate curve for each season (here: year) and where all the series share the same x-axis (here: the day of the year); see Figure 16.7.

```
cmap = plt.colormaps.get_cmap("coolwarm")
x = spokane.set_index("date").loc["1970-01-01":].reset_index()
for year in range(1970, 2022, 5):  # selected years only
    y = x.loc[x.date.dt.year == year, :]
    plt.plot(y.date.dt.dayofyear, y.temp,
             c=cmap((year-1970)/(2021-1970)), alpha=0.3,
             label=year if year % 10 == 0 else None)
avex = x.temp.groupby(x.date.dt.dayofyear).mean()
plt.plot(avex.index, avex, "g-", label="Average")  # all years
plt.legend()
plt.xlabel("Day of year")
plt.ylabel("Temperature")
plt.show()
```




Figure 16.6. Trend and seasonal pattern for the Spokane temperatures in recent years.




Figure 16.7. Seasonal plot: temperatures in Spokane vs the day of the year for 1970–2021.

Exercise 16.10 Draw a similar plot for the whole data range, i.e., 1889–2021.

Exercise 16.11 Try using `pd.Series.dt.strftime` with a custom formatter instead of `pd.Series.dt.dayofyear`.

16.3.4 Imputing missing values

Missing values in time series can be imputed based on the information from the neighbouring non-missing observations. After all, it is usually the case that, e.g., today's weather is “similar” to yesterday's and tomorrow's.

The most straightforward ways for dealing with missing values in time series are:

- *forward-fill* – propagate the last non-missing observation,
- *backward-fill* – get the next non-missing value,
- *linearly interpolate between two adjacent non-missing values* – in particular, a single missing value will be replaced by the average of its neighbours.

Example 16.12 The classic `air_quality_1973`⁷ dataset gives some daily air quality measurements in New York, between May and September 1973. Let's impute the first few observations in the solar radiation column:

```
air = pd.read_csv("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/r/air_quality_1973.csv",
    comment="#")
x = air.loc[:, "Solar.R"].iloc[:12]
pd.DataFrame(dict(
    original=x,
    ffilled=x.ffill(),
    bffilled=x.bfill(),
    interpolated=x.interpolate(method="linear")
))
```

| | original | ffilled | bffilled | interpolated |
|-------|----------|---------|----------|--------------|
| ## 0 | 190.0 | 190.0 | 190.0 | 190.000000 |
| ## 1 | 118.0 | 118.0 | 118.0 | 118.000000 |
| ## 2 | 149.0 | 149.0 | 149.0 | 149.000000 |
| ## 3 | 313.0 | 313.0 | 313.0 | 313.000000 |
| ## 4 | NaN | 313.0 | 299.0 | 308.333333 |
| ## 5 | NaN | 313.0 | 299.0 | 303.666667 |
| ## 6 | 299.0 | 299.0 | 299.0 | 299.000000 |
| ## 7 | 99.0 | 99.0 | 99.0 | 99.000000 |
| ## 8 | 19.0 | 19.0 | 19.0 | 19.000000 |
| ## 9 | 194.0 | 194.0 | 194.0 | 194.000000 |
| ## 10 | NaN | 194.0 | 256.0 | 225.000000 |
| ## 11 | 256.0 | 256.0 | 256.0 | 256.000000 |

Exercise 16.13 (*) With the `air_quality_2018`⁸ dataset:

⁷ https://github.com/gagolews/teaching-data/raw/master/t/air_quality_1973.csv

⁸ https://github.com/gagolews/teaching-data/raw/master/marek/air_quality_2018.csv.gz

1. Based on the hourly observations, compute the daily mean PM2.5 measurements for Melbourne CBD and Morwell South.

For Melbourne CBD, if some hourly measurement is missing, linearly interpolate between the preceding and following non-missing data, e.g., a PM2.5 sequence of $[\dots, 10, \text{NaN}, \text{NaN}, 40, \dots]$ (you need to manually add the NaN rows to the dataset) should be transformed to $[\dots, 10, 20, 30, 40, \dots]$.

For Morwell South, impute the readings with the averages of the records in the nearest air quality stations, which are located in Morwell East, Moe, Churchill, and Traralgon.

2. Present the daily mean PM2.5 measurements for Melbourne CBD and Morwell South on a single plot. The x-axis labels should be human-readable and intuitive.
3. For the Melbourne data, determine the number of days where the average PM2.5 was greater than in the preceding day.
4. Find five most air-polluted days for Melbourne.

16.3.5 Plotting multidimensional time series

Multidimensional time series stored in the form of an $n \times m$ matrix are best viewed as m time series – possibly but not necessarily related to each other – all sampled at the same n points in time (e.g., m different stocks on n consecutive days).

Consider the currency exchange rates for the first half of 2020:

```
eurxxx = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/eurxxx-20200101-20200630-no-na.csv",
    delimiter=",")
eurxxx[:6, :] # preview
## array([[1.6006 , 7.7946 , 0.84828, 4.2544 ],
##        [1.6031 , 7.7712 , 0.85115, 4.2493 ],
##        [1.6119 , 7.8049 , 0.85215, 4.2415 ],
##        [1.6251 , 7.7562 , 0.85183, 4.2457 ],
##        [1.6195 , 7.7184 , 0.84868, 4.2429 ],
##        [1.6193 , 7.7011 , 0.85285, 4.2422 ]])
```

This gives EUR/AUD (how many Australian Dollars we pay for 1 Euro), EUR/CNY (Chinese Yuans), EUR/GBP (British Pounds), and EUR/PLN (Polish Złotys), in this order. Let's draw the four time series; see Figure 16.8.

```
dates = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/euraud-20200101-20200630-dates.txt",
    dtype="datetime64[s]")
labels = ["AUD", "CNY", "GBP", "PLN"]
styles = ["solid", "dotted", "dashed", "dashdot"]
for i in range(eurxxx.shape[1]):
    plt.plot(dates, eurxxx[:, i], ls=styles[i], label=labels[i])
plt.legend(loc="upper right", bbox_to_anchor=(1, 0.9)) # a bit lower
plt.show()
```





Figure 16.8. EUR/AUD, EUR/CNY, EUR/GBP, and EUR/PLN exchange rates in the first half of 2020.

Unfortunately, they are all on different scales. This is why the plot is not necessarily readable. It would be better to draw these time series on four separate plots (compare the trellis plots in [Section 12.2.5](#)).

Another idea is to depict the currency exchange rates *relative* to the prices on some day, say, the first one; see [Figure 16.9](#).

```
for i in range(eurxxx.shape[1]):
    plt.plot(dates, eurxxx[:, i]/eurxxx[0, i],
             ls=styles[i], label=labels[i])
plt.legend()
plt.show()
```

This way, e.g., a relative EUR/AUD rate of c. 1.15 in mid-March means that if an Aussie bought some Euros on the first day, and then sold them three-ish months later, they would have 15% more wealth (the Euro become 15% stronger relative to AUD).

Exercise 16.14 Based on the EUR/AUD and EUR/PLN records, compute and plot the AUD/PLN as well as PLN/AUD rates.

Exercise 16.15 (*) Draw the EUR/AUD and EUR/GBP rates on a single plot, but where each series has its own⁹ y-axis.

Exercise 16.16 (*) Draw the EUR/xxx rates for your favourite currencies over a larger period.

⁹ https://matplotlib.org/stable/gallery/subplots_axes_and_figures/secondary_axis.html




Figure 16.9. EUR/AUD, EUR/CNY, EUR/GBP, and EUR/PLN exchange rates relative to the prices on the first day.

Use [data](https://www.ecb.europa.eu/stats/policy_and_exchange_rates/euro_reference_exchange_rates/html/index.en.html)¹⁰ downloaded from the European Central Bank. Add a few moving averages. For each year, identify the lowest and the highest rate.

16.3.6 Candlestick plots (*)

Consider the BTC/USD data for 2021:

```
btcusd = np.genfromtxt("https://raw.githubusercontent.com/gagolews/" +
    "teaching-data/master/marek/btcusd_ohlcv_2021.csv",
    delimiter=",")
btcusd[:6, :4] # preview (we skip the Volume column for readability)
## array([[28994.01 , 29600.627, 28803.586, 29374.152],
##        [29376.455, 33155.117, 29091.182, 32127.268],
##        [32129.408, 34608.559, 32052.316, 32782.023],
##        [32810.949, 33440.219, 28722.756, 31971.914],
##        [31977.041, 34437.59 , 30221.188, 33992.43 ],
##        [34013.613, 36879.699, 33514.035, 36824.363]])
```

This gives the open, high, low, and close (OHLC) prices on the 365 consecutive days, which is a common way to summarise daily rates.

The `mplfinance`¹¹ (`matplotlib-finance`) package defines a few functions related to the plotting of financial data. Let's briefly describe the well-known candlestick plot.

¹⁰ https://www.ecb.europa.eu/stats/policy_and_exchange_rates/euro_reference_exchange_rates/html/index.en.html

¹¹ <https://github.com/matplotlib/mplfinance>

```

import mplfinance as mpf
dates = np.arange("2021-01-01", "2022-01-01", dtype="datetime64[D]")
mpf.plot(
    pd.DataFrame(
        btcusd,
        columns=["Open", "High", "Low", "Close", "Volume"]
    ).set_index(dates).iloc[:31, :],
    type="candle",
    returnfig=True
)
plt.show()

```




Figure 16.10. A candlestick plot for the BTC/USD exchange rates in January 2021.

Figure 16.10 depicts the January 2021 data. Let's stress that it is *not* a box and whisker plot. The candlestick body denotes the difference in the market opening and the closing price. The wicks (shadows) give the range (high to low). White candlesticks represent bullish days – where the closing rate is greater than the opening one (uptrend). Black candles are bearish (decline).

Exercise 16.17 Draw the BTC/USD rates for the entire year and add the 10-day moving averages.

Exercise 16.18 (*) Draw a candlestick plot manually, without using the *mplfinance* package. Hint: `matplotlib.pyplot.fill` might be helpful.

Exercise 16.19 (*) Using `matplotlib.pyplot.fill_between` add a semi-transparent polygon that fills the area bounded between the Low and High prices on all the days.

16.4 Further reading

Data science classically deals with information that is or can be represented in tabular form and where particular observations (which can be multidimensional) are usually independent from but still to some extent similar to each other. We often treat them as samples from different larger populations which we would like to describe or compare at some level of generality (think: health data on patients being subject to two treatment plans that we want to evaluate).

From this perspective, time series are distinct: there is some dependence observed in the time domain. For instance, a price of a stock that we observe today is influenced by what was happening yesterday. There might also be some seasonal patterns or trends under the hood. For a comprehensive introduction to forecasting; see [55, 74]. Also, for data of this kind, employing statistical modelling techniques (*stochastic processes*) can make a lot of sense; see, e.g., [90].

Signals such as audio, images, and video are different because *structured randomness* does not play a dominant role there (unless it is a noise that we would like to filter out). Instead, more interesting are the patterns occurring in the frequency (think: perceiving pitches when listening to music) or spatial (seeing green grass and sky in a photo) domain.

Signal processing thus requires a distinct set of tools, e.g., Fourier analysis and finite impulse response (discrete convolution) filters. This course obviously cannot be about everything (also because it requires some more advanced calculus skills that we did not assume the reader to have at this time); but see, e.g., [87, 89].

Nevertheless, keep in mind that these are not completely independent domains. For example, we can extract various features of audio signals (e.g., overall loudness, timbre, and danceability of each recording in a large song database) and then treat them as tabular data to be analysed using the techniques described in this course. Moreover, machine learning (e.g., convolutional neural networks) algorithms may also be used for tasks such as object detection on images or optical character recognition; see, e.g., [44].

16.5 Exercises

Exercise 16.20 Assume we have a time series with n observations. What is a 1- and an n -moving average? Which one is smoother, a $(0.01n)$ - or a $(0.1n)$ -one?

Exercise 16.21 *What is the UNIX Epoch?*

Exercise 16.22 *How can we recreate the original series when we are given its `numpy.diff`-transformed version?*

Exercise 16.23 (*) *In your own words, describe the key elements of a candlestick plot.*

Changelog

Important Any [bug/typo reports/fixes](#)¹² are appreciated. The most up-to-date version of this book can be found at <https://datawranglingpy.gagolewski.com/>.

Below is the list of the most noteworthy changes.

- **2026-... (v1.1.0.9xxx)** (in progress):
 - We now mention how own function modules can be created.
 - More programming exercises.
 - Minor extensions and bug fixes.
 - Updated to Python 3.13, **numpy** 2.3, **pandas** 2.3, **matplotlib** 3.10 (amongst others).
- **2025-02-17 (v1.1.0)**:
 - New HTML theme (includes the light and dark mode).
 - Not using **seaborn** where it can easily be replaced by a few calls to the lower-level **matplotlib**, especially in the **numpy** chapters. This way, we can learn how to create some popular charts from scratch. In particular, we are now using own functions to display a heat map and a pairs plot.
 - Use **numpy.genfromtxt** more eagerly.
 - A few more examples of using *f*-strings for results' pretty-printing.
 - Minor extensions and bug fixes.
 - Updated to Python 3.11, **numpy** 2.2, **pandas** 2.2, **matplotlib** 3.10 (amongst others).
- **2023-02-06 (v1.0.3)**:
 - Numeric reference style; updated bibliography.
 - Reduce the file size of the screen-optimised PDF at the cost of a slight decrease of the quality of some figures.

¹² <https://github.com/gagolews/datawranglingpy>

- The print-optimised PDF now uses selective rasterisation of parts of figures, not whole pages containing them. This should increase the quality of the printed version of this book.
- Bug fixes.
- Minor extensions, including: `pandas.Series.dt.strftime`, more details how to avoid pitfalls in data frame indexing, etc.
- **2022-08-24 (v1.0.2):**
 - The first printed (paperback) version can be ordered from [Amazon](https://www.amazon.com/dp/0645571911)¹³.
 - Fixed page margin and header sizes.
 - Minor typesetting and other fixes.
- **2022-08-12 (v1.0.1):**
 - Cover.
 - ISBN 978-0-6455719-1-2 assigned.
- **2022-07-16 (v1.0.0):**
 - Preface complete.
 - Handling tied observations.
 - Plots now look better when printed in black and white.
 - Exception handling.
 - File connections.
 - Other minor extensions and material reordering: more aggregation functions, `pandas.unique`, `pandas.factorize`, probability vectors representing binary categorical variables, etc.
 - Final proofreading and copyediting.
- **2022-06-13 (v0.5.1):**
 - The Kolmogorov–Smirnov Test (one and two sample).
 - The Pearson Chi-Squared Test (one and two sample and for independence).
 - Dealing with round-off and measurement errors.
 - Adding white noise (jitter).
 - Lambda expressions.
 - Matrices are iterable.
- **2022-05-31 (v0.4.1):**
 - The Rules.

¹³ <https://www.amazon.com/dp/0645571911>

- Matrix multiplication, dot products.
- Euclidean distance, few-nearest-neighbour and fixed-radius search.
- Aggregation of multidimensional data.
- Regression with k -nearest neighbours.
- Least squares fitting of linear regression models.
- Geometric transforms; orthonormal matrices.
- SVD and dimensionality reduction/PCA.
- Classification with k -nearest neighbours.
- Clustering with k -means.
- Text Processing and Regular Expression chapters merged.
- Unidimensional Data Aggregation and Transformation chapters merged.
- `pandas.GroupBy` objects are iterable.
- Semitransparent histograms.
- Contour plots.
- Argument unpacking and variadic arguments (`*args`, `**kwargs`).
- **2022-05-23 (v0.3.1):**
 - More lightweight mathematical notation.
 - Some equalities related to the mathematical functions we rely on (the natural logarithm, cosine, etc.).
 - A way to compute the most correlated pair of variables.
 - A note on modifying elements in an array and on adding new rows and columns.
 - An example seasonal plot in the time series chapter.
 - Solutions to the SQL exercises added; to ignore small round-off errors, use `pandas.testing.assert_frame_equal` instead of `pandas.DataFrame.equals`.
 - More details on file paths.
- **2022-04-12 (v0.2.1):**
 - Many chapters merged or relocated.
 - Added captions to all figures.
 - Improved formatting of elements (information boxes such as *note*, *important*, *exercise*, *example*).
- **2022-03-27 (v0.1.1):**

- The first public release: most chapters are drafted, more or less.
 - Using `Sphinx` for building.
- **2022-01-05 (v0.0.0):**
 - Project started.

References

- [1] Abramowitz, M. and Stegun, I.A., editors. (1972). *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables*. Dover Publications. URL: <https://personal.math.ubc.ca/~cbm/aands/intro.htm>.
- [2] Aggarwal, C.C. (2015). *Data Mining: The Textbook*. Springer.
- [3] Arnold, B.C. (2015). *Pareto Distributions*. Chapman and Hall/CRC. DOI: 10.1201/b18141.
- [4] Arnold, T.B. and Emerson, J.W. (2011). Nonparametric goodness-of-fit tests for discrete null distributions. *The R Journal*, 3(2):34–39. DOI: 10.32614/RJ-2011-016.
- [5] Bartoszyński, R. and Niewiadomska-Bugaj, M. (2007). *Probability and Statistical Inference*. Wiley.
- [6] Beirlant, J., Goegebeur, Y., Teugels, J., and Segers, J. (2004). *Statistics of Extremes: Theory and Applications*. Wiley. DOI: 10.1002/0470012382.
- [7] Benaglia, T., Chauveau, D., Hunter, D.R., and Young, D.S. (2009). Mixtools: An R package for analyzing mixture models. *Journal of Statistical Software*, 32(6):1–29. DOI: 10.18637/jss.v032.i06.
- [8] Bezdek, J.C., Ehrlich, R., and Full, W. (1984). FCM: The fuzzy c-means clustering algorithm. *Computer and Geosciences*, 10(2–3):191–203. DOI: 10.1016/0098-3004(84)90020-7.
- [9] Billingsley, P. (1995). *Probability and Measure*. John Wiley & Sons.
- [10] Bishop, C. (2006). *Pattern Recognition and Machine Learning*. Springer-Verlag. URL: <https://www.microsoft.com/en-us/research/people/cmbishop>.
- [11] Blum, A., Hopcroft, J., and Kannan, R. (2020). *Foundations of Data Science*. Cambridge University Press. URL: <https://www.cs.cornell.edu/jeh/book.pdf>.
- [12] Box, G.E.P. and Cox, D.R. (1964). An analysis of transformations. *Journal of the Royal Statistical Society. Series B (Methodological)*, 26(2):211–252.
- [13] Bullen, P.S. (2003). *Handbook of Means and Their Inequalities*. Springer Science+Business Media.
- [14] Campello, R.J.G.B., Moulavi, D., Zimek, A., and Sander, J. (2015). Hierarchical density estimates for data clustering, visualization, and outlier detection. *ACM Transactions on Knowledge Discovery from Data*, 10(1):5:1–5:51. DOI: 10.1145/2733381.

- [15] Chambers, J.M. and Hastie, T. (1991). *Statistical Models in S*. Wadsworth & Brooks/Cole.
- [16] Clauset, A., Shalizi, C.R., and Newman, M.E.J. (2009). Power-law distributions in empirical data. *SIAM Review*, 51(4):661–703. DOI: 10.1137/070710111.
- [17] Connolly, T. and Begg, C. (2015). *Database Systems: A Practical Approach to Design, Implementation, and Management*. Pearson.
- [18] Conover, W.J. (1972). A Kolmogorov goodness-of-fit test for discontinuous distributions. *Journal of the American Statistical Association*, 67(339):591–596. DOI: 10.1080/01621459.1972.10481254.
- [19] Cramér, H. (1946). *Mathematical Methods of Statistics*. Princeton University Press. URL: <https://archive.org/details/in.ernet.dli.2015.223699>.
- [20] Dasu, T. and Johnson, T. (2003). *Exploratory Data Mining and Data Cleaning*. John Wiley & Sons.
- [21] Date, C.J. (2003). *An Introduction to Database Systems*. Pearson.
- [22] Deisenroth, M.P., Faisal, A.A., and Ong, C.S. (2020). *Mathematics for Machine Learning*. Cambridge University Press. URL: <https://mml-book.github.io/>.
- [23] Dekking, F.M., Kraaikamp, C., Lopuhaä, H.P., and Meester, L.E. (2005). *A Modern Introduction to Probability and Statistics: Understanding Why and How*. Springer.
- [24] Devroye, L., Györfi, L., and Lugosi, G. (1996). *A Probabilistic Theory of Pattern Recognition*. Springer. DOI: 10.1007/978-1-4612-0711-5.
- [25] Deza, M.M. and Deza, E. (2014). *Encyclopedia of Distances*. Springer.
- [26] Efron, B. and Hastie, T. (2016). *Computer Age Statistical Inference: Algorithms, Evidence, and Data Science*. Cambridge University Press.
- [27] Ester, M., Kriegel, H.P., Sander, J., and Xu, X. (1996). A density-based algorithm for discovering clusters in large spatial databases with noise. In: *Proc. KDD'96*, pp. 226–231.
- [28] Feller, W. (1950). *An Introduction to Probability Theory and Its Applications: Volume I*. Wiley.
- [29] Forbes, C., Evans, M., Hastings, N., and Peacock, B. (2010). *Statistical Distributions*. Wiley.
- [30] Freedman, D. and Diaconis, P. (1981). On the histogram as a density estimator: L_2 theory. *Zeitschrift für Wahrscheinlichkeitstheorie und Verwandte Gebiete*, 57:453–476.
- [31] Friedl, J.E.F. (2006). *Mastering Regular Expressions*. O'Reilly.
- [32] Gagatewski, M. (2015). *Data Fusion: Theory, Methods, and Applications*. Institute of Computer Science, Polish Academy of Sciences. DOI: 10.5281/zenodo.6960306.

- [33] Gagolewski, M. (2015). Spread measures and their relation to aggregation functions. *European Journal of Operational Research*, 241(2):469–477. DOI: 10.1016/j.ejor.2014.08.034.
- [34] Gagolewski, M. (2021). genieclust: Fast and robust hierarchical clustering. *SoftwareX*, 15:100722. URL: <https://genieclust.gagolewski.com/>, DOI: 10.1016/j.softx.2021.100722.
- [35] Gagolewski, M. (2022). stringi: Fast and portable character string processing in R. *Journal of Statistical Software*, 103(2):1–59. URL: <https://stringi.gagolewski.com/>, DOI: 10.18637/jss.v103.i02.
- [36] Gagolewski, M. (2026). *Deep R Programming*. URL: <https://deepr.gagolewski.com/>, DOI: 10.5281/zenodo.7490464.
- [37] Gagolewski, M., Bartoszek, M., and Cena, A. (2016). *Przetwarzanie i analiza danych w języku Python (Data Processing and Analysis in Python)*. PWN. in Polish.
- [38] Gagolewski, M., Bartoszek, M., and Cena, A. (2021). Are cluster validity measures (in)valid? *Information Sciences*, 581:620–636. DOI: 10.1016/j.ins.2021.10.004.
- [39] Gentle, J.E. (2003). *Random Number Generation and Monte Carlo Methods*. Springer.
- [40] Gentle, J.E. (2009). *Computational Statistics*. Springer-Verlag.
- [41] Gentle, J.E. (2020). *Theory of Statistics*. book draft. URL: <https://mason.gmu.edu/~jgentle/books/MathStat.pdf>.
- [42] Gentle, J.E. (2024). *Matrix Algebra: Theory, Computations and Applications in Statistics*. Springer.
- [43] Goldberg, D. (1991). What every computer scientist should know about floating-point arithmetic. *ACM Computing Surveys*, 21(1):5–48. URL: <https://perso.ens-lyon.fr/jean-michel.muller/goldberg.pdf>.
- [44] Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. MIT Press. URL: <https://www.deeplearningbook.org/>.
- [45] Grabisch, M., Marichal, J.-L., Mesiar, R., and Pap, E. (2009). *Aggregation Functions*. Cambridge University Press.
- [46] Grimmett, G.R. and Stirzaker, D.R. (2020). *Probability and Random Processes*. Oxford University Press.
- [47] Gumbel, E.J. (1939). La probabilité des hypothèses. *Comptes Rendus de l'Académie des Sciences Paris*, 209:645–647.
- [48] Harris, C.R. and others. (2020). Array programming with NumPy. *Nature*, 585(7825):357–362. DOI: 10.1038/s41586-020-2649-2.
- [49] Hart, E.M. and others. (2016). Ten simple rules for digital data storage. *PLOS Computational Biology*, 12(10):1–12. DOI: 10.1371/journal.pcbi.1005097.
- [50] Hastie, T., Tibshirani, R., and Friedman, J. (2017). *The Elements of Statistical Learning*. Springer-Verlag. URL: <https://hastie.su.domains/ElemStatLearn>.

- [51] Higham, N.J. (2002). *Accuracy and Stability of Numerical Algorithms*. SIAM. DOI: 10.1137/1.9780898718027.
- [52] Hopcroft, J.E. and Ullman, J.D. (1979). *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley.
- [53] Huber, P.J. and Ronchetti, E.M. (2009). *Robust Statistics*. Wiley.
- [54] Hunter, J.D. (2007). Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3):90–95.
- [55] Hyndman, R.J. and Athanasopoulos, G. (2021). *Forecasting: Principles and Practice*. OTexts. URL: <https://otexts.com/fpp3>.
- [56] Hyndman, R.J. and Fan, Y. (1996). Sample quantiles in statistical packages. *American Statistician*, 50(4):361–365. DOI: 10.2307/2684934.
- [57] Kleene, S.C. (1951). Representation of events in nerve nets and finite automata. Technical Report RM-704, The RAND Corporation, Santa Monica, CA. URL: https://www.rand.org/content/dam/rand/pubs/research_memoranda/2008/RM704.pdf.
- [58] Knuth, D.E. (1992). *Literate Programming*. CSLI.
- [59] Knuth, D.E. (1997). *The Art of Computer Programming II: Seminumerical Algorithms*. Addison-Wesley.
- [60] Kuchling, A.M. (2023). *Regular Expression HOWTO*. URL: <https://docs.python.org/3/howto/regex.html>.
- [61] Lee, J. (2011). *A First Course in Combinatorial Optimisation*. Cambridge University Press.
- [62] Ling, R.F. (1973). A probability theory of cluster analysis. *Journal of the American Statistical Association*, 68(341):159–164. DOI: 10.1080/01621459.1973.10481356.
- [63] Little, R.J.A. and Rubin, D.B. (2002). *Statistical Analysis with Missing Data*. John Wiley & Sons.
- [64] Lloyd, S.P. (1957 (1982)). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28:128–137. Originally a 1957 Bell Telephone Laboratories Research Report; republished in 1982. DOI: 10.1109/TIT.1982.1056489.
- [65] Matloff, N.S. (2011). *The Art of R Programming: A Tour of Statistical Software Design*. No Starch Press.
- [66] McKinney, W. (2022). *Python for Data Analysis*. O'Reilly. URL: <https://wesmckinney.com/book>.
- [67] Modarres, M., Kaminskiy, M.P., and Krivtsov, V. (2016). *Reliability Engineering and Risk Analysis: A Practical Guide*. CRC Press.
- [68] Monahan, J.F. (2011). *Numerical Methods of Statistics*. Cambridge University Press.

- [69] Müllner, D. (2011). *Modern hierarchical, agglomerative clustering algorithms*. arXiv:1109.2378 [stat.ML]. URL: <https://arxiv.org/abs/1109.2378v1>.
- [70] Nelsen, R.B. (1999). *An Introduction to Copulas*. Springer-Verlag.
- [71] Newman, M.E.J. (2005). Power laws, Pareto distributions and Zipf's law. *Contemporary Physics*, pages 323–351. DOI: 10.1080/00107510500052444.
- [72] Oetiker, T. and others. (2021). *The Not So Short Introduction to LaTeX 2ε*. URL: <https://tobi.oetiker.ch/lshort/lshort.pdf>.
- [73] Olver, F.W.J. and others. (2025). *NIST Digital Library of Mathematical Functions*. URL: <https://dlmf.nist.gov/>.
- [74] Ord, J.K., Fildes, R., and Kourentzes, N. (2017). *Principles of Business Forecasting*. Wessex Press.
- [75] Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., and Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830.
- [76] Poore, G.M. (2019). Codebraid: Live code in pandoc Markdown. In: *Proc. 18th Python in Science Conf.*, pp. 54–61. DOI: 10.25080/Majora-7ddc1dd1-008.
- [77] Press, W.H., Teukolsky, S.A., Vetterling, W.T., and Flannery, B.P. (2007). *Numerical Recipes. The Art of Scientific Computing*. Cambridge University Press.
- [78] Pérez-Fernández, R., Baets, B. De, and Gagolewski, M. (2019). A taxonomy of monotonicity properties for the aggregation of multidimensional data. *Information Fusion*, 52:322–334. DOI: 10.1016/j.inffus.2019.05.006.
- [79] Rabin, M. and Scott, D. (1959). Finite automata and their decision problems. *IBM Journal of Research and Development*, 3:114–125.
- [80] Ritchie, D.M. and Thompson, K.L. (1970). QED text editor. Technical Report 70107-002, Bell Telephone Laboratories, Inc. URL: <https://wayback.archive-it.org/all/20150203071645/http://cm.bell-labs.com/cm/cs/who/dmr/qedman.pdf>.
- [81] Robert, C.P. and Casella, G. (2004). *Monte Carlo Statistical Methods*. Springer-Verlag.
- [82] Ross, S.M. (2020). *Introduction to Probability and Statistics for Engineers and Scientists*. Academic Press.
- [83] Ross, S.M. (2024). *Introduction to Probability Models*. Elsevier.
- [84] Rousseeuw, P.J., Ruts, I., and Tukey, J.W. (1999). The bagplot: A bivariate boxplot. *The American Statistician*, 53(4):382–387. DOI: 10.2307/2686061.
- [85] Rubin, D.B. (1976). Inference and missing data. *Biometrika*, 63(3):581–590.
- [86] Sandve, G.K., Nekrutenko, A., Taylor, J., and Hovig, E. (2013). Ten simple rules for reproducible computational research. *PLOS Computational Biology*, 9(10):1–4. DOI: 10.1371/journal.pcbi.1003285.

- [87] Smith, S.W. (2002). *The Scientist and Engineer's Guide to Digital Signal Processing*. Newnes. URL: <https://www.dspguide.com/>.
- [88] Spicer, A. (2018). *Business Bullshit*. Routledge.
- [89] Steiglitz, K. (1996). *A Digital Signal Processing Primer: With Applications to Digital Audio and Computer Music*. Pearson.
- [90] Tijms, H.C. (2003). *A First Course in Stochastic Models*. Wiley.
- [91] Tufte, E.R. (2001). *The Visual Display of Quantitative Information*. Graphics Press.
- [92] Tukey, J.W. (1962). The future of data analysis. *Annals of Mathematical Statistics*, 33(1):1–67. URL: <https://projecteuclid.org/journalArticle/Download?urlId=10.1214%2Faoms%2F1177704711>, DOI: 10.1214/aoms/1177704711.
- [93] Tukey, J.W. (1977). *Exploratory Data Analysis*. Addison-Wesley.
- [94] van Buuren, S. (2018). *Flexible Imputation of Missing Data*. CRC Press. URL: <https://stefvanbuuren.name/fimd>.
- [95] van der Loo, M. and de Jonge, E. (2018). *Statistical Data Cleaning with Applications in R*. John Wiley & Sons.
- [96] Venables, W.N., Smith, D.M., and R Core Team. (2025). *An Introduction to R*. URL: <https://cran.r-project.org/doc/manuals/r-release/R-intro.pdf>.
- [97] Virtanen, P. and others. (2020). SciPy 1.0: Fundamental algorithms for scientific computing in Python. *Nature Methods*, 17:261–272. DOI: 10.1038/s41592-019-0686-2.
- [98] Wainer, H. (1997). *Visual Revelations: Graphical Tales of Fate and Deception from Napoleon Bonaparte to Ross Perot*. Copernicus.
- [99] Waskom, M.L. (2021). seaborn: Statistical data visualization. *Journal of Open Source Software*, 6(60):3021. DOI: 10.21105/joss.03021.
- [100] Wickham, H. (2011). The split-apply-combine strategy for data analysis. *Journal of Statistical Software*, 40(1):1–29. DOI: 10.18637/jss.v040.i01.
- [101] Wickham, H. (2014). Tidy data. *Journal of Statistical Software*, 59(10):1–23. DOI: 10.18637/jss.v059.i10.
- [102] Wickham, H., Çetinkaya-Rundel, M., and Golemund, G. (2023). *R for Data Science*. O'Reilly. URL: <https://r4ds.hadley.nz/>.
- [103] Wierzchoń, S.T. and Kłopotek, M.A. (2018). *Modern Algorithms for Cluster Analysis*. Springer. DOI: 10.1007/978-3-319-69308-8.
- [104] Wilson, G. and others. (2014). Best practices for scientific computing. *PLOS Biology*, 12(1):1–7. DOI: 10.1371/journal.pbio.1001745.
- [105] Wilson, G. and others. (2017). Good enough practices in scientific computing. *PLOS Computational Biology*, 13(6):1–20. DOI: 10.1371/journal.pcbi.1005510.
- [106] Xie, Y. (2015). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC.